

Validation

content/generated/markdown/textbook/reading-copies/validation.md

Validation Protocols

This chapter collects the observer-level and simulation-level checks that the framework must pass if it is to retain an absolute frame without contradicting established data. Its purpose is to translate foundational claims into concrete validation tasks: null tests, atomic comparisons, and frame-proxy checks.

It should be used with [Failure Criteria](#), [Constraint Ledger](#), [No-Go Theorems](#), [Detecting the Absolute Frame](#), and [Lorentz Kinematics](#).

The note is therefore a gatekeeping document. It should tell the reader what has to be reproduced, what mechanism is being claimed, and where falsification would occur.

Validation Protocols: Preferred-Frame Leakage and Frame Proxies

Complete-State and Observational Proxies

- **Complete-state diagnostic:** The $\mathbb{U}_{\text{now}}$ universe-state perspective can use the source-tagged wake-concentricity diagnostic in [Detecting the Absolute Frame](#). This is a foundational and simulation-level bookkeeping result, not an operational laboratory protocol for Physical Observers.
- **CMB rest-frame proxy:** The CMB dipole-free frame is an empirical large-scale cosmology proxy for Noether sea rest. It is not an identification of the Euclidean-void rest frame, and it does not give Physical Observers direct access to complete source-tagged wake geometry.
- **Protocol:** Compare simulation outputs against CMB-frame observational summaries only as a large-scale consistency check for Noether sea state and cosmological transport records.

Null Tests for Absolute-Frame Drift

- **Simulation Protocol:** Run a simulated Michelson-Morley experiment through a declared Noether sea state.
- **Success Criterion:** The observer-level interference pattern must remain invariant, within the declared leakage bound, as the assembly is rotated relative to the Euclidean-void rest frame.
- **Mechanism to Verify:** Check that the Noether braid naturally contracts by γ^{-1} due to its internal architrino trajectories being compressed by motion through the Noether sea state.

Precision Atomic Comparison

- **Protocol:** Compare the predicted shift in the 1S-2S Hydrogen transition for a system moving relative to the Euclidean-void rest frame vs. one at rest.
- **Bounds:** Must match experimental lack of sidereal variation to $< 10^{-16}$.

Architrino SI Base Units

This chapter examines how the modern SI system interfaces with AAA. Its purpose is to ask which defining constants might be derivable, which remain primitive, and what kinds of constant-relations the theory should eventually explain if its geometric closure program succeeds.

It should be read together with [Parameter Ledger](#), [Angular Momentum and Spin](#), [Mapping the Planck Scale](#), [Energy](#), [Particle Masses](#), and [Atomic Spectra](#).

Executive Summary

The **2019 revision of the SI** redefined all seven base units in terms of **fixed fundamental constants**, eliminating physical artifacts. This is structurally aligned with the goal of AAA: deriving observable physics from a minimal set of substrate postulates (Euclidean void, absolute time, architrino polarity-unit magnitude $\epsilon = |e|/6$, field speed c_f).

The AAA program can potentially:

1. **Derive** the numerical values of SI-defining constants from architrino geometry
2. **Explain** why certain constants are fundamental while others are emergent
3. **Predict** relationships between constants that appear independent in the Standard Model
4. **Replace** several SI constants with a smaller set of architrino parameters

The 2019 SI Revision: What Changed

The **new SI** defines all units via **exact values** of seven constants:

Constant	Symbol	Exact Value (by definition)	Defines Unit
Hyperfine transition of Cs-133	$\Delta\nu_{\text{Cs}}$	9,192,631,770 Hz	second (s)
Speed of light in vacuum	c	299,792,458 m/s	meter (m)
Planck constant	h	$6.62607015 \times 10^{-34}$ J·s	kilogram (kg)
Elementary charge	e	$1.602176634 \times 10^{-19}$ C	ampere (A)
Boltzmann constant	k_B	1.380649×10^{-23} J/K	kelvin (K)
Avogadro constant	N_A	$6.02214076 \times 10^{23}$ mol ⁻¹	mole (mol)
Luminous efficacy of 540 THz radiation	K_{cd}	683 lm/W	candela (cd)

Key insight: These SI rows are definitions, not measurements. Their exactness is a property of the unit system. A physical closure claim still has to recover the observer-level records that make those definitions useful: spectral frequencies, charge inventories, action increments, thermal energy scales, and signal propagation.

CODATA 2022 Benchmark Discipline

The 2022 CODATA constants tables add a second layer to the SI discussion. Exact SI-defining constants, adjusted constants, and derived conversion factors should not be mixed as if they carried the same evidential status.

Class	Examples	How AAA should use it
Exact SI definitions	$c, h, e, k_B, N_A, \Delta\nu_{\text{Cs}}$	Treat as unit conventions and observer-level target scales. A derivation must recover why the same convention is stable across clocks, rulers, charges, action records, and thermodynamic records.
Adjusted dimensionless or near-direct benchmarks	$\alpha, \alpha^{-1}, m_p/m_e, \text{magnetic-moment ratios}$	Use as high-pressure residual rows because they are mostly independent of arbitrary unit scale.
Adjusted dimensional benchmarks	$G, m_e c^2, m_p c^2, m_n c^2, m_\mu c^2, R_\infty$	Use only after the substrate-to-observer unit map is declared. These rows test mass, gravity, and spectral closure, but they cannot be inserted as primitive inputs.
Derived conversion factors	$\ell_P, m_P, t_P, \text{electron volt relationships, atomic-mass relationships}$	Use as consistency checks, not independent constraints, because their uncertainties inherit the constants used to construct them.

The current numerical anchors are severe in different ways. The fine-structure constant is

$$\alpha = 7.2973525643 \times 10^{-3}, \quad u_r(\alpha) \approx 1.51 \times 10^{-10}$$

while the Newtonian constant is

$$G = 6.67430 \times 10^{-11} \text{ m}^3 \text{ kg}^{-1} \text{ s}^{-2}, \quad u_r(G) \approx 2.25 \times 10^{-5}$$

Thus α is a much sharper dimensionless target than G , while Planck-unit rows such as ℓ_P, m_P , and t_P inherit roughly half of the relative uncertainty of G through square-root dependence. A Planck-alignment claim should therefore not over-read the apparent precision of derived Planck-unit numbers.

The standard uncertainty convention matters for scoring. For a measured or adjusted row X , use

$$Z_X = \frac{X_{\text{AAA}} - X_{\text{CODATA}}}{u(X_{\text{CODATA}})}$$

when the same observable has been derived from the same record. For exact SI rows, do not form a false zero-uncertainty residual; instead test whether the unit map and the adjusted rows that depend on it close simultaneously.

AAA: Fundamental Parameters

In this framework, the candidate substrate-level quantities are:

Category A: Ontological Substrate

- **Euclidean void** (no intrinsic structure)
- **Absolute time** t (linear, forward-only parameter)
- **Field propagation speed** c_f (primitive propagation speed for causal wakes)

Category B: Fundamental Entity

- **Architrino polarity-unit magnitude** $\epsilon = |e|/6$

- **Causal wake interaction kernel** (inverse-square line-of-action weighting over causal wake surfaces, with regularized coincidence handling)

Category C: Assembly Geometry (Emergent but Calculable)

- **Nested shell braid radius ratios** (inner/middle/outer scales)
- **Maximum curvature binary radius** $r_{\text{max-curv}}$ (where $v \gg c_f$)
- **Reference Noether braid density** $\rho_{\text{NS},0}$ (the normalization scale for $n(\mathbf{x}, t)$)

Everything else (masses, coupling constants, cosmological parameters) should be **derivable** from these via:

- Self-hit dynamics (non-Markovian evolution)
- Nested shell braid stability conditions (quantization)
- Noether sea coupling (emergent metric, inertia)

Primitive-to-Derived Measure Ladder

For the units program, it is useful to distinguish primitive measures from derived ones rather than treating the SI list as a flat catalog.

- **Primitive substrate inputs:** field speed c_f , architrino polarity-unit magnitude $\epsilon = |e|/6$, absolute time ordering, and the geometric closure scales that belong to stable assemblies.
- **First-order derived measures:** characteristic time, length, action, and energy scales attached to a single stable closure problem.
- **Second-order derived measures:** area, volume, velocity ratios, densities, currents, and transport coefficients built from the first-order scales.

This ladder matters because it fixes the order of derivation. The program should first identify the minimal closure scales of the substrate and only then build compound observer-level units from them. On this reading, many SI constants are not peers inside the ontology; they are bookkeeping conventions sitting at different heights in the derivation tree.

Mapping SI Constants to Architrino Physics

The Second (Time Unit) — $\Delta\nu_{\text{Cs}}$

SI Definition: The second is defined by the hyperfine transition frequency of Cesium-133:

$$1 \text{ s} = \frac{9,192,631,770}{\Delta\nu_{\text{Cs}}}$$

Architrino Interpretation:

The hyperfine transition is caused by:

- Interaction between the **outer electron's nested shell braid** (magnetic moment from its Middle Binary orbital motion at $v \approx c_f$)
- The **nuclear spin** (magnetic moment from proton/neutron Middle Binary configurations)

This is an atomic-clock validation target, not a closed spin derivation. The electron magnetic moment, nuclear spin ledger, and hyperfine coupling must inherit [Angular Momentum and Spin](#), [Atomic Structure](#), and [Atomic Spectra](#) before $\Delta\nu_{\text{Cs}}$ can be claimed from first principles.

What we must derive:

$$\Delta\nu_{\text{Cs}} = f(\text{nested shell braid geometry, } c_f, \epsilon, \text{ Noether sea coupling})$$

Challenge: The frequency is determined by:

- The Middle Binary's orbital frequency (sets the magnetic moment)
- The coupling strength between electron and atomic nucleus (mediated by Noether sea response, with photon exchange as the observer-level channel)
- The nuclear configuration (133 nucleons = complex assembly)

Pathway:

1. Calculate the electron's Middle Binary orbital frequency ω_{MB} for Cs ground state
2. Calculate the magnetic moment $\mu = \frac{\epsilon\omega_{\text{MB}}r_{\text{MB}}}{2}$ (classical analogue)
3. Calculate the nuclear spin coupling via Noether sea-mediated potential exchange
4. Derive the splitting frequency

The Meter (Length Unit) — c

SI Definition:

$$1 \text{ m} = \frac{c}{299,792,458} \text{ seconds}$$

where c is the speed of light.

Architrino Interpretation:

The speed of light c is **not fundamental**. It is the low-gradient operational speed of photon assemblies, modeled as coaxial contra-rotating pro/anti planar pairs, propagating through the Noether sea.

Key relation:

$$c_\gamma(\mathbf{x}, t) = \frac{c_f}{\chi_\gamma(\mathbf{x}, t)}, \quad \chi_\gamma(\mathbf{x}, t) = f_\gamma(\rho_{\text{NS}}(\mathbf{x}, t), n(\mathbf{x}, t), \text{Noether sea state})$$

In the low-energy limit (flat spacetime, weak Noether sea gradients):

$$c \approx c_f \quad (\text{small corrections from Noether sea refraction})$$

What we must show:

- Photons are coaxial contra-rotating pro/anti planar pairs whose bosonic/statistical behavior is recovered as a downstream closure target
- Their propagation through the Noether sea is **not instantaneous** but limited by c_f
- The effective speed c measured by operational observers (made of assemblies) matches c_f within experimental precision ($\sim 10^{-17}$ for Lorentz tests)

Candidate deviation channels:

- In strong gravitational fields (dense Noether sea): $c_\gamma < c_f$ in the photon channel (gravitational lensing, Shapiro delay)
- At Planck scales (Noether sea microstructure resolves): $c_\gamma \neq c_f$ in the photon channel (Lorentz violation signatures)

The Kilogram (Mass Unit) — h

SI Definition:

$$1 \text{ kg} = \frac{h}{(6.62607015 \times 10^{-34}) \text{ m}^2 \text{ s}^{-1}}$$

via the Kibble balance (relating mechanical power to electromagnetic power).

Architrino Interpretation:

The Planck constant h is the observer-level benchmark for a quantum of **closed-cycle action**. The [AAA](#) derivation target is to recover this scale from nested shell braid geometry and the lower recordable basin-measure scale, not to assume it as a primitive input. In the target reduction, h is related to the radian-normalized **outer-binary rotational action** by $\hbar = h/(2\pi)$:

$$L_{\text{outer}} = n\hbar = n \frac{h}{2\pi}$$

Hypothesis:

$$\hbar \overset{\text{hyp.}}{\approx} \epsilon \cdot c_f \cdot r_{\text{outer}}, \quad h = 2\pi\hbar$$

where r_{outer} is the characteristic radius of the outer binary in the hydrogen ground-state assembly. This is an internal nested shell braid action variable, not the observer-level electron orbital angular momentum quantum number ℓ of the hydrogen $1s$ state.

Derivation pathway:

- Calculate the outer-binary radius for the hydrogen ground-state assembly (energy minimization + self-hit constraints).
- Show that closed-cycle action quantization ($\oint p dq = nh$) and the equivalent radian-normalized relation ($I = n\hbar$) arise from geometric quantization of the internal binary orbit.
- Relate h and $\hbar$ to ϵ , c_f , and nested shell braid geometry.

Target relation:

$$h \propto \epsilon \cdot c_f \cdot (\text{geometric factor from nested shell braid})$$

The Ampere (Current Unit) — e

SI Definition:

$$1 \text{ A} = \frac{e}{1.602176634 \times 10^{-19}} \text{ C/s}$$

Architrino Interpretation:

The elementary charge magnitude is recovered in the observer-level bookkeeping convention:

$$|e| = 6\epsilon$$

What we must explain:

- Why only integer multiples of ϵ appear in stable observer-level electric-charge inventories (charge quantization)
- Why we observe $0, \pm|e|/3, \pm 2|e|/3, \pm|e|$ in nature, never an isolated $\pm\epsilon$ polarity unit
- Candidate answer: **confinement or dynamical suppression**. The ϵ polarity units are bound in nested shell braids (quarks) or assemblies (leptons). Isolated $\pm\epsilon$ polarity units are not observed as stable observer-level particles, so the suppression mechanism remains a closure target rather than a completed infinite-energy theorem.

The Kelvin (Temperature Unit) — k_B

SI Definition:

$$1 \text{ K} = \frac{1.380649 \times 10^{-23} \text{ J}}{k_B}$$

Architrino Interpretation:

Boltzmann's constant k_B is the conversion factor between **energy** and **temperature**. But what **is** temperature in $\mathbb{A}\mathbb{A}\mathbb{A}$?

Hypothesis:

Temperature is the **mean kinetic energy per degree of freedom** in the Noether sea bath:

$$\langle E_{\text{kinetic}} \rangle = \frac{1}{2} k_B T$$

For a neutral Noether braid assembly in the Noether sea:

- 6 degrees of freedom (3 translational + 3 rotational)
- Mean energy: $\langle E \rangle = 3k_B T$

What we must derive:

$$k_B = f(\text{neutral Noether braid assembly mass}, c_f, \text{thermal equilibrium distribution})$$

Pathway:

1. Calculate the **effective mass** of a neutral Noether braid assembly (from nested shell braid dynamics)
2. Assume **thermal equilibrium** (Maxwell-Boltzmann distribution in the Noether sea)
3. Relate the width of the velocity distribution to $k_B T$

Derivation target:

$$k_B \propto m_{\text{NS}} \cdot c_f^2 / (\text{typical thermal velocity})^2$$

The Mole — N_A

SI Definition:

$$1 \text{ mol} = \frac{N_A}{6.02214076 \times 10^{23}} \text{ entities}$$

Architrino Interpretation:

Avogadro's constant is **not fundamental**. It's a **unit conversion factor** between atomic mass units (amu) and grams.

Relation:

$$N_A = \frac{1\text{ g}}{1\text{ amu}} = \frac{1\text{ g}}{m_{\text{proton}}/12}$$

What we must derive:

- The proton mass m_p from nested shell braid geometry (3 quarks = 3 nested shell braids + gluon wake structure + Noether sea coupling)

The Candela (Luminous Intensity) — K_{cd}

SI Definition:

$$1\text{ cd} = \frac{683}{K_{\text{cd}}} \text{ lm/W at } 540\text{ THz}$$

Architrino Interpretation:

This is a **psychophysical constant**, not a physical one. It relates:

- Physical power (photons/second)
- Human perception (brightness)

The frequency 540 THz corresponds to green light ($\lambda \approx 555\text{ nm}$), where the human eye is most sensitive.

What we can say:

- Photons at 540 THz are planar-mode phase records with $\omega = 2\pi \times 540 \times 10^{12}\text{ rad/s}$; assigning that frequency to a specific Middle Binary is still a derivation target.
- The human retina’s photoreceptors (assemblies themselves) couple resonantly to this frequency
- The constant 683 lm/W is **arbitrary**—it’s a choice of units based on human biology

Summary Table: SI Constants vs AAA Parameters

SI Constant	Status in AAA	Derivation Pathway
$\Delta\nu_{\text{Cs}}$	Derivation target (open)	Hyperfine splitting from middle-binary magnetic moments
c	Operational limit near c_f	Low-gradient photon-channel speed; deviations are encoded by χ_γ
h	Derivation target (open)	Closed-cycle action quantization; equivalent outer-binary rotational-action increments in units of $\hbar$; lower recordable basin-measure scale after quantum closure
e	Recovered observer benchmark	$ e = 6\epsilon$ after choosing the observer-level electric bookkeeping normalization
k_B	Derivation target (open)	Noether sea thermal equilibrium + assembly mass
N_A	Emergent	Follows from proton mass derivation
K_{cd}	Anthropic	Human biology; not fundamental physics

Implications: Reducing the SI to Architrino Postulates

If the AAA program succeeds, we can **replace** the seven SI-defining constants with:

Candidate Substrate Inputs (Architrino SI)

1. **Architrino polarity-unit magnitude** $\epsilon = |e|/6$ (with observer charge benchmark $|e| = 6\epsilon$)
2. **Field speed** c_f (replaces c)
3. **Nested shell braid geometry parameter** (e.g., outer radius r_{outer} or max-curvature radius) (replaces h)
4. **Neutral Noether braid assembly mass** m_{NS} (replaces k_B when combined with c_f)

Everything else is intended to be derived after closure:

- $|e| = 6\epsilon$
- $c_{\text{eff}} \rightarrow c_f$ in the low-gradient Noether sea limit
- $h^{\text{target}} \stackrel{\text{def}}{=} 2\pi\epsilon \cdot c_f \cdot r_{\text{outer}}$ after the action-closure derivation, not by definition
- $k_B = f(m_{\text{NS}}, c_f)$
- $N_A = f(m_p/m_{\text{NS}})$
- $\Delta\nu_{\text{Cs}} = f(\text{Cs nested shell braid geometry})$

Result target: If the closure program succeeds, the seven SI constants reduce to **3-4 fundamental parameters**, with the rest emergent.

Tier 1 (Must Answer)

1. Derive $\hbar$ from nested shell braid geometry

- Show that Outer Binary quantization yields $L = n\hbar$
- Calculate r_{outer} for hydrogen 1s
- Predict $\hbar$ and compare to SI value

2. Confirm $c = c_f$ within bounds

- Show photon propagation through the Noether sea matches c to $< 10^{-17}$
- Identify where/how deviations appear (Planck scale, strong gravity)

3. Derive particle masses

- First derive the calibration-free A_0 reference-attractor packet described in [Particle Masses](#).
- Use that packet to extract $E_{\text{internal}}(A_0)$, $\zeta(A_0)$, and the baseline $\mathcal{M}_{\text{sea}}^{ab}$ response map before using electron, proton, or charged-lepton data as benchmarks.
- Only after the mass-map gate is fixed, test downstream predictions such as m_e , m_p , and $m_p/m_e \approx 1836$.

Tier 2 (High Priority)

4. Calculate $\Delta\nu_{\text{Cs}}$ from first principles

- Map Cs atomic structure to Noether braid assemblies
- Derive hyperfine coupling strength
- Compare to 9,192,631,770 Hz

5. Derive k_B from Noether sea thermodynamics

- Calculate neutral Noether braid assembly effective mass
- Show thermal equilibrium reproduces Maxwell-Boltzmann
- Predict k_B value

Tier 3 (Refinement)

6. Map all SM particles to nested shell braid recipes

- Create “particle cookbook” (analogous to chemical formulas)
- Show charge, spin, statistics all emerge from geometry

7. Explain fine-structure constant α

- $\alpha = \frac{e^2}{4\pi\epsilon_0\hbar c} \approx 1/137$
- In architrino terms: $\alpha = f(\epsilon, c_f, r_{\text{outer}}, \text{Noether sea})$
- Derive numerically; explain why $\alpha \ll 1$

Philosophical Payoff

If we succeed, the **2019 SI revision** will be seen as a **halfway house**:

- It eliminated physical artifacts (kilogram prototype)
- But it enshrined **7 constants** as fundamental

The **architrino revision** completes the journey:

- It eliminates **ontological constants** (replacing them with geometric consequences)
- It reduces the foundation to **3-4 substrate parameters**
- It makes all measurements traceable to **void geometry + absolute time**

The ultimate goal: A measurement system where every quantity is expressed in terms of:

- **Lengths** (in units of $c_f \cdot t$)
- **Times** (in absolute time units)
- **Polarity units** (in units of $\epsilon = |e|/6$)

No kilograms, no kelvins, no moles—just **geometry, time, and polarity bookkeeping**, with observer units recovered above that layer.

That would be a substrate-level measurement framework, with observer units recovered as derived conventions.

Parameter Ledger

This chapter is the canonical bookkeeping page for the symbols that control closure across the current **AAA** corpus. Its purpose is not to re-derive every quantity. Its purpose is to keep the roles of primitive postulates, geometric closure targets, constitutive coefficients, state variables, and observer-level benchmarks from collapsing into one another.

The central bookkeeping rule is simple: not every symbol that appears in an equation is a free parameter. Some symbols are fixed substrate inputs, some are assembly-dependent outputs, some are constitutive functions of the Noether sea, and some are measured benchmarks that the theory is supposed to recover.

Purpose

This ledger records, for each recurrent symbol:

- what kind of object it is,
- whether it is currently treated as primitive, derived, or still open,
- which chapter owns its definition,
- and which closure program is responsible for fixing it.

That distinction matters because the corpus currently spans several layers at once:

- substrate dynamics in the Euclidean void,
- assembly geometry and delay-lock structure,
- effective spacetime constitutive maps,
- and observer-level fits to standard benchmarks.

Without a ledger, those layers can silently trade symbols back and forth as if they were interchangeable. They are not.

Status Classes

Use the following classes consistently.

- **Fundamental parameter:** part of the substrate-level postulate set.
- **Regulator / convention:** introduced for regularization, nondimensionalization, or normalization; not itself an ontological observable.
- **Geometric closure target:** should be fixed by assembly geometry, delay locking, or branch selection.
- **Constitutive closure target:** effective-medium quantity that must be extracted once and then reused across observables.
- **State variable / field:** varies over space, time, or assembly; not a single global fit constant.
- **Observable benchmark:** measured output used to test the closure map.

Canonical Guardrails

Field-speed notation

The corpus uses both v and c_f for field speed in different chapters. This ledger treats

$$c_f$$

as the canonical symbol for the physical field speed, while $v = 1$ or $c_f = 1$ denotes a chapter-local nondimensionalization convention.

Parameter versus field

The following should **not** be treated as free global constants:

- $n(\mathbf{x}, t)$,
- $\rho_{\text{NS}}(\mathbf{x}, t)$,
- $\Phi_{\text{eff}}(\mathbf{x}, t)$,
- $c_{\text{eff}}(\mathbf{x})$,
- $\chi_{\text{sea}}(\mathbf{x}, t)$,
- $m_{\text{inertial}}(A)$ for a specific assembly A .

These are state variables, constitutive fields, or derived outputs. They may be controlled by a smaller parameter set, but they are not themselves independent parameters.

Benchmark versus postulate

The following observer-level quantities are closure targets, not primitive inputs:

- e ,
- $h, \hbar$,
- G ,
- $\gamma_{\text{eff}}, \beta_{\text{eff}}, \alpha_i$,
- particle masses and electroweak angles,
- observer-level redshift and expansion summaries such as $Z_X, a(t), H(t)$, and H_{eff} .

If the theory must reset them independently for each chapter, parameter closure has failed.

CODATA Benchmark Contract

The NIST/CODATA constants tables should be used as a benchmark contract, not as an extra ontology layer. The 2022 CODATA adjustment separates three different kinds of entries:

- exact SI-defining constants, whose numerical values are fixed by unit convention;
- adjusted measured constants, whose quoted standard uncertainties are experimental and theoretical benchmark widths;
- derived conversion factors, whose uncertainty follows from the constants used to construct them.

This distinction controls how residuals are formed. If a candidate closure predicts a measured dimensionless or conversion-independent quantity X , compare it to the CODATA value by

$$Z_X = \frac{X_{\text{AAA}} - X_{\text{CODATA}}}{u(X_{\text{CODATA}})}, \quad \rho_X = \frac{X_{\text{AAA}} - X_{\text{CODATA}}}{X_{\text{CODATA}}}$$

where $u(X)$ is the quoted standard uncertainty. If X is exact by SI definition, the residual is not a measurement residual. The closure test is instead whether the same substrate-to-observer unit map recovers the exact convention while also passing the adjusted measured rows that depend on it.

The uncertainty convention is also fixed. A standard uncertainty $u(y)$ is an estimated standard deviation for the result y , and the relative standard uncertainty is

$$u_r(y) = \frac{u(y)}{|y|}$$

for $y \neq 0$. When the quoted distribution is approximately Gaussian, $y \pm u(y)$ is the one-standard-uncertainty comparison interval, not a broad tolerance band to be enlarged after a fit.

Useful 2022 CODATA rows for the current closure stack are:

Quantity	CODATA 2022 value	Standard uncertainty	Ledger role
c	$299792458 \text{ m s}^{-1}$	exact	SI convention and low-gradient photon-channel benchmark; not primitive ontology unless $c_\gamma \rightarrow c_f$ is derived.
h	$6.62607015 \times 10^{-34} \text{ J Hz}^{-1}$	exact	SI convention and action benchmark; the Planck-alignment program must recover the action scale rather than fit it.
$\hbar$	$1.054571817 \dots \times 10^{-34} \text{ J s}$	exact	Radian-normalized action benchmark derived from $h/(2\pi)$ in SI units.
e	$1.602176634 \times 10^{-19} \text{ C}$	exact	Observer-level electric-charge convention; substrate polarity bookkeeping still uses ϵ and the charge-reconstruction map.
k_B	$1.380649 \times 10^{-23} \text{ J K}^{-1}$	exact	Thermodynamic unit convention; Noether sea thermodynamics must recover the energy-temperature map.
N_A	$6.02214076 \times 10^{23} \text{ mol}^{-1}$	exact	Counting convention, not a substrate particle number.
α	$7.2973525643 \times 10^{-3}$	1.1×10^{-12}	Dimensionless electromagnetic benchmark; strong test of any charge/action/signal-speed closure.
α^{-1}	137.035999177	2.1×10^{-8}	Same benchmark in inverse form; do not count both as independent residuals.
G	$6.67430 \times 10^{-11} \text{ m}^3 \text{ kg}^{-1} \text{ s}^{-2}$	$1.5 \times 10^{-15} \text{ m}^3 \text{ kg}^{-1} \text{ s}^{-2}$	Gravity-side benchmark with comparatively weak relative uncertainty $u_r \approx 2.25 \times 10^{-5}$.
$m_e c^2$	$0.51099895069 \text{ MeV}$	$1.6 \times 10^{-10} \text{ MeV}$	Mass-map benchmark after A_0 , shielding, and response-map extraction; not an input.
$m_p c^2$	938.27208943 MeV	$2.9 \times 10^{-7} \text{ MeV}$	Hadronic mass benchmark after confinement and residual-strong closure.
$m_n c^2$	939.56542194 MeV	$4.8 \times 10^{-7} \text{ MeV}$	Neutron/proton split benchmark; tests hadronic plus electromagnetic and weak-stability bookkeeping.
$m_\mu c^2$	105.6583755 MeV	$2.3 \times 10^{-6} \text{ MeV}$	Charged-lepton hierarchy benchmark after the first mass map exists.
m_p/m_e	1836.152673426	3.2×10^{-8}	Dimensionless mass-ratio benchmark for hierarchy closure.
u	$1.66053906892 \times 10^{-27} \text{ kg}$	$5.2 \times 10^{-37} \text{ kg}$	Atomic-mass conversion benchmark for nuclear and chemistry-facing rows.

Quantity	CODATA 2022 value	Standard uncertainty	Ledger role
R_∞	$10973731.568157 \text{ m}^{-1}$	$1.2 \times 10^{-5} \text{ m}^{-1}$	Spectral benchmark binding m_e , α , h , and c in the hydrogen/atomic closure stack.
ℓ_P	$1.616255 \times 10^{-35} \text{ m}$	$1.8 \times 10^{-40} \text{ m}$	Derived Planck-unit comparison dominated by G uncertainty; not independent of h , c , G .
m_P	$2.176434 \times 10^{-8} \text{ kg}$	$2.4 \times 10^{-13} \text{ kg}$	Derived Planck-unit comparison dominated by G uncertainty; not an extra fitted mass.
t_P	$5.391247 \times 10^{-44} \text{ s}$	$6.0 \times 10^{-49} \text{ s}$	Derived Planck-time comparison dominated by G uncertainty; use only after the alignment map declares its SI conversion.

The most important mining consequence is procedural: exact rows such as h , e , k_B , and c are not easier physical targets because their listed uncertainty is zero. They are exact in SI because the units are defined through them. The physical pressure comes from the adjusted and dimensionless rows, especially α , m_p/m_e , R_∞ , particle mass-energy equivalents, and G .

Naturalness and sensitivity

When a symbol is claimed as a closure output rather than a free fit, use the fine-tuning quotient

$$\text{FTQ}(p) = \frac{\Delta p/p}{\Delta_{\text{obs}}/\text{obs}}$$

as the default sensitivity diagnostic.

Here $\Delta p/p$ is the fractional perturbation of a parameter or closure output, and $\Delta_{\text{obs}}/\text{obs}$ is the resulting fractional perturbation of the observable being tested. Values $\text{FTQ}(p) > 10$ should be treated as fine-tuning pressure unless a discrete topology, symmetry, attractor basin, or measured benchmark explains the sensitivity.

Current status:

- $\epsilon = |e|/6$ is treated as a discrete polarity-unit input and an explanatory target, not as a continuous fit.
- κ is the universal coupling in the primitive acceleration law. In the bare two-body scale closure below it combines with c_f and ϵ to set length and time units rather than an independent dimensionless tuning knob, while its primitive, derived, or normalization-sensitive status in the observer-level unit map remains open.
- $\rho_{\text{NS},0}$ and related medium-density normalizations remain naturalness risks until energy shielding and cosmological closure are quantified.

Regulator versus physical pulse

The wake-width regulator η is a computational and analytic regularization, not a claim that causal wakes are fundamentally pulsed. It smooths causal wake surfaces so integrals and simulations can be evaluated with finite resolution. As $\eta \rightarrow 0$, the intended limit is the continuous path-history law, with each discrete time step in a simulation approximating the contribution from a narrow causal wake surface rather than replacing the underlying continuous emission.

Layer-I two-body scale closure

The exact bare two-body kernel has no independent dimensionless tuning constant after the regulator is removed or treated as a numerical convention. The dimensional substrate triplet

$$(c_f, \kappa, \epsilon)$$

spans the base dimensions (L, T, Q) because

$$[c_f] = \text{L T}^{-1}, \quad [\kappa] = \text{L}^3 \text{T}^{-2} \text{Q}^{-2}, \quad [\epsilon] = \text{Q}$$

It therefore defines canonical two-body units

$$Q_* = \epsilon, \quad R_* = \frac{\kappa \epsilon^2}{c_f^2}, \quad T_* = \frac{R_*}{c_f} = \frac{\kappa \epsilon^2}{c_f^3}$$

For $\tilde{\mathbf{x}} = \mathbf{x}/R_*$, $\tilde{t} = t/T_*$, and $\tilde{q}_i = q_i/\epsilon = \pm 1$, the causal constraint and bare acceleration law reduce to

$$\tilde{r}_{ij} = \tilde{t} - \tilde{t}_0$$

and

$$\frac{d^2 \tilde{\mathbf{x}}_i}{d\tilde{t}^2} = \sum_j \sum_{\tilde{t}_0 \in \tilde{\mathcal{C}}_{ij}(\tilde{t})} \sigma_{ij} \frac{|\tilde{q}_i \tilde{q}_j|}{\tilde{r}_{ij}^2 |\tilde{J}_{ij}|} \hat{\mathbf{r}}_{ij}$$

up to the separately declared regulator ratio η/R_* when a mollified surrogate is being used.

Consequently, every dimensionless output of the isolated bare two-body problem is a pure branch-geometry result: root multiplicities, branch-birth thresholds, maximum-curvature speed ratios, residual signs, and any certified radius in units of R_* . This does not certify that a stable maximum-curvature binary exists. It says that if a certified two-body branch produces such a number, that number is computed by the root ledger and stability problem rather than fitted by changing a Layer-I dimensionless constant.

Layer I: Substrate and Kernel Parameters

These symbols belong to the delayed microscopic law itself.

ID	Symbol	Class	Current status	Meaning	Primary home
K1	c_f	Fundamental parameter	Primitive	field speed of causal wake propagation	.../dynamics/master-equation.md , .../foundations/absolute-timespace.md
K2	ϵ	Fundamental parameter	Primitive	potential polarity-unit magnitude, with observer-level electric charge reconstructed from it	.../assemblies/fermions/quantum-number-mapping.md , .../assemblies/gauge-structure-emergence.md
K3	κ	Fundamental parameter or normalization-sensitive coupling	Open as primitive/normalization split; universal in the substrate acceleration law	coupling multiplying $\sigma_{ij} q_i q_j /(r_{ij}^2 J_{ij})$ in the per-hit acceleration law; because a single architrino has no primitive inertial mass, this is not an $F = ma$ coefficient; with c_f and ϵ it sets the two-body scale $R_* = \kappa \epsilon^2 / c_f^2$ rather than a Layer-I dimensionless fit constant; dimensional row $[\kappa] = \text{L}^3 \text{T}^{-2} \text{Q}^{-2}$	.../dynamics/master-equation.md , architrino-si-base-units.md , .../foundations/architrino.md
K4	η	Regulator / convention	Open but non-ontological	mollifier width used to regularize causal wake surfaces for smooth dynamics and numerics	simulations/action-energy/well-posedness-and-regularization.md , .../dynamics/master-equation.md
K5	Z_e	Regulator / convention	Convention, default $Z_e = 1$	coarse-graining / normalization factor in the substrate-to-observer charge map	.../assemblies/gauge-structure-emergence.md , .../assemblies/fermions/quantum-number-mapping.md

Layer II: Assembly-Geometry Closure Targets

These quantities belong to Noether braid architecture, shielding, branch structure, and assembly response.

ID	Symbol	Class	Current status	Meaning	Primary home
G0	A_0	Geometric closure target	Open	calibration-free neutral rest-branch Noether braid reference attractor used to derive the first mass-map outputs before particle benchmarks enter	Particle Masses , Nested Shell Braid Dynamics , Energy
G0a	$\mathcal{P}_{A_0}$	Geometric closure target	Open; compact finite-coordinate no-go recorded, branch-chart revision required before Tier 1 continuation	certificate packet tying the finite closure graph $\mathcal{G}_{A_0}$, active root ledger, quotient Floquet gap Δ_k , shielding extraction, and $\mathcal{M}_{\text{sa}}^{\text{ab}}$ response probe into one promotion sequence	simulations/a0-branch-certificate-protocol.md , simulations/a0-tier0-result-interpretation.md , .../assemblies/particle-masses.md
G1	$R_{\text{inner}}, R_{\text{middle}}, R_{\text{outer}}$	Geometric closure target	Open	characteristic radii of the nested binaries in the Noether braid	Noether Braid , Nested Shell Braid Geometry , Nested Shell Braid Dynamics
G2	$\omega_{\text{inner}}, \omega_{\text{middle}}, \omega_{\text{outer}}$	Geometric closure target	Open	characteristic binary frequencies associated with the nested radii	Nested Shell Braid Dynamics , Particle Masses
G3	R_{align}	Geometric closure target	Open, conjectural	outer-binary alignment radius in the terminal Planck-alignment map	Mapping the Planck Scale to the Nested Shell Braid Geometry
G4	$\mathcal{A}_{\text{align}}^{\text{cycle}}, I_{\text{align}}$	Geometric closure target	Open, conjectural	closed-cycle action and radian-normalized rotational-action increment of the aligned terminal mode	Mapping the Planck Scale to the Nested Shell Braid Geometry
G5	$\zeta(A)$	Geometric closure target	Open	shielding or leakage factor of assembly A , defined by far-field suppression relative to naive constituent exposure	.../dynamics/energy.md , .../assemblies/particle-masses.md
G6	α	Geometric closure target	Open	axial-frame misalignment angle used in the weak-mixing / quark-geometry program	.../assemblies/fermions/weak-mixing-angle.md
G7	ϕ_c	Geometric closure target	Open	color-sector azimuth selecting the exceptional axial-frame orientation	.../assemblies/fermions/weak-mixing-angle.md

Layer III: Constitutive Spacetime Parameters

These symbols control the handoff from the Euclidean substrate plus Noether sea to effective metric language.

ID	Symbol	Class	Current status	Meaning	Primary home
C1	$\rho_{NS,0}$	Constitutive closure target	Open	reference Noether braid density used to normalize the Noether sea	.../spacetime/emergent-metric.md , .../spacetime/proper-time-and-time-dilation.md
C2	$n(\mathbf{x}, t)$	State variable / field	Derived field	normalized Noether braid density, $n = \rho_{NS} / \rho_{NS,0}$	.../spacetime/emergent-metric.md , .../spacetime/proper-time-and-time-dilation.md
C3	$\Omega(\mathbf{x}), \xi(\mathbf{x})$	Constitutive closure target	Open	clock-channel and ruler-channel response functions in the effective metric subclass	.../spacetime/emergent-metric.md , .../spacetime/lorentz-kinematics.md
C4	$\Phi_{\text{eff}}(\mathbf{x}, t)$	State variable / field	Derived field	constitutive effective potential defined from the clock channel	.../spacetime/emergent-metric.md , .../spacetime/proper-time-and-time-dilation.md
C5	$c_{\text{eff}}(\mathbf{x}, t)$	State variable / field	Derived field	Noether sea dressed assembly-channel propagation speed used for clock/ruler closure and effective-metric comparisons, with $c_{\text{eff}} \rightarrow c_f$ in weak homogeneous conditions; separate from photon-channel speed c_γ unless Gate A closes that identification	.../spacetime/emergent-metric.md , .../spacetime/ppn-parameters.md
C5a	$\chi_{\text{sea}}(\mathbf{x}, t)$	Derived response field	Derived from c_{eff}	Noether sea delay factor, $\chi_{\text{sea}} = c_f / c_{\text{eff}}$; replaces optical refractive-index notation in Noether sea propagation maps	.../spacetime/noether-sea.md , .../spacetime/emergent-metric.md , .../spacetime/ppn-parameters.md
C6	γ_{eff}	Constitutive closure target with observable meaning	Open	first-order refraction / space-curvature coefficient in the weak-field map	.../spacetime/ppn-parameters.md
C7	C_2 or β_{eff}	Constitutive closure target with observable meaning	Open	second-order clock-channel nonlinearity entering the g_{00} expansion	.../spacetime/ppn-parameters.md
C8	$\Xi_1, \Xi_2, \Xi_3, \Xi_4$	Constitutive closure target	Open	preferred-frame leakage coefficients in the weak-field constitutive expansion	.../spacetime/ppn-parameters.md
C9	$\mathcal{M}_{\text{sea}}^{ab}$	Constitutive closure target	Open	medium-response tensor that maps shielded internal assembly energy to inertial momentum response, reducing to $h^{ab} / c_{\text{eff}}^2$ in a homogeneous isotropic Noether sea cell	.../dynamics/energy.md , .../assemblies/particle-masses.md

Layer IV: Observer-Level Benchmarks and Derived Outputs

These quantities are where closure is tested. They are not substrate inputs.

ID	Symbol	Class	Current status	Meaning	Primary home
O1	e	Observable benchmark	Derived target	elementary charge reconstructed from substrate charge and normalization map	.../assemblies/fermions/quantum-number-mapping.md , .../assemblies/gauge-structure-emergence.md
O2	$h, \hbar$	Observable benchmark / geometric target	Open	full-cycle action quantum and radian-normalized angular-momentum quantum to be related to nested shell braid alignment, orbital closure, and any lower recordable basin-measure scale derived by quantum closure	Angular Momentum and Spin, Mapping the Planck Scale to the Nested Shell Braid Geometry, Architrino SI Base Units
O3	G or G_{eff}	Observable benchmark / constitutive target	Open	effective gravitational coupling emerging from medium compliance and alignment geometry	Mapping the Planck Scale to the Nested Shell Braid Geometry, Emergent Metric
O4	$m_{\text{inertial}}(A)$	Derived output	Open	inertial mass of assembly A , extracted operationally from shielding and medium response	.../dynamics/energy.md , .../assemblies/particle-masses.md
O5	θ_W^{bare} and θ_W	Geometric target / observable benchmark	Open	bare geometric weak-mixing increment and the measured electroweak mixing angle it must eventually inform	.../assemblies/fermions/weak-mixing-angle.md , .../assemblies/gauge-structure-emergence.md
O6	$(\alpha_1, \alpha_2, \alpha_3)$	Observable benchmark	Open	standard PPN preferred-frame coefficients derived from (Ξ_1, Ξ_2, Ξ_3)	.../spacetime/ppn-parameters.md
O7	$Z_X^{E \rightarrow R}, Y_{X,E \rightarrow R}$, and $H_{\text{eff},X}$	Observer-level derived output	Open	total signed photon-frequency transfer, path-history exchange contribution, and inferred redshift-transfer slope for a declared source/receiver record; not primitive expansion parameters	.../cosmology/expansion-mechanism.md , simulations/redshift-budget-toy-model.md , reaction-cosmology-provenance-ledger.md

Canonical Relations

The ledger above is only useful if the interfaces between layers stay explicit. The following relations are the current canonical handoff points in the corpus.

1. Microscopic delayed dynamics

The regularized exact law uses the kernel-side set

$$(c_f, \epsilon, \kappa, \eta)$$

A representative regularized form is

$$\mathbf{a}_a(t) = \sum_b \kappa \sigma_{ab} |q_a q_b| \int_{-\infty}^t dt_0 \frac{\hat{\mathbf{r}}_{ab}(t; t_0)}{r_{ab}(t; t_0)^2} \delta_\eta(r_{ab}(t; t_0) - c_f(t - t_0))$$

This is the substrate-side parameter core. Any exact or numerical closure that changes these symbols chapter by chapter is not a closed theory.

2. Charge reconstruction

The current substrate-to-observer charge map is

$$|e| = 6\epsilon\sqrt{\kappa c_f} Z_e$$

with canonical normalization choice

$$Z_e = 1$$

This relation is important because it shows that the elementary charge magnitude is not presently a primitive input in the architrino ontology. It is a recovered observer-level benchmark.

This equation is a normalization-sensitive substrate-to-observer reconstruction, not a second primitive definition of ϵ . The primitive polarity convention used by the foundation and mathematics-canon pages is $\epsilon = |e|/6$ after the observer-level electric bookkeeping normalization is fixed. If $\sqrt{\kappa c_f} Z_e \neq 1$ in a dimensional convention, that factor belongs to the conversion map rather than to a separate architrino charge ontology. Closing this equivalence remains tied to the K3/K5 normalization problem.

3. Medium normalization and clock-channel potential

The constitutive spacetime layer uses

$$\rho_{\text{NS}}(\mathbf{x}, t) = \rho_{\text{NS},0} n(\mathbf{x}, t)$$

and

$$\Phi_{\text{eff}}(\mathbf{x}) = c_f^2 \ln(\Omega(\mathbf{x})\xi(\mathbf{x}))$$

Here ξ is the Noether braid envelope shape ratio, while $\Omega\xi$ is the clock-rate factor used by this exponential metric subclass after the geometry-to-clock map is fixed.

This is the cleanest current statement of the Noether sea-to-metric handoff:

$$(\delta_{ij}, n, \chi_{\text{sea}}, \Phi_{\text{eff}}, \text{stress}) \mapsto g_{\mu\nu}^{\text{eff}}$$

4. Weak-field PPN extraction

The observable weak-field coefficients are read from the constitutive map through

$$\chi_{\text{sea}}(\mathbf{x}) \equiv \frac{c_f}{c_{\text{eff}}(\mathbf{x})} = 1 - (1 + \gamma_{\text{eff}}) \frac{\Phi_N(\mathbf{x})}{c_f^2} + \mathcal{O}\left(\frac{\Phi_N^2}{c_f^4}\right)$$

and

$$\beta_{\text{eff}} = \frac{1 + 2C_2}{2}$$

Preferred-frame leakage is encoded by

$$\alpha_1 = \Xi_1, \quad \alpha_2 = \Xi_2, \quad \alpha_3 = \Xi_1 - \Xi_2 - \Xi_3$$

The zero-leakage closure condition is therefore

$$\Xi_1 = \Xi_2 = \Xi_3 = \Xi_4 = 0 \quad \Longleftrightarrow \quad \alpha_1 = \alpha_2 = \alpha_3 = 0$$

5. Mass map

The current assembly-side inertial map is

$$m_{\text{inertial}}(A) \approx \alpha_m \frac{\zeta(A) E_{\text{internal}}(A)}{c_{\text{eff}}^2}$$

with α_m fixed once by a reference assembly rather than re-fit separately for each particle.

This relation means that $m_{\text{inertial}}(A)$ is not a primitive parameter. It is an output of shielding, internal energy, and medium response.

Notation note: α_m denotes the mass-map normalization; bare α remains reserved for the measured fine-structure benchmark or for a locally declared weak-mixing branch angle, while α_i denotes PPN preferred-frame coefficients.

In a resolved Noether sea environment, this scalar relation is the homogeneous isotropic limit of the tensor response

$$p_{\text{int}}^a \approx \alpha_m \zeta(A) E_{\text{internal}}(A) \mathcal{M}_{\text{sea}}^{ab} V_{\text{cm},b}, \quad \mathcal{M}_{\text{sea}}^{ab} \rightarrow \frac{h^{ab}}{c_{\text{eff}}^2}$$

Here h^{ab} is the inverse Euclidean spatial metric on the local substrate slice. The tensor $\mathcal{M}_{\text{sea}}^{ab}$ is not a particle-specific fit parameter. It is a constitutive closure target for the Noether sea response map.

The first closure gate for this relation is the reference attractor A_0 . It must report geometry, winding, root-ledger, stability, internal-energy, shielding, medium-response, and mass-facing outputs before any observed particle mass or charged-lepton ratio is used as a benchmark. The mass-facing output is the calibration-free combination

$$\frac{\zeta(A_0) E_{\text{internal}}(A_0)}{E_0}$$

together with the unresolved constants and response-map assumptions needed to turn that dimensionless coefficient into an observer-level mass prediction.

The current compact finite-coordinate no-go is a status blocker inside $\mathcal{P}_{A_0}$, not an additional free parameter and not a benchmark input. It requires a predeclared branch-chart revision before Tier 1 continuation can be interpreted as progress toward the mass-facing output above.

6. Planck-alignment map

The current Planck-scale program uses the conjectural relations

$$\mathcal{A}_{\text{align}}^{\text{cycle}} \overset{\text{hyp.}}{\approx} h, \quad I_{\text{align}} \overset{\text{hyp.}}{\approx} \hbar, \quad 2\pi R_{\text{align}} = \ell_P$$

and the effective gravity-side alignment estimate

$$G_{\text{eff}} \equiv \frac{R_{\text{align}}^2 c_f^3}{\mathcal{A}_{\text{align}}^{\text{cycle}}}$$

These are not yet closed derivations. They are the current alignment-side targets connecting geometric closure to (h, G) .

7. Weak-mixing branch structure

The weak-mixing geometry note uses

$$\sin^2 \theta_W^{\text{bare}} = \frac{1}{4}, \quad \theta_W^{\text{bare}} = 30^\circ$$

and the discrete axial-frame branch hypothesis

$$\alpha_n = n \theta_W^{\text{bare}}$$

This means the present quark-sector use of α is a geometric branch label tied to a candidate bare electroweak increment, not yet a finished derivation of the measured weak angle.

What Is Not Yet Closed

The current corpus supports the following conservative closure assessment.

Closed enough to treat as canonical

- c_f is treated consistently as the substrate propagation speed, even when chapters temporarily write $v = 1$.
- ϵ is treated consistently as the potential polarity-unit magnitude.
- The exact bare two-body kernel admits the canonical nondimensionalization by $R_* = \kappa \epsilon^2 / c_f^2$ and $T_* = R_* / c_f$, so branch thresholds and residual equations are parameter-free once a branch chart is declared.
- $\rho_{\text{NS},0}$ is the reference density symbol for the Noether sea.
- $\Phi_{\text{eff}} = c_f^2 \ln(\Omega \xi)$ is the canonical clock-channel potential definition for the exponential metric subclass, with ξ retained as a geometry-first Noether braid shape ratio.

Still genuinely open

- whether κ is primitive, derived, or partly a normalization artifact,
- whether η should disappear entirely from physical statements after the weak limit is taken,
- whether any specific maximum-curvature binary branch exists and is stable under the full signed-root, finite-window two-body dynamics,
- the A_0 reference-attractor output packet,
- the actual nested shell braid radii/frequency ladder,
- the shielding map $\zeta(A)$ across the fermion spectrum,
- the medium-response tensor $\mathcal{M}_{\text{sea}}^{ab}$ that turns shielded internal energy into inertial and gradient response,
- the constitutive functions (Ω, ξ) and the weak-field coefficient set $(\gamma_{\text{eff}}, C_2, \Xi_i)$,
- the Planck-alignment identification of $(R_{\text{align}}, \mathcal{A}_{\text{align}}^{\text{cycle}}, I_{\text{align}}, h, \hbar, G)$,
- and the reduction of weak-mixing branch labels to a predictive electroweak closure.

Immediate Parameter-Closure Priorities

The shortest path to a better closure score is:

1. Fix the observer-level status of κ once, with an explicit statement of what part is physical coupling, what part is absorbed normalization, and how the two-body scale $R_* = \kappa \epsilon^2 / c_f^2$ enters the unit map.
2. Derive or numerically extract a reusable constitutive parameterization for (Ω, ξ) , then hold it fixed across redshift, Shapiro delay, lensing, and preferred-frame tests.
3. Resolve the A_0 branch-chart revision and accepted branch packet, then replace symbolic shielding language with an operational $\zeta(A)$ extraction protocol and a reusable $\mathcal{M}_{\text{sea}}^{ab}$ response map that can be applied to electron, quark, and neutrino assemblies without redefinition.
4. Decide whether the Planck-alignment map yields (h, G) as true outputs or only as analogy-level scaling relations.
5. Reduce the weak-mixing angle program from discrete branch suggestion to an actual minimization problem for $E_{\text{eff}}(\alpha, \phi_e)$.

Falsification Gate

Parameter closure fails if any of the following occurs:

- a symbol advertised as fundamental changes meaning across chapters,
- a constitutive coefficient must be re-fit independently for different observable classes,
- a state field such as $n(\mathbf{x}, t)$ is implicitly treated as a free global constant to rescue a calculation,
- or observer-level benchmarks such as e , G , or particle masses are matched only by introducing one-off per-sector normalizations.

In compact form, the closure target is a nonempty shared parameter set

$$\mathcal{P}_{\text{shared}} \neq \emptyset$$

where $\mathcal{P}_{\text{shared}}$ is the common substrate-plus-constitutive set that survives particle, spacetime, and quantum-side tests simultaneously.

Related Chapters

- [constraint-ledger.md](#)
- [architrino-si-base-units.md](#)
- [.../dynamics/master-equation.md](#)
- [.../dynamics/energy.md](#)
- [.../philosophy-history/theory-bridges/angular-momentum-and-spin.md](#)
- [Mapping the Planck Scale to the Nested Shell Braid Geometry](#)
- [.../spacetime/emergent-metric.md](#)
- [.../spacetime/ppn-parameters.md](#)
- [.../assemblies/fermions/weak-mixing-angle.md](#)

Reaction Ledger

This ledger records how reaction channels should account for constituent architrinos, Noether braids, axial layers, energy, momentum, charge, polarity, and path-history provenance. Its purpose is not to replace Standard Model reaction notation. Its purpose is to state what an $\Delta\Delta\Delta$ interpretation must conserve before a reaction map can be treated as more than a provisional diagram.

For radiative channels, use this ledger together with [Radiation](#). For cosmology-facing radiation and thermalization channels, use it together with [Reaction-Cosmology Provenance Ledger](#).

Scope and Status

Reaction provenance is a closure target. A channel may use standard observer notation such as $d \rightarrow u + W^-$ or $\gamma + \gamma \rightarrow e^+ + e^-$, but the AAA map is not closed until the underlying constituent ledger is explicit.

The conservative status is:

- Architrino count and polarity conservation are required constraints.
- Noether sea participation is allowed, but it must be recorded as a reactant, product reservoir, or medium-excitation channel rather than left implicit.
- W, Z, photon, and pair-production language may be retained at observer level, while the substrate map must identify the transient assembly, exchanged payload, or planar-mode nucleation event being invoked.
- Radiative, photon-capture, and sub-threshold shedding entries must attach the shared radiation event-record schema: source assembly, trigger geometry, $\delta\Theta_a$, E_{exc} , E_γ , recoil, medium excitation, polarization handoff, causal-wake ledger, and closure status.
- Any weak-channel ledger that depends on chirality, axial-frame orientation, CKM/PMNS mixing, or antineutrino routing remains provisional until the corresponding geometry is derived.
- Any reaction-level spin, helicity, polarization, or vector-channel angular-momentum entry is a downstream consumer of the angular-momentum and spin workstream. It should record what must close, not function as a local proof of that closure.

Provenance Protocol

Each reaction record should state:

1. **Observer channel:** the standard reaction label, including historical labels such as beta decay only when immediately translated into native reaction language.
2. **Active assemblies:** which incoming assemblies actually reconfigure, and which are spectators.
3. **Noether sea participation:** whether local Noether braids, neutral binaries, axial layers, or medium excitations are consumed, split, reconfigured, or returned.
4. **Constituent inventory:** total E and P counts before and after, separated into braid and axial-layer contributions where the distinction matters.
5. **Polarity and charge accounting:** how observer-level charge bookkeeping emerges from the E/P routing.
6. **Energy-momentum and angular-momentum accounting:** where kinetic energy, internal binding energy, photon assemblies, recoil, medium excitation, spin/vector ledger terms, and wake-carried angular momentum enter and exit.
7. **Path-history provenance:** which emitted causal wakes, source identities, and delayed interactions are needed to make the reaction deterministic in absolute time.
8. **Radiation event record, when applicable:** for emitted, absorbed, shifted, captured, or failed photon channels, attach the shared event fields from [Radiation](#), including E_{exc} , E_γ , recoil, medium excitation, polarization handoff, and causal-wake ledger.
9. **Hybrid Standard Model matching, when applicable:** identify the source lane for the observer-level prediction: perturbative electroweak chart, matched weak effective theory, lattice-QCD or nuclear matrix element, infrared-safe QCD observable, QED, kinetic model, or detector functional. Include the scheme, operator or observable definition, matching normalization, CKM/PMNS factor when applicable, expansion or scaling parameter, systematic remainder, and regulator-removal or continuum record when one is used.
10. **Closure status:** baseline, provisional map, derivation target, failed map, or inherited gate.

Record Template

Field	Required content
Observer channel	Standard reaction notation and native reaction label
Active assembly change	Braid and axial-layer changes for the transformed assembly
Noether sea input/output	Neutral braids, axial material, or medium excitations recruited or returned
Conserved inventory	E/P totals and charge/polarity balance
Energy-momentum and angular-momentum ledger	Internal energy, recoil, emitted assemblies, spin/vector ledger terms, wake-carried angular momentum, and medium excitation
Radiation event record, when applicable	Source assembly, source-depletion row, trigger geometry, $\delta\Theta_a$, E_{exc} , E_γ , recoil, medium excitation, polarization handoff, causal-wake ledger, photon Gate B event residual when $E_\gamma \neq 0$, and closure status
Provenance data	Source identity, emission time, causal-root branch, and local Noether sea state
Hybrid Standard Model matching, when applicable	Source lane, scheme, operator or observable, matching normalization, CKM/PMNS factor when applicable, matrix-element or factorization source, expansion or scaling parameter, systematic remainder, and regulator-removal or continuum record
Closure status	What is established, what is assumed, and what remains to derive

Residual-Routing Event-Ledger Contract

Residual-routing material enters this ledger only as a theorem-target contract. It does not by itself prove that any weak, radiative, pair-production, nuclear, or cosmology-facing reaction channel has closed. The common target is:

$$\mathcal{R}(\Gamma, \mathcal{H}, \rho_{\text{NS}}, \chi_{\text{sea}}, \dots) \longrightarrow \{B_i\} \longrightarrow \mathcal{L}_{E\mathbf{p}\mathbf{J}}$$

Here $\mathcal{R}$ is the replayable residual computed from the local assembly state, path-history ledger, Noether braid density, Noether sea delay factor, and any named sector variables. The set $\{B_i\}$ is the finite list of admissible output channels, such as retuning, bound excitation, radiation, recoil, medium heating, weak or nuclear reaction, record formation, release channel, or branch transition. The event ledger $\mathcal{L}_{E\mathbf{p}\mathbf{J}}$ is the balance object that must close after all selected outputs are named.

For a reaction attempt, the input state should be recorded as:

$$X = (\Gamma, \mathcal{H}, \rho_{\text{NS}}(\mathbf{x}, t), \chi_{\text{sea}}(\mathbf{x}, t), Z_S)$$

where Z_S denotes sector-local variables such as nuclear configuration, weak-corridor data, apparatus state, or horizon-interface boundary data when those variables control the route. A routed reaction event is a triple

$$\mathbf{e} = (X, I_e, Y_e)$$

where I_e is the selected finite channel set and Y_e lists outgoing assemblies, recoil targets, medium updates, remnant states, and provenance records. A single reaction vertex may select more than one output channel when photon output, recoil, medium update, and reaction products are simultaneous terms in one closed event.

The shared ledger object is:

$$\mathcal{L}_{E\mathbf{p}\mathbf{J}}(\mathbf{e}) = (\Delta_E, \Delta_{\mathbf{p}}, \Delta_{\mathbf{J}}, \Delta_{\text{pol}}, \Delta_{\text{arch}}, \Delta_{\text{path}}, \Delta_{\text{med}}, \Delta_{\text{rem}})(\mathbf{e})$$

Ledger closure means:

$$\mathcal{L}_{E\mathbf{p}\mathbf{J}}(\mathbf{e}) = \mathbf{0}$$

componentwise across the tuple. Nonzero physical recoil, medium heating, remnant excitation, outgoing product energy, or photon output is allowed only as a named term inside Y_e ; it is not allowed as an implicit loss.

The stronger event-balance target bundles energy, momentum, and angular momentum instead of checking photon polarization separately from the source ledger. For $\mathcal{Q} \in \{E, \mathbf{p}, \mathbf{J}\}$, define source depletion by

$$\Delta \mathcal{Q}_{\text{src}}^0 = \mathcal{Q}_{\text{src}}^- - \mathcal{Q}_{\text{src}}^+$$

A resolved radiative event closes only if

$$\Delta \mathcal{Q}_{\text{src}}^0 = \mathcal{Q}_{\gamma}^{\text{sub}} + \mathcal{Q}_{\text{recoil}}^0 + \mathcal{Q}_{\text{med}}^0 + \mathcal{Q}_{\text{wake}}^0 + \mathcal{Q}_{\text{handoff}}^0 + \mathcal{Q}_{\text{rem}}^0$$

with normalized residual

$$\Delta_{\text{evt}}^{\gamma} = \sum_{\mathcal{Q} \in \{E, \mathbf{p}, \mathbf{J}\}} \frac{\|\Delta \mathcal{Q}_{\text{src}}^0 - \mathcal{Q}_{\gamma}^{\text{sub}} - \mathcal{Q}_{\text{recoil}}^0 - \mathcal{Q}_{\text{med}}^0 - \mathcal{Q}_{\text{wake}}^0 - \mathcal{Q}_{\text{handoff}}^0 - \mathcal{Q}_{\text{rem}}^0\|}{\varepsilon_{\mathcal{Q}} + \|\Delta \mathcal{Q}_{\text{src}}^0\|}$$

The Gate B angular-momentum row is the $\mathcal{Q} = \mathbf{J}$ projection of this same identity. Let the event window be labeled by superscript 0, and let $\mathbf{J}_{\text{src}}^-$ and $\mathbf{J}_{\text{src}}^+$ be the source angular-momentum ledger before and after the event. Define

$$\Delta \mathbf{J}_{\text{src}}^0 = \mathbf{J}_{\text{src}}^- - \mathbf{J}_{\text{src}}^+$$

The photon event row is

$$\Delta \mathbf{J}_{\text{src}}^0 = \mathbf{J}_{\gamma}^{\text{sub}} + \mathbf{J}_{\text{recoil}}^0 + \mathbf{J}_{\text{med}}^0 + \mathbf{J}_{\text{wake}}^0 + \mathbf{J}_{\text{handoff}}^0 + \mathbf{J}_{\text{rem}}^0$$

Define the corresponding balance defect by

$$\mathbf{B}_{\gamma}^0 = \Delta \mathbf{J}_{\text{src}}^0 - \mathbf{J}_{\gamma}^{\text{sub}} - \mathbf{J}_{\text{recoil}}^0 - \mathbf{J}_{\text{med}}^0 - \mathbf{J}_{\text{wake}}^0 - \mathbf{J}_{\text{handoff}}^0 - \mathbf{J}_{\text{rem}}^0$$

For a Gate B-admissible photon row, helicity is the projection

$$\lambda_{\text{hel}} = \frac{\hat{\mathbf{e}} \cdot \mathbf{J}_{\gamma}^{\text{sub}}}{\hbar}, \quad \lambda_{\text{hel}} \in \{+1, -1\}$$

and the event balance bounds the projection error:

$$\left| \frac{\hat{\mathbf{e}} \cdot \mathbf{J}_\gamma^{\text{sub}}}{\hbar} - \frac{\hat{\mathbf{e}} \cdot (\Delta \mathbf{J}_{\text{src}}^0 - \mathbf{J}_{\text{recoil}}^0 - \mathbf{J}_{\text{med}}^0 - \mathbf{J}_{\text{wake}}^0 - \mathbf{J}_{\text{handoff}}^0 - \mathbf{J}_{\text{rem}}^0)}{\hbar} \right| \leq \frac{\|\mathbf{B}_\gamma\|}{\hbar}$$

The normalized event-balance residual is

$$\Delta_{\text{bal}}^\gamma = \frac{\|\Delta \mathbf{J}_{\text{src}}^0 - \mathbf{J}_\gamma^{\text{sub}} - \mathbf{J}_{\text{recoil}}^0 - \mathbf{J}_{\text{med}}^0 - \mathbf{J}_{\text{wake}}^0 - \mathbf{J}_{\text{handoff}}^0 - \mathbf{J}_{\text{rem}}^0\|}{1 + \|\Delta \mathbf{J}_{\text{src}}^0\|}$$

The denominator is understood in the normalized angular-momentum units of the event ledger. Missing source, recoil, medium, wake, handoff, or remnant rows keep the photon record provisional even when the outgoing photon substrate ledger is algebraically clean.

Provenance-Preserving Polarity Inventory

Count conservation is not enough for reaction closure. Since the ontic architrino set $\mathcal{A}$ is fixed, every serious reaction record must route identity-labeled architrinos through the event after expanding the input and output state to include any explicitly recruited or returned Noether sea content.

Let $R_e^{\text{in}} \subset \mathcal{A}$ and $R_e^{\text{out}} \subset \mathcal{A}$ denote the participating architrino identities before and after the event. A closed event must supply a bijection

$$\Pi_e : R_e^{\text{in}} \rightarrow R_e^{\text{out}}$$

such that, for every routed identity a ,

$$q_{\Pi_e(a)} = q_a, \quad q_a = \sigma_a \epsilon, \quad \sigma_a \in \{-1, +1\}$$

Equivalently, the polarity inventory vector

$$\mathbf{N}_e = (\#\{a : q_a = -\epsilon\}, \#\{a : q_a = +\epsilon\})$$

must agree before and after the event once all named reservoir terms are included. Photon assemblies, causal wakes, and corridor payloads may carry energy, momentum, angular momentum, phase, and path-history data, but they do not create new elements of $\mathcal{A}$. If a pair-production, weak, charged-pair relock, bremsstrahlung, synchrotron, or scattering record lacks Π_e or an equivalent identity-routing statement, the record remains provisional even when its net observer-level charge balances.

The contract for each serious channel is:

Contract field	Required content
Residual	Define $\mathcal{R}$ from the local state, causal-wake ledger, density field, Noether sea delay factor, and sector variables.
Threshold or separatrix	State the critical surface, basin boundary, channel boundary, or return-map condition that selects an admissible route.
Candidate channels	List the allowed routes, including radiative, recoil, medium, reaction, remnant, or record-forming terms when applicable.
Event ledger	Close E , $\mathbf{p}$, $\mathbf{J}$, charge/provenance, recoil, medium update, remnant state, architrino inventory, and identity routing where applicable.
Benchmark recovery	Name the observer-level reaction, cross-section, threshold, rate, or conservation benchmark recovered by the route.
Closure status	Mark the record as baseline, provisional map, derivation target, failed map, or inherited gate.

Promotion Criterion

A reaction record may be promoted beyond a provisional map only when all of the following conditions have been met in the same sector case:

1. **Replayable residual:** $\mathcal{R}(X)$ is computed from Γ , $\mathcal{H}$, $\rho_{\text{NS}}(\mathbf{x}, t)$, $\chi_{\text{sea}}(\mathbf{x}, t)$, and explicitly named sector variables, with no hidden sector-specific residual term.
2. **Boundary selection:** each selected channel has a stated boundary test $g_i(X, \mathcal{R}) \geq 0$, and every excluded channel required by the sector either fails its boundary test or is ruled out by a compatibility condition.
3. **Admissible output:** Y_e names all outgoing assemblies, recoil targets, medium updates, remnant states, and provenance records required by the selected channel set.
4. **Ledger closure:** $\mathcal{L}_{E\mathbf{pJ}}(e) = \mathbf{0}$ after adding the sector-required charge, polarity, architrino-inventory, identity-routing, path-history, Noether sea, and remnant rows.
5. **Benchmark compatibility:** the promoted event recovers the sector benchmark without breaking any required weak, quantum, gravity, hadronic, radiation, cosmology, conservation-law, or direct-observation acceptance gate.

This is a promotion criterion, not a completed theorem. Worked sector cases remain open until at least one channel supplies a named residual, a named threshold or separatrix, a channel decision, a complete $\mathcal{L}_{E\mathbf{pJ}}$ ledger, a benchmark recovery, and a failure diagnostic in one record. The free-neutron beta reaction, the $t \rightarrow b + W^+$ channel, radiation-coupled pair channels, and nuclear reaction examples therefore remain

provisional where their sector records still lack closed residual routing, outgoing braid provenance, angular-momentum balance, rate recovery, or quantitative benchmark closure.

Failure Modes

Failure mode	What blocks promotion
Residual replay failure	Two records with the same $(\Gamma, \mathcal{H}, \rho_{NS}, \chi_{sea}, \mathcal{Z}_S)$ produce different $\mathcal{R}$ values or different selected channel sets without an additional recorded state variable.
Boundary failure	A resolved event occurs while every required $g_i(X, \mathcal{R}) < 0$, or two mutually exclusive selected channels demand incompatible output assignments.
Ledger residual failure	After all sector-required rows are included, $\Delta_E \neq 0$, $\Delta_p \neq 0$, or $\Delta_J \neq 0$.
Inventory or provenance failure	$\Delta_{pol} \neq 0$, $\Delta_{arch} \neq 0$, or $\Delta_{path} \neq 0$ after the claimed Noether sea, corridor, source-identity, emission-time, causal-root, and branch-Jacobian records are included.
Identity-routing failure	No bijection Π_e , or equivalent identity route, maps participating input architrinos to participating output architrinos after named Noether sea reservoir terms are included.
Medium or remnant failure	$\Delta_{med} \neq 0$ or $\Delta_{rem} \neq 0$, meaning the route used medium heating, recoil, retained excitation, or remnant deformation as an implicit loss term.
Retuning failure	The same benchmark family can be recovered only by changing the residual definition, the channel boundary, or the Noether sea state variables between sector cases.
Cross-sector failure	The local route succeeds only by violating another required sector acceptance gate.

Weak-Corridor Provenance Gate

Weak reactions now require an explicit corridor-provenance stance. The current corpus supports two live possibilities:

- Transaction-payload corridor:** $W^\pm$ carries the charged triad payload and phase relation, while final-state pro/anti Noether braid material is supplied by the local Noether sea or by explicitly identified incoming assemblies.
- Provenance-carrying corridor:** $W^\pm$ carries not only the charged transaction payload but also enough pro/anti Noether braid provenance to seed some final-state lepton or antilepton braid content.

The ledger should not choose between these silently. For each serious weak record, add a row or note that states which stance is being used, which Noether braid material enters and exits, and what would falsify the accounting. This gate is coupled to the weak-coupling-triad exposure problem: the same geometry that permits left-handed charged-current docking must also determine which corridor payload can be transferred and where the outgoing lepton braids come from.

Minimum weak-channel records should therefore include:

- the active weak-coupling-triad swap,
- the corridor provenance stance,
- all Noether sea or incoming-assembly braid material used for charged lepton and neutrino outputs,
- the CKM/PMNS overlap weight when a flavor or generation branch is selected,
- and the energy, angular momentum, polarity, and path-history terms needed for deterministic replay.

Weak Reaction Case: $t \rightarrow b + W^+$ Channel

Observer-level notation:

$$t \rightarrow b + W^+, \quad W^+ \rightarrow e^+ + \nu_e$$

Native status: provisional weak-reaction provenance map.

The active quark change is an axial-layer reconfiguration. In the current assembly catalog, the top-to-bottom transition is represented as a shift from the top axial pattern to the bottom axial pattern:

$$(1E, 5P)_{\text{axial}} \rightarrow (4E, 2P)_{\text{axial}}$$

Equivalently, the active quark sector requires a $+3E, -3P$ axial exchange. In observer language this is the W^+ channel. In substrate language it is a transient payload and coupling event whose geometry, chirality selection, and energy routing still need closure.

The lepton products cannot be asserted as creation from nothing. Their braid and axial-layer material must be drawn from a local Noether sea reservoir or from explicitly identified incoming assemblies. The provisional ledger target is:

Component	Ledger requirement	Status
Top-to-bottom axial exchange	Route the $+3E, -3P$ change through a weak-channel coupling event	Provisional
Positron assembly	Identify the Noether braid and axial material used to form the charged lepton output	Provisional
Electron-neutrino assembly	Identify neutral braid and axial-layer routing, including chirality/orientation	Provisional

Component	Ledger requirement	Status
Energy-momentum	Account for quark mass difference, lepton energies, recoil, and medium excitation	Derivation target
Weak geometry	Derive the left-handed selection rule and allowed coupling operator	Derivation target

This channel should not be presented as a completed architrino derivation until the inventory table balances E/P counts, braid orientation, axial-layer routing, and energy-momentum in one consistent record.

Free Neutron Beta Reaction

Observer-level notation:

$$n \rightarrow p + e^- + \bar{\nu}_e$$

with the active quark-level comparison

$$d \rightarrow u + W^-, \quad W^- \rightarrow e^- + \bar{\nu}_e$$

Native label: free-neutron beta reaction.

The spectator structure is straightforward: one u and one d in the neutron pass through the reaction unchanged. The active channel is the second down-like assembly reconfiguring into an up-like assembly.

The axial-layer comparison is:

$$(4E, 2P)_{\text{axial}} \rightarrow (1E, 5P)_{\text{axial}}$$

So the active quark assembly sheds three E -type axial units and receives three P -type axial units. The natural provenance hypothesis is that local neutral Noether sea material supplies the compensating polarity units while the ejected E -type material participates in electron axial-layer formation.

Exposure-operator record

The controlled beta channel now has a first finite-state exposure operator in [Weak-Mixing CKM](#). The ledger record for this channel should use that operator as the geometry gate before any rate or provenance claim is made.

This gate inherits the unresolved spinor/helicity proof in [Angular Momentum and Spin](#). The blocked right-handed branch, antineutrino orientation, and weak-channel angular-momentum balance remain provisional until the weak-coupling-triad exposure geometry and the reaction-level angular-momentum ledger are derived from the same substrate proof.

Gate field	Beta-reaction record
Active assembly	One generation-I down-like quark inside the neutron
Spectators	One u and one d assembly pass through by identity
Exposure domain	$\Sigma_{\text{WCT}}^{(L)}$ on the leading, phase-matched weak-coupling triad
Gate condition	Left-handed charged-current docking with $ \Sigma_{\text{WCT}}^{(L)} = 3$ and active inventory $3E$
Blocked condition	Right-handed d channel has no charged-corridor docking in the finite-state model
Quark-side action	$A_\Sigma = 3E \rightarrow 3P$, with shielded inventory $A_{\text{sh}} = (1E, 2P)$ unchanged
Corridor payload	W^- carries the opposite transaction $\Delta A_W = 3(E - P)$, net charge $-e$
CKM weight	V_{ud} , interpreted as the same-tier weak-basis to shielding-eigenstate overlap
Provenance stance	Transaction-payload corridor unless a later derivation proves provenance-carrying corridor content is required

This record keeps the beta reaction from becoming two separate stories. The same exposed triad must explain the left-handed selection rule, supply the V_{ud} overlap domain, and identify what the W^- corridor transfers. The remaining open work is to identify the electron and antineutrino braid provenance and then attach the energy, angular momentum, recoil, and path-history terms.

The conservative ledger is:

Component	Required provenance statement	Closure status
Active $d \rightarrow u$ assembly	Route $3E$ out of the active axial layer and route $3P$ into it	Provisional map
Electron assembly	Combine the ejected $3E$ contribution with additional local Noether sea material and a suitable braid	Provisional map
Antineutrino assembly	Identify neutral braid orientation, axial-layer routing, and weak-channel phase relation	Open derivation target
Noether sea	Record every neutral braid, axial layer, or medium excitation consumed or returned	Required
Energy and angular momentum	Track mass difference, recoil, electron kinetic energy, antineutrino energy, and medium response	Required

This map supports a strong but bounded claim: beta reaction charge bookkeeping can be interpreted as local separation and rerouting of neutral Noether sea material plus active quark axial reconfiguration. It does not yet establish a full weak-interaction derivation, because chirality selection, antineutrino routing, and quantitative rate closure still belong to the weak-sector closure program.

Closure Targets

The reaction ledger needs at least four tables for each serious channel:

- 1. **Constituent inventory table:** braid and axial-layer E/P counts for every input, output, Noether sea contribution, and returned medium product.
- 2. **Energy-momentum table:** internal energy changes, kinetic output, recoil, photon assemblies, neutrino channel, and medium excitation.
- 3. **Geometry table:** axial frame, braid orientation, chirality, polarity routing, and allowed coupling/docking geometry.
- 4. **Path-history table:** causal-root branches, source identities, emission times, and local Noether sea state variables needed for deterministic replay.

Radiative or photon-coupled channels also need the shared radiation event-record table. The polarization handoff in that table remains inherited from Gate B; this ledger records the required transverse and capture/rejection fields but does not derive photon spin locally.

Validation Links

- Weak-sector geometry and chirality closure remain tied to [Quantum Number Mapping](#), [Weak Mixing Angle](#), and [Weak-Mixing CKM](#).
- Radiative and pair-production provenance should use [Synchrotron](#), [Bremsstrahlung](#), and [Reaction-Cosmology Provenance Ledger](#).
- Parameter closure belongs in [Parameter Ledger](#).

Reaction-Cosmology Provenance Ledger

This ledger connects local reaction provenance to cosmology-facing radiation, thermalization, and source-history claims. It is the bridge record for channels where synchrotron cascades, bremsstrahlung, pair production, BBN photon loading, and CMB thermalization all depend on the same underlying bookkeeping.

Use it with [Reaction Ledger](#), [Radiation](#), [Synchrotron](#), [Bremsstrahlung](#), [BBN Constraints](#), and [CMB](#).

Purpose

Cosmology-facing reaction claims need more than a source story. They need a record of what enters and exits each channel at the substrate level, and how those local channels become observer-level background quantities such as photon bath temperature, N_{eff} , light-element yields, redshift, and TT/TE/EE spectra.

This ledger separates four levels:

- **Ontology:** architrinos, Noether braids, axial layers, photon assemblies, and Noether sea state variables.
- **Reaction mechanics:** association, dissociation, planar-mode nucleation, pair production, recoil, and medium excitation.
- **Transport and thermalization:** opacity, scattering, cascade depth, diffusion, cooling, path-history redshift, and signed photon-frequency exchange.
- **Effective observables:** emissivity, light-element yield, blackbody spectrum, anisotropy, polarization, and inferred cosmological parameters.

Leap Opportunity Record

The opportunity tracked here is a possible unification of four previously separate bookkeeping problems: radiative planar-mode nucleation, pair-production provenance, BBN photon loading, and CMB thermalization. The shared claim is not that these channels are already derived from one equation. The disciplined claim is that they may need one common provenance ledger because each asks the same question at a different scale: which assemblies, Noether braid material, energy-momentum terms, and Noether sea state variables enter and exit the channel?

Current Claim Status

Claim	Bucket	Status	Decision gate
Bremsstrahlung and synchrotron both require planar-mode nucleation from assembly stress or wake concentration	Derivation-closure target	Provisional map	A common threshold condition must recover standard emissivity scalings in validated regimes
Pair production reorganizes local substrate content rather than creating charged assemblies from nothing	Ontology plus derivation-closure target	Accepted as ontology framing, open as quantitative derivation	Event records must balance architrino inventory, energy-momentum, and Breit-Wheeler rate behavior
BBN photon loading can be supplied by the same radiation and pair channels used in high-energy transport	Speculation promoted to closure target	Open	The source-zone photon ledger must preserve D, ^4He , Li, and N_{eff} constraints without per-source retuning
CMB blackbody recovery can be treated as source-to-transport-to-decoupling provenance rather than as an isolated	Derivation-closure target	Open	Thermalization depth, damping, anisotropy, polarization, and redshift handoff must all survive one shared

Claim	Bucket	Status	Decision gate
source story			parameter map

Discussion Gate

Before this bridge is promoted from ledger opportunity to mainline cosmology doctrine, the corpus needs a first quantitative record for at least one full path:

source channel $\rightarrow$ photon or pair assembly output $\rightarrow$ thermalization path $\rightarrow$ observer-level background variable

The minimal useful first path is BBN photon loading: identify a source-zone radiation channel, record its event-level provenance, propagate it through the local thermalization assumptions, and show whether it can support effective $\eta \approx 6 \times 10^{-10}$ during the deuterium bottleneck window.

Shared Provenance Fields

Field	What must be recorded	Why it matters
Architrino inventory	E/P counts, braid/axial-layer separation, and identity routing for recruited or returned substrate content	Prevents creation-from-nothing wording in pair and weak channels
Noether sea state	$\rho_{\text{NS}}(\mathbf{x}, t)$, $n(\mathbf{x}, t)$, $\chi_{\text{sea}}(\mathbf{x}, t)$, anisotropy, and excitation state	Keeps density, delay, and transport variables distinct
Radiation event record	Source assembly, trigger geometry, $\delta\Theta_\alpha$, E_{exc} , E_γ , recoil, medium excitation, polarization handoff, causal-wake ledger, and closure status	Provides the local event schema that can be propagated into source-zone, transport, and observer-level cosmology claims
Photon assembly channel	Planar-mode nucleation threshold, emitted energy, direction, polarization basis, and transverse angular-momentum ledger	Links bremsstrahlung, synchrotron, and CMB photon-bath claims
Pair channel	Incoming photon assemblies, identity-routed recruited Noether braid content, final e^+e^- assemblies, and recoil/medium excitation	Keeps pair production as association from local substrate content, not ex nihilo creation
Energy-momentum ledger	Internal energy, kinetic energy, recoil, emitted assemblies, and medium excitation	Required for observer-rate and spectrum recovery
Thermalization path	scattering depth, coupling time, cooling time, and escape time	Determines when local reactions can feed BBN or CMB background claims
Observer handoff	emissivity, opacity, redshift kernel, effective temperature, N_{eff} , and C_ℓ inputs	Keeps standard comparison variables useful without treating them as ontology

Photon Closure Gates

Photon-channel records should be sorted into three gates before they are used in cosmology-facing arguments.

The chapter-level source for the photon ontology and Gate A theorem scaffold is [Electroweak Bosons](#). This ledger records what a reaction or cosmology channel must carry forward from that scaffold before it uses photon propagation, polarization, pair production, or thermal radiation as settled input. Gate B is downstream of [Angular Momentum and Spin](#); the fields below are acceptance records, not an independent derivation of photon spin or polarization statistics. Photon-pair Bell/CHSH claims and CMB polarization-transfer claims must therefore inherit Gate B and the pair-provenance measure rather than being closed by cosmology bookkeeping alone.

Gate	Claim bucket	What the ledger must track	Closure test
Gate A: kinematics and optics	Derivation-closure target	c_f , c_γ , $\delta_\gamma \equiv 1 - c_\gamma/c_f$, planar-pair spacing d , phase frequency ω , geometric phase, and medium delay state	Recover $E_\gamma = h\nu$, $p = h/\lambda$, masslessness, no rest proper-time branch, nondispersion, and no unacceptable preferred-frame leakage
Gate B: polarization and spin	Derivation-closure target	transverse ledger orientation, analyzer basis, helicity, projection/capture geometry, accepted/rejected channel outcomes, source depletion, recoil, causal-wake, handoff, and event-balance rows	Recover exactly two transverse modes, no longitudinal mode, Malus' law, helicity ± 1 , single-photon statistics, no-signaling constraints, and $\mathcal{R}_{\gamma B}^{\text{event}}$ below tolerance
Gate C: vertices and transitions	Derivation-closure target	emission, absorption, pair production, recoil, medium excitation, transition rates, and overlap/capture probabilities	Recover QED/Maxwell limits, Breit-Wheeler thresholds and rates, blackbody behavior, Compton-like scattering, photon-photon limits, and the effective coupling scale α

These gates are not separate ontologies. They are bookkeeping filters that prevent a local photon-source story from being used as cosmology doctrine before the same event record also closes photon transport, polarization, pair conversion, and observer-level comparison variables. The shared radiation event record is the carrier for those gate handoffs; Gate B remains inherited and is not re-derived by this cosmology ledger.

Channel Map

Channel	Source document	Provenance target	Current status
Bremsstrahlung planar-mode nucleation	Bremsstrahlung	Record electron assembly energy loss, target recoil, photon assembly output, and medium excitation	Provisional map
Synchrotron planar-mode nucleation	Synchrotron	Derive photon output from curved charged-assembly transport in anisotropic Noether sea states	Provisional map

Channel	Source document	Provenance target	Current status
Breit-Wheeler pair channel	Synchrotron	Record incoming photon assemblies, recruited Noether braid content, and final e^+e^- assemblies	Derivation target
BBN photon bath	BBN Constraints	Show that pair, bremsstrahlung, synchrotron, and related channels maintain effective $\eta \approx 6 \times 10^{-10}$ during the bottleneck window	Closure target
CMB thermal spectrum	CMB	Show that source emission, transport, and thermalization produce a near-blackbody photon bath with allowed anisotropy and damping structure	Closure target
Redshift and clock handoff	Expansion Mechanism	Map photon transport through ρ_{NS} , n , χ_{sea} , and clock-rate comparison	Effective summary with open derivation
Sunyaev-Zeldovich / Compton-like frequency exchange	CMB and Radiation	Record incoming photon packet, intervening electron or medium state, outgoing frequency, recoil, medium energy change, and thermalization side effects	Calibration row and closure target

Minimum Records by Channel

Each minimum record below specializes the shared event schema in [Radiation](#). Additional cosmology variables may be added, but the source assembly, source-depletion row, trigger geometry, $\delta\Theta_a$, E_{exc} , E_γ , recoil, medium excitation, polarization handoff, photon Gate B event residual when $E_\gamma \neq 0$, causal-wake ledger, identity routing, and closure status fields remain required.

Bremsstrahlung

The minimum event record is:

$$E_{\text{exc}}^{\text{br}} = E_\gamma + \Delta E_{\text{recoil}} + \Delta E_{\text{med}} + \Delta E_{\text{rem}}$$

The provenance record must also include the source electron assembly, target assembly, source-depletion row, trigger geometry, $\delta\Theta_a$, local $\rho_{\text{NS}}(\mathbf{x}, t)$, $n(\mathbf{x}, t)$, $\chi_{\text{sea}}(\mathbf{x}, t)$, planar-mode threshold status, emitted photon assembly direction, recoil, medium excitation, causal-wake ledger, identity routing, closure status, and whether the event occurs in a regime where standard free-free emissivity remains the observer-level scaffold. Its polarization handoff inherits photon Gate B rather than deriving photon spin locally, and the record remains provisional until the event residual routes source, recoil, medium, wake, handoff, and remnant rows.

Synchrotron Emission

The event record must connect charged-assembly curvature, the effective magnetic-field map, source depletion, trigger geometry, $\delta\Theta_a$, E_{exc} , photon assembly output E_γ , recoil, medium excitation, causal-wake ledger, identity routing, and photon Gate B event residual. The closure target is to derive the standard $\nu_c \propto \gamma^2 B$ and $P_{\text{syn}} \propto U_B \gamma^2$ scalings from Noether sea anisotropy and wake-strain threshold conditions rather than fitting a separate emission rule. Synchrotron polarization records inherit Gate B, so this ledger carries the transverse handoff without proving photon helicity locally.

Pair Production

The event record must avoid creation-from-nothing wording. Incoming photon assemblies trigger association of local substrate content into e^+e^- assemblies when the observer-level threshold is satisfied. The incoming photons supply energy, momentum, polarization handoff, and trigger geometry; they do not supply new architrino identities. The incoming photons should preserve their radiation event records through the pair vertex. The pair-channel record must include:

- incoming photon assembly energies and directions,
- incoming photon polarization handoffs as inherited Gate B records,
- local Noether braid material recruited or reconfigured, including identity routing for the architrinos assigned to the final charged assemblies,
- final charged assembly inventories,
- recoil and medium-excitation terms,
- causal-wake ledger and closure status,
- and the standard-limit cross-section target.

This is the ledger distinction that ordinary absorption does not need: atomic or material capture closes the photon ledger into an existing target or medium record, while pair production closes the photon ledger and separately recruits identity-routed substrate content into new charged assemblies.

BBN Photon Loading

The BBN module needs a source-zone photon ledger. It must identify which radiation channels supply the effective photon-dominated environment and whether they preserve deuterium survival, helium clustering, and N_{eff} compatibility without per-source retuning.

Matter-Asymmetry Provenance

The observed baryon-to-photon ratio is a data-product constraint, not permission to import an external baryogenesis mechanism as doctrine. Any matter-asymmetry story used by the cosmology program must be rewritten as a reaction provenance record over a declared source window W . Let $N_B(W)$, $N_{\bar{B}}(W)$, and $N_\gamma(W)$ be the baryon, antibaryon, and photon counts after the event records have been transported to the BBN comparison surface. Define

$$\eta_B^{\text{ledger}}(W) = \frac{N_B(W) - N_{\bar{B}}(W)}{N_\gamma(W)}$$

For leptogenesis-like source routes, the ledger must also carry a neutrino/antineutrino CP-asymmetry comparison term rather than assuming the external mechanism. Source leads for this row are primary neutrino-oscillation and leptogenesis sources: long-baseline $\nu/\bar{\nu}$ transition measurements, PMNS CP-phase summaries, and baryogenesis/leptogenesis rate calculations. The comparison term is

$$\Delta_{\nu\bar{\nu}}^{\text{CP}}(E, L; \alpha, \beta) = P_{\nu_\alpha \rightarrow \nu_\beta}(E, L) - P_{\bar{\nu}_\alpha \rightarrow \bar{\nu}_\beta}(E, L)$$

where E is neutrino energy, L is baseline, and α, β label flavor channels. The source-window ledger may report $\Delta_{\nu\bar{\nu}}^{\text{ledger}}(W)$ as the event-record-weighted version of this comparison over W , but that reported value is only an input constraint on the matter-asymmetry closure. It is not an established AAA derivation of baryon excess.

The acceptance residual should be reported as

$$\mathcal{R}_{B/\gamma}(W) = \max \left(\frac{|\eta_B^{\text{ledger}}(W) - \eta_B^{\text{obs}}|}{\varepsilon_\eta}, \frac{|\Delta_{\nu\bar{\nu}}^{\text{ledger}}(W) - \Delta_{\nu\bar{\nu}}^{\text{obs}}(W)|}{\varepsilon_{\nu\bar{\nu}}}, \frac{|\Delta B_{\text{unrec}}(W)|}{\varepsilon_B}, \frac{|\Delta Q_{\text{unrec}}(W)|}{\varepsilon_Q}, \frac{|\Delta E_{\text{unrec}}(W)|}{\varepsilon_E} \right)$$

Here $\Delta_{\nu\bar{\nu}}^{\text{ledger}}$, ΔB_{unrec} , ΔQ_{unrec} , and ΔE_{unrec} are not new ontology. They are comparison or failure counters for CP-asymmetric neutrino/antineutrino transition rates, baryon-number bookkeeping, electric-charge bookkeeping, and energy balance after all declared reaction, recoil, medium, and escape channels have been included. A leptogenesis-like source model may remain in the comparison ledger only when $\mathcal{R}_{B/\gamma} \leq 1$ and the same event record also passes the BBN photon-loading and CMB thermalization checks below.

CMB Thermalization

The CMB module needs a source-to-transport-to-decoupling ledger. It must track:

- source-channel selection from SMBH-local release, medium relaxation, and conversion/dissociation pathways,
- thermalization depth and blackbody recovery,
- anisotropy and polarization transfer, with the polarization handoff inherited from Gate B,
- redshift and clock-rate handoff,
- and separation between source interpretation and the shared prediction target C_ℓ .

The thermalization-depth record is a diagnostic field, not a new substrate entity. For each modeled source-to-decoupling path, the ledger should record

$$\mathcal{D}_{\text{th}}(\nu; t_a, t_b) = \int_{t_a}^{t_b} \tau_{\text{th}}^{-1}(\nu, t) dt$$

with τ_{th}^{-1} decomposed into the specific event-recorded channels being used: planar-mode capture/release, Compton-like redistribution, pair channels, and non-radiative medium exchange. A CMB blackbody claim requires $\mathcal{D}_{\text{th}} \gg 1$ before decoupling, effective photon chemical potential driven to zero, and a post-decoupling transport map that preserves the already-generated spectrum while carrying anisotropy, polarization, damping, and redshift information.

Path Frequency Exchange

Post-emission photon frequency changes are not automatically new photon emission events. A photon packet may exchange energy with an intervening electron population, plasma, or Noether sea state and continue as the same transported packet. For each such event or coarse segment, the ledger must record incoming frequency ν^- , outgoing frequency ν^+ , the local medium state, recoil or target momentum, and the residual

$$\mathcal{R}_{\nu\text{-ex}} = \frac{|h(\nu^+ - \nu^-) + \Delta E_{\text{med}} + \Delta E_{\text{recoil}} + \Delta E_{\text{rem}}|}{\epsilon_E}$$

The same row must state whether the exchange is thermalizing, spectrally distorting, or coherently transported. A Sunyaev-Zeldovich-type boost is admissible only when the electron or medium record supplies the photon energy increase and when the side effects remain compatible with the CMB spectrum, anisotropy, polarization, and kSZ/tSZ observable rows. A depletion row is admissible only when the lost photon energy is routed into a named medium, recoil, remnant, or thermalization channel.

Closure Targets

1. **Planar-mode threshold closure:** derive a shared threshold condition for bremsstrahlung and synchrotron photon assembly output.
2. **Pair-production provenance closure:** prove that local Noether sea recruitment can satisfy architrino inventory, energy-momentum, and Breit-Wheeler rate constraints in the same event record.
3. **Photon-bath closure:** show that the relevant radiation channels can maintain BBN-compatible photon loading during the deuterium bottleneck window.
4. **Matter-asymmetry closure:** derive η_B^{ledger} from event-level reaction provenance without hidden baryon inventory, charge, or energy sources; for leptogenesis-like routes, also recover $\Delta_{\nu\bar{\nu}}^{\text{ledger}}$ from primary-source neutrino CP-asymmetry comparisons without promoting leptogenesis to doctrine.
5. **Detailed-balance closure:** derive the rate symmetry and ensemble weight relation that make emission, absorption, and stimulated terms recover Planck occupation with zero effective photon chemical potential.
6. **Blackbody closure:** show that distributed source channels plus Noether sea transport can generate and preserve the CMB blackbody spectrum within observational limits.
7. **Clock/redshift closure:** use one Noether sea state map for photon propagation, endpoint clock comparison, and redshift-distance inference.

Failure Modes

The provenance program fails for a channel if a source story cannot survive the same ledger used for reaction, transport, thermalization, and observer handoff.

Failure mode	What fails	Diagnostic consequence
Single Noether braid temperature mistake	A single excited Noether braid is treated as thermodynamically hot rather than internally excited, closure-mismatched, or metastable	Temperature is being used before an ensemble distribution or entropy-energy relation has been established
Inventory gap	Architrino inventory, Noether braid recruitment, recoil, or returned medium content cannot be balanced without unrecorded substrate creation	Pair and radiation channels cannot be promoted beyond provisional maps
Per-observable refit	The same Noether sea state variables must be re-fit independently for photon loading, blackbody recovery, damping, redshift, or growth observables	The cosmology interpretation loses its shared Noether sea state map
Standard-limit violation	Pair, Compton-like, bremsstrahlung, synchrotron, or photon propagation channels violate validated limits in regimes where those limits are already measured	The proposed substrate route fails before it can claim new deviations
Insufficient thermalization depth	$\mathcal{D}_{\text{th}}$ is too small, or its channel decomposition is not tied to event records	Source photons need not relax to a Planck bath, and a nonzero effective photon chemical potential or spectral distortion remains
Matter-asymmetry ledger failure	η_B^{ledger} cannot match the observed baryon-to-photon ratio, or $\Delta_{\nu\bar{\nu}}^{\text{ledger}}$ is imported without event-record support; the source route then relies on unrecorded baryon inventory, charge imbalance, or energy imbalance	A baryogenesis-like or leptogenesis-like source story cannot be promoted into cosmology provenance
BBN photon-loading failure	Source-zone photon production cannot preserve deuterium survival, helium clustering, lithium constraints, and N_{eff} compatibility	The BBN local-reactor mapping cannot replace the standard photon-to-baryon environment
CMB handoff failure	Blackbody precision, damping behavior, anisotropy, polarization, or TT/TE/EE coherence cannot be carried through the same transport and redshift map	CMB thermalization cannot be treated as a successful source-to-observer provenance path
Frequency-exchange ledger failure	A path segment changes photon frequency without a closed medium, recoil, remnant, or side-effect row	Redshift, blueshift, SZ, or distance-ladder claims are being used without photon provenance

Constraint Ledger

Notes collected here document the falsification criteria, ordering priorities, and supporting mechanisms for the architrino framework. Keep this page focused on observable constraints so each model version can be checked against experimental scrutiny.

Experimental Constraint Ledger and Falsification Criteria

This ledger crystallizes the measurable thresholds and theoretical guardrails that could falsify the architrino proposal. Each numbered entry combines the empirical bound, the proposed mechanism, and the explicit failure condition so that we can track how discrete experimental results shape or reject the model.

Lorentz Invariance & Preferred Frame Effects (Tier 1)

The purpose of this section is to define the combination of experimental isotropy and observational invariance that must hold if a putative absolute frame is to remain hidden. We identify the observables, derive the emergent timing/ruler behavior implied by the Noether sea, and explicitly state the tolerance beyond which the preferred frame would become perceivable.

- **Constraint** – isotropy from Michelson–Morley and resonator experiments constrains $|\Delta c/c| < 10^{-17}$ while atomic clock sidereal drift stays below 10^{-16} , keeping Lorentz-invariance leakage under the 10^{-17} falsification threshold.

- **Consolidated Requirement** – prove preferred-frame hiding: architrino assemblies must acquire Lorentz-compatible deformation and clock behavior in the Euclidean-void rest frame so no local observer can detect the Noether sea’s rest frame.
- **Observable** – local Lorentz invariance is preserved.
- **Mechanism** – assembly-based clocks/rulers must emerge with proper time τ rather than absolute time t .
- **Shared Residual** – the structural-integrity common-limit closure in [Lorentz Kinematics](#) couples this row to photon and gravitational-wave speed gates: the same branch record must make $c_{\text{mat}}^{\text{lim}}$, c_{eff} , c_γ , and c_0 agree within $O(\epsilon_{\text{LV}})$ while also producing clock/ruler deformation and two-way photon synchronization.
- **Failure Condition** – any detectable preferred-frame orientation above 10^{-17} or residual δ in $L_{\text{moving}} = L_{\text{rest}}(\gamma^{-1} + \delta)$ that exceeds 10^{-17} invalidates the theory.

Photon Time-of-Flight Dispersion Gate

High-energy transient events at cosmological distance test whether photon-channel propagation accumulates a frequency-dependent delay. The observable is a time-of-arrival residual after source-intrinsic emission lag has been modeled; it is not direct evidence for or against microscopic spatial grains by itself.

For two photon phase frequencies ω_a and ω_b emitted by the same source at redshift z , a candidate photon-channel delay is

$$\Delta t_\gamma^{\text{model}}(\omega_a, \omega_b; z) = \int_{\Gamma_z} \frac{\chi_\gamma(\omega_a, \mathbf{x}, t) - \chi_\gamma(\omega_b, \mathbf{x}, t)}{c_0} d\ell$$

Here Γ_z is the observer-level path used by the comparison, and χ_γ is the photon-channel delay factor from the same branch record used for photon synchronization. A useful residual is

$$\mathcal{R}_{\gamma\text{disp}} = \sup_{\mathcal{E}} \frac{|\Delta t_{\text{obs}} - \Delta t_{\text{src}} - \Delta t_\gamma^{\text{model}}|}{\sigma_{\Delta t}}$$

where $\mathcal{E}$ is the declared transient catalog, Δt_{src} is the modeled source lag, and $\sigma_{\Delta t}$ is the adopted timing uncertainty.

- **Constraint** – the same photon branch that recovers local Lorentz synchronization must keep $\mathcal{R}_{\gamma\text{disp}}$ below the declared catalog threshold without per-source retuning.
- **Observable** – measured arrival-time differences across photon energy or frequency bands, source-lag model, redshift, instrument timing uncertainty, and event-selection rule.
- **Validation Target** – Gate A in [Electroweak Bosons](#) must derive a nondispersive weak homogeneous photon branch rather than assume it after the fact.
- **Shared Residual** – this is the photon-channel component of the same common-limit residual defined in [Lorentz Kinematics](#); the χ_γ record cannot be repaired independently of the clock/ruler c_{eff} record.
- **Failure Condition** – a photon closure branch fails if it predicts an accumulated frequency-dependent delay in the validated band, hides that delay by changing the source-lag model event by event, or uses a different c_γ / χ_γ record from the one used in [Lorentz Kinematics](#).

The Absolute-Frame Drift Check (Lorentz Contraction Enforcement)

This entry frames the requirement that the underlying Noether sea affords a dynamical contraction mechanism to assemblies moving through the Euclidean void; without such a mechanism, assemblies would reveal their motion relative to the sea and the preferred frame would manifest.

- **Constraint** – the Noether sea must supply a dynamical closure that yields Lorentz-compatible contraction of assemblies; otherwise the model is equivalent to an untested preferred frame.
- **Failure Condition** – without contraction enforced by the Sea, preferred frame effects become measurable and falsify the theory.

Noether Sea Drag

Here we catalogue how coupling between macroscopic bodies and the Noether sea can influence orbital dynamics. The constraint ensures any additional dissipation or effective drag remains below the levels already constrained by gravitational-wave-based orbital decay measurements in general relativity.

- **Constraint** – interactions with the Noether sea must not induce orbital decay that outpaces GR’s gravitational-wave emission bounds.
- **Validation Target** – match observed orbital stability and perihelion advance within GR limits while modeling any extra coupling as a conserving medium-dressed response rather than ordinary dissipative drag.

Condensed-Matter Response Gate

Ordinary materials supply a broad recovery surface for the same assembly, electron-envelope, and Noether sea response variables. The gate is not that [AAA](#) adopts band theory as ontology. The gate is that periodic material branches recover the benchmark mathematics of bands, lattice scattering, phonons, and Hall response without per-probe retuning.

- **Constraint** – one material-branch record $\theta_{\text{mat}} = (\mathcal{B}_e, \mathcal{B}_{\text{lat}}, \rho_{\text{NS}}, n, \chi_{\text{sea}}, \mathcal{M}_{\text{sea}}^{ab})$ must recover Bloch-form bands $E_\alpha(\mathbf{k})$, effective mass tensor $(m_{\alpha,*}^{-1})^{ij} = \hbar^{-2} \partial_i \partial_j E_\alpha$, Fermi-surface or band-gap classification, reciprocal-lattice scattering $\mathbf{q} \in \Lambda^*$ with structure factor $S(\mathbf{q})$, and phonon dispersion from one declared lattice branch.
- **Hall / Topology Target** – for two-dimensional gapped branches with an effective U(1) connection, the same record must recover $\sigma_{xy} = (e^2/2\pi\hbar)C$ with integer Chern number C and ρ_{xx} below tolerance on the plateau. Fractional Hall, anyon, and Chern-Simons descriptions are recovery/comparison structures unless a local branch derivation consumes them directly.
- **No-Drag Consistency** – the ideal periodic branch must not require ordinary dissipative drag; finite τ^{-1} must be routed to disorder, vacancies, phonons, boundary exchange, heating, radiation-like shedding, or branch transition.
- **Failure Condition** – the condensed-matter branch fails if it fits band curvature, phonon stiffness, scattering peaks, Hall conductance, and transport relaxation with independent material records, if a filled band carries unlogged current or heat, if a topological plateau changes without a gap closure or branch change, or if ordinary Noether sea drag is used to explain resistance below the transport threshold in [Condensed Matter](#).

GW Speed

The propagation speed of gravitational-wave disturbances in the Noether sea must align with the measured gravitational-wave velocity, so this section records the tolerance within which new physics can coexist with GW timing data without contradicting the LIGO/Virgo baseline. The relevant benchmark is now multi-messenger rather than merely assumed: GW170817/GRB 170817A constrained the gravity-channel and photon-channel speed difference at roughly the 10^{-15} level.

- **Constraint** – gravitational waves, modeled as collective Noether sea disturbances, must satisfy the multi-messenger speed gate, with GW170817/GRB 170817A giving the reference scale

$$-3 \times 10^{-15} \lesssim \frac{v_{\text{GW}} - c_0}{c_0} \lesssim 7 \times 10^{-16}$$

Equivalently, for the weak homogeneous observer branch in which $c_\gamma \rightarrow c_0$,

$$\left| \frac{c_{\text{GW}}}{c_\gamma} - 1 \right| \lesssim 10^{-15}$$

is the order-of-magnitude ledger tolerance. Any tighter ledger tolerance adopted for a specific validation band should be stated explicitly rather than inferred from ontology.

- **Shared-Channel Requirement** – the effective gravitational-wave channel and photon channel must be derived from one Noether sea state record in the weak-field branch, as required by [Lorentz Kinematics](#). A medium-based gravity model fails this row if it lets gravitational waves and photons acquire independently tunable dressed speeds in the same region.
- **Mode and Dispersion Gate** – finite-range or medium-compliance corrections must keep accumulated dispersion, false-alarm residuals, calibration residuals, and any scalar, vector, or longitudinal gravitational-wave detector response below the residual bounds for the validated band.
- **Low-Frequency Extension** – if a cosmological-scale weakening channel claims finite-range behavior, it must also report the low-frequency residual $\mathcal{R}_{\text{GW,low}}(\theta)$ from [Gravitational Waves](#) for the declared pulsar-timing or space-interferometer band. A band not yet measured may be listed as a forecast, but it cannot be used to override the existing high-frequency speed, polarization, and dispersion gates.
- **Failure Condition** – a cosmological-scale weakening channel fails if it predicts measurable gravitational-wave dispersion, an unsuppressed non-TT mode, or a speed offset in the same regime where the weak-field metric map is supposed to recover GR.

Euclidean vs. Metric Pathing (The Refraction Mapping)

This constraint explains how apparent metric deviations (Shapiro delay and light bending) emerge from a Euclidean signalling framework endowed with a varying Noether sea delay factor χ_{sea} , which allows us to compare the emergent delay with the standard GR potential.

- **Constraint** – Shapiro delay and light bending must match GR to within PPN bounds ($|\gamma - 1| < 10^{-5}$).
- **Architrino Interpretation** – signals propagate through Euclidean space, but observer-level paths are effective travel-time extremals in the Noether sea delay map. The perceived delay or curvature arises from χ_{sea} responding to spatial variations in ρ_{NS} and related Noether sea state variables.
- **Validation Target** – map $g_{00} \approx 1 + 2\Phi/c^2$ onto the refractive slowing experienced by Noether sea signals moving through the Euclidean void with Noether sea delay.

Gravitational Time Dilation

We require that the proposed mechanical slowing induced by Noether braid density aligns quantitatively with geodetic and redshift observations such as GPS offsets and the Pound–Rebka experiment, offering a concrete mapping between the new microphysics and the classical time-dilation effects.

- **Constraint** – reproduce GPS clock offsets (38 $\mu\text{s/day}$), the Pound–Rebka redshift, and height-resolved optical-clock redshift with $\Delta\nu/\nu \approx gL/c_0^2$; this includes the approximate scales 1.1×10^{-19} across 1 mm and 3.6×10^{-17} across 33 cm near Earth’s surface.
- **Mechanism** – mechanical slowing of nested shell braid orbital frequencies couples to the local Noether braid density and Noether sea delay factor, generating the observed dilation without changing the constitutive map used for other weak-field observables.

Massive-Superposition Gravitational Distinguishability

Massive-interference experiments and precision gravity readouts jointly test whether the effective-metric channel carries enough branch information to become a which-path record. The observable is not whether spacetime is declared classical or quantum. The observable is whether two mass-density histories produce a distinguishable gravitational response before the apparatus has formed a durable record.

- **Constraint** – for two branch-level mass-density histories ρ_1 and ρ_2 , the gravitational distinguishability diagnostic

$$\mathcal{D}_{\text{grav}}(T; \theta) = \int_0^T \int_0^T \Delta h_A(t) N_{AB}^{-1}(t, t') \Delta h_B(t') dt dt'$$

with $\Delta h_A(t) = h_A(t; \rho_1, \theta) - h_A(t; \rho_2, \theta)$, must remain below the declared which-path threshold for any interference-preserving run unless a record-forming separatrix crossing and persistence window are also derived.

- **Observable** – the data products are massive-superposition coherence time, branch separation and mass-displacement history, precision-gravity response, detector noise covariance, any two-probe entanglement witness, non-gravitational coupling residuals, and the absence or presence of a durable which-path record.
- **Validation Target** – combine long-coherence interferometry with Cavendish-like, atom-interferometric, or gravitational-wave-instrument precision bounds to constrain $\mathcal{D}_{\text{grav}}$ using one effective-metric constitutive record θ ; the concrete scaffold is [Massive-Superposition Gravity Validation Packet](#).
- **Mediated-Entanglement Target** – for gravitationally induced entanglement comparisons, the same θ must generate the branch interaction phase $\Delta\Phi_{\text{ent}}$ needed for the observed witness C_{obs} while keeping $\mathcal{R}_{\text{nongrav}}$ below the isolation threshold and $\mathcal{D}_{\text{grav}}$ below the which-path threshold.
- **Failure Condition** – the measurement and spacetime branches fail jointly if the same parameter record predicts $\mathcal{D}_{\text{grav}} \gg 1$ for an interference-preserving experiment while no apparatus/environment record satisfies the record-autonomy condition in [Measurement Ontology](#).

CMB Scalar/Tensor Gate

The cosmology branch must recover the CMB scalar and tensor observables as data products before any source interpretation is promoted.

- **Constraint** – one Noether sea and assembly record must recover TT/TE/EE spectra, damping, CMB-lensing reconstruction, blackbody preservation, scalar amplitude A_s , scalar tilt n_s , acoustic phase coherence, vector-mode suppression, and the tensor bound $r \leq r_{\text{max}}$ without changing Noether sea state variables between the CMB, BBN, expansion, and growth modules.
- **Observable** – the CMB comparison residual $\mathcal{R}_{\text{CMB}}(\theta)$ defined in [CMB](#) must remain within the declared tolerance for the data release being used, and the added $\mathcal{R}_{\text{phase}}(\theta)$, $\mathcal{R}_V(\theta)$, and $\mathcal{R}_{\text{lens}}(\theta)$ gates must not require a separate medium history.
- **Smoothness Check** – the same record must also bound the effective smoothness residual $\mathcal{R}_{\text{smooth}}(\theta)$, so early-universe smoothness is tested as low observer-level gravitational free-mode content rather than assumed from an imported origin story.
- **Failure Condition** – if the framework can fit the source story only by retuning scalar power, acoustic phase, vector-mode content, CMB-lensing reconstruction, tensor contribution, blackbody recovery, or TT/TE/EE transfer independently, the cosmology closure fails at the observational layer.

Compact Dark-Sector Local-Detection Gate

Compact neutral-assembly or primordial-defect branches must face local gravitational searches as data products, not only cosmological abundance fits. For a branch record θ_A with representative mass M_A , local fraction f_A , and relative-speed distribution $p(v_{\text{rel}})$, the first rate estimate is

$$\Gamma_{\text{flyby}}(b_{\text{max}}, M_A; \theta_A) = \frac{f_A \rho_{\text{DM}}}{M_A} \pi b_{\text{max}}^2 \langle v_{\text{rel}} \rangle_{\theta_A}$$

The corresponding impulse scale on a tracked body is

$$\Delta v_{\text{test}} \simeq \frac{2GM_A}{b v_{\text{rel}}}$$

with the accepted comparison using the full ephemeris covariance rather than this estimate alone.

- **Constraint** – any claimed local compact dark-sector signal must produce an ephemeris residual $\Delta \mathbf{x}_{\text{ephem}}^\theta(t)$ above the declared ranging and model-error floor while remaining inconsistent with ordinary catalogued bodies under the same orbit-reconstruction covariance.

- **Co-Signature Check** – if the branch predicts high-energy particles, radiation, or gravitational-wave sidebands, those observables must use the same trajectory, mass, and abundance record as the ephemeris perturbation.
- **Failure Condition** – a compact dark-sector branch fails locally if it explains cosmological abundance with one mass or population record but requires a different record for ephemerides, visible-object exclusions, or high-energy null results.

Closure Program Tracking Hooks

Use this ledger as the acceptance layer for the six integrated closure programs:

Program	Primary chapters	Ledger gate
CKM holonomy closure	theory-bridges/weak-mixing-ckm.md	CKM hierarchy and CP-phase consistency with propagated uncertainty
PMNS neutral braid closure	assemblies/fermions/neutrinos.md	Oscillation pattern consistency across L/E and medium regimes
Emergent metric / PPN closure	spacetime/emergent-metric.md , spacetime/ppn-parameters.md , spacetime/proper-time-and-time-dilation.md	Lorentz leakage, PPN, redshift, Shapiro, GW-speed bounds
Non-relativistic Schrödinger + Born closure	theory-bridges/pilot-wave-character.md , quantum/wavefunction-ontology.md , theory-bridges/superposition-mechanism.md	Effective fixed-particle-number wave equation + statistical outcome consistency
Photon Gate A/B/C closure	assemblies/bosons/electroweak-bosons.md , theory-bridges/angular-momentum-and-spin.md , validation/reaction-cosmology-provenance-ledger.md , spacetime/lorentz-kinematics.md	Gate A massless nondispersive photon kinematics, Gate B polarization and squared-amplitude capture as a downstream spin/helicity ledger, and Gate C Maxwell/QED vertices, pair/radiation provenance, and α recovery
Topological spin/confinement closure	dynamics/causal-action-functional.md , assemblies/fermions/color-charge-su3.md	4π spin structure and open-vs-closed color-energy scaling

Cross-program acceptance principle:

$$\mathcal{C}_{\text{CKM}} \cap \mathcal{C}_{\text{PMNS}} \cap \mathcal{C}_{\text{PPN/GR}} \cap \mathcal{C}_{\text{QM}} \cap \mathcal{C}_{\text{Photon}} \cap \mathcal{C}_{\text{Topo}} \neq \emptyset$$

If the intersection is empty after uncertainty propagation, the integrated model version is rejected.

Failure Criteria

This chapter states the hard-stop conditions for $\mathbb{A}\mathbb{A}\mathbb{A}$. Its purpose is to distinguish ordinary incompleteness from genuine failure modes, especially where a local success in one sector cannot survive the shared closure intersection.

Its operational companions are [Validation Protocols](#), [No-Go Theorems](#), [Known Tensions](#), [Lorentz Kinematics](#), [Proper Time and Time Dilation](#), and [Absolute Time Defense](#).

Shared Closure Record

Let

$$\mathfrak{S} = \{\text{weak, quantum, gravity, hadronic, radiation, cosmology}\}$$

be the sector set. A candidate promoted closure is a record $\theta \in \mathfrak{X}$ whose shared coordinates include

$$\theta_{\text{join}} = (A, \Gamma, \mathcal{H}, \mathcal{R}, \mathcal{L}_{\text{PJ}}, \zeta, \mathcal{M}_{\text{sea}}^{ab}, \{B_i\})$$

where A is the assembly or branch family, Γ is the assembly microstate, $\mathcal{H}$ is the path-history and causal-wake ledger, $\mathcal{R}$ is the active residual family, $\mathcal{L}_{\text{PJ}}$ is the event ledger, ζ is shielding or exposure data, $\mathcal{M}_{\text{sea}}^{ab}$ is the Noether sea response object, and $\{B_i\}$ is the basin or channel partition. Sector-local coordinates $Z_S(\theta)$ record the benchmark variables, theorem assumptions, provenance rows, and tolerances used by sector S .

For each sector S , fix a gate predicate $P_S : \mathfrak{X} \rightarrow \{0, 1\}$, a benchmark map $\mathcal{B}_S : \mathfrak{X} \rightarrow \mathfrak{B}_S$, a validated benchmark region $\mathfrak{B}_S^{\text{obs}} \subseteq \mathfrak{B}_S$, a benchmark metric d_S , a tolerance ϵ_S , and a no-go pass predicate $\mathcal{G}_S : \mathfrak{X} \rightarrow \{0, 1\}$. Define the distance from a benchmark point to the validated region by

$$\text{dist}_{d_S}(b, \mathfrak{B}_S^{\text{obs}}) = \inf_{b' \in \mathfrak{B}_S^{\text{obs}}} d_S(b, b')$$

The sector acceptance set is the mathematical subset

$$\mathcal{C}_S = \{\theta \in \mathfrak{X} : P_S(\theta) = 1, \quad \text{dist}_{d_S}(\mathcal{B}_S(\theta), \mathfrak{B}_S^{\text{obs}}) \leq \epsilon_S, \quad \mathcal{G}_S(\theta) = 1\}$$

The shared acceptance intersection is

$$\mathcal{C}_{\text{AAA}} = \bigcap_{S \in \mathfrak{S}} \mathcal{C}_S$$

A closure attempt survives the validation gate only as an element of $\mathcal{C}_{\text{AAA}}$. A sector result that lies in one $\mathcal{C}_S$ but in no element of the full intersection remains a local result rather than a promoted AAA closure.

Residual-Bearing Criticism

A proposed failure claim must name the coordinate in θ_{join} , the sector predicate P_S , the benchmark distance, the no-go predicate $\mathcal{G}_S$, or the residual family $\mathcal{R}$ that it changes. Generic skepticism that leaves the closure record and every residual unchanged is not a closure-blocking condition. It may remain a comparison concern, but it does not promote to a validation failure until it moves an existing gate.

Let q be a proposed criticism of a candidate record θ . The notation $P_S(\theta; q)$, $\mathcal{B}_S(\theta; q)$, and $\mathcal{G}_S(\theta; q)$ means that the corresponding sector gate has been re-evaluated after applying the claimed change. Then q can block promotion only if

$$[\exists S \in \mathfrak{S} : P_S(\theta; q) = 0] \vee [\exists S \in \mathfrak{S} : \text{dist}_{d_S}(\mathcal{B}_S(\theta; q), \mathfrak{B}_S^{\text{obs}}) > \epsilon_S] \vee [\exists S \in \mathfrak{S} : \mathcal{G}_S(\theta; q) = 0]$$

This rule does not make the validation suite less severe. It prevents a residual-bearing closure record from being rejected by a criticism that has not identified which accepted observable, mathematical consistency condition, or no-go assumption has actually changed.

Null-Result Residual for Added Channels

When a closure attempt predicts channels outside the validated Standard Model and GR-facing benchmark set, those channels must be tested against null results before the record can be promoted. Let $\mathfrak{C}_\theta^{\text{new}}$ be the set of predicted additional channels for a candidate record θ : unstable baryon channels, new charged or neutral partners, extra gauge or transport modes, preferred-frame leakage channels, or other non-baseline outputs that would have produced an observed rate, cross-section, lifetime shift, branching ratio, dispersion, or anisotropy. For each channel e , let $O_e(\theta) \geq 0$ be the predicted observable and O_e^{max} the accepted upper bound in the comparison regime. Define

$$\mathcal{R}_{\text{null}}(\theta) = \sup_{e \in \mathfrak{C}_\theta^{\text{new}}} \left[\log \frac{O_e(\theta)}{O_e^{\text{max}}} \right]_+, \quad [x]_+ \equiv \max(x, 0)$$

A promoted record must satisfy

$$\mathcal{R}_{\text{null}}(\theta) = 0$$

using the same shared coordinates θ_{join} that recover the positive benchmarks. A channel may avoid this gate only by being outside the validated comparison domain, by being an exactly unobservable gauge redundancy, or by being proven absent in the accepted branch family. It is not enough to add a large symmetry, partner family, hidden transport dimension, or unstable reaction corridor and then tune it below every bound with sector-specific parameters.

For symmetry-container comparisons, the extra-sector test is part of the positive claim rather than a later cleanup. If a larger algebra, hidden sector, or partner family is invoked to explain one observed pattern, every non-baseline channel it brings into the tested domain must either be exactly redundant, absent in the accepted branch family, or routed through the same $\mathcal{R}_{\text{null}}^{\text{op}}$ record that recovered the observed pattern.

Operational Null-Result Ledger

For audits and simulations, the same condition should be expanded into a channel ledger rather than left as a single symbol. Let θ_+ denote the record used for the positive Standard-Model, GR, quantum, and cosmology benchmarks, and let θ_e denote the record used to suppress a predicted non-baseline channel e . Define the shared-record split

$$\Delta_{\text{shared}}(e; \theta) = \text{dist}_{\text{shared}}(\pi_{\text{shared}}\theta_e, \pi_{\text{shared}}\theta_+)$$

where π_{shared} keeps the common Noether sea, assembly, weak-exposure, metric, and provenance coordinates consumed by both the positive benchmark and the null channel. The operational audit residual is

$$\mathcal{R}_{\text{null}}^{\text{op}}(\theta) = \sup_{e \in \mathfrak{C}_\theta^{\text{new}}} \left(\left[\log \frac{O_e(\theta)}{O_e^{\text{max}}} \right]_+ + \lambda_{\text{split}} \Delta_{\text{shared}}(e; \theta) \right)$$

The original promotion condition is recovered by requiring $\mathcal{R}_{\text{null}}^{\text{op}}(\theta) = 0$. This form rejects a second failure mode: a channel can be numerically hidden but still fail because its suppression uses a different shared record from the one that fit the observed sector.

Added-channel family	Example observable $O_e(\theta)$	Null data product	Same-record requirement
Mirror matter or added charged partners	production cross-section, branching ratio, stable relic abundance	collider exclusions, precision electroweak fits, cosmological abundance bounds	the axial-layer and gauge-representation record that yields observed fermions must also exclude the partner branch
Superpartners or large symmetry partners	missing-energy rate, partner mass threshold, coupling strength	collider missing-energy and resonance searches	partner absence must follow from the accepted branch family, not from an independent mass threshold

Added-channel family	Example observable $O_e(\theta)$	Null data product	Same-record requirement
Proton-instability or baryon-violating corridors	$\Gamma_p(\theta)$ or forbidden nuclear transition rate	proton-lifetime and rare-event limits	the same color/topology and reaction-provenance ledger used for hadrons must suppress the channel
Extra gauge bosons or gauge modes	resonance rate, precision-contact term, long-range force strength	collider, fifth-force, and precision-scattering bounds	the effective gauge residual must recover $U(1)_Y \times SU(2)_L \times SU(3)_c$ without an unsuppressed added mode
Hidden transport or extra propagation modes	dispersion, birefringence, scalar/vector gravitational-wave response	photon, gravitational-wave, and timing residuals	the same Noether sea response map must set clock, signal, and metric channels
Sterile or neutral partner branches	mixing angle, ΔN_{eff} , relic abundance, free-streaming scale	oscillation, BBN, CMB, and structure-formation bounds	the neutral-sector Hamiltonian and cosmology record must be shared
Preferred-frame leakage channels	two-way anisotropy, clock drift, PPN preferred-frame coefficients	resonator, atomic-clock, solar-system, and gravitational-wave timing bounds	the Lorentz-closure map must suppress leakage without retuning clock, ruler, or signal coefficients

For the hidden-transport family, free-space birefringence is a direct null-result specialization rather than a new ontology. If $v_+(\omega, \hat{\mathbf{k}}; \theta)$ and $v_-(\omega, \hat{\mathbf{k}}; \theta)$ are the two physical photon-polarization propagation speeds extracted from the same record θ , define

$$\mathcal{R}_{\text{biref}}(\theta) = \sup_{\omega, \hat{\mathbf{k}}} \left| \frac{v_+(\omega, \hat{\mathbf{k}}; \theta) - v_-(\omega, \hat{\mathbf{k}}; \theta)}{c_0} \right|$$

The photon/effective-metric record can be promoted only when $\mathcal{R}_{\text{biref}}(\theta) \leq \epsilon_{\text{biref}}$ in the declared weak homogeneous regime and when the same θ also supplies the clock, ruler, signal, and metric coefficients used for the positive GR-facing benchmarks. If birefringence is numerically hidden by switching to a different channel record than the one used for lensing, Shapiro delay, spectra, or photon synchronization, $\mathcal{R}_{\text{null}}^{\text{op}}$ fails even if the split is individually small.

Null-Result Ownership Matrix

The following matrix assigns each recurring null-result family to the corpus homes that should carry the positive derivation and the absence proof. The owner document does not need to reproduce every experimental limit; it must state the observable $O_e(\theta)$, name the comparison bound O_e^{max} , and route the channel through $\mathcal{R}_{\text{null}}^{\text{op}}$ when the channel is predicted.

Channel family	Observable vector	Bound symbol	Primary owner	Supporting gates
Mirror matter / added charged fermions	$(\sigma_{\text{prod}}, B_{\text{vis}}, \Omega_{\text{relic}})$	$O_{\text{mirror}}^{\text{max}}$	Quantum Number Mapping	Gauge Symmetries, Known Tensions
Superpartners / symmetry partners	$(\sigma_{\text{miss}}, m_{\text{partner}}, B_{\text{cascade}})$	$O_{\text{partner}}^{\text{max}}$	Gauge Symmetries	Theory Differentials, No-Go Theorems
Proton-instability corridors	$(\Gamma_p, B_{p \rightarrow e^+ \pi^0}, B_{p \rightarrow \bar{\nu} K^+})$	Γ_p^{max}	Color Charge SU(3)	Reaction Ledger, Known Tensions
Extra gauge bosons / gauge modes	$(\sigma_{Z'}, \sigma_{W'}, g_{\text{new}}, \Delta_{\text{contact}})$	$O_{\text{gauge}+}^{\text{max}}$	Gauge Symmetries	Gauge Structure Emergence, Electroweak Bosons
Hidden transport / extra propagation modes	$(\Delta v/c, \omega_{\text{disp}}, h_{\text{scalar}}, h_{\text{vector}})$	$O_{\text{transport}}^{\text{max}}$	Constraint Ledger	Observer Framework, PPN Parameters
Sterile / neutral partner branches	$(\theta_{\text{mix}}, \Delta N_{\text{eff}}, \Omega_{\nu R}, \lambda_{\text{fs}})$	$O_{\text{sterile}}^{\text{max}}$	Neutrinos	Dark Matter, CMB
Preferred-frame leakage	$(\Delta_{\text{tw}}, \delta \nu/\nu, \alpha_1, \alpha_2, \alpha_3)$	$O_{\text{LV}}^{\text{max}}$	Lorentz Kinematics	PPN Parameters, Constraint Ledger

For proton-instability corridors, convert every current partial-mean-life lower limit τ_c^{min} into a channel-rate ceiling

$$\Gamma_{p,c}^{\text{max}} = \frac{1}{\tau_c^{\text{min}}}$$

The current benchmark scale is already severe: PDG 2024 summaries give $\tau/B(p \rightarrow e^+ \pi^0) > 2.4 \times 10^{34} \text{ yr}$ and proton neutrino/kaon modes near $5.9 \times 10^{33} \text{ yr}$ at 90% confidence. These numbers are comparison anchors, not permanent constants; a closure packet should cite the current experimental source when the hadronic gate is evaluated.

Sector Acceptance Sets

Sector	Predicate $P_S(\theta) = 1$	Benchmark condition	Falsifier
$\mathcal{C}_{\text{weak}}$	One weak-coupling-triad exposure record $\mathcal{E}_{\text{weak}}(A) = Q_{\text{weak}}[\Pi_{\text{weak}} \mathcal{L}_A]$ supplies V-A, CKM/PMNS overlap, and weak-corridor provenance without redefining Π_{weak} , Q_{weak} , or the exposed domain.	$\mathcal{B}_{\text{weak}}(\theta)$ lies in the observed charged-current handedness, mixing, and provenance region within ϵ_{weak} .	Right-handed charged-current coupling is not strongly suppressed in the validated regime, or the weak exposure domain changes between chirality, mixing, and provenance.
$\mathcal{C}_{\text{quantum}}$	A transfer operator or return map $\mathcal{T}_{\Delta t}$, basin partition $\{B_i\}$, invariant or metastable measure μ_* , and detector kernel produce $p_i = \mu_*(B_i)$ from Γ and $\mathcal{H}$ without assigning probabilities as an external rule.	$\mathcal{B}_{\text{quantum}}(\theta)$ lies in the Born-rule, Bell/CHSH/Tsirelson/GHZ/Hardy, Leggett-Garg temporal-correlation, detector-record, and no-signaling benchmark region within $\epsilon_{\text{quantum}}$.	The validated regime gives non-Born weights, a classical-axis linear-correlation failure, untracked temporal-measurement disturbance, superluminal signal transfer, or a detector kernel not derived from the recorded causal state.
$\mathcal{C}_{\text{gravity}}$	One Noether sea response map $\mathcal{M}_{\text{sea}}^{ab}$ supplies clock, ruler, effective signal-speed, weak-field	$\mathcal{B}_{\text{gravity}}(\theta)$ lies in the redshift, Shapiro-delay, lensing, orbital, gravitational-wave-speed, PPN, and preferred-frame bound region within $\epsilon_{\text{gravity}}$.	Clock, ruler, signal, or metric coefficients must be tuned independently, ordinary dissipative drag appears in stable

Sector	Predicate $P_S(\theta) = 1$	Benchmark condition	Falsifier
	metric, and PPN channels without changing coefficients per observable.		motion, or preferred-frame leakage exceeds the recorded bounds.
$\mathcal{C}_{\text{hadronic}}$	An accepted branch family A , exposure quotient, color/topology ledger, residual strong channel set, and $\mathcal{L}_{E\mathbf{p}\mathbf{J}}$ close confinement, quark mass, baryon-stability, and nuclear-binding rows.	$\mathcal{B}_{\text{hadronic}}(\theta)$ lies in the confinement, quark-hierarchy, proton-stability, deuteron, saturation, and alpha-like benchmark region within $\epsilon_{\text{hadronic}}$.	The sector predicts generic fast proton decay, unphysical nuclear binding signs, missing color/topology closure, or an unbalanced architrino / Noether braid inventory.
$\mathcal{C}_{\text{radiation}}$	A radiation residual $\mathcal{R}_\Theta$ selects admissible channels from $\{B_i\}$ and closes $\mathcal{L}_{E\mathbf{p}\mathbf{J}}$ with photon output, recoil, medium update, signed photon-frequency exchange, non-radiative remnant, or reaction rows explicitly recorded.	$\mathcal{B}_{\text{radiation}}(\theta)$ lies in the Larmor/Lienard, bremsstrahlung, synchrotron, pair-threshold, Compton-like, SZ-like transfer, and blackbody benchmark region within $\epsilon_{\text{radiation}}$.	Any benchmark requires per-observable retuning, untracked energy loss or gain, a missing recoil/provenance row, a free longitudinal photon mode, or a blackbody fit not tied to the event ledger.
$\mathcal{C}_{\text{cosmology}}$	One source, transport, signed photon-frequency-transfer, thermalization, and clock-rate record uses the same $\rho_{\text{NS}}(\mathbf{x}, t)$, $n(\mathbf{x}, t)$, $\chi_{\text{sea}}(\mathbf{x}, t)$, $\mathcal{M}_{\text{sea}}^{\text{ab}}$, and reaction provenance ledger across local source channels and observer-level cosmology.	$\mathcal{B}_{\text{cosmology}}(\theta)$ lies in the BBN, CMB blackbody, damping, anisotropy, polarization handoff, redshift-budget, $H(z)$, BAO, and growth benchmark region within $\epsilon_{\text{cosmology}}$.	BBN photon loading, CMB thermalization, redshift handoff, frequency-exchange closure, or structure growth requires unbalanced substrate creation, unlogged photon energy transfer, per-source retuning, or Noether sea variables incompatible with local reaction / radiation ledgers.

Promotion Lemma

For sector S , let $\pi_S : \mathfrak{X} \rightarrow \mathfrak{X}_S$ be the projection that keeps the sector- S coordinates and shared coordinates consumed by that sector. For a local sector result $c \in \mathfrak{X}_S$, define the extension fiber

$$\text{Ext}_S(c) = \{\theta \in \mathcal{C}_{\mathbb{A}\mathbb{A}\mathbb{A}} : \pi_S(\theta) = c\}$$

Lemma. A local sector result c is promotable through the validation gate if and only if $c \in \pi_S(\mathcal{C}_S)$ and

$$\text{Ext}_S(c) \neq \emptyset$$

Proof route: if c is promoted, the promoted record must retain the sector- S result and pass every sector gate, so it is an element of $\text{Ext}_S(c)$. Conversely, any $\theta \in \text{Ext}_S(c)$ is a shared closure record whose sector- S projection equals c and whose weak, quantum, gravity, hadronic, radiation, and cosmology predicates all pass; therefore the local result has survived the validation gate. If the fiber is empty, the result is blocked by at least one sector predicate, benchmark region, no-go record, or failure condition.

Incompatibility Witnesses

A local claim c imposes a constraint subset $I(c) \subseteq \mathfrak{X}$ consisting of all closure records that preserve the claim's definitions, coefficients, ledger rows, and effective-limit assumptions. For a target sector T , define the constrained target set

$$\mathcal{C}_T \mid c = \mathcal{C}_T \cap I(c)$$

An incompatibility witness from sector S to sector T is the object

$$W_{S \rightarrow T}(c) = (c, T, I(c), P_T, \mathcal{B}_T, \mathfrak{B}_T^{\text{obs}}, d_T, \epsilon_T, \mathcal{G}_T, \delta_T(c))$$

where

$$\delta_T(c) = \epsilon_T - \inf_{\theta \in I(c), \Pr(\theta)=1, \mathcal{G}_T(\theta)=1} \text{dist}_{d_T}(\mathcal{B}_T(\theta), \mathfrak{B}_T^{\text{obs}})$$

The witness empties the target gate when $\mathcal{C}_T \mid c = \emptyset$. It damages the target gate when $\mathcal{C}_T \mid c \neq \emptyset$ but $\delta_T(c)$ removes a required tolerance margin, forces a hidden sector-specific parameter split, or leaves a required ledger row undefined.

Witness class	Imposed local claim c	Target effect	Failure code
Weak-domain split	$I(c)$ requires distinct weak exposure domains for V-A, CKM/PMNS, and weak-corridor provenance.	$\mathcal{C}_{\text{weak}} \mid c = \emptyset$ because P_{weak} requires one weak-coupling-triad exposure record.	weak.hidden_domain_split
Gravity coefficient split	$I(c)$ requires separate clock, ruler, signal, and PPN coefficients not derived from one $\mathcal{M}_{\text{sea}}^{\text{ab}}$.	$\mathcal{C}_{\text{gravity}} \mid c = \emptyset$ if the split is needed for benchmark recovery.	gravity.hidden_tuning
Radiation-cosmology split	$I(c)$ fits blackbody recovery with $\chi_{\text{sea}}^{\text{CMB}}(\mathbf{x}, t)$ incompatible with the BBN or local radiation event ledger.	$\mathcal{C}_{\text{cosmology}} \mid c = \emptyset$ or $\delta_{\text{cosmology}}(c) < 0$.	cosmology.incompatible_transport_limit
Quantum signal leak	$I(c)$ recovers Bell correlations through a detector kernel that transfers controllable signals outside the causal-wake ledger.	$\mathcal{C}_{\text{quantum}} \mid c = \emptyset$ and the same record damages $\mathcal{C}_{\text{gravity}}$ through preferred-frame leakage.	quantum.signal_transfer
Event-ledger omission	$I(c)$ routes radiation, reaction, measurement, or strong-field release without a required E , $\mathbf{p}$, $\mathbf{J}$, polarity, provenance, medium, or remnant row.	The target sector using that event has no admissible $\mathcal{L}_{E\mathbf{p}\mathbf{J}}$ completion.	event.missing_ledger_row

Witness class	Imposed local claim c	Target effect	Failure code
Null-result violation	$I(c)$ predicts a non-baseline channel $e \in \mathcal{E}_\theta^{\text{new}}$ with $O_e(\theta) > O_e^{\text{max}}$ in a tested regime.	The relevant sector gate may fit its positive benchmark, but the shared closure record fails $\mathcal{R}_{\text{null}}(\theta) = 0$.	null.observed_absence_violation

Testable Failure Modes

Failure mode	Mathematical test	Routed workstream
Empty intersection	$\mathcal{C}_{\text{AAA}} = \emptyset$ or $\text{Ext}_S(c) = \emptyset$ for a proposed local promotion.	Known Tensions, Closure Scorecard
Hidden tuning	A shared variable or map has sector-specific values $p_S \neq p_T$ with no recorded state variable, or the same benchmark family is recovered only by changing $\Pi_S, Q_S, \mathcal{R}, \{B_i\}$, the branch-chart revision record, equality map, root-coordinate split, $\mathcal{M}_{\text{sea}}^{\text{ab}}$, $\rho_{\text{NS}}(\mathbf{x}, t)$, or $\chi_{\text{sea}}(\mathbf{x}, t)$ between cases. Branch-chart revisions selected after residual inspection rather than declared from branch geometry fail this test.	Parameter Ledger, Constraint Ledger
Null-result violation	$\mathcal{R}_{\text{null}}(\theta) > 0$ for a predicted added channel in a validated comparison regime.	Known Tensions, Constraint Ledger
Missing conservation/provenance field	$\mathcal{L}_{\text{EPI}}(e)$ has an undefined or nonzero required row after all claimed outputs, recoil, medium updates, remnants, polarity / charge, architrino inventory, source identity, emission time, causal-root branch, and branch-Jacobian records are included.	Reaction Ledger, Reaction-Cosmology Provenance Ledger
Benchmark-only fitting	A target benchmark in $\mathfrak{B}_S^{\text{obs}}$ is used as an input to $\mathcal{L}_A, \Pi_S, Q_S, \mathcal{R}, \{B_i\}$, a branch-chart revision, an equality map, a root-coordinate split, or $\mathcal{M}_{\text{sea}}^{\text{ab}}$ rather than as an output of a replayable closure record.	Particle Masses, Measurement Ontology, Radiation
Incompatible effective limits	Two sectors require asymptotic maps whose overlap is empty, for example incompatible weak-field metric limits, photon / radiation limits, blackbody / BBN transport limits, or quantum no-signaling / gravity causal limits.	Known Tensions, General Relativity, Cosmology Ontology

Preferred-Frame Hiding Stop Condition

- Hard wall:** If the Euclidean-void rest frame is detectable by any physical experiment, for example a Michelson-Morley-type null test, at $\Delta c/c > 10^{-17}$, the theory fails.
- Required compensation:** Moving assemblies must acquire the Lorentz-compatible deformation and clock laws, $L_{\parallel} = L_0/\gamma$ and $T = \gamma T_0$, from delayed causal closure and Noether sea response rather than from kinematic postulates.
- Coefficient closure:** Clock, ruler, signal, and metric response coefficients must suppress two-way anisotropy and other preferred-frame leakage to the validated bounds. A qualitative contraction story is not sufficient.
- Dissipative drag:** If the Noether sea induces ordinary drag that slows cosmological bodies without a conserving medium-dressed response mechanism, the theory is falsified.

Critical Stop Conditions

- c_f variance:** If field speed varies in the true void, the theory fails.
- Noether sea drag:** If the Noether sea causes orbital decay or secular kinetic-energy loss through ordinary dissipative drag, rather than a reversible medium-dressed inertial response, the theory fails.
- Lorentz leakage:** If absolute motion affects atomic spectra above 10^{-17} , the theory fails.
- Empty shared intersection:** If quantitative development makes $\mathcal{C}_{\text{AAA}} = \emptyset$, the present implementation is rejected even if individual sector chapters remain locally suggestive.

No Go Theorems

This chapter classifies the formal obstruction results that act as validation filters for **AAA**. A no-go theorem is not useful here as a decorative citation. It is useful only when its assumptions, conclusion, and replacement burden can be recorded against a candidate closure.

The operational companion is [Failure Criteria](#). That page defines the shared closure intersection. This page defines how a theorem enters one sector gate: directly as a rejection condition, as an assumption mismatch, as a replacement constraint, or as an irrelevant comparison.

Applicability Record

For a no-go family G , let $\mathcal{A}_G$ be its assumption set and let

$$\sigma_{\theta,G} : \mathcal{A}_G \rightarrow \{\text{accepted, rejected, replaced, effective, absent}\}$$

record the **AAA** stance toward each assumption in the candidate record θ . The applicability class is

$$\text{app}(G, \theta) \in \{\text{direct, assumption mismatch, replacement constraint, irrelevant comparison}\}$$

The class is `direct` when the theorem’s assumptions are accepted or effective in the tested regime and its conclusion applies as a rejection condition. The class is `assumption mismatch` when a required assumption is rejected or absent and the theorem does not by itself supply a validated replacement burden. The class is `replacement constraint` when an assumption is rejected or replaced but the theorem protects a validated behavior that the candidate record must recover by `AAA` objects. The class is `irrelevant comparison` when G shares no benchmark variable, conservation condition, or effective limit with the local claim under test.

Applicability Map

No-go family	Applicability class	Assumption status	Replacement constraint or falsifier
Bell/CHSH/Tsirelson, including GHZ and Hardy subbenchmarks	replacement constraint	Bell local-causality, ordinary common-cause screening, Markov screening, or context-independent local value assumptions are not substrate assumptions when $\mathcal{H}$ and detector response are retained; no-signaling, validated correlation bounds, GHZ perfect-correlation products, and Hardy zero/positive probability patterns remain benchmark constraints.	Derive pair provenance, detector kernels, Born weights, no-signaling, Tsirelson-compatible correlations, GHZ product signs, and Hardy event margins from $\mathcal{T}_{\Delta t}$, $\{B_i\}$, and μ_+ . Record reconstruction is not sufficient unless the induced joint record measure also passes the Bell, no-signaling, measurement-independence, factorization-residual, GHZ parity, and Hardy-event gates. Failure occurs if the model reduces to the classical-axis linear-correlation mode, uses controllable superluminal transfer, treats final records as an explanation without deriving their tested joint distribution, lets the declared common-past record screen the wings into a Bell-local product law, assigns context-independent local values across GHZ contexts, or erases Hardy’s zero-probability constraints while claiming the positive event.
Kochen-Specker / noncontextual operator values	replacement constraint	A context-independent value assignment to every self-adjoint observable is not a substrate assumption. Effective operator values exist only after a preparation, apparatus kernel, coarse-graining, and record channel are declared. The protected benchmark is the quantum contextuality pattern: commuting context products, compatible shared marginals, and the absence of a global noncontextual value map in validated regimes.	Derive context-indexed apparatus records $r_{O,C} = R_{O,C}(\Phi_{\mathcal{C}}^{\text{tot}}(\Gamma_0; \mathcal{K}_C))$ from one substrate flow, while recovering the declared context product constraints and shared-observable marginals. Failure occurs if the closure silently assigns substrate values to all effective operators, changes the target state per context, or recovers contextuality only by making apparatus records inconsistent across overlapping calibrated contexts.
Pusey-Barrett-Rudolph quantum-state reality theorem	replacement constraint	Preparation independence and ontic-state overlap assumptions are not substrate axioms; the wavefunction is observer-level bookkeeping rather than a primitive physical field. The protected benchmark is stronger: independently prepared systems must have declared preparation records, product or non-product provenance status, and the standard state-discrimination statistics.	A candidate wavefunction account must state whether its substrate preparation measure factorizes for independently prepared systems and must expose any provenance correlation needed to avoid the theorem. Failure occurs if the model treats overlapping effective wavefunctions as harmless while also accepting product preparation independence and the PBR measurement statistics, or if it evades the theorem by hiding unrecorded correlations between supposedly independent preparation devices.
Leggett-Garg temporal-correlation inequalities	replacement constraint	Macroscopic realism per se and noninvasive measurability are not substrate axioms. A measurement in <code>AAA</code> is a physical apparatus-target coupling, so temporal readouts may disturb later basin dynamics; the protected benchmark is the observed sequential-correlation data together with an explicit disturbance ledger.	A candidate measurement account must declare the apparatus kernels used at each time, recover the tested temporal correlators, and report whether earlier probes perturb later record statistics. Failure occurs if the model asserts a definite macro-trajectory with noninvasive readout while accepting a Leggett-Garg violation, or if it explains the violation only by untracked apparatus disturbance rather than a declared record-channel residual.
Frauchiger-Renner / Wigner-friend observed-observer consistency	replacement constraint	The standard no-go setup assumes that quantum state descriptions can be applied to other theory-users, that one observer may import another observer’s certified certainty, and that one declared record channel cannot certify mutually exclusive outcomes. <code>AAA</code> rejects an external classical-observer cut, but it also rejects importing another observer’s conclusion without a physical record channel, access region, apparatus kernel, and boundary-data model.	A measurement closure that includes observed Physical Observers must derive every imported statement from the same substrate flow, record-autonomy test, and finite communication channel used for ordinary apparatus records. Failure occurs if a model needs a hidden external observer, lets a Physical Observer import certainty without a durable record, treats an unbuildable reference/readout setup as a completed experiment, or allows two mutually exclusive outcomes to be certified inside one declared record channel. If the reference or readout channel cannot satisfy the physical record criteria, the thought experiment is blocked by realizability rather than promoted into ontology.
Groenewold-van Hove / global quantization map	replacement constraint	A global quantization map from all classical observables $C^\infty(M)$ to Hilbert-space operators, preserving every Poisson bracket as a commutator, is not a substrate assumption. The protected benchmark is narrower: in validated quantum regimes, the selected observer-level observables must recover the tested commutator algebra on the calibrated record domain.	Derive an admissible observable set from the same coarse-graining, apparatus kernel, retained path-history data, and record window used for the effective operator model, then bound the quantization-domain residual in Quantum Operator Mapping . Failure occurs if a closure claims bracket-to-commutator recovery for all smooth classical functions, uses a choice of polarization or representation as hidden ontology, or changes the observable domain per benchmark without recording the physical apparatus and coarse-graining that justify the restriction.
Lorentz invariance and preferred-frame tests	direct	Observer-level clock, ruler, two-way signal, PPN, and spectral bounds apply directly to any candidate effective metric or transport map.	Bound ϵ_{LV} , $\Delta_{lw}(\beta)$, PPN parameters, spectra, and gravitational-wave-speed differences within recorded limits. Failure occurs when absolute motion is detectable above the accepted thresholds.
Spin-statistics / exchange	replacement constraint	Local Lorentz-QFT axioms are not fundamental substrate assumptions, but matter stability and exchange classes are validated effective constraints.	Derive the ordered-frame lift, 4π spinor behavior, and bosonic/fermionic exchange classes from Noether braid topology and angular-momentum ledger. Failure occurs if the lift cannot separate fermionic and bosonic closure classes.

No-go family	Applicability class	Assumption status	Replacement constraint or falsifier
CPT theorem / local relativistic QFT assumptions	replacement constraint	Local relativistic QFT assumptions are not substrate assumptions for absolute time, Euclidean void, and delayed causal wakes. This includes local field operators, microcausal commutation structure, fundamental Poincare symmetry, and a Lorentz-invariant vacuum as primitive assumptions. The protected benchmarks remain observer-level particle/antiparticle mass degeneracy, charge-conjugate reaction bookkeeping, neutral-meson and lepton-sector CPT bounds, Lorentz-leakage bounds, and the absence of unobserved baryon/lepton channels.	Recover the tested CPT-facing benchmarks from architrino polarity, pro/anti assembly mapping, delayed dynamics, effective Lorentz closure, and the existing null-result ledger. A candidate record should publish a residual vector such as $\mathcal{R}_{\text{CPT}}(\theta) = (\Delta m_{\eta\bar{\eta}}, \Delta q_{\eta\bar{\eta}}, \Delta \Gamma_{\text{conj}}, \epsilon_{\text{LV}}, \mathcal{R}_{\text{null}})$ and show that each component stays within the declared experimental or closure bound. Failure occurs if the record hides rejected local-QFT assumptions inside the proof, predicts CPT-violating mass or reaction asymmetries above bounds, or restores the symmetry only by adding untracked channels outside $\mathcal{R}_{\text{null}}$.
Exact global architrino flips or permutations	assumption mismatch with replacement constraint when effective indistinguishability is claimed	Substrate architrinos are provenance-bearing entities with path-history and causal-wake records. A global flip, polarity reassignment, or label permutation is not exact unless it preserves those records and all causal-root relations, not merely the instantaneous exposed properties.	State whether the symmetry is a kernel/background symmetry, a full-history symmetry on a special state, or an effective coarse-grained equivalence. Effective exchange, gauge, flavor, or charge bookkeeping may be used only after the suppressed provenance data and replacement recovery target are named. Failure occurs if a closure treats provenance-suppressed interchangeability as substrate identity, or if an effective symmetry claim cannot recover the validated observer-level degeneracies, conservation laws, and exchange classes.
Coleman-Mandula / gauge unification constraints	assumption mismatch with replacement constraint when effective scattering is claimed	Exact Lorentz-invariant analytic S-matrix assumptions are not substrate assumptions for delayed absolute-time dynamics. Compact internal symmetry, unitarity, positive-energy particle states, and effective gauge-sector factorization become benchmarks when Standard-Model-facing scattering or mixing is claimed. A pre-effective symmetry container may evade the theorem's literal hypotheses only before observer-level spacetime, scattering states, and gauge factors have been recovered; after that recovery, the same record must reproduce the validated factorization and may not use mixed spacetime/internal generators to create observed-sector shortcuts.	State which assumptions are effective, recover compact internal gauge behavior in the tested regime, and derive gauge-like symmetries without contradicting observed factorization. Failure occurs if a claimed unification predicts forbidden effective-sector mixing, hides added channels outside $\mathcal{R}_{\text{null}}$, uses gauge covariance as an unexplained fit, or suppresses non-baseline sectors with a record different from the positive recovery record.
Weinberg-Witten-like obstructions	assumption mismatch with replacement constraint when emergent photon or gravity language is claimed	Lorentz-covariant conserved stress-tensor assumptions of the theorem are not fundamental substrate assumptions for Noether sea state and assembly closures. Photon and gravity claims must still recover the validated effective channels.	Keep photon and metric objects as medium/assembly closures with explicit domain limits. Failure occurs if the record claims a fundamental Lorentz-covariant composite photon/graviton while also denying the theorem's assumptions, or if effective limits cannot be recovered.
Boundary-Hamiltonian / kinematic-locality constraints on emergent gravity	replacement constraint when emergent gravity, boundary unitarity, or black-hole information claims are made	In generally covariant gravity comparisons, the Hamiltonian can be a boundary term, and Marolf-style arguments show that non-linear gravity is not straightforwardly recovered from a kinematically local theory with independently commuting bulk observables. AAA does not accept local QFT operator algebras, boundary Hamiltonians, or asymptotic boundary observables as substrate primitives, but the protected benchmark remains: effective gravity must carry unitary observer-level information accounting without freezing local dynamics or treating local horizon entanglement as a sharply defined substrate observable.	A candidate record must replace the rejected assumptions with finite boundary wake data, declared reference resources, access-region limits, and a Noether sea continuation map that recovers both local effective dynamics and boundary-accessible bookkeeping. Failure occurs if the model claims emergent GR from purely local commuting substrate variables, hides all bulk dynamics behind a boundary algebra, or treats horizon-crossing correlations as lost or recovered without a declared Physical Observer access model.
Global-GR underdetermination and observationally indistinguishable spacetime results	replacement constraint when a global cosmology, horizon, or effective-metric claim is promoted from observer records	Lorentzian manifold ontology, global spacetime extension classes, and model-class maximality assumptions are not substrate assumptions. The protected benchmark is methodological: rich local records and local-property preservation do not by themselves license a unique global reconstruction.	A promoted global claim must state the Physical Observer access region, data-product projection, local-induction assumptions, and ambiguity residual that make the claim invariant across admissible closure records. Failure occurs if a cosmology or strong-field packet treats a fitted FLRW, de Sitter, extension, or horizon interpretation as final ontology merely because it reproduces the observer-accessible data, or if it changes the admissible model class to obtain determinism or uniqueness without recording that assumption as part of the closure.
Massive-gravity and finite-range-gravity obstructions	replacement constraint when large-scale gravity modification is claimed	Fundamental massive-graviton and Lorentzian spin-2 assumptions are not AAA substrate assumptions. The protected constraints remain local GR recovery, bounded physical energy, stable mode counting, low-energy positivity bounds where the effective comparison domain accepts their assumptions, de Sitter or cosmological background bounds, and gravitational-wave polarization, speed, and dispersion limits.	A medium-response closure must recover the GR limit in validated regimes, keep perturbation energy bounded below after gauge and effective redundancies are removed, pass the accepted positivity tests for any claimed low-energy effective scattering or response map, and prevent extra scalar or longitudinal gravitational-wave modes or finite-range drift from exceeding observational bounds. Failure occurs if large-scale weakening is obtained only by allowing ghost-like negative-energy modes, positivity-violating effective coefficients, order-one solar-system deviations, or unconstrained gravitational-wave dispersion or polarization.

No-go family	Applicability class	Assumption status	Replacement constraint or falsifier
AdS/CFT, island, replica-wormhole, string, or loop-quantum-gravity comparison constraints	irrelevant comparison unless a specific tested benchmark is imported	These frameworks are comparison tools unless the local packet imports a precise entropy, unitarity, horizon, or observational condition as a gate.	No acceptance burden is created by analogy alone. A burden is created only by a named benchmark such as area-scaling entropy, Page-curve-compatible accounting, horizon regularity, or direct compact-object data.

Black-hole CPT comparisons are handled by the CPT row, not by importing global mirror-boundary ontology. If a horizon-interface packet uses CPT or thermal-equilibrium language, it must publish the corresponding formation/release balance residual in the black-hole chapter and keep $\mathcal{R}_{\text{CPT}}(\theta)$ within the tested particle-sector bounds. A record fails if it restores apparent balance only by adding untracked release channels, spectator species outside $\mathcal{R}_{\text{null}}$, or a second state record for horizon entropy.

Cosmic Bell tests sharpen the Bell row by converting measurement-independence leakage into an observationally bounded residual. When detector settings are chosen from distant photons, quasars, or other causally screened sources, a candidate closure may not rely on an untracked common cause linking those settings to the pair-preparation record. Such a route must be recorded as nonzero Δ_{MI} and compared against the experimental setting-source covariance bound rather than hidden inside pair provenance.

The GHZ and Hardy subbenchmarks sharpen the Bell row by removing any reliance on a single CHSH average. For a calibrated three-party GHZ setup, the four product contexts $\mathcal{C}_{\text{GHZ}} = \{XXX, XYY, YXY, YYX\}$ carry signs χ_C whose product is -1 , while any context-independent local value assignment makes the product $+1$. A compact record residual is

$$\Delta_{\text{GHZ}} = \max_{C \in \mathcal{C}_{\text{GHZ}}} [1 - \chi_C E_\theta(C)]_+$$

where $E_\theta(C)$ is the product expectation for the declared apparatus context and $[x]_+ \equiv \max(x, 0)$. For a Hardy setup with binary observables U_i, D_i , use the zero-probability constraints and positive Hardy event as a margin:

$$\Delta_{\text{Hardy}} = [P_\theta(D_1 = 1, D_2 = 1) - P_\theta(U_1 = 1, U_2 = 1) - P_\theta(D_1 = 1, U_2 = 0) - P_\theta(U_1 = 0, D_2 = 1)]_+$$

A useful Bell-family closure must make Δ_{GHZ} small on the perfect-correlation contexts, produce the positive Hardy margin where the experiment requires it, and still keep Δ_{MI} and Δ_{NS} inside tolerance. These are validation targets for the joint record measure, not new ontology.

The Kochen-Specker row includes the Mermin-Peres magic square as a preferred compact benchmark when a candidate operator map claims contextuality recovery. In that subcase the six commuting row/column contexts carry product signs $\chi_C \in \{+1, +1, +1, +1, +1, -1\}$. The closure must derive context-indexed apparatus records that satisfy those products and preserve shared marginals while refusing a global noncontextual value map. A proof that only assigns prewritten substrate values to all effective operators fails the parity check: each observable appears twice, so the product of all assigned values is $+1$, whereas the benchmark product signs multiply to -1 .

The Pusey-Barrett-Rudolph row is a preparation-independence audit, not a license to ignore independent preparation. For two declared preparations P_A and P_B , with substrate preparation measures $\rho_A(\lambda_A|P_A)$, $\rho_B(\lambda_B|P_B)$, and joint measure $\rho_{AB}(\lambda_A, \lambda_B|P_A, P_B)$, define

$$\Delta_{\text{PI}} = D_{\text{TV}}(\rho_{AB}(\lambda_A, \lambda_B|P_A, P_B), \rho_A(\lambda_A|P_A)\rho_B(\lambda_B|P_B))$$

If a candidate avoids the theorem by allowing $\Delta_{\text{PI}} > 0$, that residual must be tied to a physical shared-provenance, boundary-data, or apparatus-coupling record. Otherwise it is an untracked preparation correlation. The useful closure target is therefore two-part: recover the PBR state-discrimination statistics in the declared record channel while reporting whether the substrate preparation measure factorizes. If both the PBR measurement statistics and preparation independence are accepted in the same domain, overlapping effective wavefunction descriptions cannot be treated as a harmless epistemic overlap.

The Leggett-Garg row protects temporal correlation data without importing macrorealism as ontology. For dichotomic records $q_i \in \{-1, +1\}$ at times t_i , define

$$C_{ij} = \sum_{q_i, q_j = \pm 1} q_i q_j P_\theta(q_i, q_j | \mathcal{K}_i, \mathcal{K}_j), \quad K_{\text{LG}} = C_{12} + C_{23} - C_{13}$$

Macrorealism plus noninvasive measurability gives $K_{\text{LG}} \leq 1$ for this sign convention. The AAA replacement burden is not to accept noninvasive readout, but to declare the disturbance residual

$$\Delta_{\text{NIM}} = \sup_{i < j} D_{\text{TV}}(P_\theta(q_j | \mathcal{K}_j), P_\theta(q_j | \mathcal{K}_i, \mathcal{K}_j))$$

recover the observed K_{LG} -type statistics, and state whether the violation is carried by ordinary record-forming apparatus coupling, weak-probe disturbance, or a still-unclosed measurement model. A result that leaves Δ_{NIM} implicit has not converted the Leggett-Garg comparison into a usable validation gate.

For finite-range gravity comparisons, positivity bounds should be treated as an effective-domain filter, not as imported ontology. Let $E_{\text{min}}^{\text{phys}}(\theta)$ denote the lowest physical perturbation energy after gauge and redundant variables are removed, and let $\Pi_a(\theta)$ denote the low-energy positivity

functionals whose signs are fixed by the accepted comparison theorem for the declared scattering or response domain. A compact residual for a candidate large-scale weakening record is

$$\mathcal{R}_{\text{range}}(\theta) = w_{\text{GR}} \mathcal{R}_{\text{GR}}(\theta) + w_E \left[\frac{-E_{\text{min}}^{\text{phys}}(\theta)}{\epsilon_E} \right]_+^2 + w_{\text{pos}} \sum_a \left[\frac{-\Pi_a(\theta)}{\epsilon_{\text{pos},a}} \right]_+^2 + w_{\text{pol}} \frac{\mathcal{P}_{\text{extra}}}{\mathcal{P}_{\text{TT}}} + w_{\text{disp}} \int_{B_{\text{CW}}} \left| \frac{\partial^2 \omega_\theta}{\partial k^2} \right|^2 d \log f + w_{\text{cos}} \mathcal{R}_{\text{shared}}(\theta)$$

where $[x]_+ \equiv \max(x, 0)$. The record is useful only if one shared Noether sea response map can make this residual small. A result that passes local GR tests by changing the energy, positivity, polarization, dispersion, or cosmology record separately is not a promoted closure.

Use in Validation

A candidate closure record must name the no-go family it touches and fill the applicability record before the result can be promoted. If $\text{app}(G, \theta) = \text{direct}$, the theorem's conclusion is a hard rejection condition. If $\text{app}(G, \theta) = \text{replacement constraint}$, the rejected assumption does not remove the burden; it only changes the object that must carry the validated behavior.

The no-go record therefore becomes one component of the sector predicate $\mathcal{G}_S(\theta)$ used in [Failure Criteria](#). A result that passes a local benchmark but evades the relevant theorem by changing assumptions without supplying the replacement constraint is not a closure result.

Known Tensions

This chapter is the pressure ledger for the present repo state. Its purpose is to collect the unresolved burdens that matter most for closure without mixing them with vague future ideas or low-stakes wishlist items.

Purpose

This chapter is the pressure ledger for [AAA](#). It collects the places where the framework is not yet closed, where the present derivation stack is thinner than the claim it supports, or where current observations impose a hard quantitative burden that the repo has not yet fully carried.

This page is not a dumping ground for vague uncertainty. Each tension should identify:

- the issue,
- why it matters,
- the current repo status,
- the closure target,
- and the failure condition.

Severity Scale

- **Tier 1:** could directly falsify the present architecture if not resolved.
- **Tier 2:** does not immediately kill the architecture, but blocks a serious Standard-Model or GR-level closure claim.
- **Tier 3:** important downstream completion issue, but not yet the main credibility gate.

Pressure Ledger

Tier	Issue	Why it matters	Current repo status	Closure target	Failure condition
1	Weak V–A selection rule	The weak interaction must distinguish left-chiral fermions from right-chiral ones.	quantum-number-mapping.md gives a geometric lock-out story, and weak-mixing-ckm.md now identifies this as part of the shared weak-coupling-triad exposure problem, but no operator derivation is complete.	Derive a docking or coupling operator that exposes the weak-coupling triad for left-handed charged-current coupling, hides it for right-handed charged-current coupling, and then reuses the same domain for CKM/PMNS overlap and weak-reaction provenance.	If right-handed neutrino or right-handed charged-fermion coupling to W is not strongly suppressed in the same regime, or if the exposure domain must be redefined separately for mixing and provenance, the current weak-sector picture fails.
1	Preferred-frame leakage	The ontology has absolute time and a medium, so observer-level Lorentz hiding must be quantitative.	The requirement is clear in constraint-ledger.md , and Lorentz Kinematics now states the moving-assembly coefficient targets plus the translating-binary residual test, but the full attractor proof is not complete.	First solve the translating two-body branch and test $T_u/T_l = \gamma_f$ and $L_0/L_\perp = 1/\gamma_f$ on the same causal-root ledger; then show that nested shell braid clocks, rulers, and signal transport suppress measurable preferred-frame effects below current experimental bounds through coupled shape, clock, and two-way anisotropy closure.	Any robust preferred-frame signal above the recorded bounds, a non-Lorentzian binary residual that cannot be traced to a controlled branch feature, or any need to tune clock and ruler coefficients independently falsifies the observer-level spacetime closure.
1	Born-rule derivation	Quantum replacement claims are not credible without a basin-measure or equivalent statistical closure.	wavefunction-ontology.md and measurement-ontology.md fix the ontology; quantum-operator-mapping.md now states the finite-time invariant-measure, thermodynamic ensemble consistency, and	Derive outcome weights from deterministic basin measures in the same regime that yields the effective wave equation, show that the same finite-window measure projects to the	If the deterministic closure produces a non-Born weighting in validated regimes, if Born weights and thermodynamic summaries require incompatible measures, or if the

Tier	Issue	Why it matters	Current repo status	Closure target	Failure condition
			admissible quantization-domain targets, but the derivation is still open.	thermodynamic summaries used for apparatus irreversibility, decoherence, and record formation, and restrict effective operators to a physically declared observable domain rather than a global quantization of all classical functions.	operator map requires ad hoc observable-domain changes per benchmark, the current quantum story fails.
1	Weak-field GR recovery	Redshift, Shapiro delay, lensing, and orbital tests must come from one constitutive map.	The interface now exists in general-relativity.md and ppn-parameters.md , but the shared fit is incomplete.	Produce one reusable parameter set for the weak-field metric map.	If different observables require incompatible constitutive coefficients, the emergent-metric program fails.
2	Low-energy quantum-gravity EFT recovery	Quantized metric methods are not AAA ontology, but their long-distance effective predictions are fixed by known low-energy degrees of freedom.	general-relativity.md and emergent-metric.md state the classical weak-field map; they need an explicit observer-level GR-EFT recovery gate.	Recover the standard long-distance quantum correction to the Newtonian potential using the same weak-field constitutive record that supports PPN, redshift, Shapiro delay, lensing, and gravitational-wave speed.	If the calculable low-energy quantum correction requires an independent coefficient set, spacetime closure is incomplete even if the classical observables are matched.
2	Parameter non-closure	Too many symbols remain geometric promises rather than fixed quantities.	parameter-ledger.md now organizes them, but most are still open.	Close κ , the mass prefactor, the metric constitutive coefficients, and the weak-mixing datum without per-observable retuning.	If the same symbol has to be re-fit independently across chapters, the closure claim weakens sharply.
2	Null-result closure for added channels	A unification claim can fail even while matching known positive benchmarks if it predicts extra channels that experiments have not seen.	failure-criteria.md now defines $\mathcal{R}_{\text{null}}(\theta)$ for predicted non-baseline channels, but the main sector ledgers have not all routed their null-result bounds through that residual. The concrete comparison cases are mirror matter, superpartners, proton-instability channels, extra gauge bosons, hidden transport modes, sterile or neutral partner branches, and preferred-frame leakage channels.	For every added partner family, unstable baryon channel, extra gauge or transport mode, preferred-frame leakage channel, or other non-baseline output, compute $O_e(\theta)$ and show $O_e(\theta) \leq O_e^{\text{max}}$ from the same shared closure record used for the positive benchmarks. A symmetry container that includes the Standard Model as a subcase passes only when the added channels are proven absent, exactly redundant, or below bounds by the same branch record that recovers the observed sector.	If unobserved channels are hidden only by sector-specific masses, thresholds, compactification-like assumptions, or disconnected suppression factors, the framework has reproduced the failure pattern of overextended unification rather than closing it.
2	Thermodynamic-gravity closure	If the metric is an emergent equation of state, the repo needs more than constitutive rhetoric.	emergent-metric.md now states the Noether sea-first picture, defines a local-horizon residual $\mathcal{R}_{\text{thermo}}(\theta)$, and links the proof scaffold to Thermodynamic Residual ; black-holes.md frames horizon entropy as a block-density count over horizon-compatible reduced Noether braid closure labels. No run has yet driven the residual small from a simulated Noether sea record.	Show that the Noether sea admits an area-scaling entropy channel $S_H = k_B \log \mathcal{B}_H $ whose local coefficient is recovered as a block entropy density, a local Rindler/Unruh recovery in the appropriate limit, a Jacobson-style $dQ = T_H dS$ residual for boundary-wake data, Page-curve-compatible information release through horizon-interface channels, and a controlled nonequilibrium regime where distinctive departures are predicted.	If GR-like recovery requires thermodynamic language but the Noether sea cannot supply area scaling, local horizon temperature, a shared stress/entropy/temperature record, Page-curve-compatible information accounting, or a coherent nonequilibrium boundary, the present gravity interpretation loses depth and may be mislocated.
2	Reaction-cosmology provenance closure	The local-reaction story and the cosmology-source story now meet at photon loading, pair production, signed photon-frequency exchange, and thermalization.	reaction-cosmology-provenance-ledger.md defines the shared ledger, including path-frequency exchange, but no full source-to-background path has been closed.	Produce one conserved provenance path from a radiation, pair, or Compton/SZ-like transfer channel through thermalization to a BBN or CMB observable, using the same Noether sea state variables throughout.	If BBN photon loading, CMB blackbody recovery, or redshift-budget reconstruction requires unbalanced substrate creation, unlogged photon energy transfer, per-source retuning, or incompatible thermalization assumptions, the local-recycling cosmology branch fails.
2	Shared cosmology state closure	Dark-energy, H_0 , S_8 , CMB, BBN, BAO, weak-lensing, redshift-budget, and pre-BBN comparison claims all consume overlapping Noether sea state variables.	cosmology-ontology.md , dark-energy.md , and hubble-s8-tensions.md now state the shared-state requirement; inflation-model.md , BBN-constraints.md , structure-formation.md , and gravitational-waves.md now route pre-BBN branch projections through the same record; simulations/cosmology-shared-residual-fit.md supplies the first mock residual-packet scaffold; and dark-energy.md now gives a thermodynamic Λ_{eff} conjugacy target, but no empirical joint residual fit exists.	Produce one θ_{sea} and projection family that keeps SN, BAO, CMB, WL, RSD, BBN, H_0 , S_8 , signed path-frequency-transfer rows, pre-BBN branch projections, and stochastic-background bounds inside tolerance without per-pipeline retuning; if Λ_{eff} is treated thermodynamically, derive it as a conjugate to an effective observer-level four-volume functional of the same θ_{sea} .	If distance, growth, early-universe, calibration, path-frequency-transfer, pre-BBN branch, stochastic-background, or thermodynamic- Λ_{eff} observables require incompatible Noether sea state records, the cosmology branch has hidden the tension rather than closed it.

Tier	Issue	Why it matters	Current repo status	Closure target	Failure condition
2	Radiation Gate C benchmark closure	Radiation must recover standard electromagnetic and QED-like benchmarks before Noether sea-dependent deviations or cosmology source claims are credible.	radiation.md now carries a classified closure-target ledger, with channel scaffolds in bremsstrahlung.md , synchrotron.md , and reaction-cosmology-provenance-ledger.md , but no unified Gate C derivation is complete.	Close Larmor/Lienard recovery, free-free emissivity, synchrotron $\gamma^2 B$ and power scaling, pair thresholds, Compton-like scattering, and blackbody detailed balance through one event record, while treating free photon polarization as a Gate B handoff only.	If any benchmark requires per-observable retuning, violates validated limits, or derives free photon polarization outside Gate B, radiation Gate C does not close.
2	CKM / PMNS quantitative closure	Flavor mixing cannot remain only qualitative if the framework claims Standard-Model replacement.	PMNS oscillation formulas exist; CKM geometry has an overlap/holonomy scaffold and is now tied to the same weak-coupling-triad exposure route as V–A and reaction provenance.	Derive one geometric overlap map for quark and lepton mixing from the exposed weak-coupling-triad domain, shielding eigenstates, and near-photon neutral-sector Hamiltonian, then test it against CKM and PMNS data.	If no stable geometry reproduces the observed hierarchy and phases, or if the CKM/PMNS definitions require a different weak-basis domain from the V–A operator, the present mixing architecture is incomplete at best.
2	Quark mass map	The quark catalog is in place, but the mass hierarchy is still not quantitative.	quarks.md closes structure, not masses.	Produce a first-pass mass map for u, d, c, s, t, b from shielding and internal-energy accounting.	If the hierarchy cannot be reproduced even at scaling level, generation-by-shielding is in trouble.
2	Spin / statistics closure	The framework repeatedly appeals to spinor and bosonic/fermionic behavior.	There is now a decent 4π story, but not a formal closure proof.	Derive the ordered-frame history-lift map cleanly enough to justify spin- $\frac{1}{2}$ and associated statistics sectors.	If the topology cannot distinguish fermionic and bosonic closure classes, several assembly claims lose their footing.
2	Baryon stability and baryon-number status	Proton stability is a major empirical constraint and a major theoretical claim.	The color chapter gives a topological argument, but the quantitative baryon-number status remains open. Current PDG 2024 comparison bounds already put representative partial mean lives at $\tau/B(p \rightarrow e^+ \pi^0) > 2.4 \times 10^{34}$ yr and proton neutrino/kaon modes near 5.9×10^{33} yr at 90% confidence, so this is an active null-result gate rather than a qualitative concern.	Show whether proton stability is exact, exponentially protected, or only effective in a quantified regime, and route every predicted baryon-violating corridor through $\Gamma_{p\bar{c}}^{\max} = 1/\tau_c^{\min}$ in Failure Criteria .	If the theory predicts generic fast proton decay, or suppresses it by a sector-local parameter not tied to the same color/topology and reaction-provenance ledger that recovers hadron structure, the hadronic sector is not viable.
3	Nuclear binding closure	The residual strong-force story must eventually recover nuclear phenomenology beyond pions-as-metaphor.	nuclear-binding.md now gives a first effective interface, but no fitted nuclear map.	Recover at least deuteron binding, saturation, and alpha-like enhancement in one coherent effective model.	If even the sign and scaling of nuclear binding cannot be stabilized, the hadronic coarse-graining is inadequate.
3	Condensed-matter branch recovery	Materials provide dense, precise tests of whether electron-envelope, lattice, phonon, and Noether sea response variables remain one record instead of becoming probe-specific fits.	condensed-matter.md now states Bloch-band, effective-mass, Fermi-surface, diffraction, phonon, Hall, and topological-response residuals; constraint-ledger.md records the corresponding response gate. No derivation yet computes these objects from the master equation or a settled material branch.	Recover Bloch form, reciprocal-lattice scattering, phonon dynamical matrices, effective mass tensors, Fermi-surface or band-gap classification, Hall sign/plateaux, and no-drag transport from one declared material branch and Noether sea state record.	If band curvature, lattice stiffness, diffraction peaks, Hall response, and transport relaxation require independent response maps, or if resistance is explained by ordinary Noether sea drag below the transport threshold, the material-response program has split from the main ontology.
3	Strong-field / black-hole closure	Strong-field claims are distinctive and therefore risky.	The alignment framing exists, but the predictive map is not yet broad.	Derive concrete departures near the alignment regime while preserving weak-field success.	If the strong-field story contradicts weak-field closure or observed compact-object data, it must be revised.

Highest-Leverage Cluster

The top credibility cluster is:

1. weak V–A,
2. preferred-frame hiding,
3. Born-rule emergence,
4. weak-field GR recovery.

Those four form the present hard gate because each one touches a major validated pillar of modern physics:

- electroweak structure,
- Lorentz hiding,
- quantum statistics,
- and relativistic gravity phenomenology.

A cross-cutting null-result discipline now sits over that cluster. A proposed closure may not buy unification by adding hidden sectors whose only role is to disappear below proton-stability, collider, precision-symmetry, preferred-frame, or cosmology bounds. The residual $\mathcal{R}_{\text{null}}(\theta)$ in [Failure Criteria](#) must vanish using the same shared record that passes the positive recovery gates.

The strict implementation is to treat each elegant symmetry package as a joint positive-and-absence record: the same θ that recovers the Standard-Model-facing rows must also set each extra-channel observable $O_e(\theta)$ to zero or below its validated bound.

If those remain open, the framework can still be a promising substrate program, but not yet a closed replacement architecture.

Interdependence Map

Several tensions are linked and should not be treated as isolated tasks.

Weak sector cluster

The weak-selection problem, right-handed neutrino stance, CKM/PMNS closure, weak-corridor provenance, and the quark misalignment parameter α all belong to the same electroweak geometry stack. The current synthesis is that these are readouts of one weak-coupling-triad exposure problem: axial-frame branch selection determines what can be exposed, the V–A operator determines which handedness can dock, the overlap integrals determine mixing weights, and the reaction ledger determines where the corridor payload and outgoing Noether braid provenance enter and exit. A clean derivation of one should now constrain the others rather than leaving them as independent stories.

The neutrino branch of this cluster has four empirical decision handles: the lightest-neutrino mass, the mass sum $\sum_i m_i$, neutrinoless double-beta limits or detection, and any evidence for a sterile or right-handed singlet. These data products should decide between the current minimal near-photon neutral-pair stance, a sterile ν_R branch, or a lepton-number-violating provenance channel. They should not be used to rewrite the charged-fermion axial-layer rule or to import a sterile dark-matter interpretation before the PMNS, reaction, BBN, CMB, and structure-formation gates are simultaneously satisfied.

A useful benchmark-only sharpening is the package $m_{\text{lightest}} \rightarrow 0$, $\sum_i m_i \approx 0.06 \text{ eV}$, suppressed neutrinoless double-beta rate, and any sterile or right-handed singlet behaving as cold collisionless matter only after the neutral-sector and cosmology gates close. These values should be treated as discriminator targets, not as adopted ontology: they can rank the neutral-lepton branches, but they cannot bypass the PMNS Hamiltonian, reaction provenance, BBN, CMB, structure-formation, and null-result residuals.

Quantum cluster

Superposition, measurement, Born-rule emergence, and Bell/nonlocality closure are one package. A good ontology chapter without a basin-measure derivation is progress, but not endpoint closure.

Penrose-Diosi gravitational-collapse tests are an external benchmark for the same finite-time threshold-resolution burden, not an adopted ontology. The comparison uses the gravitational self-energy of the difference between two mass distributions,

$$\Delta E_G \sim \frac{G}{2} \iint \frac{(\rho_1 - \rho_2)(\mathbf{x})(\rho_1 - \rho_2)(\mathbf{y})}{\|\mathbf{x} - \mathbf{y}\|} d^3x d^3y$$

with the collapse-time estimate

$$\tau_G \sim \frac{\hbar}{\Delta E_G}$$

The useful comparison pressure is the tension between local free-fall equivalence and linear superposition when the two branches carry measurably different mass distributions. The validation burden is to compare τ_{meas} for massive-superposition records, including BEC-scale spatial superpositions of roughly 10^9 to 10^{10} atoms, against τ_G and against ordinary environmental decoherence, while preserving the AAA claim that branch selection is finite-time threshold resolution rather than fundamental gravitational collapse. Any collapse variant that predicts persistent spontaneous heating must also pass low-background and compact-object heating bounds before it can serve even as a comparison baseline.

Spacetime cluster

Preferred-frame hiding, redshift, Shapiro delay, lensing, gravitational-wave speed, and the long-distance quantum correction to Newtonian gravity are all readouts of the same observer-level constitutive map. Thermodynamic-gravity closure belongs in the same cluster because area scaling, local horizon temperature, and nonequilibrium breakdown define whether the constitutive picture is merely suggestive or genuinely explanatory. Low-energy quantum-gravity EFT is kept here as a recovery benchmark, not as a commitment that the effective metric is microscopic ontology. These issues rise or fall together.

Reaction-cosmology cluster

Radiative planar-mode nucleation, pair-production provenance, BBN photon loading, CMB blackbody recovery, signed path-frequency exchange, and redshift handoff form one closure cluster when cosmology is read through SMBH-local recycling and Noether sea transport. A local source story is not enough; the same provenance record must carry architrino inventory, energy-momentum, thermalization depth, photon-frequency transfer, and observer-level comparison variables without changing the Noether sea state map between channels.

Pre-BBN comparison branches belong to this same cluster. They can add value only as stress tests on the shared record: light-element yields, N_{eff} , CMB acoustic and lensing products, matter power, and stochastic gravitational-wave bounds must all be projections of the same Noether sea history. If the branch is kept alive by independent hiding assumptions, it is a null-result failure rather than a productive extension.

Radiation benchmark checks

Radiation Gate C closure is validated only if the same event record passes the following classified checks. These checks are not alternate ontologies; they are benchmark recoveries that prevent source-channel language from outrunning the photon and reaction ledgers.

Check	Class	Required validation	Failure signal
Radiative event record	ontology	Record routed closure residuals, planar-mode photon output when present, non-photon shedding channels, recoil, local Noether sea state, and conservation ledgers.	Radiation is treated as primitive field emission or as untracked energy loss.
Larmor/Lienard recovery	derivation target	Recover $P \propto \ \mathbf{a}\ ^2$ in the weak nonrelativistic limit and the Larmor/Lienard observer-level power/angular behavior after clock conversion.	Low-speed power is not quadratic in acceleration, or relativistic recovery needs a separate fit.
Bremsstrahlung emissivity	derivation target	Recover $d\sigma/dk$, screening/form-factor corrections, $\epsilon_{\nu}^{\text{eff}} \propto Z^2 n_e n_i T^{-1/2} e^{-h\nu/(k_B T)} g_{\text{ff}}$, and $\epsilon_{\text{ff}} \propto Z^2 n_e n_i T^{1/2}$ in LTE.	Cross-section and emissivity require incompatible Noether sea variables or plasma-specific hidden fits.
Synchrotron $\gamma^2 B$ scaling	derivation target	Recover $\nu_c \propto \gamma^2 B$, $P_{\text{syn}} \propto U_B \gamma^2$, and cooling breaks from one effective magnetic-state map.	The γ^2 frequency scaling is absent, or the B map changes between curvature and emission.
Pair thresholds	derivation target	Recover $s \geq 4m_e^2 c^4$ and the angle-dependent photon-photon threshold while preserving architrino inventory and pair provenance.	Pair channels imply creation from nothing, wrong thresholds, or unbalanced Noether braid recruitment.
Compton-like scattering	derivation target	Recover the Compton shift, Thomson limit, Klein-Nishina correction, recoil, and outgoing photon provenance in one Gate C vertex.	The channel becomes phenomenological frequency loss without a closed recoil and photon ledger.
Aharonov-Bohm phase	derivation target	Recover a relative phase proportional to enclosed magnetic flux while the local force channel on the interferometer arms vanishes, using the same effective U(1) connection and photon/action ledger as the rest of Gate C.	The phase requires a local force on the arms, an independent phase fit, or a literal gauge-potential ontology rather than a derived effective connection.
Blackbody recovery	derivation target	Recover Planck occupation, zero effective photon chemical potential, thermalization depth, damping, anisotropy, polarization handoff, and redshift handoff without retuning the Noether sea map.	The CMB or thermal branch needs unbalanced photon loading, per-observable retuning, or incompatible transport assumptions.
Free photon polarization boundary	derivation target	Use Gate B records for transverse modes, helicity, Malus' law, and analyzer statistics; radiation and cosmology pages may only consume that handoff.	Any radiation channel derives free photon polarization locally, adds a free longitudinal mode, or treats Gate B as already proven.
Noether sea-dependent deviations	speculation	State a benchmark-preserving limit and a measurable residual before promoting any $p_{\text{NS}}(\mathbf{x}, t)$, $\chi_{\text{sea}}(\mathbf{x}, t)$, anisotropy, or threshold-floor effect.	A deviation is used to repair a failed standard recovery or is fitted separately per observable.

Ontology Watchlist

The foundational ontology hub now keeps only stable commitments. Open questions that used to live there are tracked here or in the relevant branch chapters:

- **Deterministic branch selection:** close the rule for active causal roots, weighted sums, phase-sensitive thresholding, and basin selection in [Master Equation](#) and [Binary Dynamics](#). The working hypothesis remains deterministic multistability, with apparent randomness coming from chaotic sensitivity to microstate and wake history.
- **Polarity unit and coupling scale:** explain $\epsilon = |e|/6$ and close κ through [Parameter Ledger](#), [Architrino SI Base Units](#), and the charge-mapping chapters. The unresolved question is whether six-site nested shell braid organization derives the ϵ unit, and whether κ is related to ϵ , c_f , $\hbar$, or Planck-alignment quantities rather than being independently postulated.
- **Quantum ontology:** keep wavefunction status, decoherence, and Born-rule recovery in [Wavefunction Ontology](#), [Measurement Ontology](#), and the Born-rule tension above. Decoherence still needs a stance on whether its irreversibility is fundamental in the Noether sea environment or practical because reversal is dynamically inaccessible to Physical Observers.
- **Symmetry and conservation:** close CPT stance, baryon-number status, and proton-stability regime through the particle and interaction chapters. The unresolved CPT issue is treated as a replacement constraint in [No-Go Theorems](#): the standard proof assumes local relativistic QFT, while this framework uses absolute time and delayed substrate dynamics, so the corpus must preserve tested CPT-facing observables without importing those assumptions as ontology.
- **Cosmological history:** keep beginning/eternity and initial-condition questions in [Cosmology Ontology](#), [Expansion Mechanism](#), and related cosmology modules. If the background is eternal, the theory still owes a large-scale homogeneity and isotropy account, including a scale-neutral residual comparing dimensionless pair-separation distributions across large windows; if it has an initialization boundary, it owes an architrino-distribution account.
- **Unification claim:** treat “all forces from nested shell braid geometry and Noether sea dynamics” as a closure program, not as a primitive ontology statement. The qualitative structure exists across interaction chapters, but quantitative derivations remain the acceptance gate.

Acceptance Principle

The framework should be judged by the intersection of its surviving closure sets:

$$\mathcal{C}_{\text{weak}} \cap \mathcal{C}_{\text{quantum}} \cap \mathcal{C}_{\text{gravity}} \cap \mathcal{C}_{\text{hadronic}} \cap \mathcal{C}_{\text{radiation}} \cap \mathcal{C}_{\text{cosmology}} \neq \emptyset$$

If that intersection becomes empty after quantitative work is done, the present implementation is rejected even if many individual chapters remain suggestive. The detailed sector predicates, benchmark tolerances, and promotion-fiber test are recorded in [Failure Criteria](#).

Related Chapters

- [constraint-ledger.md](#)
- [closure-scorecard.md](#)
- [.../assemblies/fermions/quantum-number-mapping.md](#)
- [.../spacetime/general-relativity.md](#)
- [.../quantum/measurement-ontology.md](#)

Massive Superposition Gravity

This packet turns massive-superposition gravity experiments into concrete validation targets. It belongs to the observable and inference layer: the task is to preserve the branch mass histories, coherence data, detector response, entanglement data, and record criteria without importing any external collapse ontology or quantum-metric ontology.

Related homes are [Measurement Ontology](#), [Observer Framework](#), and [Constraint Ledger](#).

Comparison Boundary

The packet may use external classical-quantum gravity proposals as comparison pressure, but only at the level of observables and inference. The comparison rows are:

External comparison	Retained pressure	AAA use	Not imported
Oppenheim-style classical-quantum gravity	A classical or effective gravity readout must not reveal branch information while the quantum branch description still shows interference.	Bound $\mathcal{D}_{\text{grav}}$, constrain N_{AB} , and require a Physical Observer record before treating gravity-side branch information as a measurement.	Stochastic-metric ontology, fundamental collapse, external terminology, or the claim that gravity must remain classical at the substrate level.
Gravitationally induced entanglement	Two isolated massive probes can acquire branch-dependent correlations through gravity alone.	Require the same effective-metric record θ to generate the branch interaction phase and to keep which-path leakage below the retained weak-probe threshold.	Constructor-theory doctrine, Q-number terminology, fundamental graviton ontology, or the claim that spacetime geometry itself has been prepared in superposition.

Every averaged quantity in this packet is a run-record summary. A covariance matrix, branch expectation value, or correlation function may be used only after the Physical Observer access region, detector channel, boundary-data model, and persistence criterion have been declared. It may not be promoted into a primitive gravity state or collapse mechanism merely because it appears in a successful inference pipeline.

Experiment-Family Classification

Different laboratory proposals enter this packet at different levels. The classification below keeps the observable pressure while preventing passive phase tests, active branch-mass tests, and mediated-entanglement tests from being treated as one result.

Experiment family	Retained observable	Packet status	Interpretation guardrail
guided/free-fall atom-interferometer phase tests	fitted cubic-time phase coefficient $\hat{\beta}_{T^3}$, fringe visibility, and control-phase record	passive external-field phase benchmark	Confirms or constrains the weak-field phase map; does not by itself test active self-gravity or fundamental collapse.
BEC, solid, nanoparticle, nanodiamond, membrane, or cantilever massive-superposition tests	branch mass histories ρ_1, ρ_2 , visibility $\mathcal{V}(T)$, τ_{mess} , ΔE_G , and $\mathcal{D}_{\text{grav}}$	active branch-mass-history benchmark	Tests whether finite-time threshold resolution, ordinary decoherence, and Penrose-Diosi-like collapse scales remain quantitatively distinguishable.
two-probe gravitationally induced entanglement tests	cross-branch phase $\Delta\Phi_{\text{ent}}$, entanglement witness $\mathcal{C}_{\text{obs}}$, and non-gravitational residual $\mathcal{R}_{\text{dongrav}}$	mediated-entanglement benchmark	Tests the shared gravity-side constitutive record without importing fundamental graviton ontology or a quantum-metric substrate.

The packet should classify a run by the strongest observable it actually carries. A passive phase benchmark may constrain θ for later active-mass tests, but it cannot be used as evidence that gravity has or has not selected a branch. Conversely, an active branch-mass run that loses visibility must still show a record-forming separatrix crossing before the loss is interpreted as measurement rather than uncontrolled environmental decoherence.

Observable Target

The target experiment compares two branch-level mass-density histories over a coherence window T :

$$\rho_1(\mathbf{x}, t), \quad \rho_2(\mathbf{x}, t)$$

The branch pair is interference-preserving only if the apparatus and environment have not produced an autonomous which-path record. The gravitational or effective-metric channel therefore becomes a constraint through the response difference

$$\Delta h_A(t) = h_A(t; \rho_1, \theta) - h_A(t; \rho_2, \theta)$$

where A labels the resolved detector response channel and θ is the shared effective-metric constitutive record.

The which-path diagnostic is

$$\mathcal{D}_{\text{grav}}(T; \theta) = \int_0^T \int_0^T \Delta h_A(t) N_{AB}^{-1}(t, t'; \theta) \Delta h_B(t') dt dt'$$

Here N_{AB} is the observer-level covariance decomposed in [Observer Framework](#). It summarizes unresolved deterministic boundary histories and calibrated detector/environment residuals; it is not an ontological randomness postulate.

Minimal Response Model

A concrete first packet can use a displaced normalized mass packet. Let φ_σ be normalized by

$$\int_{\Sigma_t} \varphi_\sigma(\mathbf{x}) d^3x = 1$$

For branch separation $\mathbf{d}(t)$ around center $\mathbf{x}_0(t)$, set

$$\begin{aligned} \rho_1(\mathbf{x}, t) &= m \varphi_\sigma\left(\mathbf{x} - \mathbf{x}_0(t) - \frac{\mathbf{d}(t)}{2}\right), \\ \rho_2(\mathbf{x}, t) &= m \varphi_\sigma\left(\mathbf{x} - \mathbf{x}_0(t) + \frac{\mathbf{d}(t)}{2}\right). \end{aligned}$$

Let $G_A(t, s; \mathbf{x}; \theta)$ be the detector response kernel implied by the same effective-metric constitutive record used for redshift, Shapiro delay, lensing, gravitational-wave speed, and, when the record is extrapolated to compact sources, horizon-scale ring/shadow imaging. The branch response is

$$h_A(t; \rho_k, \theta) = \int_0^t \int_{\Sigma_s} G_A(t, s; \mathbf{x}; \theta) \rho_k(\mathbf{x}, s) d^3x ds$$

Therefore

$$\Delta h_A(t) = \int_0^t \int_{\Sigma_s} G_A(t, s; \mathbf{x}; \theta) [\rho_1(\mathbf{x}, s) - \rho_2(\mathbf{x}, s)] d^3x ds$$

When $\|\mathbf{d}(t)\|$ is small relative to the packet scale,

$$\rho_1(\mathbf{x}, t) - \rho_2(\mathbf{x}, t) = -m d^i(t) \partial_i \varphi_\sigma(\mathbf{x} - \mathbf{x}_0(t)) + O(\|\mathbf{d}(t)\|^3)$$

so the leading branch response is

$$\Delta h_A(t) \approx -m \int_0^t d^i(s) \int_{\Sigma_s} G_A(t, s; \mathbf{x}; \theta) \partial_i \varphi_\sigma(\mathbf{x} - \mathbf{x}_0(s)) d^3x ds$$

This gives the first closure equation: a mass displacement history should map to a predicted detector-channel separation before any interpretive claim about classical or quantum spacetime is introduced.

Mediated Entanglement Comparison

A complementary massive-superposition test asks whether two independently prepared massive probes can become entangled through the gravity-side channel while non-gravitational couplings are suppressed or bounded. This is a positive branch-phase benchmark, not a new ontology. The observable is the final two-probe correlation record, together with the calibration record showing that electromagnetic, spin-spin, thermal, and apparatus cross-talk channels are too small to account for the effect.

Let the two probes be A and B , with branch labels $a, b \in \{+, -\}$ and branch mass histories $\rho_A^a(\mathbf{x}, t)$ and $\rho_B^b(\mathbf{x}, t)$. The same weak-field constitutive record θ used for redshift, Shapiro delay, lensing, PPN, gravitational-wave speed, compact-source ring/shadow extrapolations, and $\mathcal{D}_{\text{grav}}$ must determine the branch interaction energy

$$U_{ab}^{\text{eff}}(t; \theta) = -G_{\text{eff}}(\theta) \int_{\Sigma_a} \int_{\Sigma_t} \frac{\rho_A^a(\mathbf{x}, t) \rho_B^b(\mathbf{y}, t)}{\|\mathbf{x} - \mathbf{y}\|} d^3x d^3y + O(c_0^{-2})$$

The branch phase is then

$$\Phi_{ab}(T; \theta) = \frac{1}{\hbar} \int_0^T U_{ab}^{\text{eff}}(t; \theta) dt$$

Local branch phases can be absorbed into the one-probe descriptions. The entangling invariant is the cross-branch phase combination

$$\Delta\Phi_{\text{ent}}(T; \theta) = \Phi_{++}(T; \theta) + \Phi_{--}(T; \theta) - \Phi_{+-}(T; \theta) - \Phi_{-+}(T; \theta)$$

For the ideal equal-amplitude two-branch packet, a first witness target is

$$C_{\text{GIE}}(T; \theta) = \left| \sin \frac{\Delta\Phi_{\text{ent}}(T; \theta)}{2} \right|$$

This formula is an observer-level benchmark. It does not say that the Euclidean void is quantized, that the effective metric is fundamental, or that a graviton field is the native substrate. It says that the same gravity-side constitutive record must produce the branch phase that standard low-energy descriptions would attribute to gravitational mediation.

The comparison is meaningful only when the non-gravitational residual is bounded. Let $\mathcal{R}_{\text{nongrav}}$ collect calibrated electromagnetic, spin-spin, Casimir, thermal, vibration, and apparatus cross-talk contributions to the same entanglement witness. A run can be used as a gravity-side validation target only if

$$\mathcal{R}_{\text{nongrav}} \leq \varepsilon_{\text{iso}}$$

with ε_{iso} declared by the apparatus class and retained alongside the covariance record N_{AB} .

Input Record Schema

The packet is evaluated on an explicit run record:

Field	Symbol	Required content
branch mass histories	ρ_1, ρ_2	normalized mass-density histories on Σ_t over $0 \leq t \leq T$
branch separation	$\mathbf{d}(t)$	center or multipole separation history with declared packet width σ
apparatus/environment record	$\mathcal{A}_{\text{rec}}$	record variable, persistence window, environmental coupling channels, and ordinary decoherence estimate
gravity response kernel	$G_A(t, s; \mathbf{x}; \theta)$	detector response derived from the same effective-metric constitutive record used in weak-field gravity
mediated-entanglement phase	$\Delta\Phi_{\text{ent}}$	cross-branch phase predicted from ρ_A^a, ρ_B^b and the shared constitutive record θ
non-gravitational residual	$\mathcal{R}_{\text{nongrav}}$	calibrated bound on non-gravity channels that could create the observed correlation
covariance decomposition	N_{AB}	detector noise, unresolved boundary-wake terms, environmental residuals, and calibration residuals
visibility data	$\mathcal{V}(T)$	observed or predicted interference visibility over the run
entanglement data	C_{obs}	measured or predicted two-probe entanglement witness in the retained readout basis
record criteria	$R, \Sigma, T_{\text{rec}}$	Physical Observer record variable, separatrix, and persistence threshold

No row may be filled by changing the weak-field metric record after the positive gravity benchmarks have already been fit. The same θ must be replayable through redshift, Shapiro delay, lensing, PPN, gravitational-wave speed, compact-source ring/shadow extrapolations, and this massive-superposition packet.

Evaluation Protocol

- Normalize the branch histories.** Verify $\int_{\Sigma_t} \rho_k(\mathbf{x}, t) d^3x = m$ for each branch and each resolved time slice, or record the known mass exchange with the apparatus ledger.
- Compute the response difference.** Use one kernel $G_A(t, s; \mathbf{x}; \theta)$ to compute $h_A(t; \rho_1, \theta)$, $h_A(t; \rho_2, \theta)$, and $\Delta h_A(t)$.
- Assemble the covariance.** Build $N_{AB} = N_{AB}^{\text{det}} + N_{AB}^{\text{env}} + N_{AB}^{\text{wake}} + N_{AB}^{\text{cal}}$, with each term either derived from the apparatus model or bounded by calibration data.
- Evaluate distinguishability.** Compute $\mathcal{D}_{\text{grav}}(T; \theta)$ and compare it with ε_{wp} .
- Evaluate record formation.** Compute τ_{meas} , Δ_{rec} , and the persistence window from the measurement chapter's record criteria.
- Evaluate mediated entanglement when present.** If the run is a two-probe mediated-entanglement experiment, compute $\Delta\Phi_{\text{ent}}$, C_{GIE} , and $\mathcal{R}_{\text{nongrav}}$ from the same run record.
- Classify the run.** Use the same output record to assign one of four statuses:

Status	Conditions	Interpretation
weak-probe	$\mathcal{D}_{\text{grav}} \leq \varepsilon_{\text{wp}}$ and no durable record forms	gravitational response is too weak to act as a which-path record
mediated-entangling	$C_{\text{GIE}} \geq C_{\text{obs}} - \varepsilon_C$, $\mathcal{R}_{\text{nongrav}} \leq \varepsilon_{\text{iso}}$, $\mathcal{D}_{\text{grav}} \leq \varepsilon_{\text{wp}}$, and no durable which-path record forms	the branch phase is strong enough to account for the entanglement witness while the gravity-side readout remains below record threshold

Status	Conditions	Interpretation
record-forming	$\mathcal{D}_{\text{grav}} > \varepsilon_{\text{wp}}$, $\tau_{\text{meas}} < T$, and Δ_{rec} stays below threshold through T_{rec}	the apparatus/environment has formed an autonomous record
falsifying	$\mathcal{D}_{\text{grav}} \gg 1$ while visibility remains high and no record-autonomy criterion is met	the effective-metric response overproduces observable which-path information

For a white-noise readout approximation, $N_{AB}(t, t') = S_{AB}\delta(t - t')$, the distinguishability reduces to

$$\mathcal{D}_{\text{grav}}(T; \theta) = \int_0^T \Delta h_A(t) S_{AB}^{-1} \Delta h_B(t) dt$$

This special case is the first numerical target because it turns the validation packet into a finite time-series calculation once m , σ , $\mathbf{d}(t)$, G_A , and S_{AB} are supplied.

Worked Acceleration Bound

A first sanity bound can use a single acceleration readout channel before introducing a full detector geometry. Suppose the branch displacement is bounded by $\|\mathbf{d}(t)\| \leq d_0$, the detector is at distance R from the branch center with $d_0 \ll R$, and the weak-field map satisfies $G_{\text{eff}}(\theta) \rightarrow G$ in the tested regime. The branch acceleration difference is bounded by

$$|\Delta h(t)| \leq \frac{2G_{\text{eff}}(\theta)M d_0}{R^3}$$

For a white acceleration readout covariance $N(t, t') = S_a\delta(t - t')$, the distinguishability obeys

$$\mathcal{D}_{\text{grav}}(T; \theta) \leq \frac{4G_{\text{eff}}^2(\theta)M^2 d_0^2 T}{R^6 S_a}$$

With benchmark values

$$M = 10^{-14} \text{ kg}, \quad d_0 = 10^{-6} \text{ m}, \quad R = 10^{-3} \text{ m}, \quad T = 1 \text{ s}$$

and an aggressive acceleration-noise amplitude

$$S_a^{1/2} = 10^{-15} \text{ m s}^{-2} / \sqrt{\text{Hz}}$$

the bound is

$$\mathcal{D}_{\text{grav}} \lesssim 1.8 \times 10^{-12} \left(\frac{M}{10^{-14} \text{ kg}} \right)^2 \left(\frac{d_0}{10^{-6} \text{ m}} \right)^2 \left(\frac{10^{-3} \text{ m}}{R} \right)^6 \left(\frac{T}{1 \text{ s}} \right) \left(\frac{10^{-15} \text{ m s}^{-2} / \sqrt{\text{Hz}}}{S_a^{1/2}} \right)^2$$

For a which-path threshold of order unity, this run is deep in the weak-probe class. Solving the same bound for the mass needed to reach $\mathcal{D}_{\text{grav}} \sim \varepsilon_{\text{wp}}$ gives

$$M_{\text{crit}} \approx \frac{R^3}{2G_{\text{eff}}(\theta)d_0} \sqrt{\frac{\varepsilon_{\text{wp}} S_a}{T}}$$

or, in the same benchmark geometry,

$$M_{\text{crit}} \approx 7.5 \times 10^{-9} \text{ kg } \varepsilon_{\text{wp}}^{1/2} \left(\frac{R}{10^{-3} \text{ m}} \right)^3 \left(\frac{10^{-6} \text{ m}}{d_0} \right) \left(\frac{S_a^{1/2}}{10^{-15} \text{ m s}^{-2} / \sqrt{\text{Hz}}} \right) \left(\frac{1 \text{ s}}{T} \right)^{1/2}$$

This is not a new ontology or an experimental forecast. It is a scale check: for ordinary mesoscopic masses, gravity-side which-path leakage is negligible unless the branch mass, separation, proximity, coherence time, or readout sensitivity moves by many orders of magnitude. A full detector calculation should replace the scalar factor $2/R^3$ with the tensor response in the Minimal Response Model above.

Acceptance Criteria

For an interference-preserving run, the metric or gravity-side readout must satisfy

$$\mathcal{D}_{\text{grav}}(T; \theta) \leq \varepsilon_{\text{wp}}$$

For a mediated-entanglement run, the same record must also satisfy

$$C_{\text{GIE}}(T; \theta) \geq C_{\text{obs}} - \varepsilon_C, \quad \mathcal{R}_{\text{nongrav}} \leq \varepsilon_{\text{iso}}$$

This combined gate preserves the observable without overclaiming the interpretation: the run tests whether the retained gravity-side constitutive record can generate the observed branch correlation while avoiding premature which-path record formation.

If a which-path record is claimed instead, the measurement chapter's record criteria must also hold:

$$\tau_{\text{meas}} < T, \quad \sup_{t \in [\tau_{\text{meas}}, \tau_{\text{meas}} + T_{\text{rec}}]} \Delta_{\text{rec}}(t; k) \leq \varepsilon_{\text{rec}}$$

The failure condition is strict. If $\mathcal{D}_{\text{grav}} \gg 1$ while interference visibility remains high and no record-autonomy condition is satisfied, the effective-metric response has overproduced observable which-path information.

The same θ must also remain compatible with the gravity-side ledger: redshift, Shapiro delay, lensing, PPN parameters, gravitational-wave speed, dispersion, detector-mode bounds, and compact-source ring/shadow extrapolations. A parameter set that fits the massive-superposition channel only by changing the weak-field metric record is not a valid closure.

Simulation Target

The minimal simulation target is the map

$$\mathcal{S}_{\text{grav}} : (m, \sigma, \mathbf{d}(t), T, G_A, N_{AB}, R, \Sigma, \rho_A^a, \rho_B^b) \mapsto (\mathcal{D}_{\text{grav}}, \mathcal{V}(T), \Delta\Phi_{\text{ent}}, C_{\text{GIE}}, \tau_{\text{meas}}, \Delta_{\text{rec}})$$

The inputs are the branch mass scale, packet width, separation history, coherence window, detector response kernel, covariance decomposition, record variable, separatrix, and two-probe branch histories when present. The outputs are the gravitational distinguishability, interference visibility, entangling phase, mediated-entanglement witness, finite measurement time, and record-autonomy residual.

The worked acceleration bound supplies the first analytic $\mathcal{D}_{\text{grav}}$ estimate. The mediated-entanglement comparison supplies the first branch-phase target. Full packet closure still requires one numerical or analytic instance that computes the retained outputs from a shared constitutive record and reports whether the branch pair is weak-probe, mediated-entangling, record-forming, or falsifying.

Validation Simulations

Architrino

This note records the minimum tier-1 simulation tests that should be passed before any strong self-hit or non-Markovian claims are trusted numerically. Its purpose is narrow: establish provenance-resolved propagation, baseline diagnostics, and a workable history-buffer strategy before moving to richer dynamics.

The file is therefore an implementation-facing checklist rather than a general theory chapter. It should be read as a gate on simulation credibility.

Tier-1 Mandatory Unit Tests (Before Self-Hit Claims)

Provenance-resolved propagation test

Implement 1-architrino and 2-architrino setups with $\mathbb{U}_{\text{now}}$ sensors arranged on causal rings:

- Verify causal isochron propagation at c_f
- Verify correct arrival ordering and phase behavior (per kernel)
- Verify numerical stability of t_{emit} inversion as $\Delta t \rightarrow \Delta t/2$
- Produce provenance tables showing correct emitter_id values and emission times

Baseline diagnostics

- Energy/momentum bookkeeping (as defined by the model) must be stable under refinement
- Cross-integrator comparison required for the above propagation test

Grid-Based History Strategy

1. **Problem:** Infinite memory cost for particle-based history in self-hit regimes.
2. **Solution:** Use the $\mathbb{U}_{\text{now}}$ Grid as the history buffer. Store potential magnitude/gradients at grid nodes.
3. **Algorithm:** When an architrino requires its self-potential from $t - \Delta t$, query the **grid node** closest to where the particle was, rather than indexing the particle list.
4. **Deliverable:** Prove convergence of this grid-based history against analytic causal isochrons.

Grid-Based History

- **Memory Strategy:** Use the fixed grid to store potential history.
- **Lookup:** Query grid nodes for history potential values (Order(1) lookup) rather than querying particle history (Order(N)).
- **Validation:** Verify causal isochron propagation and phase ordering on the grid.

Convergence Tests

This chapter defines the convergence standard for simulations that include self-hit structure and other delayed-memory effects. Its role is to specify which observables are checked, which refinement ladders are required, and what pass/fail thresholds count as numerical control rather than artifact.

Because self-hit dynamics are especially prone to fake structure under poor time or history resolution, this document should be read as a validation gate rather than as optional numerical hygiene.

All convergence claims in this chapter are finite-window claims. Passing the gates below validates the declared observables on the analysis window, with the stated detector set, history horizon, and regulator choices. It does not decide unbounded reachability questions for the full delayed dynamics; those would require a separate theorem about the global flow rather than a stronger convergence plot.

Convergence in Non-Markovian (Self-Hit) Dynamics

Scope and default observable set

For each claim, compute convergence on a fixed analysis window $W = [t_a, t_b]$ and detector set $\{x_k\}$ using:

- $\Phi(x_k, t)$
- $\|\nabla \Phi(x_k, t)\|$
- self-hit event rate $\lambda_{\text{self}}(x_k)$
- key invariant drift (e.g., normalized energy drift) ϵ_E

Comparison metrics (required)

For any observable Y on two runs A (coarser) and B (finer), define

$$E_{\text{rel}}(Y; A, B) \equiv \frac{\|R(Y_B) - Y_A\|_{L^2(W, \{x_k\})}}{\|R(Y_B)\|_{L^2(W, \{x_k\})} + \varepsilon_0}, \quad \varepsilon_0 = 10^{-12}$$

Here R is restriction of the finer run to the coarser sampling grid.

For provenance distributions of solved t_{emit} , define:

$$D_W \equiv \frac{W_1(P_A, P_B)}{\text{IQR}(P_B) + \varepsilon_0}, \quad D_{JS} \equiv \text{JSD}(P_A \| P_B)$$

where W_1 is 1-Wasserstein distance and JSD is Jensen-Shannon divergence.

For delayed source-state interpolation, the run must declare an order- q history interpolation operator I_h^q . On a fixed analysis window W , define

$$E_{\text{hist}}(S_\eta; \Delta h, \Delta h/2; W) = \frac{\left(\sum_{m \in W} \|I_{\Delta h/2}^q S_\eta(t_{\text{emit}, m}) - I_{\Delta h}^q S_\eta(t_{\text{emit}, m})\|^2 w_m \right)^{1/2}}{\left(\sum_{m \in W} \|I_{\Delta h/2}^q S_\eta(t_{\text{emit}, m})\|^2 w_m \right)^{1/2} + \varepsilon_0}$$

For nonsmooth state-dependent delay windows, define the jump residual rows

$$\mathcal{D}_{\text{jump}} = \{(\xi_a, k_a, \ell_a, \xi_{\pi(a)}, R_{\text{jump}, a})\}, \quad R_{\text{jump}, a} = \frac{|t_{0, \ell_a}(\xi_a) - \xi_{\pi(a)}|}{\max(\Delta t, \Delta h, \eta/c_f, \varepsilon_0)}$$

Required refinements with pass/fail thresholds

1. Temporal refinement (Δt and $\Delta t/2$, plus $\Delta t/4$ for order check):

- Pass if $E_{\text{rel}}(\Phi) \leq 0.02$, $E_{\text{rel}}(\|\nabla \Phi\|) \leq 0.03$, and $|\Delta \lambda_{\text{self}}|/\lambda_{\text{self}} \leq 0.05$.
- Estimated observed order:

$$p_{\text{obs}}(Y) = \log_2 \frac{E_{\text{rel}}(Y; \Delta t, \Delta t/2)}{E_{\text{rel}}(Y; \Delta t/2, \Delta t/4)}$$

Require $p_{\text{obs}} \geq 0.8$ for at least one primary field channel (Φ or $\|\nabla \Phi\|$).

2. History-resolution refinement (history step halved or interpolation order increased):

- Pass if $E_{\text{rel}}(\Phi) \leq 0.02$, $E_{\text{rel}}(\|\nabla \Phi\|) \leq 0.03$.
- Provenance stability mandatory: $D_W \leq 0.05$ and $D_{JS} \leq 0.02$.
- Delayed-source interpolation stability mandatory whenever delayed states are evaluated from stored history: $E_{\text{hist}} \leq \tau_{\text{hist}}$ with τ_{hist} declared before the run.

3. Spatial refinement (grid/particle resolution increase):

- Pass if $E_{\text{rel}}(\Phi\text{-map}) \leq 0.03$ and $E_{\text{rel}}(\nabla\Phi\text{-map}) \leq 0.05$.
- Self-hit counts and stability-window boundaries must satisfy relative shift ≤ 0.05 .

4. Cross-integrator validation (e.g., symplectic vs RK with matched resolution):

- Pass if $E_{\text{rel}}(\Phi) \leq 0.03$, $E_{\text{rel}}(\|\nabla\Phi\|) \leq 0.05$.
- Provenance agreement must satisfy $D_W \leq 0.08$ and $D_{JS} \leq 0.03$.
- The cross-integrator report must name solver family, interpolation policy, solver residual controls, and event/restart handling. If the compared runs select different active-root identities or transition statuses, the claim fails even if observable plots are close.

5. Continuum moment refinement when a run promotes a coarse PDE, kinetic moment, or Noether sea transport equation:

- Pass if the retained density/current channel satisfies

$$E_{\text{rel}}(R_{\rho}^{\text{cg}}) \leq 0.03, \quad E_{\text{rel}}(R_P^{\text{cg}}) \leq 0.05, \quad E_{\text{rel}}(R_E^{\text{cg}}) \leq 0.05$$

- The moment-closure residual must decrease under temporal, history, and spatial refinement. A continuum plot is not promotion evidence if the next unresolved moment grows or if the memory-current residual is absorbed into fitted constants.

6. Stochastic and response refinement when a run adds Langevin, Fokker-Planck, or fluctuation-response summaries:

- For the first two moments of any declared distribution $P(z, t)$, require agreement with direct event-root ensembles:

$$E_{\text{rel}}(\langle z \rangle) \leq 0.03, \quad E_{\text{rel}}(\text{Cov}(z)) \leq 0.05$$

- If a diffusion tensor $D^{ij}(z)$ is inferred from jump or ledger increments, require it to remain positive semidefinite on the retained domain and stable under refinement.
- If a response kernel χ_{AB} is promoted, require the causal dispersion residual $\mathcal{R}_{\text{KK}}(\chi_{AB}) \leq 0.05$ on the declared frequency band and require any fluctuation-dissipation residual to be reported from the same record.

7. Revised branch-coordinate model selection when a run changes a reduced branch coordinate, chart partition, or residual basis before rerun:

- The proposed coordinate must declare its source fields, equality map, symmetry quotients, and excluded locked keys before any coefficient fit or rerun.
- The selection report must include a held-out residual check, and it must include a phase-origin check whenever the coordinate uses an observation-phase split.
- The design must remain overdetermined after quotienting, with $N_{\text{eq}} > N_{\text{coef}}$ or $R_{\text{df}} > 0$. Report

$$R_{\text{df}} = \frac{N_{\text{eq}} - N_{\text{coef}}}{N_{\text{eq}}}, \quad \frac{\text{tr } H}{N_{\text{eq}}} \leq \frac{1}{2}, \quad \max_i H_{ii} \leq \frac{1}{2}$$

or an explicitly justified equivalent if a linear hat matrix H is not available.

- Branch identity must persist under temporal refinement, history-window refinement, regulator refinement when a regulator is used, and root-ledger refinement. A coordinate that only improves the fitted residual while changing the active branch identity fails model selection.

Machine-checkable convergence output

Every promoted claim must emit `convergence_table.csv` with one row for each required gate: temporal refinement, history-resolution refinement, history-interpolation refinement when delayed states are reconstructed from stored history, spatial refinement, cross-integrator validation, regulator ladder when used, transition-window refinement when a fold-layer or active-root status transition is claimed, and negative control. Each row records the two run identifiers being compared, the restricted observable channel, $E_{\text{rel}}(\Phi)$, $E_{\text{rel}}(\|\nabla\Phi\|)$, D_W , D_{JS} , E_{hist} when applicable, p_{obs} , active-root mismatch, self-hit or stability-window shift, transition-window status, pass/fail status, and failure code.

For continuum or stochastic promotions, append rows for moment-closure, distribution-moments, diffusion-tensor, causal-response, and fluctuation-dissipation when those channels are claimed. These rows must include the artifact hash of the direct event-root run and the artifact hash of the reduced continuum or stochastic run being compared.

For field-theory or continuum-limit promotions, the packet must also declare the scaling-limit datum: regulator family, scaling trajectory, volume or window trajectory when relevant, test-observable class, observable maps from the regulated state to the promoted variables, normalization and mixing rules for composite observables, convergence topology, positivity or reconstruction condition when the claim uses a quantum-field analogue, and the artifact hashes for every regulated run consumed by the limit. Without this datum, a finite-regulator trend is a diagnostic, not a promoted continuum claim.

If the promoted claim invokes an Osterwalder-Schrader-like or Wightman-like field-theory reconstruction, the packet must identify the full reconstruction package it is borrowing: positivity, covariance or symmetry, locality or support condition, vacuum-sector or clustering condition, test-function space, regularity and growth control, and the target reconstructed object. Reflection positivity alone is not enough to promote a regulated numerical family into a local quantum-field analogue.

For revised branch-coordinate promotions, append rows for branch-coordinate-source, branch-coordinate-heldout, branch-coordinate-phase-origin when applicable, branch-coordinate-design, and branch-identity-refinement. These rows must include the artifact hash of the predeclared coordinate packet and the rerun candidate that consumes it.

The regulator row must include each promoted observable Y and the value of

$$E_\eta(Y; \eta, \eta/2) = \frac{\|R(Y_{\eta/2}) - Y_\eta\|_{L^2(W, \{x_k\})}}{\|R(Y_{\eta/2})\|_{L^2(W, \{x_k\})} + 10^{-12}}$$

It also records whether active root-ledger entries match between η and $\eta/2$ after matching source, receiver, root class, and branch status. A convergence plot is not promotion evidence unless the table row containing the plotted quantity is present and tied to the campaign artifact hash.

Regulator extrapolation fits must report the fitted observable, the regulator ladder, the assumed asymptotic form, excluded points if any, stability under fit-window changes, endpoint or singular-window controls when they affect the extrapolation, and a negative-control observable. A fit that behaves smoothly but has no declared observable map, topology, normalization, volume or window estimate when relevant, remainder bound, or independent continuum reconstruction remains below theorem-grade evidence.

Negative control (null test, mandatory)

Run at least one intentionally wrong model choice (wrong history kernel, wrong c_f , or perturbed emission-time solver).

Pass condition for the *pipeline* (not the null run): the null run must break expected invariants by a clear margin, with at least one of:

- invariant drift increase by $\geq 5\times$ relative to the validated run,
- provenance instability $D_W > 0.10$ or $D_{JS} > 0.05$,
- stability-window shift > 0.10 .

If the null run still passes the convergence gates above, treat the claim as numerically unvalidated.

Global acceptance rule

A claim is numerically validated only if all applicable refinement gates pass and the null test fails as required. Conditional gates such as revised branch-coordinate model selection apply only when the claim changes the reduced coordinate, chart partition, or residual basis before rerun.

Perspective

This framework appears to fit a surprising breadth of phenomena not because of any single novelty, but because a small set of simple, mutually reinforcing structural decisions is doing most of the heavy lifting. Two widely discussed choices—reduction to $+\epsilon$ and $-\epsilon$ architrino polarities and choosing $\epsilon = |e|/6$ —help with parsimony and observer-level charge bookkeeping, but the outsized wins come from how delayed line-of-action action, Jacobian-weighted causal flux, and same-source causal-root branches conspire to produce stability, scale selection, and emergent “magnetic-like” behavior without ever invoking right-hand-rule cross products.

Historically, general relativity and quantum mechanics are extraordinarily successful as effective theories that summarize large classes of phenomena. We position this neoclassical, delayed line-of-action model as a simpler dynamical substrate whose coherent assemblies recover GR/QM-like phenomenology in appropriate coarse-grained, slow/weak, or phase-locked limits.

We work throughout in units with primitive wake speed $c_f = 1$; per-hit accelerations are directed along $\hat{\mathbf{r}}$, weighted by the causal Jacobian, and superpose linearly.

Delayed Emission on Jacobian-Weighted Isochrons

- What we assume:
 - Sources emit potential on expanding causal isochrons with surface density $\propto 1/r^2$, represented distributionally by $\delta(r - \tau)$ with $\tau = t - t_0$.
 - Each causal hit is directed along $\hat{\mathbf{r}}$ from the source history point to the receiver, with received magnitude weighted by the branch Jacobian.
- Why it matters:
 - Gauss-like behavior follows immediately ($1/r^2$ on causal wake fronts).
 - Moving systems automatically generate tangential components in the receiver’s frame due to path-history geometry and causal-flux bunching: the “aim point” is in the past, and source motion enhances or suppresses active branches through the Jacobian. Orbital

and vortex-like patterns emerge from delay, not from any $B \propto \mathbf{v} \times \mathbf{E}$ construction.

- Consequence:
 - Many “magnetic” phenomenologies (circulation, axial vortices, flux tubes) can be reproduced as kinematic consequences of delayed, Jacobian-weighted line-of-action pushes. There is no right-hand rule, no cross products, just geometry, flux weighting, and time delay.
-

Constant per-wavefront emission

- What we assume:
 - Emission cadence and per-wavefront amplitude are constant at the source.
 - Why it matters:
 - Simplifies calibration and emphasizes that stability and scale selection arise from delay and self-interaction. Receiver motion influences instantaneous power via $\mathbf{F} \cdot \mathbf{v}$ through the radial component v_r , while source motion modulates the received force magnitude through the Jacobian.
 - With η -mollification ($\delta \rightarrow \delta_\eta$), the calculation can define Φ_η and verify $\Delta E_k = -\Delta U$ on resolved intervals while still taking $\eta \rightarrow 0$ for sharp impulses.
-

Self-Hit Root Onset

- What we assume:
- Same-source self-hit is accepted only when the root equation

$$\mathcal{C}_{aa}(t) = \{ s < t : \|\mathbf{x}_a(t) - \mathbf{x}_a(s)\| = c_f(t - s) \}$$

is nonempty and the active root passes the transversality/Jacobian floor. A speed excursion above c_f is a necessary warning condition for simple nontrivial roots, not a sufficient criterion.

- Self-hits are always repulsive (like-on-like).
 - Why it matters:
 - This nonlinearity is the core stabilizer. Strictly sub-field-speed interval history rules out nontrivial self-hit roots on that interval, while super-field-speed curved history can open an internal, strong, repulsive channel that balances or overtakes inward trends.
 - Scale selection emerges: the balance of delayed attraction with self-repulsion defines a smallest sustainable orbital radius d_0 and a fastest natural frequency, yielding a canonical time unit t_0 .
-

Superposition with isochrons and η -regularization

- What we assume:
 - All wake contributions superpose linearly at the level of distributions (isochrons add).
 - We use a narrow Gaussian isochron δ_η when continuous-time derivatives are needed.
 - Why it matters:
 - Locality: inverse-square geometric weighting together with finite-speed branch selection makes nearby coherent roots dominant, but infinite populations still require an explicit cutoff, screening rule, cancellation estimate, sampled mean field, or principal-value/mean-field subtraction.
 - Bookkeeping: with δ_η , standard ODE solvers can integrate numerically; with δ , the analysis can reason about impulses and events. Both views agree in the $\eta \rightarrow 0$ limit for integrals over resolved intervals.
-

Assembly grammar -> nested shell braid and flux tubes

- What we assume:
 - Binary orbits are the base motif; binaries can nest with wide scale separation; a nested shell braid is dynamically robust.
 - Polar regions of fast binaries host persistent axial structures (vortex-like loci in the delayed wake geometry), which couple between assemblies.
- Why it matters:

- Color-like structure arises naturally from three internal binaries: distributing axial architrinos across three axes creates three distinguishable, yet symmetric, configurations.
- Flux-tube-like coupling is not a particle exchange but a persistent geometric linkage between polar vortices—consistent with confinement-like phenomenology without invoking a separate gauge field.

Charge quantization at $\epsilon=|e/6|$

- What we assume:
 - The architrino polarity magnitude is ϵ , so observer-level quark electric charges are integers of ϵ .
- Why it matters:
 - Observed quark fractions ($\pm 1/3$, $\pm 2/3$ of e) become $\pm 2\epsilon$ and $\pm 4\epsilon$ integers in the natural unit. This removes “fractionality” at the fundamental level and simplifies assembly rules and conservation statements.

Consequences that explain the “fit”

- Stability without fine-tuned potentials:
 - The $\|\mathbf{v}\| = c_f$ switch and delay geometry set operating points and prevent singular collapse.
- Scale emergence:
 - d_0 and t_0 arise from dynamics; they are not postulated rulers and clocks but attractors of the binary system.
- Shielding and apparent inertia:
 - Fast internal motion produces far-zone cancellation; the tiny residual wake signature of a coherent assembly behaves like inertial mass in interactions with the outside.
 - Magnetism without magnetism:
 - Tangential effects and axial structures appear as a corollary of path-history plus Jacobian-weighted line-of-action per-hit action. No cross products required.

What the model explicitly does not use

- No Lorentzian spacetime metric at the fundamental level (background is absolute time + Euclidean space; emergent cones are effective, not kinematic).
- No right-hand-rule magnetism or $\mathbf{v} \times \mathbf{B}$ forces; every per-hit action is along $\hat{\mathbf{r}}$.
- No gauge field inventory beyond architrino causal wakes; interaction carriers are the geometry of delayed isochrons and their couplings.

Validation and next steps (concrete)

1. Far-field cancellation and the zero-potential axis
 - Compute the time-averaged multipole expansion of a high-frequency binary; show leading terms cancel along the rotation axis and decay rapidly off-axis.
 - Observable: a “quiet line” (near-zero net potential) threading the binary.
2. Scale selection for d_0 and t_0
 - With $\delta \rightarrow \delta_\eta$, compute the mean inward attraction from the partner versus the mean outward self-repulsion across one orbit; the fixed point defines d_0 and the maximum orbital frequency $2\pi/t_0$.
 - Prediction: the same d_0 appears across binaries with the same ϵ and c_f , independent of initial conditions after sufficient relaxation.
3. Energy consistency across the $\|\mathbf{v}\| = c_f$ transition
 - Use Φ_η to evaluate U and verify $\Delta E_k = -\Delta U$ across events that cross the self-hit onset boundary; in the $\eta \rightarrow 0$ limit, impulses integrate to the same work.
4. Numerical recipe (robust, minimal assumptions)
 - For each receiver time t : (i) root-find causal emission times t_0 for all sources (and self), (ii) discard non-physical roots ($H(0) = 0$, handle $r = 0$ by symmetry), (iii) sum $a_{ot \leftarrow o}(t; t_0)$, (iv) integrate velocity and position with an event-aware scheme. Use ϵ -thickening for smooth integration when needed.

Comparisons and falsifiable edges

- Classical E&M:

- Replace Maxwell + Lorentz force with delayed, radial-only action; predict the same far-zone radiation patterns for coherent assemblies but different near-zone dynamics when $\|\mathbf{v}\| \approx c_f$ or self-hits occur.
- QCD phenomenology:
 - Confinement-like behavior emerges from polar-vortex coupling; falsifiable via constraints on hadron breakup channels and energy distributions if the coupling geometry is perturbed.
- Inertia/apparent mass:
 - Predicts context-dependent inertia from shielding; assemblies in different internal phases could exhibit small, measurable variations in response to identical external effective fields.

Open Closure Questions

- Exact analytic forms for d_0 and t_0 in the symmetric binary with the canonical modulation.
- Rigorous conditions for uniqueness/multiplicity of causal roots in accelerated motion and their contribution to stability.
- Statistical mechanics of many-body wake structures: when and how do coherent, Lorentz-consistent effective cones emerge from moving-assembly deformation, clock/ruler retuning, and Noether sea response, and with what characteristic speed relative to the declared branch speed c_* ?

Plain language summary: radial hits, time delay, constant per-wavefront amplitude, and self-hit for fast movers are enough to produce stable orbits, natural rulers and clocks, shielding that looks like inertia, and “magnetic-like” structures without right-hand-rule magnetism.

Effective observables and states (quantum-like layer)

Premise: single-hit information is sparse. At an instant, a receiver learns only (i) the net magnitude of the push and (ii) an unoriented line of action through its current position. The $\mathbb{U}_{\text{now}}$ universe-state perspective can include the full source-tagged emission ledger as complete-state bookkeeping, but a local receiver or Physical Observer cannot infer that hidden ledger from a single hit.

- Emission ledger (microstate): the set of tuples $(t_0, \mathbf{s}_j(t_0), \mathbf{v}_j(t_0), q_j)$ over all sources j that causally affect the receiver.
- Observational map: ledgers map to histories of hits $\{A(t_k), L(t_k)\}$ across one or more receivers and over time.
- Observational equivalence: two ledgers are equivalent if they induce indistinguishable hit histories at the chosen resolution (including mollifier width η , temporal sampling, and receiver geometry).
- Coarse-grained PDE observables (Method 1):
 - Number density $n(\mathbf{x}, t)$: count-per-volume of architrinos.
 - Polarity density $\rho(\mathbf{x}, t)$: net $+\epsilon - \epsilon$ per unit volume; natural source term in continuum PDE variants.
 - Energy density $\mathcal{E}(\mathbf{x}, t)$: local kinetic + potential energy density for validation and conservation checks.
 - Use: these fields are the natural inputs/targets for grid-based PDE runs and for validating event-driven simulations in aggregate.

Observability axioms:

- A1 Single-hit observables are magnitude A and an unoriented line L ; orientation along L , source identity, distance r , and emitter speed $\|\mathbf{v}_{\text{em}}\|$ are not individually observable at an instant.
- A2 All practical observables are functionals of hit histories across time and receivers; unique micro inversion is generically impossible.
- A3 An effective “state” is a probability measure over observationally equivalent ledger classes, updated as new hits arrive.

Bayesian operational stance:

- State update = conditioning on new hit histories; active interventions (changing receiver geometry/filters) alter future histories and thus the posterior over ledger classes.

Plain language: a receiver never sees the full ledger of who emitted what; it sees only a time series of push magnitudes and lines. The appropriate language is therefore statistical over micro-histories that fit those pushes.

$\mathbb{U}_{\text{now}}$ Note: Limits of Perfect Clocks and Frames

Absolute time and Euclidean frames remove coordinate ambiguity (synchronization and alignment) but not physical ambiguity:

- Sign/side ambiguity: attraction from $+\epsilon$ on one side vs repulsion from $-\epsilon$ on the opposite side along the same line remain indistinguishable at an instant.
- Baseline distance scaling plus Jacobian modulation: $A \propto 1/(r^2|J|)$; emitter speed affects both the timing of causal roots and the received per-hit amplitude through the branch Jacobian.
- Collinear superposition: several sources along the same unoriented line can sum to the same instantaneous A and L .

- Self-hit aliasing: self-intersections can mimic external sources along L .
- Surrogate location recast: any instantaneous hit may be recast to a stationary surrogate source placed somewhere along L with an adjusted emission time; useful for inference and visualization, but it does not resolve the sign/side ambiguity or fix distance without temporal data.

Consequence: embedded observers and synthetic detector records must reason statistically over ledger classes. The $\mathbb{U}_{\text{now}}$ universe-state perspective can compare those classes against the complete ledger, but the observer-accessible data remain many-to-one; “quantum-like” observability is not a contradiction but a necessity.

Single-source multi-hit nuance vs universal superposition

Even for a single source, the receiver cannot be sure that a given shove did not come from multiple distinct emission times $t_0 \in \mathcal{C}_{\text{obj}}(t)$ on that same source. When $\|\mathbf{v}_j\| > 1$ or the source trajectory curves, several roots of $r = v(t - t_0)$ can occur and arrive in close succession along the same unoriented line of action, contributing separate per-hit pushes that are locally indistinguishable as to origin.

However, this is not the dominant practical difficulty. The governing issue is global superposition: at any instant the net field is the linear sum of contributions from all architrinos in the universe whose causal isochrons intersect the receiver “now.” While inverse-square surface dilution and Jacobian weighting usually make nearby sources dominate, the mapping from the universal emission ledger to observed hit histories remains vastly many-to-one. Consequently, inference must be temporal, statistical, and multi-view, not a frame-perfect instantaneous inversion.

Operational noncommutativity and contextuality (emergent)

Measurement procedures are interventions that condition future hit histories:

- Let F, G be experimental contexts (e.g., planar-mode analyzers, path blockers, timing gates). Because they modify trajectories and thus the set of future causal roots, their composition generally satisfies $F \circ G \neq G \circ F$ at the level of observed statistics.
- Contextuality: the distribution over ledger classes that best explains data depends on which filters were applied and in what order; the outcomes are context-dependent without invoking microscopic cross-product forces.

Plain language: a present intervention changes which pushes will be recorded later; doing A then B is not generally the same as doing B then A .

Interference and amplitude-squared from planar-mode overlap

Linear superposition at the isochron level plus coherent geometry yields interference-like patterns in aggregates:

- Photon planar-mode ledgers from multiple sources add linearly at the effective-amplitude level; a detector that integrates over a small time window and area effectively accumulates a complex amplitude A_{mode} from coherent sub-bundles.
- Intensity emerges as an overlap norm proportional to $|A_{\text{mode}}|^2$ under time/ensemble averaging of phase-like structure encoded by path histories.
- Polarization example (already used): Malus’s law arises as a geometric projection of a planar mode’s transverse ledger onto an analyzer axis, giving $\cos^2 \theta$ transmission without right-hand-rule magnetism.

Plain language: aligned planar-mode records add, misaligned ones cancel, and the recorded strength scales like the square of the pattern overlap.

Reconstruction Under Information Bounds

Instantaneous inversion is ill-posed; reconstruction is temporal, multi-view, and prior-guided:

- Multi-receiver geometry: use separated receivers to triangulate unoriented lines at the same t ; intersecting rays yield two-sided candidate loci.
 - Time-series constraints: track $L(t)$ and timing-derived $r(t)$ proxies; curvature and rotation of L constrain source paths.
 - Active probing: vary receiver motion/filters to sample different roots and break degeneracies.
 - Priors: charge inventories, speed bounds, assembly templates (e.g., binaries, planar-mode statistics) shrink the hypothesis space.
 - Estimation: run Bayesian filters or particle sets over ledger classes; update with each hit; report identifiability and uncertainty, not single-point “sources.”
-

Worked micro-to-effective examples

- Two-planar-mode interference:
 - Setup: two coherent photon planar modes reach a screen. The observed intensity pattern is the squared norm of their geometric overlap along the screen, set by relative phase encoded in path history.

- Which-way intervention: inserting a context that disrupts one planar mode's coherence changes the ledger classes and removes the overlap term, flattening the pattern.
- Polarization analyzer:
 - The analyzer projects the planar mode's transverse ledger onto its axis; transmission $\propto \cos^2 \theta$ follows immediately from geometric projection.
- Sequential filters (order matters):
 - Two non-parallel analyzers $F(\theta_1)$ and $G(\theta_2)$ applied in different orders yield different transmitted patterns because they recondition future causal roots differently: $F \circ G \neq G \circ F$.

Falsifiable edges and tests (observability-focused)

- Context order test: demonstrate order-dependent transmission with sequential analyzers on coherent planar modes; quantify the asymmetry $F \circ G$ vs $G \circ F$.
- Planar-mode interference robustness: map how partial decoherence (deliberate jitter in source paths) suppresses the overlap term; compare to predicted $|A|^2$ decay with coherence length.
- Multi-receiver triangulation under ambiguity: show that two-sided localization from unoriented lines plus time series reduces, but does not eliminate, sign/side and distance-speed degeneracies—matching Step 9 limits.
- Bell-type correlation target (open): assess whether planar-mode phase models with absolute time can reproduce observed $\cos(2\theta)$ correlations across separated analyzers without hidden cross-product forces; treat Tsirelson-like bounds as a stringent benchmark.

Plain language: we can test the framework by checking order effects, interference weakening when we scramble coherence, and how much multiple receivers really help; reproducing quantum correlations is the toughest, and we flag it as an explicit target.

README

This note is the launch overview for the simulation branch. It explains the common simulation frame, the role of the virtual $\mathbb{U}_{\text{now}}$ universe-state perspective, and the separation between raw microstate logging and detector-level synthetic observables.

Use it as the top orientation document before reading the more specialized [Simulation Run Protocols](#), convergence checks, action-energy notes, the [A₀ branch certificate protocol](#), the [A₀ Tier 0 result interpretation](#), the [Nested Shell Braid Action-Increment Protocol](#), the [Retuning-Map Toy Model](#), the [Cosmology Shared Residual Fit Protocol](#), the [Redshift-Budget Toy Model](#), the [Static Response Vector Toy Model](#), and the [Hydrogen I_N Spectral Coefficient Row Toy Scan](#).

Simulation Frame: Virtual $\mathbb{U}_{\text{now}}$ Perspective

- All tiers are implemented in the absolute Euclidean frame (fixed x,y,z; absolute t).
- The simulator effectively plays the role of the $\mathbb{U}_{\text{now}}$ universe-state perspective by integrating the master equation and maintaining $S(t)$.
- Raw outputs are $\mathbb{U}_{\text{now}}$ -style (fields, provenance, microstate summaries).
- “What experiments see” is generated by post-processing: embed detector assemblies with worldlines $\mathbf{X}_{\text{det}}(t)$, compute $\tau_{\text{det}}(t)$, and generate detector-like logs.

Checklist per tier:

- What $\mathbb{U}_{\text{now}}$ records (Φ , $\nabla\Phi$, Noether sea variables, provenance)
- How to compute physical observables (τ , redshift, lensing proxies, public gravitational-wave packet outputs: detector strain, phase, event energy ledger, timing residuals, and provenance)
- Convergence requirements for each output type
- For the first mass-map target, how the A₀ Tier 0/Tier 1 branch certificate is separated from later energy, shielding, and Noether sea response interpretation
- For Tier 0 rows, how active roots, raw roots, excluded near-zero self roots, residual semantics, and promotion gates should be read before any Tier 1 continuation

Simulation Frame and the $\mathbb{U}_{\text{now}}$ universe-state perspective

All simulation tiers are implemented in the absolute frame:

- **Spatial frame:** fixed Cartesian grid in the Euclidean void, (x, y, z) constant in time.
- **Temporal frame:** global absolute time t , advanced in discrete steps Δt .
- **Microdynamics:** architrino positions and velocities updated according to the master equation; potentials propagated at speed c_f .

From the code's perspective, we are always the $\mathbb{U}_{\text{now}}$ **universe-state perspective**:

- We know $S(t)$ (all architrinos, all assemblies) at each time step.
- We can compute fields and Noether sea state anywhere in the domain.

To connect to experiment:

- We embed **model detectors** (assembly worldlines) in this frame.
- We compute:
 - What fields they experience along their paths,
 - How their internal clocks tick (τ vs t),
 - What signals they register (arrival times, redshifts, intensity patterns).
- Synthetic observables are derived from these detector responses, not from raw $S(t)$ directly.

This enforces a clean separation between:

- Fundamental dynamics in the absolute frame (what the simulation integrates),
- Emergent observational physics (what real experiments would see).

Run Protocols

This chapter defines the mandatory runtime protocol for simulations carried out in the absolute-frame implementation of the theory. Its role is to standardize the frame, logging requirements, provenance bookkeeping, metadata, and acceptance gates so results from different runs can be compared and audited coherently.

The opening gives the top-level simulation rule set; the later sections unpack the absolute-frame interpretation and the required $\mathbb{U}_{\text{now}}$ instrumentation in more detail.

Master Simulation Protocol (Absolute Frame)

1. **Coordinate Anchor:** All simulations run on a fixed Cartesian grid chosen as the coordinate scaffold for the Euclidean void. $\text{Grid}[x][y][z]$ is a chart address, not an intrinsic label in the void.
2. **Clock Rate:** The simulator uses a global Time counter (absolute t). No relativistic scaling is applied to the integration step itself.
3. $\mathbb{U}_{\text{now}}$ **universe-state interface:** Every run must instantiate an array of fixed virtual sensors to log Φ and $\nabla\Phi$ at declared absolute-frame grid addresses.
4. **Noether sea Initialization:** Low-excitation Noether sea runs must pre-populate the grid with a lattice of coupled pro/anti Noether braids to simulate Noether sea influence on test particles.
5. **Convergence:** Δt refinement must be accompanied by “History Resolution” refinement to ensure self-hit calculations are numerically stable.
6. **Campaign Packet:** Any run used for a proof certificate, branch-certificate gate, or promoted validation claim must emit a machine-checkable packet rather than only plots or summaries.

Simulation Campaign Object

Every promoted numerical claim is carried by a campaign object, not by an isolated plot or best-fit table:

$$\mathcal{C}_{\text{sim}} = (\text{id}, S_\eta, \mathcal{G}_h, \Delta t, \eta, I_h^q, \mathcal{L}_{\text{root}}, \mathcal{T}_\eta, \mathcal{R}_{\text{branch}}, \Pi_{\mathbb{U}_{\text{now}}}, \mathcal{E}_{\text{conv}}, \mathcal{F})$$

Here id fixes the run identifier and source commit, S_η is the regularized state history, $\mathcal{G}_h$ is the spatial and history mesh, Δt is the absolute-time step, $\eta > 0$ is the causal-wake regularization width, I_h^q is the declared order- q history interpolation operator, $\mathcal{L}_{\text{root}}$ is the causal-root ledger, $\mathcal{T}_\eta$ is the transition-record family for fold-layer, separator, or active-root status windows, $\mathcal{R}_{\text{branch}}$ is the named branch-residual vector, $\Pi_{\mathbb{U}_{\text{now}}}$ is the provenance log, $\mathcal{E}_{\text{conv}}$ is the convergence-measure vector, and $\mathcal{F}$ is the finite failure-code set.

When a campaign is used for a continuum, field-theory, or regulator-removal claim, it must also attach an extraction map: the regulated observables, test windows, volume or window trajectory when relevant, normalization and mixing rules, convergence topology, positivity or reconstruction condition when applicable, and the artifact hashes for the regulator ladder. If independent methods or benchmarks are used, the packet must expose their normalization conventions and error envelopes before comparing coordinates. These fields tell reviewers exactly what is claimed to survive the finite run and what remains only a regulator-level diagnostic.

For a QFT-like reconstruction claim, the campaign must also state the presentation being targeted, such as Wightman data, Osterwalder-Schrader data, a local observable net, or a weaker named comparison. The packet must then list the hypotheses required by that presentation rather than using generic terms such as `continuum field` or `reconstructed field`.

The state history is

$$S_\eta(t) = \{(\mathbf{x}_i(t), \mathbf{v}_i(t), q_i)\}_{i=1}^N, \quad S_{\eta,t}(\theta) = S_\eta(t + \theta), \quad \theta \in [-h, 0]$$

A Tier 1 packet must state whether this history is evaluated in $C^1([-h, 0])$, $W^{1,\infty}([-h, 0])$, or a stricter history class. A missing history class is an incomplete artifact, because the delayed source-state evaluation cannot be audited without it.

The mesh and interpolation record is

$$\mathcal{G}_h = (\Omega_h, \Delta x, \{x_k\}_{k=1}^K, \Theta_h, \Delta h, \text{bc})$$

where $\Omega_h \subset \mathbb{R}^3$ is the Euclidean-void computational domain, $\{x_k\}$ are the fixed $\mathbb{U}_{\text{now}}$ sample points, $\Theta_h \subset [-h, 0]$ is the stored path-history mesh, Δh is the history resolution, and bc records boundary conditions. The interpolation operator I_h^q is part of the packet; delayed source states cannot be reconstructed by an implicit or undocumented lookup rule.

Executable Diagnostic Contract

A campaign that disciplines a proof certificate must reduce its numerical status to predeclared scalar diagnostics. The default diagnostic vector is

$$\mathcal{D}_{\text{exec}} = (D_{\text{branch}}, D_{\text{ref}}, D_{\text{ord}}, D_{\text{hist}}, D_{\text{space}}, D_{\text{cross}}, D_{\text{prov}}, D_{\text{cons}}, D_{\eta}, D_{\text{jump}})$$

where every component is a ratio with passing threshold 1. The component meanings are:

Component	Required role
D_{branch}	largest branch residual divided by its declared tolerance
D_{ref}	temporal refinement residual for Φ , $\ \nabla\Phi\ $, and self-hit rate
D_{ord}	observed-order gate for the retained primary field channel
D_{hist}	history-resolution, interpolation, and provenance-distribution gate
D_{space}	spatial refinement and self-hit stability-window gate
D_{cross}	cross-integrator agreement with matching branch identity
D_{prov}	$\mathbb{U}_{\text{now}}$ causal-provenance residual
D_{cons}	energy, momentum, and angular-momentum drift gate
D_{η}	regulator-dependence gate for promoted observables
D_{jump}	jump or branch-transition residual for nonsmooth windows

The Tier 1 acceptance predicate is

$$\text{Accept}_1(\mathcal{C}_{\text{sim}}) \iff R_0 \in \text{Candidate}_1, \quad \max_a \mathcal{D}_{\text{exec},a} \leq 1, \quad \Delta_{\text{root}}(\Delta t, \Delta t/2) = 0$$

$$\Delta_{\text{root}}(\Delta h, \Delta h/2) = 0, \quad \Delta_{\eta, \text{root}} = 0, \quad \text{NullFail} = 1, \quad \text{Artifacts} = 1$$

Here NullFail = 1 means the negative control violates at least one required null-test margin, and Artifacts = 1 means every required artifact exists with a content hash and source commit.

Failure routing is deterministic. Missing required artifacts, source commits, pre-run tolerances, or hashes route to artifact_incomplete. Changing a promoted observable, tolerance, branch label, or regulator ladder after output inspection routes to hidden_tuning. Unstable active-root identity routes to branch_root_instability; failed refinement routes to mesh_nonconvergence; failed provenance routes to provenance_discontinuity; failed conservation routes to conservation_drift; failed regulator rows route to regulator_dependence; and exit from the admissible η continuation set routes to eta_continuation_failure.

Proof-Certificate Handoff

The proof-to-simulation handoff for a finite certificate is

$$\mathbf{H}_{\text{proof} \rightarrow \text{sim}} = (\text{certificate_id}, S_{\eta,0}, W, \Lambda, \mathcal{L}_{\text{root}}^{\text{expected}}, \tau_{\text{branch}}, \tau_{\text{conv}}, \tau_{\eta}, \text{Null}, \text{Outputs})$$

It names the source certificate, initial history, analysis window, branch label, expected active-root classes, branch tolerances, convergence tolerances, regulator ladder, negative-control mutation, and required output channels before the run starts.

The simulation-to-proof handoff is

$$\mathbf{H}_{\text{sim} \rightarrow \text{proof}} = (\text{artifact_hashes}, \mathcal{L}_{\text{root}}^{\text{matched}}, \mathcal{T}_{\eta}, \mathcal{R}_{\text{branch}}, \mathcal{E}_{\text{conv}}, \mathcal{D}_{\text{exec}}, \text{failure_code}, \text{promotion_status})$$

A proof packet may cite a simulation only through this handoff. It must state whether every expected active root was matched under Δt , Δh , and η refinement, which residual component controls the verdict, and which artifact contains each value.

Lemma (simulation-promotion criterion). Let Q be a priority-theory claim whose variables are contained in $\mathcal{C}_{\text{sim}}$, and let R_1 be a Tier 1 continuation of a Tier 0 candidate R_0 . If R_0 satisfies the Tier 0 acceptance criteria, R_1 satisfies the Tier 1 acceptance criteria, the negative control fails as required, and

$$\max_a \frac{\mathcal{E}_{\text{ref},a}}{\tau_{\text{ref},a}} \leq 1, \quad \max_a \frac{\mathcal{E}_{\text{prov},a}}{\tau_{\text{prov},a}} \leq 1, \quad \max_a \frac{\mathcal{E}_{\text{cons},a}}{\tau_{\text{cons},a}} \leq 1$$

$$\max_a \frac{\mathcal{R}_{\text{branch},a}}{\tau_{\text{branch},a}} \leq 1, \quad \max_Y \frac{E_\eta(Y)}{\tau_{\eta,Y}} \leq 1$$

with all tolerances declared before the run, then the result may be promoted from numerical candidate to simulation-supported priority claim for Q . This lemma does not convert a simulation-supported priority claim into an analytic theorem; it authorizes proof-program routing only with artifact hashes and the failure-code ledger attached.

A_0 Branch-Certificate Protocol

The first mass-map target has a specialized protocol in [A₀ Branch Certificate Protocol](#), with Tier 0 row semantics summarized in [A₀ Tier 0 Result Interpretation](#). That protocol separates four stages:

1. Tier 0 algebraic branch search for finite root-ledger candidates.
2. Tier 1 $\eta > 0$ delayed-dynamics continuation and Floquet diagnostics.
3. Tier 2 internal-energy and shielding extraction.
4. Tier 3 Noether sea response tensor probes.

A rerun after a finite-coordinate no-go must include the predeclared branch-chart revision record; residual-selected coordinates, locked keys promoted into branch geometry, or benchmark-derived inputs invalidate the packet as hidden fitting.

After the current compact scalar-basis no-go, an A_0 rerun must also predeclare the corrected one-period branch-equation basis, the non-circular carrier correction if used, the residual-balance ledger, held-out residual rule, and failure code before it can proceed to Δ_k or η -ladder persistence.

No simulation run should report $\zeta(A_0)$, $E_{\text{internal}}(A_0)$, or $\mathcal{M}_{\text{sea}}^{ab}$ as accepted outputs unless the preceding branch-certificate gates have passed.

Cosmology Shared-Residual Protocol

The first cosmology-facing validation scaffold is [Cosmology Shared Residual Fit Protocol](#). It specializes the campaign-packet rule to the shared dark-energy and cosmology calibration gate. The packet tests whether SN, BAO, CMB, weak-lensing, redshift-space-distortion, BBN, and pre-BBN branch residuals can consume one θ_{sea} without per-observable retuning.

No cosmology packet should report a promoted dark-energy, H_0 , S_8 , BBN, CMB, or growth closure unless its ordinary residuals and cross-family projection penalty are both inside declared tolerances.

Public Gravitational-Wave Benchmark Protocol

A public gravitational-wave benchmark packet tests the effective gravitational-radiation limit against versioned open strain and parameter-estimation records. The packet is not evidence for a fundamental metric ripple in the Euclidean void. It is an observer-level validation object: the $\Delta\Delta\Delta$ simulation must predict detector strain, phase, event-ledger energy balance, and any photon/gravity timing residual through its Noether sea response map and then compare those predictions to public artifacts.

The packet object is

$$\mathcal{C}_{\text{GW}} = (\text{event_id}, \text{catalog}, \text{event_version}, \mathcal{D}, \mathcal{S}_h, \mathcal{P}_{\text{PE}}, \mathcal{P}_{\text{wave}}, \mathcal{Q}_{\text{det}}, \mathcal{L}_{\text{EpJ}}, \mathcal{R}_{\text{GW}}, \Pi_{\text{wave}}, \mathcal{F})$$

Here $\mathcal{D}$ names the detectors, $\mathcal{S}_h$ names the strain files, $\mathcal{P}_{\text{PE}}$ names posterior-sample and parameter-estimation records, $\mathcal{P}_{\text{wave}}$ names the waveform-family or numerical-relativity provenance, $\mathcal{Q}_{\text{det}}$ carries calibration, data-quality, injection-mask, down-sampling, and glitch-treatment records, $\mathcal{L}_{\text{EpJ}}$ is the event conservation ledger, $\mathcal{R}_{\text{GW}}$ is the residual vector, and Π_{wave} maps each fitted or plotted sample back to public artifacts.

The residual vector is

$$\mathcal{R}_{\text{GW}} = (R_h, R_\phi, R_E, R_J, R_{c_g}, R_{\text{det}}, R_{\text{PE}}, R_{\text{prov}})$$

R_h compares whitened or otherwise declared detector strain on the predeclared analysis window; R_ϕ compares unwrapped inspiral-merger phase on the declared frequency band; R_E checks source masses, remnant mass, radiated energy, recoil, ejecta or heat-channel terms, and boundary exchange in one conservation ledger; R_J checks angular-momentum accounting when the packet claims spin or recoil closure; R_{c_g} is used only for multimessenger timing rows; and the final three residuals are provenance-completeness checks.

For a multimessenger row,

$$R_{c_g} = \frac{\Delta t_{\text{obs}} - \Delta t_{\text{src}}}{D_L/c_\gamma}, \quad \Delta t_{\text{obs}} = t_\gamma - t_{\text{GW}}$$

The intrinsic source-emission delay Δt_{src} must be declared before fitting the gravity-channel speed. A packet fails as hidden tuning if it absorbs photon/gravity timing into an undeclared source delay, changes the analysis band after inspecting residuals, substitutes a cleaned strain product without recording a new provenance row, or changes waveform family after comparing to the data.

The minimum artifact list is `event.json`, `strain_files.json`, `detector_quality.json`, `parameter_estimation.json`, `waveform_provenance.json`, `analysis_window.json`, `strain_residuals.csv`, `phase_residuals.csv`, `energy_ledger.csv`, `speed_residual.json` when applicable, `artifact_hashes.json`, and `failure_report.md`. For long binary-neutron-star inspirals the packet must also include a glitch/cleaning row, a low-frequency cutoff row, and a reason if any detector is excluded from a visible-strain comparison. For short binary-black-hole benchmarks the packet must include an inspiral-merger-ringdown window, detector arrival-time comparison, and ringdown handoff row.

The normalized public-data diagnostic is

$$\mathcal{D}_{\text{GW}} = (D_h, D_\phi, D_E, D_{c_g}, D_{\text{det}}, D_{\text{PE}}, D_{\text{prov}})$$

with

$$D_h = \frac{R_h}{\tau_h}, \quad D_\phi = \frac{R_\phi}{\tau_\phi}, \quad D_E = \frac{R_E}{\tau_E}, \quad D_{c_g} = \frac{|R_{c_g}|}{\tau_{c_g}}$$

D_{det} , D_{PE} , and D_{prov} are binary completeness ratios whose value is 0 only when detector masks/calibration, parameter-estimation release metadata, and artifact hashes are all present. A packet can support a promoted gravitational-wave claim only if

$$\max_a \mathcal{D}_{\text{GW},a} \leq 1$$

and the public-data provenance row was fixed before waveform comparison.

The first benchmark triad is:

Packet row	Required public-data role	Failure routed if missing
GW150914_short_bbh	Short inspiral-merger-ringdown strain, two-detector arrival timing, radiated-energy ledger, numerical-relativity waveform provenance, and ringdown handoff	artifact_incomplete or hidden_tuning
GW170817_long_bns	Long inspiral strain, three-detector timing, glitch/cleaning provenance, chirp-mass phase benchmark, and parameter-estimation waveform-family record	provenance_discontinuity or mesh_nonconvergence
GW170817_GRB_speed	Photon/gravity timing residual with luminosity distance, observed delay, and intrinsic source-emission lag nuisance	hidden_tuning or conservation_drift

This public benchmark packet is a success marker under the existing simulation provenance and conservation gates, not a new gate family. Its value is that public strain, parameter-estimation samples, waveform provenance, and multimessenger timing make strong-field radiation tests replayable without importing GR waveform success as [AAA](#) ontology.

Tier 0 / Tier 1 Campaign Packet

Tier 0 and Tier 1 results are accepted only through an auditable campaign packet. The packet must include the source commit, pre-run tolerances, root ledger, branch residual vector, convergence table, η ladder when a regulator claim is made, declared history interpolation, failure report, and artifact hashes. When a run crosses a fold-layer, separator, or active-root status transition, the packet must also include transition records for that window.

The minimum Tier 0 packet contains `campaign.json`, `mesh.json`, `state_vector.json`, `root_ledger.json`, `branch_residuals.json`, `candidate_rows.csv`, `failure_codes.md`, and `promotion_gate.md`. For corrected branch-equation reruns, `branch_residuals.json` must include the branch-native basis, predeclared coefficient rule, held-out residual rule, and pass/fail value for the residual-balance row. The minimum Tier 1 packet adds `run_metadata.json`, U_{how} provenance data, `history_interpolation.json`, `convergence_table.csv`, `eta_ladder.csv`, `conservation_ledger.csv`, `cross_integrator_report.md`, `negative_control_report.md`, `failure_report.md`, and `promotion_lemma_check.md`. If a Tier 1 run claims a branch transition, it also emits `transition_records.json` with the status, regularization route, transition-window scale, root-ledger rows, and promoted observables for each transition window.

The `cross_integrator_report.md` artifact must name the solver family, delayed interpolation polynomial or reconstruction rule, nonlinear solve residuals when implicit stages are used, small-delay or vanishing-delay encounters, and event or restart handling. Cross-integrator agreement is evidence only when the branch identity and transition records match, not merely when plotted observables are close.

A Tier 1 packet supports a proof or validation claim only when the branch residuals, convergence checks, provenance checks, conservation checks, regulator-dependence checks, and negative control all pass with tolerances declared before the run. If any promoted scalar, root count, branch label, stability gap, or tolerance is selected after inspecting output, the packet fails as hidden tuning.

Run Protocol: Absolute-Frame + $\mathbb{U}_{\text{now}}$ Logging

Absolute frame rule

All simulations integrate dynamics in the absolute Euclidean frame:

- Fixed Cartesian coordinates (x,y,z) in a chosen scaffold representing the Euclidean void
- Global absolute time t with step Δt
- No relativistic time dilation applied to the integration clock (proper time is derived only in post-processing)

Void and Noether Sea Terminology (Simulation-Facing)

- “Euclidean void” = the fixed spatial container represented by the chosen coordinate chart / grid indices
- “Noether sea” = coupled pro/anti braids instantiated as objects or response variables in the void

Mandatory $\mathbb{U}_{\text{now}}$ universe-state perspective ($\mathbb{U}_{\text{now}}$) grid

Every run must instantiate $\mathbb{U}_{\text{now}}$ sensors:

- $\mathbb{U}_{\text{now}}$ grid definition: chart points/worldlines, spacing, bounds, boundary conditions
- Logged channels (minimum): Φ , $\nabla\Phi$
- Optional: Noether sea state variables (for example, ρ_{NS} and alignment metrics)
- Provenance tables: receiver_id, t , emitter_id, t_{emit} , contribution_strength when feasible

Causal wake surface bookkeeping requirement

When a potential wake surface intersects a $\mathbb{U}_{\text{now}}$ sensor or contributes to $\Phi(x, t)$, the code must:

- Solve for emission time t_{emit} using $\|x - x_{\text{emitter}}(t_{\text{emit}})\| = c_f(t - t_{\text{emit}})$
- Record emitter identity plus t_{emit} (provenance logging)

Metadata (required)

Each run must store:

- c_f , kernel parameters, Δt , integrator name/order, tolerances
- history-window/compression settings (if any)
- initial conditions seed
- version hash / commit id

Acceptance gate

No major physical claim is accepted without:

- $\mathbb{U}_{\text{now}}$ logs
- Δt convergence
- history-resolution convergence
- cross-integrator comparison (for critical results)

$\mathbb{U}_{\text{now}}$ universe-state perspective Implementation & Grid Protocols

1. **Grid Initialization:** All simulations run on a rigid Cartesian grid chosen as the coordinate scaffold for the **Euclidean void**. The grid is pre-loaded with a lattice of coupled Noether braids to instantiate the Noether sea.
2. **Fiducial Sensor Array:** Instantiate a grid of virtual sensors at fixed chart locations (x, y, z) . Each records Φ and $\nabla\Phi$.
3. **Causal Time Lookup:** When a causal isochron intersects a sensor, the simulator uses the grid history to “look back” to the emitter’s position at t_{history} .
4. **Logging Standard:** All runs must log $\mathbb{U}_{\text{now}}$ channels (Φ , $\nabla\Phi$, provenance tables) to allow cross-run convergence auditing.

$\mathbb{U}_{\text{now}}$ universe-state perspective Grid

- **Grid:** Initialize rigid Cartesian Grid[x] [y] [z] as the chosen chart for the Euclidean void.
- **Sea Initialization:** Pre-load the grid with coupled Noether braids for low-excitation Noether sea runs.
- **Logging:** Record Φ and $\nabla\Phi$ at fixed nodes ($\mathbb{U}_{\text{now}}$ universe-state sensors).
- **Time:** Global step Δt (absolute time).

A0 Branch Certificate Protocol

This protocol defines the simulation-facing handoff for the A_0 reference attractor described in [Particle Masses](#), [Nested Shell Braid Dynamics](#), and [Energy](#). It specializes the general [Simulation Run Protocols](#) to the first neutral rest-branch nested shell braid mass-map target.

The protocol does not treat A_0 as a particle label. It treats A_0 as a calibration-free branch certificate problem: find a finite, stable, multi-scale causal-root ledger before energy, shielding, Noether sea response, or mass comparisons enter.

Master-Equation Handoff Boundary

If a run consumes a master-equation branch-chart object $\mathfrak{B}(\Gamma, \mathcal{S}; h, \eta, \epsilon_c)$, the consumed data must remain branch-certificate data: active roots, inactive gaps, Jacobian floor, memory depth, returned-section residual, section stability, and the refinement schedule that preserves the same branch identity. These fields may support Tier 0 and Tier 1 certification only.

The same packet must keep downstream extraction fields separate. `energy_ledger`, `far_field_shielding`, `medium_response`, and `mass_summary` remain not-computed until their tiers pass. A run fails the handoff if $\zeta(A_0)$, $\mathcal{L}_{\text{aniso}}$, or $\mathcal{M}_{\text{sea}}^{ab}$ changes under root-ledger refinement, inactive-gap refinement, history-window extension, or controlled η refinement while the branch label and quotient row are claimed to be unchanged.

Certificate Packet Schema

An auditable A_0 branch certificate should preserve one top-level packet shape across all tiers. Fields that are not computed at a given tier must remain present with an explicit status, role, and note rather than disappearing from the packet.

Field	Required content	Promotion role
metadata	run identifier, code or derivation version, source commit, integrator, tolerances, η , sampling schedule, and history-window rule	makes the packet reproducible
sea_cell	u_{sea}^i , G_{grad} , n , χ_{sea} , declared c_* , and boundary conditions	fixes the homogeneous Noether sea cell and prevents mixing c_f with c_{eff}
branch_label	layer windings, inter-layer closure integers, handedness, carrier ellipticity, and active root-branch summary	identifies the branch being certified
z_lambda	quotient-coordinate row z_Λ : ε_{IM} , ε_{MO} , T_I/T_M , T_M/T_O , δ_M , layer ellipticities, G_{em} , χ_N , H_I , H_M , H_O , Φ_{rel} , removed gauges $SO(3)$, S_k^1 , Γ_Λ , branch class $[\Lambda]$, and quotient-degeneracy status	records the reduced moduli coordinate rather than an unquotiented carrier representative
branch_chart_revision	conditional pre-rerun record for any revised reduced branch coordinate, including source fields, equality map, equation and coefficient counts, held-out residual rule, phase-origin rule when a phase split is used, symmetry or quotient behavior, locked-key exclusion, benchmark exclusion, and <code>accepted_history_boundary: false</code>	prevents residual-selected coordinates or post-fit added columns from masquerading as branch geometry
state_vector	six architrino labels, polarities, reduced geometry, frequencies, phase offsets, carrier chart, history segment, and center gauge	gives the reduced Noether braid state vector
closure_system	active variables, causal-root equations, layer phase closure, inter-layer closure, center-gauge closure, speed-ordering inequalities, and tolerances	ties closure labels to equations rather than only to names
root_ledger	active and raw partner, self, and inter-layer root classes with delays, branch Jacobians, separator flags, parity events, and excluded near-zero self roots separated	verifies finite causal-root bookkeeping
term_classification	terms assigned to averaging, locking, and leakage channels, with measured or derived residual size	prevents internal corrections from being hidden before promotion
residuals	complete branch-row residual surface $\mathcal{R}_{A_0}$, with $\mathcal{R}_{\text{state}}$, $\mathcal{R}_{\text{root}}$, $\mathcal{R}_{\text{phase}}$, $\mathcal{R}_E$, $\mathcal{R}_{\text{drift}}$, $\mathcal{R}_{\text{speed}}$, $\mathcal{R}_{\text{avg}}$, $\mathcal{R}_{\text{lock}}$, $\mathcal{R}_{\text{leak}}$, and $\mathcal{R}_{\text{Floquet}}$, each with value, tolerance, status, role, and note fields	gives a machine-checkable promotion surface with later-tier omissions explicit
residual_values	numeric mirror of $\mathcal{R}_{A_0}$ values, with Tier 0 omissions recorded as null rather than hidden	gives scripts a stable audit surface without erasing row semantics
Delta_k	Δ_k value, status, role, nonpositive-gap failure code, and note; Tier 0 emits null with <code>not_computed_in_tier0</code>	keeps the Floquet handoff visible before Tier 1 computes the return map
stability	monodromy or finite-difference return map, excluded symmetry modes, non-symmetry Floquet multipliers, and the computed Δ_k once Tier 1 exists	separates integer closure from attractor stability
group_velocity_anisotropy	$\mathbf{V}_{\text{cm}}$, declared c_* , β_* , envelope ratio, forward/backward delay ratio, tensor $\mathcal{A}_{\text{gv}}^{ij}$, refinement status, and whether the entry is rest residue, small-velocity response, or probe-induced drift	keeps motion-induced deformation separate from shielding leakage
energy_ledger	sign-resolved kinetic content, interaction terms, wake/history terms, layer totals E_I , E_M , E_O , $E_{\text{internal}}(A_0)$, and action per closed cycle	supplies the unshielded energy reservoir after Tier 1 passes
far_field_shielding	extraction radii, angular grid, selected wake channel, $\mathcal{L}(\hat{\mathbf{R}})$, naive constituent sum, leading isotropic projection, $\zeta(A_0)$, $\mathcal{L}_{\text{aniso}}$, and convergence status	turns shielding into an extracted far-field quantity after Tier 1 passes
medium_response	acceleration probes, gradient probes, extracted $\mathcal{M}_{\text{sea}}^{ab}$ baseline, symmetric tensor part, antisymmetric residue, and response anisotropy	compatibility field for testing Noether sea inertial and gravitational response after shielding passes
mass_summary	$\zeta(A_0)E_{\text{internal}}(A_0)/E_0$, unresolved constants, response-map assumptions, and explicitly excluded particle benchmarks	records only calibration-free mass-facing output

Field	Required content	Promotion role
certificate_gates	pass/fail/not-computed gates for quotient nondegeneracy, scale separation, speed ordering, phase closure, carrier residuals, root residual, active root-ledger stability, active separator-root handling, near-zero self-root handling, residual semantics, Floquet handoff, and Tier 0 continuation	controls promotion between branch search, attractor, shielding, and response claims
failure_code	reason the row or packet failed, or the next allowed promotion status	prevents failed packets from being read as mass-map results

The residuals field is the complete branch-row surface

$$\mathcal{R}_{A_0} = (\mathcal{R}_{\text{state}}, \mathcal{R}_{\text{root}}, \mathcal{R}_{\text{phase}}, \mathcal{R}_E, \mathcal{R}_{\text{drift}}, \mathcal{R}_{\text{speed}}, \mathcal{R}_{\text{avg}}, \mathcal{R}_{\text{lock}}, \mathcal{R}_{\text{leak}}, \mathcal{R}_{\text{Floquet}})$$

Tier 0 may compute only part of this surface. The row must still emit every component. Missing later-tier components use explicit `not_computed_in_tier0` status, null value, null tolerance when no tolerance exists yet, a promotion role, and a note that names the tier responsible for computing the entry.

The group-velocity anisotropy entry uses the reduced centered covariance of the six-worldline state. With

$$\mathbf{C}_{A_0}(t) = \frac{1}{6} \sum_{a \in A_0} \mathbf{s}_a(t)$$

define

$$D_{A_0}^{ij}(\mathbf{V}_{\text{cm}}) = \left\langle \sum_{a \in A_0} (s_a^i - C_{A_0}^i) (s_a^j - C_{A_0}^j) \right\rangle_{T_k}$$

$$Q_{A_0}^{ij} = \frac{D_{A_0}^{ij}}{h_{mn} D_{A_0}^{mn}}, \quad \mathcal{A}_{\text{gv}}^{ij} = Q_{A_0}^{ij} - \frac{1}{3} h^{ij}$$

This tensor measures motion-induced or probe-induced Noether braid deformation. It is not the same object as the far-field leakage residue $\mathcal{L}_{\text{aniso}}$, which is extracted from cycle-averaged wake coefficients in Tier 2.

Tier 0: Algebraic Branch Search

Tier 0 is a reduced branch-search pass. It samples diagnostic carrier charts, solves delayed root equations on those charts, classifies internal terms, and emits candidate rows. It does not claim a physical attractor.

Required inputs:

- homogeneous Noether sea cell with $u_{\text{sea}}^i = 0$, $G_{\text{grad}} = 0$, $n = 1$, $\chi_{\text{sea}} = 1$, and primitive wake speed c_f ;
- layer labels $\ell \in \{I, M, O\}$ and polarity labels $\sigma \in \{+, -\}$;
- scale ratios $\varepsilon_{IM} = R_I/R_M$ and $\varepsilon_{MO} = R_M/R_O$;
- speed offsets enforcing $s_I > c_f$, $s_M \approx c_f$, and $s_O < c_f$;
- candidate handedness tuple and carrier ellipticity;
- $\eta > 0$, sampling resolution, and history-window rule.

Required outputs:

Output	Meaning
branch_label	layer windings, inter-layer closure integers, handedness, and active root-branch summary
closure_labels	declared T_k , winding integers, inter-layer closure integers, and active root classes
z_lambda	reduced quotient-coordinate row z_Λ , including radius ratios, period ratios, δ_M , layer ellipticities, plane Gram data $G_{\ell m}$, χ_N , handedness labels, phase-offset quotient status, removed gauges, branch class $[\Lambda]$, and quotient_degenerate
state_vector	reduced geometry, frequencies, phase offsets, carrier chart, and center gauge
closure_system	active causal-root, phase-closure, inter-layer-closure, center-gauge, and speed-ordering equations used by the row
root_ledger	active and raw partner, self, and inter-layer root counts with delays, branch Jacobians, separator flags, and excluded near-zero self roots separated
term_classification	terms assigned to averaging, locking, and leakage channels
residuals	every component of $\mathcal{R}_{A_0}$, each with value, tolerance, status, role, and note fields; $\mathcal{R}_E$ and $\mathcal{R}_{\text{Floquet}}$ are explicit Tier 0 omissions unless supplied by a later diagnostic
residual_values	numeric value mirror for the same $\mathcal{R}_{A_0}$ components, with omitted components recorded as null
Delta_k	Δ_k status object; Tier 0 sets value to null and status to <code>not_computed_in_tier0</code> until Tier 1 constructs the monodromy or finite-difference return map
group_velocity_anisotropy	rest-branch residue if computed, or an explicit not-computed Tier 0 status; no Tier 0 row may use this as shielding evidence

Output	Meaning
certificate_gates	pass/fail/not-computed gates for quotient coordinates, scale separation, speed ordering, phase closure, carrier residuals, root residual, active root ledger, active separator roots, near-zero self roots, residual vector semantics, Δ_k , and Tier 0 continuation
failure_code	reason the row failed, or candidate if it survives Tier 0

Tier 0 passes only if at least one row has a finite causal-root ledger, nondegenerate quotient coordinates, retained scale separation, correct speed ordering, bounded carrier residuals, no unclassified separator term, and a complete residual surface. Passing Tier 0 only authorizes Tier 1 continuation.

Tier 0 Failure-Code Enum

The row-level failure_code field is a machine-readable enum. The accepted values are:

Code	Trigger	Promotion consequence
candidate	all Tier 0 promotion gates pass	row may seed Tier 1 continuation only
quotient-degenerate	z_Λ has degenerate plane-normal Gram or orientation data after quotienting global rotations	reject the row as a reduced moduli coordinate
scale-separation-collapse	radius or period ratios violate the declared separated-scale Tier 0 regime	reject the row or widen the scan only as a controlled scale-separation test
speed-order-collapse	$\mathcal{R}_{\text{speed}}$ fails the intended $s_I > c_f$, $s_M \approx c_f$, $s_O < c_f$ ordering	reject the row before attractor continuation
phase-closure-open	$\mathcal{R}_{\text{phase}}$ fails layer winding closure over T_k	reject the row until integer closure is restored
carrier-residual-open	$\mathcal{R}_{\text{state}}$ or $\mathcal{R}_{\text{drift}}$ fails the Tier 0 carrier chart tolerance	reject the row as an unclosed diagnostic carrier
root-residual-open	$\mathcal{R}_{\text{root}}$ fails on candidate active causal-root branches	reject the row until active roots solve within tolerance
averaging-residual-open	$\mathcal{R}_{\text{avg}}$ fails its declared averaging tolerance	keep the term in the branch equations or reject the row
locking-residual-open	$\mathcal{R}_{\text{lock}}$ fails its declared locking tolerance	keep the near-separator or resonance term in Tier 1 or reject the row
separator-singularity-unresolved	active near-separator roots exceed the configured allowance without a locking continuation rule	reject the row until separator handling is explicit
near-zero-self-root-excluded	excluded near-zero self roots exceed the configured allowance under $H(0) = 0$	reject the row until a positive-delay self branch or regularized fold-layer rule exists
root-ledger-instability	the active causal-root ledger is empty or lacks partner, self, or inter-layer classes	reject the row as a finite-ledger failure
nonpositive-floquet-gap	Tier 1 computes $\Delta_k \leq 0$	reject the branch as a non-attractor even if integer closure holds

At Tier 0, nonpositive-floquet-gap appears only as the reserved Delta_k.failure_code_if_nonpositive and certificate_gates.floquet_gap.failure_code, because Tier 0 does not compute Δ_k .

Near-Zero Self Roots

Tier 0 must distinguish raw self-root sightings from active self-hit branches. A self root at the configured near-zero delay threshold is recorded in the raw ledger but excluded from the active ledger as an instantaneous self-kick artifact under the convention $H(0) = 0$.

Such a root may not count as self-hit closure merely because a fold-layer diagnostic preserves the locked self-root keys. The current fold-layer row is a transition candidate only; it promotes after a corrected one-period branch-equation attempt passes the declared residual surface, with Δ_k and η -ladder persistence still downstream.

The reader-facing interpretation of these rows is in [A₀ Tier 0 Result Interpretation](#).

Tier 1: $\eta > 0$ Continuation

Tier 1 promotes a surviving Tier 0 row into direct delayed dynamics with the regularized wake kernel still active. It must preserve the absolute-frame logging standard.

Required checks:

1. direct evolution over at least one declared T_k ;
2. root-ledger stability under Δt and history-window refinement;
3. persistence of averaging, locking, and leakage classifications;
4. no secular center drift after symmetry modes are removed;
5. monodromy or finite-difference return-map estimate with symmetry modes quotiented;
6. positive non-symmetry Floquet gap $\Delta_k > 0$;
7. convergence under the standards in [Convergence Tests](#);

8. a Floquet or monodromy report stating whether the state-dependent delay derivative term was included in the variational operator;
9. `transition_records.json` whenever the run crosses a fold-layer, separator, or active-root status transition.

Branch-Chart Revision Checkpoint

If a Tier 1 diagnostic or corrected carrier attempt reaches a finite-coordinate no-go and proposes a revised branch chart, the revision is admissible only as a pre-rerun record. The proposed reduced coordinate z_Λ^* or finer branch partition μ^* must be declared from branch geometry, causal-root data, quotient-row data, or corrected carrier state before residual fitting. It may not be selected from residual-sign binning, particle benchmarks, fitted weights, or post-fit cancellation.

The pre-rerun record must report `coordinate_source_fields`, `equality_map`, `equation_count`, `coefficient_count`, held-out residual checks, phase-origin checks when a phase split is used, locked-key exclusions, symmetry quotients, benchmark exclusions, and `accepted_history_boundary: false`. The design must remain overdetermined after quotienting, for example by satisfying $N_{\text{eq}} > N_{\text{coef}}$ or $R_{\text{dif}} > 0$, and the same branch identity must survive the refinement checks in [Convergence Tests](#).

Such a row is a revision candidate only. A branch-chart checker may authorize only a new Tier 1 rerun path; it does not accept history. If the checker rejects the packet for a hidden fit split, inadequate degrees of freedom, or held-out residual failure, then the compact-coordinate no-go remains a controlled chart failure. If the checker passes, the branch still requires corrected one-period residuals, quotient-row identity, monodromy or Δ_k , and η -ladder persistence with the same branch identity.

Tier 1 passes only if the same branch remains stable before any $\eta \rightarrow 0^+$ extrapolation.

Corrected One-Period Branch-Equation Boundary

The current fold-layer-locked compact fixture is a controlled negative result, not an accepted attractor and not a broad falsification of the A_0 program. The direct one-period runner can preserve the locked self-root keys in $\mathcal{R}_{\text{lock}}$ without solving the branch equations: state return, root closure, phase closure, speed ordering, center drift, and energy-like speed closure still fail. The scalar branch-native relation basis over B_{self} , B_{partner} , and B_{inter} leaves a relative acceleration residual far above the Tier 1 tolerance, so the next allowed rerun must predeclare either a non-circular carrier correction $\mathbf{d}_\ell(t)$ or a richer branch-native interaction basis before residual fitting.

For a declared period window $W = [t_0, t_0 + T_k]$, the corrected carrier has the form

$$\mathbf{x}_{a,\ell}^*(t) = \mathbf{x}_{a,\ell}^{(0)}(t) + \mathbf{d}_\ell(t), \quad \mathbf{d}_\ell(t + T_k) = \mathbf{d}_\ell(t), \quad \langle \mathbf{d}_\ell \rangle_W = 0$$

The one-period residual is

$$\mathcal{R}_{1\text{per}} = \frac{\left(\int_W \sum_a \left\| \mathbf{a}_a^{\text{ME}}(t; \mathbf{d}) - \sum_{B \in \{B_{\text{self}}, B_{\text{partner}}, B_{\text{inter}}\}} \alpha_B \mathbf{A}_{a,B}(t; \mathbf{d}) \right\|^2 dt \right)^{1/2}}{\left(\int_W \sum_a \left\| \mathbf{a}_a^{\text{ME}}(t; \mathbf{d}) \right\|^2 dt \right)^{1/2} + \varepsilon_0}$$

The rerun may proceed toward monodromy only if

$$\mathcal{R}_{1\text{per}} \leq 0.02$$

with $\mathbf{d}_\ell(t)$, the basis terms $\mathbf{A}_{a,B}$, the coefficient rule for α_B , and any held-out interval declared before fitting. A scalar-basis no-go is therefore a chart or basis failure; it does not become an attractor failure unless every admissible corrected carrier and branch-native basis inside the declared search class fails the same residual boundary.

Tier 2: Energy and Shielding

Tier 2 begins only after Tier 1 passes. It computes the internal-energy ledger and far-field shielding extraction described in [Energy](#). The required outputs are:

- E_I , E_M , E_O , and $E_{\text{internal}}(A_0)$;
- interaction and wake/history bookkeeping with no double counting;
- far-field wake coefficients $\mathcal{L}(\hat{\mathbf{R}})$ over extraction radii and angular grids;
- the naive constituent sum $\mathcal{L}_{\text{naive}}$ and the leading isotropic projection $\Pi_0 \mathcal{L}$;
- $\zeta(A_0)$ from the leading isotropic projection;
- anisotropic leakage $\mathcal{L}_{\text{aniso}} = (1 - \Pi_0) \mathcal{L}$ retained as a separate tensor or channel list;
- convergence status under extraction radius, angular resolution, Δt , history-window, and η refinement.

Tier 2 fails if particle masses, charged-lepton ratios, electron radius, or measured α enter as inputs.

Tier 3: Medium-Response Probe

Tier 3 begins only after Tier 2 passes. It applies small acceleration and gradient probes to the accepted branch and extracts the homogeneous baseline for $\mathcal{M}_{\text{sea}}^{ab}$. The probe must report whether the acceleration and gradient channels share the same shielded-energy coefficient to first

order, and it must report response anisotropy separately from both $\mathcal{A}_{\text{gv}}^{ij}$ and $\mathcal{L}_{\text{aniso}}$.

Runtime Artifact

The first reduced Tier 0 artifact is `scripts/mass-map/a0-tier0-branch-search.mjs`, with default grid `scripts/mass-map/a0-tier0-default-grid.json`. It is an algebraic branch-search scaffold, not a production simulator. Its required role is to emit candidate rows with parameter choices, quotient-coordinate rows, carrier diagnostics, root ledgers, term classifications, residual surfaces, Δ_k handoff status, leakage placeholders, certificate gates, and failure codes matching this protocol.

The Tier 1 handoff scaffold is `scripts/mass-map/a0-tier1-continuation-scaffold.mjs`. It consumes Tier 0 JSON rows and emits the $\eta > 0$ continuation contract, reduced-coordinate chart, symmetry-quotiented monodromy plan, Δ_k acceptance boundary, and required output artifact list. It is not a delayed-dynamics solver and cannot certify the branch without a later Tier 1 run.

The companion audit is `scripts/audit-a0-mass-map-promotion.mjs`. It scans `AAA` and mass-map priority prose for premature statements that treat $\zeta(A_0)$, $E_{\text{internal}}(A_0)$, or $\mathcal{M}_{\text{sea}}^{ab}$ as accepted before the Tier gates pass.

Acceptance Boundary

The A_0 branch is not an attractor until Tier 1 passes. It is not a mass-map result until Tier 2 passes. It is not an inertial-response result until Tier 3 passes. A reported group-velocity anisotropy tensor is a deformation diagnostic, not a shielding extraction and not a substitute for the Noether sea response probe.

Synthetic Observables

This note defines the canonical logging standard for the virtual $\mathbb{U}_{\text{now}}$ perspective and explains how those logs are turned into detector-like synthetic observables. Its purpose is to keep the separation clear between exact simulation bookkeeping and the post-processed quantities that stand in for what a physical observer would measure.

The file therefore serves as both a data-contract note and an observer-interface note for the simulation stack.

$\mathbb{U}_{\text{now}}$ Logging Standard

Purpose

Define a canonical $\mathbb{U}_{\text{now}}$ universe-state perspective ($\mathbb{U}_{\text{now}}$) used across all simulation tiers for logging, diagnostics, and synthetic datasets. $\mathbb{U}_{\text{now}}$ is not physically realizable; it is a bookkeeping operator acting on the full microstate.

Definition

A $\mathbb{U}_{\text{now}}$ is defined by:

- Fixed Euclidean sample points or worldlines $P = \{x_k\}$ in a declared coordinate scaffold on $\mathbb{R}^3$
- Access to the full state $S(t) = \{(x_i(t), v_i(t), q_i, \dots)\}$ for all architrinos
- Output channels:
 - Local potential $\Phi(x_k, t)$
 - Local gradient $\nabla\Phi(x_k, t)$ (force proxy)
 - Optional local Noether sea state variables (e.g., ρ_{NS} , alignment/orientation metrics)
 - Causal wake surface provenance/event tags: for each received contribution at (x_k, t) , record emitter_id together with t_{emit} , satisfying $\|x_k - x_{\text{emitter}}(t_{\text{emit}})\| = c_f(t - t_{\text{emit}})$
 - Photon packet provenance when a radiation channel is declared: source event, path segment, before/after frequency, recoil or medium-energy exchange, remnant row, and signed exchange residual
 - Optional finite-window operator diagnostics for declared reconstructed channels $\mathbf{Y}_\eta$, including Gauss, Stokes, and wake-surface normalization residuals

Minimal synthetic products

- Time series: $\Phi(t)$, $\nabla\Phi(t)$ at fixed points (“stationary detectors”)
- Snapshot field maps: $\Phi(x, t_0)$, $\nabla\Phi(x, t_0)$ over grids at fixed t_0
- Provenance tables: receiver_id, t , emitter_id, t_{emit} , contribution_strength
- Propagation diagnostics: arrival-time distributions, dispersion tests, effective c_{eff} estimates
- Coarse kinetic moments when a continuum reduction is claimed: density, current, momentum-current tensor, energy-flux vector, and memory-current residuals derived from the same event-root records
- Stochastic summaries when a noise model is claimed: drift vector, diffusion tensor, first two distribution moments, and direct ensemble comparison against event-root histories
- Reaction-diffusion probes when pattern or front language is claimed: front speed, unstable-mode band, selected wavelength, and conservation or source ledger for each reaction term

- Jet/outflow source products when a collimated release or working surface is claimed: beam radius, head radius, bow-shock speed, Mach number, jet-to-ambient density ratio, knot spacing, cooling ratio, synthetic line map, synthetic synchrotron map, inverse-Compton map, polarization fraction, and polarization angle
- Cosmology-facing photon products when redshift is inferred: total Z_X , endpoint/source/launch/path decomposition, signed path-frequency exchange $Y_{X,\text{path}}$, packet-cadence stretch, flux factors, and image-sharpness diagnostics

Mapping: $\mathbb{U}_{\text{now}}$ data $\rightarrow$ Physical observables

Synthetic observables must be generated by post-processing $\mathbb{U}_{\text{now}}$ logs with a model of a *physical* observer (assembly clock/detector):

1. Extract local Noether sea state along detector worldline $X_{\text{det}}(t)$
2. Compute derived detector clock time τ_{det} via the declared clock map $d\tau = F(\text{Noether sea state}, v_{\text{det}}, \Phi, \nabla\Phi, \dots) dt$ from [Proper Time and Time Dilation](#)
3. Generate detector-like outputs:
 - clock readings $\tau(t)$
 - photon arrival times and frequency shifts, with signed exchange rows separated from endpoint cadence and launch geometry
 - inferred “geodesics” (effective paths) from travel-time minimization through the Noether sea effective signal speed c_{eff}

Validation checks (must pass)

- **Causality residual (per record m):**

$$\rho_m \equiv \frac{||x_k - x_{i_m}(t_{\text{emit},m})|| - c_f(t_m - t_{\text{emit},m})}{\max(c_f \Delta t, \varepsilon_r)}, \quad \varepsilon_r = 10^{-12}$$

Pass if at least 99.9% of records satisfy $\rho_m \leq 10^{-2}$ and $\max_m \rho_m \leq 5 \times 10^{-2}$.

- **Temporal ordering check:**

$$\theta_m \equiv \frac{t_{\text{emit},m} - t_m}{\Delta t}$$

Pass if fraction with $\theta_m > 10^{-9}$ is $\leq 10^{-6}$.

- **Cross-integrator invariance:**

For any channel Y use

$$E_{\text{rel}}(Y; A, B) \equiv \frac{\|R(Y_B) - Y_A\|_{L^2}}{\|R(Y_B)\|_{L^2} + 10^{-12}}$$

Pass if $E_{\text{rel}}(\Phi) \leq 0.03$ and

$E_{\text{rel}}(\|\nabla\Phi\|) \leq 0.05$.

- **Finite-window Gauss/Stokes residuals:** for any declared reconstructed vector channel $\mathbf{Y}_\eta$ on Σ_t , use

$$R_G[V, t; \mathbf{Y}_\eta] \equiv \frac{|\int_{\partial V} \mathbf{Y}_\eta \cdot \hat{\mathbf{n}} dS - \int_V \nabla \cdot \mathbf{Y}_\eta dV|}{\int_{\partial V} |\mathbf{Y}_\eta \cdot \hat{\mathbf{n}}| dS + \int_V |\nabla \cdot \mathbf{Y}_\eta| dV + 10^{-12}}$$

and

$$R_S[S, t; \mathbf{Y}_\eta] \equiv \frac{|\oint_{\partial S} \mathbf{Y}_\eta \cdot d\mathbf{x} - \int_S (\nabla \times \mathbf{Y}_\eta) \cdot \hat{\mathbf{n}} dS|}{\oint_{\partial S} |\mathbf{Y}_\eta \cdot d\mathbf{x}| + \int_S |(\nabla \times \mathbf{Y}_\eta) \cdot \hat{\mathbf{n}}| dS + 10^{-12}}$$

Pass if both residuals are $\leq 2 \times 10^{-2}$ on resolved windows and decrease under spatial refinement. These are diagnostics on reconstructed continuum channels, not claims that the channel is substrate ontology.

- **Distributional wake-surface normalization:** for emitted wake surface m with source strength q_m , delay $\tau_m = t - t_{\text{emit},m}$, and radial annulus $R_- \leq r_m \leq R_+$ around the emission point, use

$$Q_{m,\eta}^{\text{ann}} = q_m H(\tau_m) \int_{R_-}^{R_+} \delta_\eta(r_m - c_f \tau_m) dr_m$$

and

$$R_{N,m} \equiv \frac{\left| \int_{R_- \leq r_m \leq R_+} \rho_{m,\eta}(t, \mathbf{x}) dV - Q_{m,\eta}^{\text{ann}} \right|}{|q_m| + 10^{-12}}$$

Pass if at least 99.9% of emitted wake surfaces satisfy $R_{N,m} \leq 10^{-2}$ and the maximum resolved-window residual is $\leq 5 \times 10^{-2}$.

- **Photon-frequency exchange closure:** when a photon packet changes frequency during transport, the logged before/after frequencies must close with medium, recoil, and remnant rows:

$$R_{\nu\text{-ex},m} = \frac{|h(\nu_m^+ - \nu_m^-) + \Delta E_{\text{med},m} + \Delta E_{\text{recoil},m} + \Delta E_{\text{rem},m}|}{\epsilon E}$$

A cosmology-facing redshift or blueshift product may consume this row only after the residual is reported with the same photon provenance used for arrival-time, flux, and image-sharpness outputs.

- **Operator consistency across PDE and event-root runs:** after resampling the event-root reconstruction onto the PDE grid, define

$$\Delta \mathbf{Y}_\eta \equiv \mathbf{Y}_\eta^{\text{PDE}} - R(\mathbf{Y}_\eta^{\text{root}}), \quad E_{\text{op}}(V, S, t) \equiv \max\{R_G[V, t; \Delta \mathbf{Y}_\eta], R_S[S, t; \Delta \mathbf{Y}_\eta]\}$$

Pass if $E_{\text{op}} \leq 0.03$ on the declared validation windows and decreases under temporal/history/spatial refinement.

- **Curvilinear-coordinate hygiene:** finite-window residuals must use the coordinate weights and operator formulas of the declared Euclidean scaffold. In spherical coordinates (r, θ, φ) ,

$$w_V = r^2 \sin \theta \Delta r \Delta \theta \Delta \varphi, \quad w_{S_R} = R^2 \sin \theta \Delta \theta \Delta \varphi$$

and a radial channel must evaluate $\nabla \cdot (F_r \hat{\mathbf{r}})$ as

$$\frac{1}{r^2} \frac{\partial}{\partial r} (r^2 F_r)$$

not as $\partial_r F_r$. Fail the run if the coordinate scaffold does not declare these weights.

- **Provenance distribution agreement:** for t_{emit} distributions, require

$$D_W \equiv \frac{W_1(P_A, P_B)}{\text{IQR}(P_B) + 10^{-12}} \leq 0.08, \quad D_{JS} \equiv \text{JSD}(P_A \| P_B) \leq 0.03$$

- **Kinetic-moment closure:** for any promoted continuum observable, compute the direct event-root moments

$$\rho_{\text{dir}}, \quad \mathbf{j}_{\text{dir}}, \quad \Pi_{\text{dir}}^{ij}, \quad \mathbf{J}_{e,\text{dir}}$$

and compare them with the reduced continuum reconstruction. The reduced channel must report

$$R_{\text{mom}} = \max_Y \frac{\|Y_{\text{cg}} - R(Y_{\text{dir}})\|_{L^2(W)}}{\|R(Y_{\text{dir}})\|_{L^2(W)} + 10^{-12}}$$

where Y ranges over the retained density, current, momentum-current, and energy-flux channels. Pass if $R_{\text{mom}} \leq 0.05$ and the omitted memory-current residual decreases under refinement.

- **Drift-diffusion reconstruction:** if a Fokker-Planck or Langevin surrogate is emitted, estimate drift and diffusion from increments,

$$u^a(z) = \lim_{\Delta t \rightarrow 0} \frac{\langle \Delta z^a \rangle_z}{\Delta t}, \quad D^{ab}(z) = \lim_{\Delta t \rightarrow 0} \frac{\langle \Delta z^a \Delta z^b \rangle_z}{2\Delta t}$$

The synthetic distribution must match direct event-root ensembles in $\langle z \rangle$ and $\text{Cov}(z)$ before higher stochastic claims are trusted. Higher cumulants may differ from the surrogate unless a separate closure row has been declared.

- **Reaction-diffusion and pattern probes:** when a reduced scalar or multi-channel field y obeys

$$\partial_t y^a = D^{ab} \Delta y_b + F^a(y) + R_{\text{rd}}^a$$

the packet must report the fixed points, the linearized growth matrix, the unstable wavenumber band if one exists, and the front-speed estimate if a traveling-front claim is made. For two-channel pattern claims, the Turing-style gate is that the homogeneous fixed point is stable before diffusion and that the diffusion-shifted linear operator has a declared finite unstable band. Without those rows, visual pattern formation is not a validated synthetic observable.

- **Jet/outflow head and radiation probes:** when a simulation claims an astrophysical jet, outflow, knot chain, or working surface, the synthetic packet must compare the logged event-root dynamics to the observer-level jet-head and radiation benchmarks. For a supersonic head with jet speed v_j , beam radius R_j , head radius R_h , density ratio $\eta_j = \rho_j / \rho_a$, and $a_h = (R_j / R_h)^2$, the bow-shock speed target is

$$v_{\text{bs, std}} = v_j \left[1 + (\eta_j a_h)^{-1/2} \right]^{-1}$$

The head residual is

$$R_{\text{head}} = \left| \frac{v_{\text{bs, map}}}{v_{\text{bs, std}}} - 1 \right|$$

For radiative shocks, also report

$$\mathcal{R}_{\text{cool}} = \frac{t_{\text{cool}}}{t_{\text{dyn}}}, \quad t_{\text{dyn}} \sim \frac{\ell_j}{v_j}$$

and route the synthetic emission to thermal line/free-free rows when $\mathcal{R}_{\text{cool}} \ll 1$, or to adiabatic, particle-acceleration, synchrotron, inverse-Compton, cocoon, and lobe rows when $\mathcal{R}_{\text{cool}} \gg 1$. If a pulsed inlet is declared with period P , the knot spacing should report

$$R_{\text{knot}} = \left| \frac{\Delta x_{\text{knot}}}{v_j P} - 1 \right|$$

up to projection, cooling, and deceleration corrections named in the packet. For synchrotron-bearing jets, the same run must emit I_{ν}^{syn} , I_{ν}^{IC} , Π_{ν} , and ψ_{ν} maps from one effective $B_{\text{eff}} \leftrightarrow \mathcal{V}_{\text{NS}}$ reconstruction. These are observer-level comparison variables; passing them does not promote MHD fields into substrate ontology.

- **Convergence triad:** must pass temporal/history/spatial gates from [convergence-tests.md](#), including null-test failure.

Failure mode

If any of the quantitative checks above fail (or if the null test does not fail), treat $\mathbb{U}_{\text{now}}$ outputs as numerically unreliable for self-hit claims until thresholds are met.

$\mathbb{U}_{\text{now}}$ as Standard Probe

1. **Definition:** A $\mathbb{U}_{\text{now}}$ universe-state perspective ($\mathbb{U}_{\text{now}}$) is a non-physical bookkeeping operator acting on the full microstate $S(t)$.
2. **Synthetic Observables:**
 - **Raw Data:** Time series of $\Phi(x, t)$ at fixed points.
 - **Post-Processing:** To simulate a physical detector, we act on the raw data by integrating the derived clock time τ of a “clock assembly” moving through the $\mathbb{U}_{\text{now}}$ grid.
3. **Separation of Concerns:** This explicitly separates Ontology (simulation state/ $\mathbb{U}_{\text{now}}$ data) from Phenomenology (synthetic detector data).

Virtual Sensor & Data Extraction

- **Virtual Sensor:** Implementation of the $\mathbb{U}_{\text{now}}$ universe-state perspective. Samples potential/gradient at fixed coordinates.
- **Post-Processing:** Convert Virtual Sensor data (Ground Truth) into Physical Observer data (what a moving clock measures).
- **Provenance:** Track emitter ID and emission time for every potential contribution at a grid point.

Branch / Quantum

A0 Tier 0 Result Interpretation

This note explains how to read the first reduced A_0 branch-search artifact. It is a companion to the [A0 Branch Certificate Protocol](#), the general [Simulation Run Protocols](#), and the convergence standards in [Convergence Tests](#).

The Tier 0 artifact is not an attractor proof. It is a certificate-facing filter that decides whether a reduced carrier chart is disciplined enough to seed Tier 1 $\eta > 0$ continuation. Its output should be read together with the mass thesis in [Particle Masses](#), the energy ledger definitions in [Energy](#), the dynamics baseline in [Nested Shell Braid Dynamics](#), and the closure bookkeeping in [Parameter Ledger](#).

Output Status

The runtime artifact is `scripts/mass-map/a0-tier0-branch-search.mjs`. It emits rows with six separate layers of interpretation:

Output layer	Meaning	Promotion role
<code>z_lambda</code>	Quotient-coordinate row z_λ after removing global rotations, the common closed-cycle phase gauge, and allowed branch-preserving chart relabelings	Decides whether the row can be read as a reduced moduli coordinate rather than a raw carrier representative
<code>root_ledger</code>	Active and raw causal-root counts by source relation, with excluded instantaneous self-root counts separated from active roots	Decides whether the carrier chart has a finite active partner, self, and inter-layer ledger

Output layer	Meaning	Promotion role
residuals and residual_values	The complete $\mathcal{R}_{A_0}$ row surface, plus a numeric mirror where Tier 0 omissions remain null	Prevents a numerical value, a diagnostic placeholder, and a later-tier obligation from being confused
Delta_k	Δ_k handoff object with null value and not_computed_in_tier0 status until Tier 1 builds the return map	Keeps Floquet stability from being silently omitted or treated as a Tier 0 result
certificate_gates	Pass/fail/not-computed status for the Tier 0 promotion checks	Decides whether the row may seed Tier 1 continuation
failure_code	One machine-readable row code, or candidate when the row survives Tier 0	Gives scripts and readers the same rejection reason

A row with status `tier0_continuation_ready` may seed Tier 1. A row with status `tier0_rejected` does not seed Tier 1 until the failing gate is resolved. Neither status accepts an attractor, computes $\zeta(A_0)$, validates $E_{\text{internal}}(A_0)$, or derives $\mathcal{M}_{\text{sea}}^{ab}$.

The same boundary applies when a compact finite-coordinate chart or coarse branch split fails. Such a failure means the proposed reduced coordinate did not earn a continuation run; it does not by itself falsify the broader A_0 branch program. A branch-chart checker can authorize only a new Tier 1 rerun path after the coordinate source, equality map, fit degrees of freedom, held-out residuals, phase-origin handling when relevant, and benchmark exclusions are declared before fitting. It does not create accepted history, and it does not convert Tier 0 readiness into an attractor claim.

Quotient-Coordinate Row

The emitted `z_lambda` object is the row-level representation of z_Λ . It records the reduced coordinate after quotienting away global rotations, the common S_k^1 phase gauge, and allowed discrete relabelings Γ_Λ that preserve polarity assignment, layer roles, speed ordering, and causal-root branch class.

z_lambda entry	Row semantics
schema	version marker for the quotient-coordinate row
radius_ratios	ϵ_{IM} and ϵ_{MO}
period_ratios	T_I/T_M and T_M/T_O , so time-scale separation is checked alongside radius separation
delta_M	middle-layer speed offset $(s_M - c_f)/c_f$
ellipticity and ellipticity_status	layer ellipticity data and whether Tier 0 used a shared scalar chart
plane_gram	G_{tm} values for the quotient-reduced binary-plane normals
orientation_class	χ_N , the triple product, and a nondegenerate or degenerate status
handedness	H_I, H_M, H_O layer handedness labels
phase_offset_quotient	Φ_{rel} status after removing the common S_k^1 phase origin; Tier 0 currently emits a gauge-fixed zero-offset representative and marks the quotient basis not_computed_in_tier0
branch_class and branch_class_status	$[\Lambda]$ data from winding integers, inter-layer closure, active and raw root classes, and excluded roots; Tier 0 marks the representative as not yet a canonical discrete quotient
removed_gauges	declared gauge removals: $SO(3)$, S_k^1 , and Γ_Λ
quotient_degenerate	Boolean failure surface for quotient-degenerate

The quotient row is not a new dynamical assumption. It is the coordinate audit that prevents a raw carrier chart, a gauge choice, and a branch class from being mistaken for three independent pieces of physics.

Near-Zero Self-Root Policy

The Tier 0 scanner distinguishes raw self-root sightings from active self-hit branches. A raw self root whose delay lies at the configured near-zero threshold is recorded but excluded from the active ledger as `excluded_instantaneous_self_kick`.

This policy follows the canonical convention $H(0) = 0$: an instantaneous self-kick is not an active causal hit. The exclusion is conservative. It does not prove that no nearby regularized fold-layer branch exists; it says only that the diagnostic carrier has not yet supplied a positive-delay self-root branch that can be promoted.

The current fold-layer diagnostic can preserve locked self-root keys as a transition candidate, but it does not by itself accept self-hit closure. A fold-layer row promotes only after a corrected one-period branch-equation attempt passes the declared residual surface; until then, Δ_k and η -ladder persistence remain downstream obligations.

Residual Semantics

The emitted `residuals` object is the complete branch-row residual surface

$$\mathcal{R}_{A_0} = (\mathcal{R}_{\text{state}}, \mathcal{R}_{\text{root}}, \mathcal{R}_{\text{phase}}, \mathcal{R}_E, \mathcal{R}_{\text{drift}}, \mathcal{R}_{\text{speed}}, \mathcal{R}_{\text{avg}}, \mathcal{R}_{\text{lock}}, \mathcal{R}_{\text{leak}}, \mathcal{R}_{\text{Floquet}})$$

Each entry carries value, tolerance, status, role, and note fields. The companion `residual_values` object mirrors only the values; omitted Tier 0 components remain null rather than disappearing.

The Tier 0 residual surface deliberately includes entries that are not computed at Tier 0:

Residual	Emitter key	Tier 0 interpretation
$\mathcal{R}_{\text{state}}$	state	Carrier-chart return mismatch over one declared period
$\mathcal{R}_{\text{root}}$	root	Active root defect on candidate causal-root branches
$\mathcal{R}_{\text{phase}}$	phase	Integer layer-winding mismatch
$\mathcal{R}_E$	energy	Not computed at Tier 0; Tier 1 or Tier 2 must supply a regularized energy/history functional
$\mathcal{R}_{\text{drift}}$	drift	Centering check for the diagnostic chart; Tier 1 must retest under direct delayed dynamics
$\mathcal{R}_{\text{speed}}$	speed	Sign-aware violation of the intended $s_I > c_f$, $s_M \approx c_f$, $s_O < c_f$ ordering
$\mathcal{R}_{\text{avg}}$	avg	Diagnostic size of terms claimed to average out
$\mathcal{R}_{\text{lock}}$	lock	Diagnostic fraction or defect of selected locking terms
$\mathcal{R}_{\text{leak}}$	leak	Far-field leakage placeholder, not a shielding extraction
$\mathcal{R}_{\text{Floquet}}$	Floquet	Not computed at Tier 0; Tier 1 must construct the monodromy diagnostic

This makes the residual vector complete as an audit surface without pretending that Tier 0 has done Tier 1 or Tier 2 work.

Floquet Handoff

The `Delta_k` object is the Tier 0 handoff for Δ_k . Tier 0 does not construct the monodromy operator, so the emitted value is null, the status is `not_computed_in_tier0`, and the role is `tier1_required`. The reserved failure code is `nonpositive-floquet-gap`, which applies only after Tier 1 computes $\Delta_k \leq 0$.

The same handoff appears in `certificate_gates.floquet_gap` with status `not_computed_in_tier0`. This is a positive omission rule: Tier 0 must show that Floquet stability remains open, not leave the field absent.

Certificate Gates and Failure Codes

The Tier 0 `certificate_gates` object names the promotion checks directly:

Gate	Meaning
<code>quotient_coordinates</code>	z_A must be nondegenerate after global rotations are removed
<code>scale_separation</code>	radius and period ratios must remain inside the declared separated-scale regime
<code>speed_ordering</code>	$s_I > c_f$, $s_M \approx c_f$, and $s_O < c_f$ must hold within tolerance
<code>phase_closure</code>	layer winding closure over T_k must hold
<code>carrier_residuals</code>	state return and center drift residuals must remain bounded
<code>root_residual</code>	active causal-root defects must remain within tolerance
<code>active_root_ledger</code>	partner, self, and inter-layer active root classes must all be present
<code>active_separator_roots</code>	active near-separator roots must have an explicit continuation rule or remain below allowance
<code>near_zero_self_roots</code>	near-zero self roots remain excluded under $H(0) = 0$ and may not count as active self hits
<code>residual_vector_semantics</code>	every residual component must carry value, tolerance, status, role, and note fields
<code>floquet_gap</code>	Δ_k is not computed at Tier 0 and must be computed in Tier 1
<code>tier0_continuation</code>	only rows whose row-level code is candidate may seed Tier 1

The row-level `failure_code` enum preserves the existing Tier 0 codes and reserves the new quotient and Floquet codes:

Code	Meaning
<code>candidate</code>	the row survives Tier 0 and may seed Tier 1 only
<code>quotient-degenerate</code>	the quotient-coordinate row is degenerate after gauge removal
<code>scale-separation-collapse</code>	radius or period ratios collapse the declared separated-scale regime
<code>speed-order-collapse</code>	sign-aware speed ordering fails
<code>phase-closure-open</code>	integer layer-winding closure fails
<code>carrier-residual-open</code>	carrier return or drift residuals fail
<code>root-residual-open</code>	active causal-root residuals fail

Code	Meaning
averaging-residual-open	terms claimed to average out exceed their declared tolerance
locking-residual-open	selected locking terms exceed their declared tolerance
separator-singularity-unresolved	active near-separator roots lack an accepted handling rule
near-zero-self-root-excluded	excluded instantaneous self roots block Tier 0 promotion
root-ledger-instability	the active root ledger is empty or lacks partner, self, or inter-layer classes
nonpositive-floquet-gap	Tier 1 computes $\Delta_k \leq 0$

Promotion Boundary

Tier 0 can only answer a finite branch-search question: does this reduced carrier chart have an active root ledger, controlled chart residuals, and no unresolved near-zero self-root obstruction?

It cannot answer the attractor question, because that requires Tier 1 direct delayed dynamics and a positive non-symmetry Floquet gap $\Delta_k > 0$. It cannot answer the mass-map question, because that requires Tier 2 energy and shielding extraction. It cannot answer the inertial-response question, because that requires Tier 3 acceleration and gradient probes for $\mathcal{M}_{\text{sea}}^{ab}$.

The safe reading is therefore:

Tier 0 pass $\implies$ eligible for Tier 1 continuation

not

Tier 0 pass $\implies$ accepted A_0 attractor

This boundary is the main protection against premature mass-map promotion.

Bell-Family Record-Measure Harness

This protocol gives the Bell-family residuals in [No-Go Theorems](#) their first executable scaffold. It is not a closure proof. It is a probability-table harness that checks whether a proposed record table preserves the standard benchmark shape before any claim is made about deriving that table from architrino dynamics, pair provenance, detector kernels, and finite-time basin measures.

The immediate target is discipline. A model that fits one Bell average can still fail GHZ parity, Hardy zero/positive-event structure, no-signaling, or measurement independence. The harness therefore evaluates CHSH, GHZ, Hardy, no-signaling, measurement-independence, and observed factorization residuals in one packet.

Runtime Artifact

Run:

```
node scripts/quantum/bell-family-residual-harness.mjs --pretty
```

To inspect one case:

```
node scripts/quantum/bell-family-residual-harness.mjs --scenario ghz_local_value_table --pretty
```

To inspect the candidate-fixture intake path:

```
node scripts/quantum/bell-family-residual-harness.mjs \
  --candidate scripts/quantum/product-screened-axis-candidate.json \
  --pretty
```

The script emits JSON with one row per scenario:

Field	Meaning
metadata.source	whether the run used built-in scenarios or a candidate JSON fixture
metadata.candidate_path	candidate fixture path when metadata.source is candidate
id	stable scenario identifier
classification	benchmark or negative_control
source_protocol	declared source construction for candidate fixtures, when supplied
source_record_count	number of retained source records in a candidate fixture

Field	Meaning
metrics.chsh	CHSH expectations, S , local-bound excess, and Tsirelson excess
metrics.ghz	GHZ product-context expectations and Δ_{GHZ} residual
metrics.hardy	Hardy zero-term probabilities and positive-event margin
metrics.no_signaling	maximum one-party marginal drift under remote setting changes
metrics.measurement_independence	total-variation drift of declared provenance labels across settings
metrics.observed_factorization	total-variation distance between the observed joint table and the product of its observed marginals
metrics.product_screening	total-variation distance between the emitted table and a declared Bell-local product-screening reconstruction
gates	pass/fail records for the residuals that apply to the scenario
witness_tags	non-failure tags such as <code>bell.chsh_local_bound_violated</code>
failure_codes	stable failure codes such as <code>bell.signal_transfer</code>

Residual Object

For a two-party CHSH table with binary outcomes $a, b \in \{-1, +1\}$, the harness computes

$$E(x, y) = \sum_{a, b = \pm 1} ab P(a, b | x, y)$$

and the convention

$$S = E(A_0, B_0) - E(A_0, B_1) + E(A_1, B_0) + E(A_1, B_1)$$

The gate reports both the local-bound excess

$$\Delta_{\text{CHSH}} = [|S| - 2]_+$$

and the Tsirelson excess

$$\Delta_{\text{Ts}} = [|S| - 2\sqrt{2}]_+$$

For GHZ, the script uses the context signs in [Bell's Theorem](#):

$$\mathcal{C}_{\text{GHZ}} = \{XXX, XYY, YXY, YYX\}, \quad \prod_{C \in \mathcal{C}_{\text{GHZ}}} \chi_C = -1$$

and computes

$$\Delta_{\text{GHZ}} = \max_{C \in \mathcal{C}_{\text{GHZ}}} [1 - \chi_C E(C)]_+$$

For Hardy, it computes the positive margin

$$\Delta_{\text{Hardy}} = [P(D_1 = 1, D_2 = 1) - P(U_1 = 1, U_2 = 1) - P(D_1 = 1, U_2 = 0) - P(U_1 = 0, D_2 = 1)]_+$$

No-signaling is evaluated as the maximum one-party marginal drift between contexts that keep that party's setting fixed:

$$\Delta_{\text{NS}}^i = \sup_{s_i, \mathbf{s}_{-i}, \mathbf{s}'_{-i}} \sum_{r_i} |P(r_i | s_i, \mathbf{s}_{-i}) - P(r_i | s_i, \mathbf{s}'_{-i})|$$

Measurement-independence leakage is represented by a declared provenance label distribution in each context:

$$\Delta_{\text{MI}} = \sup_{\mathbf{s}} D_{\text{TV}}(\rho_{\text{prov}}(\Pi | \mathbf{s}), \rho_{\text{prov}}(\Pi | \mathbf{s}_0))$$

where $\mathbf{s}_0$ is the packet baseline. A real closure packet should replace this toy provenance distribution with the pair-provenance ledger described below.

For generated pair-provenance cases, the harness also checks whether the emitted table is exactly reconstructed by a Bell-local product-screening form:

$$\Delta_{\text{screen}} = \sup_{\mathbf{s}} D_{\text{TV}} \left(P_{\theta}(\mathbf{r} | \mathbf{s}), \int_{\Pi} \prod_i K_i(r_i | s_i, \Pi) d\rho_{\text{prov}}(\Pi) \right)$$

Here $\Delta_{\text{screen}} = 0$ is not a success for Bell closure. It means the proposed table has collapsed back into the screened common-cause model excluded by the Bell-family gate. A closure candidate must avoid that collapse while still keeping Δ_{MI} and Δ_{NS} within tolerance.

Generated Pair-Provenance Path

The first generated path is a deliberately failing local-axis model. It declares a finite pair-provenance grid

$$\Pi_{AB}^{(N)} = \{(\phi_k, \phi_k + \pi, w_k)\}_{k=1}^N, \quad w_k = \frac{1}{N}$$

and two local deterministic apparatus kernels:

$$K_A(a|A_i, \Pi_k) = \mathbf{1}[a = \text{sgn} \cos(A_i - \phi_k)]$$

$$K_B(b|B_j, \Pi_k) = \mathbf{1}[b = \text{sgn} \cos(B_j - \phi_k - \pi)]$$

The generated table is then

$$P_{\text{gen}}(a, b|A_i, B_j) = \sum_k w_k K_A(a|A_i, \Pi_k) K_B(b|B_j, \Pi_k)$$

This is a useful negative control because it has explicit pair provenance, explicit local kernels, clean no-signaling, and clean measurement independence, but it still reaches only the classical-axis correlation. The product-screening residual is zero by construction, so the `product_screening_escape` gate must fail with `bell.product_screening_collapse`.

The candidate-reader path now makes that obstruction inspectable from a declared source-record fixture rather than only from built-in tables. The current fixture `scripts/quantum/product-screened-axis-candidate.json` supplies eight explicit source records, local deterministic response tables, normalized source weights, and four CHSH contexts. It is not a positive Bell candidate. It is a compact negative control showing that explicit provenance can still reduce to Bell-local product screening unless the completed record law supplies a stronger joint record-basin measure.

Built-In Scenarios

Scenario	Role	Expected signal
<code>chsh_quantum_singlet</code>	benchmark	$ S = 2\sqrt{2}$, no-signaling passes, measurement independence passes
<code>local_classical_axis</code>	negative control	classical-axis response reaches only the local CHSH bound
<code>separable_pair_measure</code>	negative control	independent outcomes produce no Bell-family structure
<code>generated_pair_provenance_screened_axis</code>	negative control	generated pair provenance and local kernels collapse to Bell-local product screening
<code>setting_dependent_provenance</code>	negative control	CHSH table is present, but $\Delta_{\text{MI}} > 0$
<code>signaling_box</code>	negative control	one-party marginals change under remote setting changes
<code>ghz_product_benchmark</code>	benchmark	GHZ product signs match with $\Delta_{\text{GHZ}} = 0$
<code>ghz_local_value_table</code>	negative control	context-independent local values fail GHZ parity
<code>hardy_no_signaling_margin</code>	benchmark	Hardy margin is positive while no-signaling passes
<code>hardy_local_forbidden_event</code>	negative control	the positive Hardy event is cancelled by a forbidden event and no-signaling also fails

These scenarios are deliberately small. The goal is to catch wiring errors, sign errors, and invalid escape routes before a larger Master-Equation packet consumes the residuals.

Proof Scaffold Boundary

The harness encodes a useful obstruction:

$$P_{\theta}(\mathbf{r}|\mathbf{s}) = \int_{\Pi} \prod_i K_i(r_i|s_i, \Pi) d\rho_{\text{prov}}(\Pi)$$

is still a Bell-local product form when $d\rho_{\text{prov}}(\Pi)$ is independent of the settings and Π is a complete common-past screen. Such a model cannot pass CHSH, GHZ, and Hardy as a family. A successful `AAA` closure must therefore derive a stronger object:

$$P_{\theta}(\mathbf{r}|\mathbf{s}) = \mu_{*,T}^{(n)}(B_{\mathbf{r}}^{\mathbf{s}})$$

where $B_{\mathbf{r}}^{\mathbf{s}}$ is the record-basin subset for the declared preparation, pair or multiplet provenance, local apparatus kernels, coarse-graining, and record window. This is the same measurement discipline used in [Measurement Ontology](#), but lifted from single-assembly basin weights to a Bell-family joint record measure.

The native proof packet must supply:

1. a pair-provenance ledger Π_{AB} or multiplet ledger Π_{ABC} ;
2. local apparatus kernels derived from the Stern-Gerlach-like or photon-analyzer channel;
3. one finite-window measure $\mu_{*,T}^{(n)}$ on the retained joint record manifold;
4. a compression audit showing why the completed record law does not reduce to Bell-local product screening;
5. no-signaling and measurement-independence residuals evaluated on the same packet.

The single-assembly Stern-Gerlach response in [Angular Momentum and Spin](#) is a prerequisite, not the Bell proof itself. Bell-family closure starts only after the pair-provenance measure and the joint record basins are explicit.

Acceptance Boundary

Passing this harness means only that the residual calculations and negative controls behave as expected. It does not validate ~~AAA~~ quantum closure.

A future closure packet becomes promotable only if:

1. the probability tables are generated from declared substrate variables rather than written by hand;
2. Δ_{MI} and Δ_{NS} remain within tolerance;
3. CHSH, GHZ, and Hardy benchmarks are evaluated together;
4. the same $\mu_{*,T}^{(n)}$ also agrees with the record and repeated-frequency discipline in [Quantum Operator Mapping](#);
5. the product-screening audit does not collapse the completed hidden-variable record into $\int_{\Pi} \prod_i K_i d\rho_{\text{prov}}$;
6. failure cases are reported when the model reduces to classical-axis response, separable pair measure, product-screened pair provenance, context-independent GHZ values, forbidden Hardy events, setting-dependent provenance, or signaling marginals.

Nested Shell Braid Action-Increment Protocol

This protocol defines the simulation-facing test for deriving or falsifying the one-cycle action increment used by the quantum closure program. It specializes [Simulation Run Protocols](#) and [Convergence Tests](#) to the question left open by [Nested Shell Braid Dynamics](#), [Dyadic Resonance Lock](#), [Angular Momentum and Spin](#), and [Mapping the Planck Scale](#).

The target is narrow. The run must compute the smallest accepted Master-Equation projected action increment from candidate nested shell braid branch transitions whose stability rows pass. It may compare the resulting scale to the observer-level h , $\hbar$ benchmark after the computation. It may not insert $\hbar$ as an input step size.

Closure Question

The action-angle bridge in [Angular Momentum and Spin](#) states the conditional theorem target:

$$\Delta I_i = \hbar \implies \Delta \Gamma_{\text{cell}} = h^n$$

for n record-facing action-angle channels. This protocol tests the missing premise. It asks whether accepted nested shell braid dynamics select a positive increment ΔI_* such that

$$h_{\text{AAA}} = 2\pi \Delta I_*$$

matches the observer-level Planck constant benchmark.

Passing this protocol would not complete quantum theory. It would only promote the action-increment step from bookkeeping convention to candidate derived output.

Accepted Transition Class

Let B_q and $B_{q'}$ denote candidate nested shell braid branch states with passed stability rows, layer radii, frequencies, speeds, plane normals, active causal-root ledger, and wake ledger. A candidate accepted transition belongs to

$$\mathcal{T}_{\text{acc}} = \emptyset$$

unless both endpoint packets first satisfy branch-certificate eligibility: matching ledger identity, matching active-root convention, positive Jacobian floors, declared inactive-root or tail status, $\Delta_{\mathbf{k}} > 0$, conservation pullback on the same rows, and refinement records sufficient to keep the endpoint status stable. Before that eligibility is supplied, a run may report diagnostics or rejected endpoint packets, but it may not promote `candidate_action_increment` or `candidate_h_recovery`.

When endpoint eligibility has been established, the accepted transition class is

$$\mathcal{T}_{\text{acc}} = \{B_q \rightarrow B_{q'} : \Delta_{\mathbf{k}} > 0, \mathcal{R}_{\text{phase}} \leq \tau_{\text{phase}}, \mathcal{R}_E \leq \tau_E, \mathcal{R}_P \leq \tau_P, \mathcal{R}_J \leq \tau_J, \Delta N_{\text{self}} \in 2\mathbb{Z}, \mathcal{R}_{\text{root}} \leq \tau_{\text{root}}\}$$

The tolerances τ_{phase} , τ_E , τ_P , τ_J , and τ_{root} must be declared before the run. The transition is not accepted merely because it improves a fit to h .

Plain language: only stable, conservation-accounted, root-accounted branch changes are allowed to vote on the action increment.

Master-Equation Increment

For each candidate transition, compute the layer torque integrals and wake boundary term directly from the delayed dynamics. With transaction axis $\hat{\mathbf{a}}$, the projected increment is

$$\Delta I_{\text{ME}} = \hat{\mathbf{a}} \cdot \left(\sum_{\ell \in \{I, M, O\}} \int_{t_i}^{t_f} \mathbf{T}_\ell(s) ds + \Delta \mathbf{L}_{\text{wake}, \partial} \right)$$

Here $\mathbf{T}_\ell$ is the layer torque reconstructed from causal-wake forces on the architrinos in layer ℓ , and $\Delta \mathbf{L}_{\text{wake}, \partial}$ is the angular momentum still carried across the chosen core boundary at the end of the transition window.

Branch-Chart Conservation Pullback

The projected action increment is a diagnostic until the exact nonlocal Noether charges close on the same live-ledger branch chart. For each accepted transition, pull the normalized delayed-interior characteristic-tail increments from [Master Equation](#) back to the candidate branch rows and report

$$\begin{aligned} \mathcal{E}_{\text{tot}}^{(\eta)} &= K_\mu + E_{\text{wake,eff}}^{(\eta)}, & \mathcal{P}_{\text{tot}}^{(\eta)} &= \mathbf{P}_{\text{mech}} + \mathbf{P}_{\text{wake,eff}}^{(\eta)} \\ \mathcal{J}_{\text{tot}}^{(\eta)} &= \mathbf{J}_{\text{mech}} + \mathbf{J}_{\text{wake,eff}}^{(\eta)} \end{aligned}$$

The residuals $\mathcal{R}_E$, $\mathcal{R}_P$, and $\mathcal{R}_J$ are the normalized window changes of these three totals after subtracting the declared Euler-residual and endpoint-leakage terms. They must use the same branch rows as the root ledger, force residual, and ΔI_{ME} calculation. A work-integral energy reconstruction or torque projection may be reported as a diagnostic, but it does not replace the exact wake-history pullback.

The candidate increment floor is

$$\Delta I_* = \inf_{B_q \rightarrow B_{q'} \in \mathcal{T}_{\text{acc}}} |\Delta I_{\text{ME}}(B_q \rightarrow B_{q'})|$$

with required positivity condition

$$0 < \Delta I_* < \infty$$

The benchmark comparison is

$$\delta_h = \left| \frac{2\pi \Delta I_* - h}{h} \right|$$

Cluster and Stability Residuals

Because a single transition can be a numerical accident, the packet must scan a family of branch transitions with passed stability rows. For a selected class $\mathcal{C} \subset \mathcal{T}_{\text{acc}}$, report

$$\delta_I(\mathcal{C}) = \frac{\text{std}_{\mathcal{C}}(\Delta I_{\text{ME}})}{|\text{mean}_{\mathcal{C}}(\Delta I_{\text{ME}})| + \varepsilon_0}, \quad \varepsilon_0 = 10^{-12}$$

Also report the Floquet basin-robustness gap

$$\Delta_{\mathbf{k}} = 1 - \max_{i \notin G} \|\mu_i(\mathbf{k})\|$$

for each endpoint branch and each transition continuation.

The action-increment claim is numerically meaningful only when δ_I is small, $\Delta_{\mathbf{k}} > 0$, and the phase, energy, and root residuals remain below their predeclared tolerances across refinement.

Required Packet Files

The minimum campaign packet contains:

File	Required contents
campaign.json	source commit, protocol version, run ids, integrator, tolerances, declared benchmark policy, and whether $h, \hbar$ entered only after the Master-Equation increment was computed
branch_pairs.csv	each $B_q \rightarrow B_{q'}$ row, branch labels, integer windings, inter-layer closure integers, transition window, and inclusion/exclusion status
state_vectors.json	pre/post layer radii, frequencies, speeds, plane normals, phase offsets, source channel, transaction axis, and mechanical endpoint charges
root_ledger_before_after.json	partner, self, and inter-layer roots before and after transition, with delays, action-level g, u , Jacobians, separator flags, and ΔN_{self}

File	Required contents
torque_integrals.csv	diagnostic $\int \mathbf{T}_I dt$, $\int \mathbf{T}_M dt$, $\int \mathbf{T}_O dt$, $\Delta \mathbf{L}_{\text{wake},\partial}$, and projection onto $\hat{\mathbf{a}}$
action_increment_rows.csv	ΔI_{ME} , absolute value, cluster id, accepted/rejected status, and failure code
energy_ledger.csv	$\sum_{\ell} \int \omega_{\ell} dI_{\ell}$, ΔE_{wake} , ΔE_{coupl} , exact $E_{\text{wake,eff}}^{(\eta)}$, diagnostic U if used, and $\mathcal{R}_E$
conservation_pullback.csv	branch-chart id, cut/window id, η , ϵ_c , h , endpoint convention, ν_J , inactive-gap minimum, h_{mem} , K_{μ} , $E_{\text{wake,eff}}^{(\eta)}$, $\mathbf{P}_{\text{mech}}$, $\mathbf{P}_{\text{wake,eff}}^{(\eta)}$, $\mathbf{J}_{\text{mech}}$, $\mathbf{J}_{\text{wake,eff}}^{(\eta)}$, $\mathcal{R}_E$, $\mathcal{R}_P$, $\mathcal{R}_J$, and verdict
phase_closure_residuals.csv	layer and inter-layer phase closure residuals, winding labels, and tolerance status
floquet_report.json	monodromy or finite-difference return map, excluded symmetry modes, multipliers, and $\Delta_{\mathbf{k}}$
cluster_summary.json	ΔI_* , class means, class standard deviations, δ_I , h_{AAA} , δ_h , and promotion status
convergence_table.csv	the convergence rows required by Convergence Tests , including active-root mismatch and stability-window shift
negative_control_report.md	null runs and the invariant, provenance, or stability channel they break
promotion_gate.md	final pass/fail statement and the strongest claim the packet authorizes

Promotion Gates

A packet may promote candidate_action_increment only if all of the following pass:

1. $h, \hbar$ are absent from the simulated equations of motion and accepted-transition selection, except as post-run benchmark labels.
2. Both endpoint packets satisfy branch-certificate eligibility on matching ledger identity and active-root convention.
3. At least one transition class has $0 < \Delta I_* < \infty$.
4. Endpoint branches and transition continuations have $\Delta_{\mathbf{k}} > 0$ after symmetry modes are removed.
5. Phase closure, root residuals, energy residuals, momentum residuals, and angular-momentum residuals pass the predeclared tolerances.
6. δ_I is below the predeclared cluster tolerance.
7. The temporal, history-resolution, spatial, cross-integrator, and negative-control checks from [Convergence Tests](#) pass.
8. The packet reports δ_h honestly, whether or not the benchmark match is good.

Only a packet that also has small δ_h may promote candidate_h_recovery. A packet with a positive and stable ΔI_* but poor δ_h promotes only a derived action increment that does not recover the measured Planck benchmark.

Failure-Code Enum

Code	Trigger
input-hbar-contamination	the run seeded transition size, branch selection, or tolerances from $\hbar$ before computing ΔI_{ME}
no-positive-increment-floor	accepted transitions accumulate arbitrarily small nonzero ΔI_{ME}
multi-cluster-action-scale	multiple stable increment clusters appear with no derived reason to choose one
nonpositive-floquet-gap	an endpoint branch or transition continuation has $\Delta_{\mathbf{k}} \leq 0$
phase-closure-open	layer or inter-layer closure residuals exceed tolerance
root-ledger-instability	active roots change under refinement or the self-hit parity condition fails
energy-ledger-open	$\mathcal{R}_E$ exceeds tolerance or the wake/root energy channel is unaccounted
conservation-pullback-open	$\mathcal{R}_P$ or $\mathcal{R}_J$ exceeds tolerance, or the exact Noether pullback uses different rows than the root ledger or force residual
convergence-fail	required convergence or cross-integrator gates fail
negative-control-fail	the intentionally wrong model still passes the packet gates
benchmark-mismatch	h_{AAA} is stable but fails the declared h benchmark tolerance

Interpretation

This protocol preserves the level distinction. A passing action-increment packet would support the action-cell step used by [Wavefunction Ontology](#). It would not by itself derive the Born rule, spin statistics, Bell correlations, photon polarization, or observer-level orbital quantum numbers. Those remain downstream closure targets.

Retuning-Map Toy Model

This protocol documents the first arithmetic fixture for the cadence-scale retuning map introduced in [Nested Shell Braid Dynamics](#). The fixture is not a delayed-dynamics proof. It replays the constrained branch bookkeeping for an accepted $\Delta A_{\text{cyc}} = \pm \hbar$ transaction and reports whether the resulting increment can be treated as a same-branch retuning.

The purpose is narrow: turn the retuning scaffold into a machine-readable packet that outputs $(\Delta \nu_N, \Delta R_I, \Delta R_M, \Delta R_O, \Delta \lambda, \Delta \xi)$ and the corresponding first estimate for the cadence-space current J_{ν} .

Runtime Artifact

Run the default mock packet with:

```
node scripts/nested-shell-braid/retuning-map-toy-model.mjs --pretty
```

The script consumes:

```
scripts/nested-shell-braid/retuning-map-mock.json
```

and emits one result row per scenario. The packet is dimensionless: action increments are in units of h , speeds are compared to the declared c_f , and radius/cadence changes are reported as logarithmic increments plus reconstructed component changes.

Replay Equation

On branch chart q , the toy state is

$$\mathbf{y}_q = (\ln \nu_I, \ln \nu_M, \ln \nu_O, \ln R_I, \ln R_M, \ln R_O, \ln \lambda, \ln \xi)^T$$

Given a positive semidefinite retuning-cost matrix $\mathbf{K}_q^{\text{ret}}$, the fixture solves

$$\Delta \mathbf{y}_{q,\sigma} = \arg \min_{\Delta \mathbf{y}} \frac{1}{2} \Delta \mathbf{y}^T \mathbf{K}_q^{\text{ret}} \Delta \mathbf{y}$$

subject to

$$DA_{\text{cyc},q}[\Delta \mathbf{y}] + \Delta A_{\text{wake}} = \sigma h$$

and the declared linearized branch constraints. The layer-speed diagnostics are then checked through

$$\Delta \ln s_\ell = \Delta \ln R_\ell + \Delta \ln \nu_\ell, \quad \ell \in \{I, M, O\}$$

The script applies the ordinary nested shell braid speed gates:

$$s'_I > c_f, \quad |s'_M - c_f| \leq \epsilon_M c_f, \quad s'_O < c_f$$

The representative Noether braid cadence increment is

$$\Delta \ln \nu_N = w_I \Delta \ln \nu_I + w_M \Delta \ln \nu_M + w_O \Delta \ln \nu_O, \quad w_I + w_M + w_O = 1$$

For a local rate density r_σ of accepted σ transactions per braid, the first current estimate is

$$J_\nu = \sum_{\sigma=\pm 1} f_N r_\sigma \Delta \nu_N^{(q,\sigma)} + O((\Delta \nu_N)^2 \partial_\nu f_N)$$

Input Packet

Each scenario supplies:

Field	Meaning
reference_state	baseline $R_I, R_M, R_O, \lambda, \xi, \nu_N, s_I, s_M, s_O, c_f, \epsilon_M$
representative_cadence_weights	weights w_I, w_M, w_O used to extract $\Delta \nu_N$
compliance_diagonal	diagonal version of $\mathbf{K}_q^{\text{ret}}$
action_gradient_h_per_log	linearized $DA_{\text{cyc},q}$ row in h units per log variable
constraints	linearized branch constraints, each with coefficients and target
f_N	local Noether braid cadence-state distribution value
partial_nu_f_N	local slope used only to estimate the higher-order current remainder
transactions	accepted or control σ transactions with wake action increment and local rate density

This fixture intentionally starts with a diagonal compliance matrix. A later branch packet can replace it with a full matrix once the linearized return map supplies off-diagonal coupling.

Output Diagnostics

The fixture reports:

Output field	Meaning
status	candidate only when constraints and speed gates pass
delta_y	solved logarithmic retuning vector
retuning_components	$(\Delta\nu_N, \Delta R_I, \Delta R_M, \Delta R_O, \Delta\lambda, \Delta\xi)$
constraint_residual_max	largest absolute residual in the declared linear constraints
speed_gates	post-retuning checks for inner, middle, and outer layer speed regimes
J_nu.contribution	$f_{Nr,\sigma} \Delta\nu_N^{(q,\sigma)}$ for the transaction
net_J_nu.value	sum of transaction contributions in the scenario
net_J_nu.higher_order_estimate	magnitude estimate for the omitted $O((\Delta\nu_N)^2 \partial_\nu f_N)$ term

Expected Mock Behavior

The default mock packet has two rows.

Scenario	Expected behavior
same_branch_plus_minus_balance	Plus and minus one- h retunings both pass the speed gates. Unequal local rates leave a small signed current, net_J_nu.value near 0.0017019.
middle_hinge_violation_control	The linear action constraint solves, but the middle layer leaves the declared hinge tolerance. The row fails with middle-hinge-violation.

These numbers are fixture expectations only. They validate arithmetic, packet shape, branch-gate reporting, and the current estimate. They do not validate a physical Noether braid branch.

Failure Reading

The first failure modes are concrete:

Diagnostic pattern	Meaning
nonzero constraint_residual_max above tolerance	the declared linearized branch constraints are not actually solved
middle-hinge-violation	the retuning cannot be treated as a same-regime middle-hinge update
inner-speed-regime-crossing or outer-speed-regime-crossing	the transaction crosses a speed-regime boundary
large higher-order current estimate	the continuum current requires smaller steps, narrower bins, or a higher-order transport model
candidate branch with missing physical return-map source	the fixture is arithmetic only and must be replaced by a delayed-dynamics branch packet before promotion

A promotable retuning packet must eventually replace the mock compliance matrix with a return-map-derived $\mathbf{K}_q^{\text{ret}}$, preserve the same causal-root ledger, and keep the speed gates attached to the same branch state that supplies $\Delta\nu_N$.

Metric / Observer

Static Response Vector Toy Model

This protocol documents the first replay fixture for the weak static response vector used in the Γ_N geometry extraction target. It is a small arithmetic gate for the endpoint row in [Proper Time and Time Dilation](#) and the Shapiro-delay coefficient in [PPN Parameters](#).

The fixture is not an empirical PPN fit. Its purpose is to keep the clock cadence row, the clock-rate row, and the signal-delay coefficient from being silently blended while the $\mathbf{A}\mathbf{A}\mathbf{A}$ constitutive response is still being derived.

Runtime Artifact

Run the default mock packet with:

```
node scripts/spacetime/static-response-vector-toy-model.mjs --pretty
```

The script consumes:

```
scripts/spacetime/static-response-vector-mock.json
```

and emits one result row per scenario.

Replay Equations

For a weak static endpoint cell, write

$$\ln n = a_n \frac{U}{c_0^2}, \quad \ln \chi_{\text{sea}} = a_\chi \frac{U}{c_0^2}, \quad \ln \lambda = a_\lambda \frac{U}{c_0^2}, \quad \ln \frac{R_{\text{core}}}{R_{\text{core},0}} = a_R \frac{U}{c_0^2}$$

The cadence-stretch row must satisfy

$$b_n a_n + b_\chi a_\chi + b_\lambda a_\lambda + b_R a_R = 1$$

while the inverse clock-rate row must satisfy

$$\omega_n a_n + \omega_\chi a_\chi + \omega_\lambda a_\lambda + \omega_R a_R = -1$$

The row-inverse condition checks

$$b_i + \omega_i = 0$$

for $i \in \{n, \chi, \lambda, R\}$.

The Shapiro-delay neighbor supplies

$$a_\chi^{\text{sig}} = 1 + \gamma_{\text{eff}}$$

so the shared clock/signal delay residual is

$$\Delta_\chi^{\text{clk-sig}} = a_\chi - a_\chi^{\text{sig}}$$

The branch is shared-delay closed only when $\Delta_\chi^{\text{clk-sig}} = 0$ within the declared tolerance.

The same arithmetic also exposes the lensing/dynamics equality burden used by dark-sector comparisons. In the weak static row, the signal-deflection channel is closed only when the spatial-compliance response gives

$$\gamma_{\text{eff}} = 1, \quad a_\chi^{\text{sig}} = 2$$

while the clock/dynamical endpoint row still satisfies the cadence and inverse-clock equations below. A response vector that changes the dynamical acceleration but leaves $a_\chi^{\text{sig}} \neq 2$ is a split clock/signal branch: it may fit rotation curves or hydrostatic motion, but it cannot yet claim the lensing mass equality required by cluster and galaxy-galaxy weak-lensing tests.

Minimal Shared-Delay Packet

The first admissible static endpoint packet is the shared scalar delay response specialization of the equations above. Define

$$A_\chi \equiv 1 + \gamma_{\text{eff}}$$

The minimal response vector is

$$(a_n, a_\chi, a_\lambda, a_R) = (0, A_\chi, 0, 0)$$

with cadence row

$$(b_n, b_\chi, b_\lambda, b_R) = (0, A_\chi^{-1}, 0, 0)$$

and inverse clock-rate row

$$(\omega_n, \omega_\chi, \omega_\lambda, \omega_R) = (0, -A_\chi^{-1}, 0, 0)$$

For the GR-matching branch, this gives $A_\chi = 2$, $a_\chi = 2$, $b_\chi = 1/2$, and $\omega_\chi = -1/2$. The `shared_delay_clean_gr_branch` row in the mock packet is exactly this replay. The `density_scale_compensated_branch` row samples the remaining compensated family, where nonzero a_n , a_λ , or a_R are allowed only if the same cadence row still satisfies $b_n a_n + b_\chi a_\chi + b_\lambda a_\lambda + b_R a_R = 1$ and the inverse row remains $\omega_i = -b_i$.

Pressure Bridge

Pressure-response packets can feed the same fixture after their anisotropic terms are separated from the isotropic static projection. For a pressure row r , the bridge uses

$$\delta \mathbf{g}_r^P = (\delta \ln n, \delta \ln \chi_{\text{sea}}, \delta \ln \lambda, \delta \ln R)_r$$

and checks the pressure version of the cadence row:

$$\widehat{\delta \ln \Gamma_{N,r}^P} = b_n \delta \ln n_r + b_\chi \delta \ln \chi_{\text{sea},r} + b_\lambda \delta \ln \lambda_r + b_R \delta \ln R_r$$

The pressure cadence residual is

$$\mathcal{R}_{\Gamma,r}^P = \widehat{\delta \ln \Gamma_{N,r}^P} - \delta \ln \Gamma_{N,r}$$

The inverse clock-rate row must also close:

$$\mathcal{R}_{C,r}^P = (\omega_n \delta \ln n_r + \omega_\chi \delta \ln \chi_{\text{sea},r} + \omega_\lambda \delta \ln \lambda_r + \omega_R \delta \ln R_r) + \delta \ln \Gamma_{N,r}$$

When `derive_response` is `gamma_normalized`, the fixture also forms a normalized static-equivalent response vector

$$a_i^{P \rightarrow \Gamma} = \frac{\delta g_i^P}{\delta \ln \Gamma_N}$$

This normalization makes pressure rows replayable by the same endpoint arithmetic, but it does not convert pressure loading into a gravitational PPN branch. The `gamma_eff_sweep` diagnostic is only an algebraic comparison against $a_\chi^{\text{sig}} = 1 + \gamma_{\text{eff}}$; a pressure-normalized value that closes for some formal γ_{eff} is not a solar-system Shapiro result.

Anisotropic pressure entries, such as $\Delta \Pi^{\parallel - \perp}$ or deviatoric strain, must be either projected out before the isotropic static row is evaluated or carried in `anisotropic_residuals`. The isotropic Γ_N row must not absorb directional pressure response as a hidden scalar coefficient.

Input Packet

Each scenario supplies:

Field	Meaning
<code>gamma_eff</code>	PPN Shapiro-delay coefficient through $a_\chi^{\text{sig}} = 1 + \gamma_{\text{eff}}$
<code>gamma_eff_sweep</code>	optional list of trial γ_{eff} values for the shared-delay diagnostic
<code>response</code>	static weak-potential response vector $(a_n, a_\chi, a_\lambda, a_R)$
<code>pressure_bridge</code>	optional pressure row used to derive a normalized static-equivalent response vector
<code>cadence_row</code>	cadence-stretch coefficients $(b_n, b_\chi, b_\lambda, b_R)$ for $\ln \Gamma_N$
<code>clock_rate_row</code>	inverse clock-rate coefficients $(\omega_n, \omega_\chi, \omega_\lambda, \omega_R)$
<code>expect_shared_delay</code>	whether the scenario is expected to satisfy $\Delta_\chi^{\text{clk-sig}} = 0$
<code>tolerance</code>	optional scenario-level residual tolerance

Output Diagnostics

The fixture reports:

Output field	Meaning
<code>diagnostics.a_chi_sig</code>	signal-delay coefficient fixed by the PPN Shapiro map
<code>diagnostics.delta_chi_clk_sig</code>	shared clock/signal delay residual
<code>diagnostics.gamma_eff_sweep</code>	optional sweep of shared-delay residuals over trial γ_{eff} values
<code>diagnostics.endpoint_sum</code>	cadence-stretch row sum
<code>diagnostics.endpoint_residual</code>	endpoint residual relative to 1
<code>diagnostics.clock_rate_sum</code>	inverse clock-rate row sum
<code>diagnostics.clock_rate_residual</code>	clock-rate residual relative to -1
<code>diagnostics.row_inverse_residuals</code>	coefficient-by-coefficient residuals $b_i + \omega_i$
<code>diagnostics.pressure_bridge</code>	optional pressure-row replay of $\mathcal{R}_T^P$, $\mathcal{R}_C^P$, and effective-speed identity

These diagnostics turn the first-order response vector into an executable closure object. A later constitutive simulation can replace the mock response values with measured $(a_n, a_\chi, a_\lambda, a_R)$ rows while keeping the same gate.

Expected Mock Behavior

The default mock packet has five rows.

Scenario	Expected behavior
<code>shared_delay_clean_gr_branch</code>	Passes with $\gamma_{\text{eff}} = 1$, $a_\chi = 2$, and $b_\chi = 0.5$.
<code>density_scale_compensated_branch</code>	Passes with nonzero density, scale, and core-radius responses while preserving the endpoint and row-inverse constraints.
<code>split_clock_signal_delay_branch</code>	Fails shared-delay closure even though its endpoint and clock-rate rows close arithmetically.
<code>underclosed_clock_row</code>	Fails the endpoint and clock-rate sums while satisfying the shared-delay residual.
<code>pressure_bridge_fe_cr_toy_isotropic_projection</code>	Passes the pressure-projected cadence and clock-rate rows using the Fe/Cr toy isotropic projection, while correctly reporting that its pressure-normalized $a_\chi^{P \rightarrow \Gamma} = 0.6$ is not the GR-matching Shapiro branch.

The two failing rows are intentional failure witnesses. They show that a model can fit the static clock row while violating shared delay, or satisfy shared delay while underclosing the endpoint row. The pressure bridge row is a third kind of witness: it demonstrates that a pressure packet can close the isotropic Γ_N arithmetic while still remaining outside the gravitational PPN interpretation.

Compensated-Family Validation Result

The executable separates three claims that should not be collapsed.

First, the minimal shared-delay row passes the weak GR endpoint:

$$(a_n, a_\chi, a_\lambda, a_R) = (0, 2, 0, 0), \quad (b_n, b_\chi, b_\lambda, b_R) = \left(0, \frac{1}{2}, 0, 0\right)$$

Second, the density/scale-compensated row also passes the endpoint and inverse-row checks:

$$(a_n, a_\chi, a_\lambda, a_R) = (0.25, 2, -0.1, 0.05), \quad (b_n, b_\chi, b_\lambda, b_R) = (0.4, 0.4, -0.5, 1)$$

because

$$0.4(0.25) + 0.4(2) + (-0.5)(-0.1) + 1(0.05) = 1$$

This is an admissibility witness for the compensated static family, not a derivation of those numbers.

Third, the Fe/Cr pressure bridge falsifies the χ_{sea} -only shared row for the toy isotropic pressure projection. The pressure-normalized response is

$$\mathbf{a}^{P \rightarrow \Gamma} = (0, 0.6, 0, 0)^T$$

so no single χ_{sea} coefficient can satisfy both

$$b_\chi(2) = 1, \quad b_\chi(0.6) = 1$$

The current validation status is therefore conditional. Nonzero static endpoint coefficients a_n , a_λ , and a_R are not required by the endpoint row itself. They become necessary only if an independent branch record, such as hydrogen spectral refinement or pressure-response replay, supplies non- χ_{sea} response that must share the same Γ_N row. The next proof obligation is to replace the toy nonzero entries with branch-derived density, envelope-scale, or R_{core} response rather than treating them as fit parameters.

The hydrogen spectral toy scan may replay this compensated row as a scaffold, but that replay is not evidence that the gravitational endpoint has acquired nonzero a_n , a_λ , or a_R . Those entries become promotable only when the hydrogen branch or another declared branch derives the same component split for the same Noether sea cell.

Hydrogen Gamma_N Spectral Coefficient Row Toy Scan

This protocol is the first proof/simulation packet for the hydrogen spectral coefficient row $\mathbf{b}_N^{\text{spec}}$. Its purpose is narrow: constrain the row that extracts Γ_N for the hydrogen spectral channel without fitting a separate clock factor to each line.

The packet depends on the clock/rate convention in [Proper Time and Time Dilation](#) and the hydrogen line-set benchmark in [Atomic Spectra](#). It keeps the cadence-stretch factor and the observer frequency multiplier separate:

$$C_{N,H}^{(\ell)} = \left(\Gamma_{N,H}^{(\ell)}\right)^{-1}$$

Runtime Artifact

Run the default executable packet with:

```
node scripts/spacetime/hydrogen-gamma-n-spectral-row-toy-scan.mjs --pretty
```

The script consumes:

```
scripts/spacetime/hydrogen-gamma-n-spectral-row-mock.json
```

and emits one result row per scenario. The packet also keeps one mock passing shared-row case and intentional failure witnesses for direct cadence multiplication, per-line row fitting, endpoint-row violation, and response-record mismatch.

The default packet now begins with `hydrogen_rydberg_static_response_scaffold`. That scenario is not a completed hydrogen derivation, but it is the first theory-bearing input scaffold: the line labels are ordinary hydrogen transitions with recovered principal labels, the executable derives normalized Rydberg line factors, the envelope gaps declare one shared line-inferred cadence stretch, the $\mathbf{g}_{N,H}^{(\ell)}$ entries preserve the density/delay/scale/core split, and the static response vector is inherited from the static response packet rather than retuned inside the spectral scan.

Theory-Bearing Input Scaffold

The scaffold uses the line factors from the hydrogen Rydberg benchmark. For each line object, the executable reads the recovered labels `principal_n_a` and `principal_n_b` and forms

$$\Lambda_{ab} = \frac{1}{n_b^2} - \frac{1}{n_a^2}$$

The record-level `frequency_scale` represents the normalized $R_{\text{H}C_{\gamma,0}}$ comparison scale. In the first scaffold it is set to one, so the executable derives

$$\nu_{a \rightarrow b}^{\text{obs},(\ell)} = \Lambda_{ab}$$

The record-level `line_inferred_ln_Gamma_N` then supplies the line-inferred cadence stretch used to derive the replay envelope gap:

$$\frac{E_{\text{env}}^{(\ell)}(a) - E_{\text{env}}^{(\ell)}(b)}{h} = e^{0.001} \Lambda_{ab}$$

so every selected line infers

$$\ln \widehat{\Gamma}_{N,H}^{(\ell)}(a, b) = 0.001$$

The accepted scaffold row is the density/scale-compensated static-response row

$$\mathbf{b}_N^{\text{spec}} = (0.4, 0.4, -0.5, 1, 1)$$

with static response vector

$$(a_n, a_\chi, a_\lambda, a_R) = (0.25, 2, -0.1, 0.05)$$

It satisfies the endpoint constraint because

$$0.4(0.25) + 0.4(2) + (-0.5)(-0.1) + 1(0.05) = 1$$

The two admissible spectral records keep different component splits while preserving the same row prediction:

$$\mathbf{g}_{N,H}^{(A)} = (0.0005, 0.002, 0.0002, 0, 0.0001)^T, \quad \mathbf{g}_{N,H}^{(B)} = (0.0007, 0.0018, 0.0001, 0, 0.00005)^T$$

and

$$\mathbf{b}_N^{\text{spec}} \cdot \mathbf{g}_{N,H}^{(A)} = \mathbf{b}_N^{\text{spec}} \cdot \mathbf{g}_{N,H}^{(B)} = 0.001$$

This makes the packet stronger than a free mock arithmetic witness, but still below a constitutive hydrogen derivation. It checks that a declared row inherited from the static response packet can control several hydrogen line labels across two admissible records without collapsing n and χ_{sea} or fitting a separate coefficient row to each transition.

The scaffold still has a limited claim level. It derives the observer-frequency and envelope-gap entries from the Rydberg line-factor equation and a declared shared cadence stretch, but it does not derive the hydrogen envelope gaps from the master dynamics, does not derive the static response vector, and does not assign real observer frequencies. Its job is to make those inputs explicit and replaceable while keeping the coefficient-row scan executable.

Hydrogen Spectral Residual Separation

The row scan uses the Rydberg principal-label factor as its leading benchmark, but real hydrogen spectroscopy is not exhausted by that factor. The source-level comparison stack separates at least five corrections that must not be hidden inside Γ_N :

Channel	Standard benchmark role	Packet treatment
reduced mass	replaces m_e by $m_e M / (m_e + M)$ in the leading Coulomb spectrum	declared input to the envelope gap, not a per-line row fit
fine structure	relativistic kinetic, spin-orbit with Thomas-precession factor, and Darwin/contact terms split levels at order $(Z\alpha)^4$	later correction residual, not part of the shared cadence row
hyperfine structure	nuclear spin and magnetic moment couple to electron spin/orbital channels	apparatus/source-branch residual unless explicitly modeled
Lamb-type shift	QED photon-field correction splitting Dirac-degenerate levels	external QED recovery residual
finite nuclear structure	nuclear size, magnetic distribution, and quadrupole effects alter small-radius states	envelope/source-model residual

For a line $a \rightarrow b$, write the declared comparison gap as

$$\Delta E_{\text{H}}^{(\ell)}(a, b) = \Delta E_{\text{Ryd}}(a, b) + \Delta E_{\text{fs}}(a, b) + \Delta E_{\text{hfs}}(a, b) + \Delta E_{\text{Lamb}}(a, b) + \Delta E_{\text{nuc}}(a, b) + \Delta E_{\text{rem}}(a, b)$$

The current toy scaffold sets the correction terms to zero by declaration and therefore tests only the shared-row handling of the leading Rydberg factor. A non-toy packet must report a residual-separation check

$$\mathcal{R}_{\text{H,res}}^{(\ell)} = \max_{(a,b) \in \mathcal{L}_{\text{H}}^0} \frac{\left| \Delta E_{\text{H}}^{(\ell)}(a, b) - \sum_{c \in \{\text{Ryd, fs, hfs, Lamb, nuc}\}} \Delta E_c(a, b) \right|}{\varepsilon_{\text{rem}}(a, b)} \leq 1$$

This prevents the coefficient scan from passing by absorbing known spectral physics into the cadence-stretch row. It also fixes the degeneracy burden: the leading Coulomb target must recover the n^2 orbital degeneracy before correction channels split it, while the fine-structure channel may depend on j and the hyperfine channel may depend on nuclear-spin records.

Compensated-Row Readout

The current scaffold makes the compensated-family test explicit. The accepted split-record row is

$$\mathbf{b}_N^{\text{spec}} = (0.4, 0.4, -0.5, 1, 1)$$

with

$$\mathbf{g}_{N,\text{H}}^{(A)} = (0.0005, 0.002, 0.0002, 0, 0.0001)^T, \quad \mathbf{g}_{N,\text{H}}^{(B)} = (0.0007, 0.0018, 0.0001, 0, 0.00005)^T$$

The refinement difference satisfies

$$\mathbf{b}_N^{\text{spec}} \cdot (\mathbf{g}_{N,\text{H}}^{(B)} - \mathbf{g}_{N,\text{H}}^{(A)}) = 0$$

so both records give the same $\ln \Gamma_{N,\text{H}} = 0.001$ while preserving separate n , χ_{sea} , λ , and R_{core} entries. By contrast, the shared-delay-only control row

$$\left(0, \frac{1}{2}, 0, 1, 0 \right)$$

predicts a refinement mismatch of -0.0001 on record B in the default scaffold. This is the current hydrogen validation result: atom-local refinement can falsify the minimal row when the accepted response record changes component split, but the scaffold does not yet require nonzero gravitational endpoint coefficients a_n , a_λ , or a_R unless a constitutive hydrogen branch derives the same split from the static endpoint response.

Input Variables

Each toy packet supplies one weak-homogeneous hydrogen line set $\mathcal{L}_{\text{H}}^0$ and one or more admissible resolution records $\ell \in \mathcal{I}_{\text{spec}}^{\text{atom}}$. For each record, the packet declares:

Variable	Meaning
$\Theta_{\text{H,spec}}^{(\ell)}$	shared hydrogen channel ledger used to extract the envelope gaps and local Noether sea response
$\mathcal{L}_{\text{H}}^0$	chosen isolated hydrogen transitions $a \rightarrow b$ with recovered labels
$E_{\text{env}}^{(\ell)}(a) - E_{\text{env}}^{(\ell)}(b)$	envelope gap from the same spectral channel record
$\nu_{a \rightarrow b}^{\text{obs},(\ell)}$	observer-level frequency used only after the clock-rate conversion is declared
$\mathbf{g}_{N,\text{H}}^{(\ell)}$	shared clock-facing deformation record for the line set
$\varepsilon_{\Gamma}, \Delta_{\Gamma}^{\text{tol}}$	line-inferred cadence-stretch denominator floor and tolerance
$\varepsilon_{\text{row}}, \Delta_{\text{row}}^{\text{tol}}$	coefficient row denominator floor and row-stability tolerance
$\mathcal{R}_{\Gamma,\text{H}}^{\text{spec},(\ell)}$	declared higher-order residual budget, not a fitted clock row

The deformation record is the one used by the hydrogen clock/rate target:

$$\mathbf{g}_{N,\text{H}}^{(\ell)} = \left(\ln n_{\text{H}}^{(\ell)}, \ln \chi_{\text{sea,H}}^{(\ell)}, \ln \lambda_{\text{H}}^{(\ell)}, -\ln \xi_{\text{H}}^{(\ell)}, \ln \frac{R_{\text{core,H}}^{(\ell)}}{R_{\text{core,0}}^{(\ell)}} \right)^T$$

For each line, the packet also forms the line-inferred cadence stretch

$$\widehat{\Gamma}_{N,\text{H}}^{(\ell)}(a, b) = \frac{E_{\text{env}}^{(\ell)}(a) - E_{\text{env}}^{(\ell)}(b)}{h\nu_{a \rightarrow b}^{\text{obs},(\ell)}}$$

This inferred value is a diagnostic readout. It is not a permission to fit a separate Γ_N or coefficient row to the transition.

Coefficient Constraints

The spectral row has the same component order as the Γ_N extraction target:

$$\mathbf{b}_N^{\text{spec}} = (b_n^{\text{spec}}, b_\chi^{\text{spec}}, b_\lambda^{\text{spec}}, 1, b_R^{\text{spec}})$$

The fixed fourth entry is the inherited Lorentz-branch constraint $b_\xi = 1$. The remaining entries must satisfy the static weak-field endpoint constraint when evaluated on the same static response vector used by the clock row:

$$b_n^{\text{spec}} a_n + b_\chi^{\text{spec}} a_\chi + b_\lambda^{\text{spec}} a_\lambda + b_R^{\text{spec}} a_R = 1$$

within the declared endpoint tolerance. If the packet also supplies the inverse clock-rate row ω^{spec} , then it must satisfy

$$\omega_i^{\text{spec}} = -b_i^{\text{spec}}$$

for $i \in \{n, \chi, \lambda, R\}$. A branch may additionally impose shared clock/signal delay only by declaring the same condition used in the static response vector packet:

$$a_\chi = 1 + \gamma_{\text{eff}}$$

The spectral coefficient row is therefore a constrained row inherited from clock closure. It is not a spectral nuisance parameter and not a per-line normalization constant.

Minimal Toy Scan

The minimal scan is a finite grid over the four free entries $(b_n^{\text{spec}}, b_\chi^{\text{spec}}, b_\lambda^{\text{spec}}, b_R^{\text{spec}})$ after setting $b_\xi = 1$.

1. Reject every row that violates the endpoint constraint

$$|b_n^{\text{spec}} a_n + b_\chi^{\text{spec}} a_\chi + b_\lambda^{\text{spec}} a_\lambda + b_R^{\text{spec}} a_R - 1| > \Delta_{\text{row}}^{\text{tol}}$$

2. For each remaining row and resolution record, compute

$$\ln \Gamma_{N,H}^{\text{row},(\ell)} = \mathbf{b}_N^{\text{spec}} \cdot \mathbf{g}_{N,H}^{(\ell)}$$

3. Compare the row prediction to every line-inferred cadence stretch:

$$\mathcal{E}_\Gamma^{(\ell)}(a, b; \mathbf{b}_N^{\text{spec}}) = \ln \widehat{\Gamma}_{N,H}^{(\ell)}(a, b) - \ln \Gamma_{N,H}^{\text{row},(\ell)}$$

4. Across refinement records, require the accepted row to keep the same predicted clock-rate conversion after the envelope-gap convergence budget is removed:

$$\mathcal{E}_{\text{ref}}(\ell, \ell'; \mathbf{b}_N^{\text{spec}}) = \ln \Gamma_{N,H}^{\text{row},(\ell)} - \ln \Gamma_{N,H}^{\text{row},(\ell')}$$

The scan output is the accepted coefficient row set

$$\mathcal{B}_H^{\text{spec}} = \{\mathbf{b}_N^{\text{spec}} \mid \text{endpoint, line-set, and refinement residuals pass}\}$$

This set may be a point, a bounded interval family, or empty. A bounded family is still useful because it constrains the coefficient row without assigning a separate row to each spectral line.

Pass Condition

The toy scan passes when $\mathcal{B}_H^{\text{spec}}$ is nonempty and every accepted row satisfies

$$\max_{\ell, (a,b) \in \mathcal{L}_H^0} \frac{|\mathcal{E}_\Gamma^{(\ell)}(a, b; \mathbf{b}_N^{\text{spec}})|}{\left| \ln \widehat{\Gamma}_{N,H}^{(\ell)}(a, b) \right| + \varepsilon_\Gamma} \leq \Delta_\Gamma^{\text{tol}}$$

with the refinement check

$$\max_{\ell, \ell'} \frac{|\mathcal{E}_{\text{ref}}(\ell, \ell'; \mathbf{b}_N^{\text{spec}})|}{\left| \ln \Gamma_{N,H}^{\text{row},(\ell)} \right| + \varepsilon_{\text{row}}} \leq \Delta_{\text{row}}^{\text{tol}}$$

The stronger extraction claim requires the diameter of $\mathcal{B}_H^{\text{spec}}$ to shrink under additional independent hydrogen records or under a constitutive response calculation for $(a_n, a_\chi, a_\lambda, a_R)$. The first packet does not require that stronger claim; it only requires that a shared constrained row survive the line set.

This is not yet the full promotion gate. That gate requires $\mathbf{g}_{N,H}^{(\ell)}$, $E_{\text{env}}^{(\ell)}(a) - E_{\text{env}}^{(\ell)}(b)$, $\nu_{a \rightarrow b}^{\text{obs},(\ell)}$, and $(a_n, a_\chi, a_\lambda, a_R)$ to be extracted from one declared hydrogen spectral channel record and the same Noether sea cell, with recoil, hyperfine structure, photon-channel propagation, and source-branch effects carried outside Γ_N unless they are in the declared residual budget.

Hydrogen Γ_N Certificate Boundary

A deterministic hydrogen row is a certificate rather than only a nonempty accepted-row set. The certificate object is

$$\mathcal{C}_H^\Gamma = \left(\Theta_{H,\text{spec}}^{(\ell)}, \mathcal{L}_H^0, \mathbf{g}_{N,H}^{(\ell)}, \Delta E_{\text{env}}^{(\ell)}, \nu_{\text{obs}}^{(\ell)}, \mathbf{a}^G, \mathbf{b}_N^{\text{spec}}, \boldsymbol{\tau} \right)$$

where $\mathbf{a}^G = (a_n, a_\chi, a_\lambda, a_R)$ is the static Noether sea response row for the same cell and $\boldsymbol{\tau}$ collects the declared tolerances.

The certificate residual vector is

$$\mathcal{R}_H^\Gamma = \left(b_\xi^{\text{spec}} - 1, \mathbf{b}_{N,\text{stat}}^{\text{spec}} \cdot \mathbf{a}^G - 1, \mathcal{R}_{\text{line}}, \mathcal{R}_{\text{ref}}, \mathcal{R}_{H,\text{res}} \right)$$

with

$$\mathcal{R}_{\text{line}} = \max_{\ell, (a,b)} \frac{\left| \ln \widehat{\Gamma}_{N,H}^{(\ell)}(a,b) - \mathbf{b}_N^{\text{spec}} \cdot \mathbf{g}_{N,H}^{(\ell)} \right|}{\left| \ln \widehat{\Gamma}_{N,H}^{(\ell)}(a,b) \right| + \varepsilon_\Gamma}$$

and

$$\mathcal{R}_{\text{ref}} = \max_{\ell, \ell'} \frac{\left| \mathbf{b}_N^{\text{spec}} \cdot \left(\mathbf{g}_{N,H}^{(\ell)} - \mathbf{g}_{N,H}^{(\ell')} \right) \right|}{\left| \mathbf{b}_N^{\text{spec}} \cdot \mathbf{g}_{N,H}^{(\ell)} \right| + \varepsilon_{\text{row}}}$$

Here $\mathbf{b}_{N,\text{stat}}^{\text{spec}} = (b_n^{\text{spec}}, b_\chi^{\text{spec}}, b_\lambda^{\text{spec}}, b_R^{\text{spec}})$ is the four-entry static endpoint subrow. The packet passes only if every component of $\mathcal{R}_H^\Gamma$ is within its declared tolerance and all packet inputs share the same provenance ledger $\Theta_{H,\text{spec}}^{(\ell)}$ and the same static Noether sea cell. Otherwise it fails with the first violated row: provenance, b_ξ , endpoint, line-set, refinement, or residual separation.

Failure Tests

The packet must include intentional failing rows or records for the following cases:

Failure test	Required failure
direct cadence multiplication	using Γ_N instead of $C_N = \Gamma_N^{-1}$ in the observer-frequency comparison fails the line-set residual
per-line row fit	allowing $\mathbf{b}_N^{\text{spec}}(a,b)$ makes isolated lines pass but fails the shared-row condition
collapsed density/delay variable	replacing (n, χ_{sea}) by one scalar fails when the packet contains density-delay split records
endpoint-row violation	a row that fits the line set but violates $b_n a_n + b_\chi a_\chi + b_\lambda a_\lambda + b_R a_R = 1$ is rejected
residual overuse	hiding recoil, hyperfine structure, photon-channel propagation, or unresolved source-branch effects inside $\mathcal{R}_{\Gamma,H}^{\text{spec},(\ell)}$ beyond the declared budget fails
response-record mismatch	changing $\mathbf{g}_{N,H}^{(\ell)}$ between lines after $\mathcal{L}_H^0$ is chosen fails
spectral-correction collapse	absorbing fine-structure, hyperfine, Lamb-type, reduced-mass, or nuclear-size corrections into $\mathbf{b}_N^{\text{spec}}$ fails once the correction channels are declared

These failure tests keep the spectral row tied to the shared clock/rate map. They also separate the proof obligations: the envelope calculation owns the line gaps, the clock-row calculation owns Γ_N and C_N , and the photon-channel event record owns emission and absorption propagation.

Output Diagnostics

The executable packet reports:

Output field	Meaning
diagnostics.accepted_rows	candidate rows that satisfy $b_\xi = 1$, the endpoint constraint, the line-set residual, and the refinement residual
diagnostics.response_record_mismatch_pass	whether every line used the shared $\mathbf{g}_{N,H}^{(\ell)}$ record for its resolution
diagnostics.per_line_spoof	whether each line could be made to pass by some row even though no shared row passes
diagnostics.row_results[].diagnostics.endpoint_residual	residual for $b_n a_n + b_\chi a_\chi + b_\lambda a_\lambda + b_R a_R = 1$
diagnostics.row_results[].diagnostics.line_residuals	line-by-line values of $\mathcal{E}_\Gamma^{(\ell)}(a,b; \mathbf{b}_N^{\text{spec}})$
diagnostics.row_results[].diagnostics.line_residuals[].line_factor_Lambda_ab	derived or declared hydrogen line factor Λ_{ab}
diagnostics.row_results[].diagnostics.line_residuals[].envelope_gap_over_h	declared or derived envelope gap divided by h

Output field	Meaning
diagnostics.row_results[].diagnostics.line_residuals[].observed_frequency	declared or derived observer frequency used in the cadence-stretch readout
diagnostics.row_results[].diagnostics.refinement_residuals	resolution-pair residuals for the shared row prediction

The packet succeeds only when its declared expectations are met. A failure witness should therefore have status: "fail" but expectation_status: "pass" when it fails for the intended reason.

Thermodynamic Residual

This protocol turns the local-horizon target in [Emergent Metric](#) into a validation scaffold. It does not assume that gravity is thermodynamic at the substrate level. It tests whether one Noether sea state and observer-channel record can supply the three observer-level quantities used in the Jacobson comparison: boundary entropy, local temperature, and boost-energy flux.

The protocol is a proof-and-simulation target, not an empirical claim. A successful packet would show that the same record that recovers weak-field ADM/Cartan and PPN behavior also makes the local Clausius residual small in the equilibrium comparison regime.

Minimal Record

For each Physical Observer O , effective-horizon patch $\partial\Omega$, and finite analysis window $W = [t_a, t_b]$, the packet must declare one shared record θ with the following content.

Channel	Required content	Failure prevented
Noether sea state	$n(\mathbf{x}, t)$, $\rho_{\text{NS}}(\mathbf{x}, t)$, $\chi_{\text{sea}}(\mathbf{x}, t)$, u_{sea}^i , e^a_i , γ_{ij} , and N on the relevant region	fitting entropy, flux, and metric response with separate Noether sea states
Physical Observer	worldline, clock-rate record, access region, reference resources, and observer acceleration a_O derived from the metric channel	importing an external observer or a free Rindler frame
Boundary patch	$\partial\Omega$, effective patch area $A_{\partial\Omega}^{\text{eff}}$, orientation, and signed crossing convention	hiding the area comparison in an undefined horizon surface
Boundary wake labels	retained label set $\mathcal{B}_{\partial\Omega}^{(O)}(\theta; W)$ with emitter identity, emission time, receiver or sensor identity, reception time, channel, and persistence criterion	counting unrecorded or inaccessible microstates
Flux projection	either $T_{\mu\nu}^{\text{eff}}(\theta)$ on the patch or a declared discrete estimator from the same causal-wake and provenance logs	fitting dQ independently of the record
Gates	predeclared ϵ_{thermo} , ϵ_A , ϵ_E , convergence tolerances, and negative controls	selecting tolerances after seeing the output

Boundary Count and Area Slope

The observer-accessible boundary label set is

$$\mathcal{B}_{\partial\Omega}^{(O)}(\theta; W) = \{b : b \text{ is a retained boundary-wake label crossing } \partial\Omega \text{ during } W \text{ and readable by } O\}$$

The first entropy estimator is the microcanonical count

$$\widehat{S}_{\partial\Omega}^{(O)}(\theta; W) = k_B \log \left| \mathcal{B}_{\partial\Omega}^{(O)}(\theta; W) \right|, \quad \left| \mathcal{B}_{\partial\Omega}^{(O)}(\theta; W) \right| > 0$$

This finite count is a packet estimator, not the final horizon-interface coefficient. For coefficient recovery, a row should be interpreted as a finite-block sample of the block-density target

$$\widehat{s}_U^{(O)}(\theta; W) = \frac{1}{|U|} \log \left| \mathcal{B}_U^{(O)}(\theta; W) \right|$$

where U is the declared connected patch block and $\mathcal{B}_U^{(O)}$ retains only labels accessible to the same observer record. The large-block target is $\widehat{s}_U^{(O)} \rightarrow 1/4$ after boundary corrections, not a literal one-patch cardinality.

Area scaling is a recovery target, not a definition. Compare neighboring patches or refinements with the same observer and record:

$$\mathcal{R}_A^{(O)} = \frac{\left| \frac{\Delta \widehat{S}_{\partial\Omega}^{(O)}}{\Delta A_{\partial\Omega}^{\text{eff}}} - \frac{k_B}{4A_{\text{align}}} \right|}{\frac{k_B}{4A_{\text{align}}} + \epsilon}$$

Passing this subgate means the retained logarithmic label count has the target area slope in the relevant equilibrium regime. It does not yet prove Page-curve recovery or black-hole endpoint closure.

Temperature and Flux

The local temperature comparison must be derived from the observer-channel acceleration:

$$\hat{T}_U^{(O)} = \frac{\hbar a_O}{2\pi k_B c_0}, \quad a_O^2 = \gamma_{ij} a_O^i a_O^j$$

The continuum flux estimator is

$$\widehat{dQ}_{\partial\Omega}^{(O)}(\theta; W) = \int_W \int_{\partial\Omega} T_{\mu\nu}^{\text{eff}}(\theta) \xi^\mu d\Sigma^\nu$$

When the run has not constructed a continuum $T_{\mu\nu}^{\text{eff}}$, the packet may use a discrete estimator, but only if every term comes from the same boundary-wake and observer record:

$$\widehat{dQ}_{\partial\Omega, \text{disc}}^{(O)}(\theta; W) = \sum_{b \in \mathcal{B}_{\partial\Omega}^{(O)}(\theta; W)} \sigma_b E_b^{(O)} \omega_b^{(O)}$$

Here σ_b is the signed crossing convention, $E_b^{(O)}$ is the observer-level energy assigned by the same channel that builds $T_{\mu\nu}^{\text{eff}}$, and $\omega_b^{(O)}$ is the declared quadrature or coarse-graining weight.

The measured local-horizon residual is then

$$\widehat{\mathcal{R}}_{\text{thermo}}^{(O)} = \frac{\left| \widehat{dQ}_{\partial\Omega}^{(O)} - \hat{T}_U^{(O)} d\widehat{S}_{\partial\Omega}^{(O)} \right|}{\left| \widehat{dQ}_{\partial\Omega}^{(O)} \right| + \hat{T}_U^{(O)} \left| d\widehat{S}_{\partial\Omega}^{(O)} \right| + \varepsilon}$$

Conservation and Same-Record Gate

The thermodynamic comparison is not allowed to pass by sacrificing local observer-level conservation. The packet must also report

$$\mathcal{R}_{E, \partial\Omega}^{(O)} = \frac{\left| \Delta E_{\Omega}^{(O)}(\theta; W) + \widehat{dQ}_{\partial\Omega}^{(O)}(\theta; W) \right|}{\left| \Delta E_{\Omega}^{(O)}(\theta; W) \right| + \left| \widehat{dQ}_{\partial\Omega}^{(O)}(\theta; W) \right| + \varepsilon}$$

A local-horizon packet passes only when

$$\widehat{\mathcal{R}}_{\text{thermo}}^{(O)} \leq \epsilon_{\text{thermo}}, \quad \mathcal{R}_A^{(O)} \leq \epsilon_A, \quad \mathcal{R}_{E, \partial\Omega}^{(O)} \leq \epsilon_E$$

and the same θ also satisfies the weak-field metric gates relevant to the run. A packet that fits $\widehat{S}$, $\hat{T}_U$, and $\widehat{dQ}$ with independent records fails even if each scalar looks plausible by itself.

Free-Energy and Response Consistency

The same record should also support the near-equilibrium free-energy direction when such a channel is claimed. Let the packet declare a coarse state $z(\theta; t)$, entropy estimator $\widehat{S}_z$, energy estimator $\widehat{E}_z$, and local temperature $\widehat{T}_z$ built from the same observer and Noether sea record. Define

$$\widehat{F}_z = \widehat{E}_z - \widehat{T}_z \widehat{S}_z$$

On a relaxation window with no declared external work, the free-energy residual is

$$\widehat{\mathcal{R}}_F^{(O)} = \frac{\left[\Delta_W \widehat{F}_z - W_{\text{ext}, z}^{(O)} \right]_+}{\left| \Delta_W \widehat{F}_z \right| + \left| W_{\text{ext}, z}^{(O)} \right| + \varepsilon}$$

The gate is optional unless the packet uses free-energy minimization, order-parameter relaxation, or Landau-Ginzburg language. If invoked, it must pass with the same θ that supplies $\widehat{\mathcal{R}}_{\text{thermo}}^{(O)}$.

If the packet includes stochastic or fluctuation claims, it must report a response/noise residual rather than fitting noise independently. For a declared observable pair (A, B) , use the measured fluctuation spectrum $S_{AB}^{(O)}(\omega)$ and the dissipative response $\chi_{AB}^{\prime\prime(O)}(\omega)$:

$$\widehat{\mathcal{R}}_{\text{FD}}^{(O)}(A, B) = \frac{\left\| S_{AB}^{(O)}(\omega) - \mathcal{F}_{\widehat{T}_z} \left(\chi_{AB}^{\prime\prime(O)}(\omega) \right) \right\|_\omega}{\left\| S_{AB}^{(O)} \right\|_\omega + \left\| \mathcal{F}_{\widehat{T}_z} \left(\chi_{AB}^{\prime\prime(O)} \right) \right\|_\omega + \varepsilon}$$

Here $\mathcal{F}_{\widehat{T}_z}$ is the packet's declared classical or quantum fluctuation-dissipation map. This check is a same-record discipline for equilibrium response. It does not assert that Noether sea dynamics is fundamentally stochastic.

Proof Route

The proof route has four controlled steps.

1. Show that $\mathcal{B}_{\partial\Omega}^{(O)}(\theta; W)$ is stable under temporal, spatial, and history-resolution refinement for the fixed observer and patch.
2. Show that $\Delta\hat{S}/\Delta A_{\partial\Omega}^{\text{eff}}$ converges to $k_B/(4A_{\text{align}})$ in the equilibrium local-horizon regime.
3. Show that the flux estimator from the same θ satisfies $d\hat{Q} = \hat{T}_U d\hat{S} + O(\epsilon_{\text{thermo}})$ while $\mathcal{R}_{E,\partial\Omega}^{(O)}$ remains small.
4. Use the existing ADM/Cartan handoff to show that the same record recovers the weak-field observer metric. Only after this step may the Jacobson comparison be promoted from analogy to a native recovery route for the effective Einstein equation.

Failure Codes

Failure code	Meaning
thermo-label-coverage-open	the packet does not record enough boundary-wake labels to define $\mathcal{B}_{\partial\Omega}^{(O)}$
thermo-area-scaling-open	$\hat{S}$ scales with volume, history length, or patch choice rather than $A_{\partial\Omega}^{\text{eff}}$
thermo-temperature-split-open	$\hat{T}_U$ requires an acceleration or clock channel not present in the metric record
thermo-flux-split-open	$d\hat{Q}$ is fitted from a stress or energy record not used by the observer metric
thermo-residual-open	$\widehat{\mathcal{R}}_{\text{thermo}}^{(O)}$ exceeds the declared tolerance
thermo-conservation-open	$\mathcal{R}_{E,\partial\Omega}^{(O)}$ exceeds tolerance
thermo-ppn-split-open	the local-horizon residual passes only for a record that fails the weak-field ADM/Cartan or PPN gates
thermo-negative-control-open	a declared negative control still passes the local-horizon packet

Negative Controls

A promoted packet must include at least three null runs:

1. randomize or drop a declared fraction of boundary-wake labels, which should break either area scaling or conservation;
2. replace a_O with a constant temperature parameter, which should fail the same-record temperature test;
3. compute flux with an independently fitted stress record, which should be rejected as a split-record pass.

If these null runs still pass, the residual is not measuring thermodynamic closure.

Runtime Artifact

The first scaffold is:

```
node scripts/gravity/thermodynamic-residual.mjs --pretty
```

It consumes:

```
scripts/gravity/thermodynamic-residual-mock.json
```

and emits a JSON result with this shape:

Output field	Meaning
observations	computed label counts or finite-block samples, entropy change, local temperature, flux, area residual, thermodynamic residual, conservation residual, same-record checks, and weak-field gate checks
negative_controls	declared null runs and whether any passed when they should have failed
totals.max_area_residual	largest area-scaling residual across local-horizon rows
totals.max_thermodynamic_residual	largest $\widehat{\mathcal{R}}_{\text{thermo}}^{(O)}$ across rows
totals.max_conservation_residual	largest $\mathcal{R}_{E,\partial\Omega}^{(O)}$ across rows
gates	label coverage, same-record temperature, same-record flux, area scaling, thermodynamic residual, conservation, weak-field same-record, and negative-control gates
failure_code	null on pass, otherwise the first failed thermodynamic-residual gate

The mock packet is deliberately dimensionless. It uses $k_B = \hbar = c_0 = A_{\text{align}} = 1$ so the packet shape can be inspected by hand before any real Noether sea simulation supplies physical units, observer records, and boundary-wake provenance.

This runtime should not be expanded into a large fixture family unless it protects a live derivation. Its main value is to keep the theory honest at the handoff point where a candidate Noether sea record claims to supply entropy, temperature, flux, and weak-field metric recovery together. Until such a record exists, additional passing and failing fixtures are lower value than deriving the record itself.

Promotion Boundary

Passing this protocol would establish a local equilibrium recovery route for thermodynamic gravity language. It would not by itself close black-hole information release, strong-field endpoint regularity, Page-curve recovery, or cosmological horizon thermodynamics. Those remain separate validation packets that may consume the same boundary-label and same-record discipline.

Cosmology Residuals

Cosmology Shared Residual Fit Protocol

This protocol turns the shared calibration gate in [Dark Energy](#) into a first machine-checkable validation scaffold. Its purpose is narrow: test whether supernova, BAO, CMB, weak-lensing, redshift-space-distortion, BBN, and pre-BBN comparison packets can consume one shared Noether sea state record without silently replacing the state per observable family.

This is not a cosmological parameter fit and not an empirical claim. The first runtime artifact is a mock packet that fixes the object shape, residual accounting, projection-penalty semantics, gates, and failure codes that a real survey-facing packet must later populate.

Residual Object

Let

$$\mathcal{X}_{\text{cos}} = \{\text{SN}, \text{BAO}, \text{CMB}, \text{WL}, \text{RSD}, \text{BBN}, \text{PREBBN}\}$$

For each family $X \in \mathcal{X}_{\text{cos}}$, the packet records a residual vector r_X , a covariance object C_X , nuisance/calibration context ν_X , and a projection $\Pi_X \theta_{\text{sea}}$ of the shared Noether sea state record into that family. The scaffold computes

$$\mathcal{R}_X = r_X(\theta_{\text{sea}}, \nu_X)^T C_X^{-1} r_X(\theta_{\text{sea}}, \nu_X)$$

and the cross-family projection penalty

$$\mathcal{P}_{XY} = \sum_{a \in K_X \cap K_Y} w_a ((\Pi_X \theta_{\text{sea}})_a - (\Pi_Y \theta_{\text{sea}})_a)^2$$

where K_X is the set of shared comparison coordinates reported by family X , and w_a is a declared dimensionless weight. The packet-level residual is

$$\mathcal{R}_{\text{shared}} = \sum_{X \in \mathcal{X}_{\text{cos}}} \mathcal{R}_X + \lambda \sum_{X < Y} \mathcal{P}_{XY}$$

A low value of the first term alone is insufficient. The second term is the split-ontology guard: it rejects a fit that keeps each observable close to its benchmark only by assigning mutually incompatible projections of θ_{sea} .

The first empirical packet should keep the leading standard comparison objects visible inside the residual vectors:

$$\begin{aligned} r_{\text{SN/BAO}} &\supset \left(\frac{d_L^\theta(z) - d_L^{\text{obs}}(z)}{\sigma_{d_L}}, \frac{D_M^\theta(z)/r_d^\theta - (D_M/r_d)^{\text{obs}}}{\sigma_{D_M/r_d}}, \frac{H^\theta(z)r_d^\theta - (Hr_d)^{\text{obs}}}{\sigma_{Hr_d}} \right) \\ r_{\text{CMB}} &\supset \left(\frac{\Delta T_{\text{bb}}^\theta}{\epsilon_{\text{bb}}}, \frac{C_\ell^\theta - C_\ell^{\text{obs}}}{\sigma_{C_\ell}}, \frac{C_L^{\phi\phi, \theta} - C_L^{\phi\phi, \text{obs}}}{\sigma_{C_L^{\phi\phi}}} \right) \\ r_{\text{growth}} &\supset \left(\frac{f\sigma_8^\theta(z, k) - f\sigma_8^{\text{obs}}(z, k)}{\sigma_{f\sigma_8}}, \frac{P^\theta(k, z) - P^{\text{obs}}(k, z)}{\sigma_P} \right) \end{aligned}$$

and r_{BBN} should retain D/H, Y_p , lithium, η , and ΔN_{eff} rows. These are data-product coordinates, not ontology claims. They make the shared packet check luminosity distance, BAO rulers, blackbody preservation, CMB lensing, growth, and BBN yield recovery before any Noether sea state interpretation is promoted.

Redshift-facing packets must now expose the signed photon-frequency transfer row rather than treating redshift as a primitive expansion coordinate. For a line or photon family X , retain

$$r_{\nu\text{-path}} \supset \left(\frac{Z_X^\theta - Z_X^{\text{obs}}}{\sigma_Z}, \frac{Y_{X, \text{path}}^\theta - Y_{X, \text{cal}}^{\text{obs}}}{\sigma_Y}, \frac{\mathcal{R}_{\nu\text{-ex}}^\theta}{\epsilon_{\nu\text{-ex}}} \right)$$

where Z_X is the total logarithmic redshift budget, $Y_{X, \text{path}}$ is the signed path-history exchange contribution, and $Y_{X, \text{cal}}^{\text{obs}}$ is any declared calibration row such as a Sunyaev-Zeldovich or kinematic-Sunyaev-Zeldovich frequency-shift packet. This row does not add a separate cosmology gate. It prevents a shared-state fit from hiding path-frequency exchange inside $H(z)$, distance modulus, or CMB temperature calibration.

The source-mined empirical packet should retain the following benchmark families without turning them into separate gates:

Family	Required packet content	Shared-state overlap
CMB_PLANCK_LAMBDA	Planck/LAMBDA frequency-map and component-separation provenance, TT/TE/EE spectra, likelihood choice, CMB lensing map or bandpower provenance, foreground and beam nuisance context	theta_star, r_d, omega_b, omega_c, tau, A_s, n_s, CMB_lensing, blackbody
CMB_ACT	ACT DR6 high- ℓ spectra or likelihood rows, ACT lensing bandpowers, covariance, foreground model context, SZ/kSZ frequency-exchange provenance when used	CMB_lensing, small_scale_damping, foreground_context, growth_projection, frequency_exchange
BAO_DESI	DESI tracer label, effective redshift, isotropic or anisotropic BAO vector, covariance, likelihood or chain provenance	r_d, D_M, D_H, D_V, H_eff, theta_acoustic
SN_SH0ES_PANTHEON	Pantheon+ light-curve and covariance provenance, redshift convention, calibration/standardization context, Cepheid/SN ladder anchor context, local H_0 row when used	D_L, H_eff_ladder, clock_endpoint, path_history, frequency_exchange, calibration_context
WL_RSD_DES	DES weak-lensing/clustering data vector, shear calibration, photo- z calibration, covariance, DESI RSD rows when present	S_8, f_sigma_8, CMB_lensing, growth_response, noether_sea_coupling
EUCLID_PUBLIC	Public release identifier, image/catalogue/mask/photo- z readiness products, covariance readiness note	mask_context, photo_z_context, shape_context, future_growth_projection

As of 2026-05-19, EUCLID_PUBLIC is a readiness row, not a cosmology-constraint row. A packet may use Euclid Q1-style products to test mask, catalogue, image, spectroscopy, and photo- z bookkeeping, but it must not count Euclid as a successful weak-lensing or clustering cosmology residual until a public cosmology release supplies the relevant data vector and covariance.

For empirical packets, the BAO row should use the explicit anisotropic/isotropic vector

$$\mathbf{r}_{\text{BAO},i} = \mathbf{C}_{\text{BAO},i}^{-1/2} \left[\begin{pmatrix} D_M^\theta(z_i)/r_d^\theta \\ D_H^\theta(z_i)/r_d^\theta \\ D_V^\theta(z_i)/r_d^\theta \end{pmatrix}_{\text{kept}} - \begin{pmatrix} (D_M/r_d)_i^{\text{obs}} \\ (D_H/r_d)_i^{\text{obs}} \\ (D_V/r_d)_i^{\text{obs}} \end{pmatrix}_{\text{kept}} \right]$$

where kept means the subset reported by the survey bin. This avoids pretending that isotropic BAO bins contain independent radial and transverse information. The SN/local-ladder row should analogously keep the distance-modulus and local-slope rows separate:

$$\mathbf{r}_{\text{SN}/H_0} = \left(\mathbf{C}_\mu^{-1/2} [\boldsymbol{\mu}^\theta - \boldsymbol{\mu}^{\text{obs}}], \frac{H_{\text{eff,ladder}}^\theta - H_{0,\text{ladder}}^{\text{obs}}}{\sigma_{H_0}}, \frac{\Delta_{\text{cal}}^\theta}{\sigma_{\text{cal}}} \right)$$

The CMB row should preserve spectra and lensing as separate but overlapping checks:

$$\mathbf{r}_{\text{CMB}} = \left(\mathbf{C}_\ell^{-1/2} [\mathbf{C}_{\ell,\text{TTTEEE}}^\theta - \mathbf{C}_{\ell,\text{TTTEEE}}^{\text{obs}}], \mathbf{C}_{\phi\phi}^{-1/2} [\mathbf{C}_L^{\phi\phi,\theta} - \mathbf{C}_L^{\phi\phi,\text{obs}}], \frac{\theta_*^\theta - \theta_*^{\text{obs}}}{\sigma_{\theta_*}}, \frac{\Delta T_{\text{bb}}^\theta}{\epsilon_{\text{bb}}} \right)$$

The overlap key CMB_lensing must appear in both CMB and growth-facing projections whenever lensing is used. Otherwise a packet can accidentally fit CMB spectra with one projection and weak-lensing or clustering with another, which is exactly the split-ontology failure this protocol is meant to catch.

Dark-sector comparison packets should also retain the linear/nonlinear split exposed by scalar-fluid and MOND-like hybrid models:

$$r_{\text{DM,split}} \supset \left(\frac{w_{\text{lin}}^\theta - w_{\text{lin}}^{\text{CDM}}}{\sigma_w}, \frac{(c_{s,\text{lin}}^2)^\theta - (c_s^2)^{\text{CDM}}}{\sigma_{c_s^2}}, \frac{v_c^\theta(r, E_{\text{gal}}) - v_c^{\text{obs}}(r, E_{\text{gal}})}{\sigma_{v_c}}, \frac{\Delta_{\text{BTFR}}^\theta(M_b, v_f, E_{\text{gal}})}{\sigma_{\text{BTFR}}}, \frac{\text{RAR}^\theta(g_{\text{bar}}, E_{\text{gal}}) - \text{RAR}^{\text{obs}}(g_{\text{bar}})}{\sigma_{\text{RAR}}}, \frac{a_*^\theta(E) - a_*^{\text{obs}}(E)}{\sigma_{a_*}}, \frac{f_*^\theta}{f_*^{\text{obs}}} \right)$$

Here w_{lin} and $c_{s,\text{lin}}^2$ are comparison coordinates for CDM-like linear loading, while $v_c(r)$, Δ_{BTFR} , RAR , $a_*(E)$, and $f_*(E)$ are nonlinear acceleration-response coordinates. A dimensionless BTFR residual can be recorded as

$$\Delta_{\text{BTFR}}^\theta \equiv \frac{G_N M_b^{\text{obs}} a_*^\theta(E_{\text{gal}})}{(v_f^\theta)^4} - 1$$

with v_f the retained flat-curve velocity and M_b the retained baryonic mass. The environment label E is not a new ontology coordinate; it is the observable context carried in ν_X . For these rows it should include at least M_{halo} , z_{vir} , σ_v , T_{eff} , the baryon profile, and, for mergers, the declared ratio v_{inf}/c_s when the comparison template supplies a sound-speed coordinate. The low-acceleration galaxy comparison may be expressed as

$$g_{\text{obs}}^\theta(r, E_{\text{gal}}) = g_{\text{bar}}(r) + g_{\text{med}}^\theta(r, E_{\text{gal}})$$

where g_{med}^θ is only the Noether sea response projection being tested against a MOND-like comparison residual. To make the galaxy-vs-cluster split measurable, the same packet should evaluate $a_*(E)$ and $f_*(E)$ at both E_{gal} and E_{cl} . Passing the galaxy rotation-curve, BTFR, and RAR rows while failing the cluster rows below is not promotable as a shared-state success. These rows are not a request to add a new fundamental scalar-fluid ontology. Their purpose is to prevent a packet from fitting CMB and matter power data with one effective dark component while fitting galaxy, cluster, and merger accelerations with a separately tuned Noether sea law.

For cluster-facing rows, include the hydrostatic/lensing equality packet

$$r_{\text{cl}} \supset \left(\frac{T_{\text{ICM}}^\theta(r) - T_{\text{ICM}}^{\text{obs}}(r)}{\sigma_T}, \frac{P_{\text{SZ}}^\theta(r) - P_{\text{SZ}}^{\text{obs}}(r)}{\sigma_P}, \frac{\Phi_{\text{lens}}^\theta(r) - \Phi_{\text{lens}}^{\text{obs}}(r)}{\sigma_{\Phi_{\text{lens}}}}, \frac{\Phi_{\text{dyn}}^\theta(r) - \Phi_{\text{dyn}}^{\text{obs}}(r)}{\sigma_{\Phi_{\text{dyn}}}}, \frac{\gamma_{\text{eff}}^\theta(r) - 1}{\sigma_\gamma}, \frac{d_{\text{lens-gal}}^\theta - d_{\text{lens-gal}}^{\text{obs}}}{\sigma_{d,\text{lg}}}, \frac{d_{\text{lens-gas}}^\theta - d_{\text{lens-gas}}^{\text{obs}}}{\sigma_{d,\text{lgas}}} \right)$$

This row is a success marker under the existing shared-state gate, not a new standalone gate. It records whether the same Noether sea state packet can recover cluster gas temperature, SZ pressure, lensing potential, dynamical potential, and Bullet-like lensing/galaxy/gas peak separation without changing the acceleration law between observables.

Merger-facing rows may be attached to the same cluster or dark-sector observable family when the packet claims regime-dependent behavior:

$$r_{\text{merge}} \supset \left(\frac{t_{\text{merge}}^\theta(v_{\text{inf}}/c_s) - t_{\text{merge}}^{\text{obs}}}{\sigma_t}, \frac{\Delta_{\text{fric}}^\theta(v_{\text{inf}}/c_s) - \Delta_{\text{fric}}^{\text{obs}}}{\sigma_{\text{fric}}}, \frac{\mathcal{I}_{\text{int}}^\theta(v_{\text{inf}}/c_s) - \mathcal{I}_{\text{int}}^{\text{obs}}}{\sigma_{\mathcal{I}}}, \frac{N_{\text{vort}}^\theta(R) - N_{\text{vort}}^{\text{obs}}(R)}{\sigma_N} \right)$$

The ratio v_{inf}/c_s distinguishes low-dissipation pass-through encounters from high-dissipation encounters in comparison templates that provide c_s . The coordinate $\mathcal{I}_{\text{int}}$ is a declared shell or interference-morphology statistic for high-relative-speed mergers, and $N_{\text{vort}}(R)$ is included only when the comparison template predicts vortex-like substructure measurable through lensing over projected radius R . Cold-atom or other laboratory analogue simulations can supply provenance for these dimensionless template variables, but visual analogy is not a substitute for astronomical residual rows under the shared-state packet.

Packet Schema

The runtime packet should preserve this shape even when a later empirical packet replaces the mock values:

Field	Required content	Promotion role
metadata	run identifier, source commit when available, input provenance, fit family, and declared comparison level	makes the packet reproducible
required_families	required observable families, defaulting to SN, BA0, CMB, WL, RSD, BBN, and PRE_BBN	prevents cherry-picking a subset of cosmology constraints
theta_sea	shared dimensionless state record used by all projections	names the single Noether sea state candidate under test
observables	one row per family with residual vector, covariance, nuisance/calibration note when available, and projection coordinates	supplies $\mathcal{R}_X$ and $\Pi_X \theta_{\text{sea}}$
projection_weights	dimensionless weights w_a for common projection coordinates	makes the split penalty explicit rather than rhetorical
lambda	nonnegative coefficient multiplying the projection penalty	controls how strongly shared-state incompatibility is penalized
thresholds	predeclared maxima for ordinary residuals, raw projection penalty, shared residual, and projection overlap	prevents post-fit gate selection
gates	pass/fail records for coverage, residual total, projection penalty, projection overlap, and total shared residual	turns the comparison into an auditable decision surface
failure_code	null on pass, otherwise the first failed gate	gives follow-up work a stable repair target

The current mock packet uses normalized comparison coordinates such as `H_norm`, `w_eff`, `n`, `chi_sea`, `G_growth`, `Y_BBN`, `Delta_N_eff`, `lambda_fs`, `Omega_GW`, `Z_total`, `Y_path`, and `frequency_exchange_residual`. These are not new ontology. They are dimensionless placeholders for observer-level expansion, equation-of-state, normalized Noether braid density, Noether sea delay, growth-response, BBN-yield, relativistic-species, free-streaming, stochastic-gravitational-wave, total redshift-budget, path-frequency-transfer, and exchange-ledger comparison channels.

Pre-BBN Branch Packet

The `PRE_BBN` row is the runtime version of the comparison gate defined in [Inflation Model](#), [BBN Constraints](#), [Structure Formation](#), and [Gravitational Waves](#). It represents one declared branch X per packet. Multiple candidate branches should be compared by running separate packets or by building an explicitly documented aggregate row, not by hiding several branches inside one unlabeled residual.

The pre-BBN residual vector should preserve the observable/data-product split:

$$r_{\text{PREBBN}} = \left(\frac{\|\Delta \mathbf{Y}_{\text{BBN}}^X\|}{\epsilon_{\text{BBN}}}, \frac{\|\Delta C_\ell^X\|}{\epsilon_{\text{CMB}}}, \frac{\|\Delta P_X(k, z)\|}{\epsilon_{\text{growth}}}, \sup_f \frac{\Omega_{\text{GW}}^X(f)}{\Omega_{\text{GW}}^{\text{max}}(f)} \right)$$

The projection keys should include the ordinary shared cosmology coordinates plus branch-facing coordinates such as `Delta_N_eff`, `lambda_fs`, and `Omega_GW`. The packet passes this subgate only when the ordinary residual $\mathcal{R}_{\text{PREBBN}}$ is small and the projection penalty shows that the same θ_{sea} is being consumed by BBN, CMB, growth, and gravitational-wave comparisons.

Frame-Split Measurement Recipe

The `cosmology.frame_split` witness is the directional subgate for the same shared-state problem. It asks whether the rest-frame correction used for CMB inference can coexist with matter dipoles, supernova residual directionality, BAO anisotropy, and local H_0 scatter without giving each family its own hidden frame.

The required frame families are

$$\mathcal{F}_{\text{frame}} = \{\text{CMB}, \text{MD}, \text{SN}, \text{BAO}, H_0\}$$

where MD denotes matter-dipole catalogues such as radio, infrared, quasar, or galaxy-count samples. Each row must report a measured three-vector $\mathbf{y}_i$, an expected three-vector $\mathbf{m}_i(\theta_{\text{frame}})$ from the declared common frame model, a covariance object C_i , calibration or mask context ν_i , and a projection $\Pi_i \theta_{\text{frame}}$ onto shared frame coordinates.

The context ν_i must distinguish observational provenance from physical residuals. At minimum it should identify the sky mask or footprint, foreground or component-separation recipe when relevant, beam or transfer-function correction, redshift-bin and selection function, standardization or calibration model, covariance construction, and any simulation, mock-catalogue, or machine-learning training source used to estimate significance. These entries do not add another cosmology gate; they prevent a frame residual from being promoted when the mismatch is actually a reduction-pipeline or training-prior artifact.

The preprocessing rules are:

- CMB: $\mathbf{y}_{\text{CMB}} = \mathbf{D}_{\text{CMB}}$ and $\mathbf{m}_{\text{CMB}}$ is the same dipole vector in the declared coordinate convention.
- Matter dipoles: for catalogue X , $\mathbf{y}_{\text{MD},X} = \mathbf{D}_X$ and

$$\mathbf{m}_{\text{MD},X} = K_X(\alpha_X, x_X) \mathbf{D}_{\text{CMB}} + \mathbf{F}_X(\theta_{\text{frame}}, \nu_X)$$

where K_X is the catalogue kinematic amplification factor and $\mathbf{F}_X$ is the allowed non-kinematic directional residual from the shared frame state and survey context.

- Supernovae: $\mathbf{y}_{\text{SN}}(z_b)$ is the fitted distance-modulus dipole in redshift bin z_b , after standardization and host-environment bookkeeping; $\mathbf{m}_{\text{SN}}(z_b)$ is the corresponding shared-frame prediction.
- BAO: $\mathbf{y}_{\text{BAO}}(z_b)$ is the anisotropic BAO-scale dipole or lowest retained directional harmonic in bin z_b ; $\mathbf{m}_{\text{BAO}}(z_b)$ is the shared-frame prediction in the same basis.
- Local H_0 : $\mathbf{y}_{H_0}(z_b)$ is the directional local-ladder or low-redshift inferred- H scatter vector; $\mathbf{m}_{H_0}(z_b)$ is the shared-frame prediction after the same peculiar-velocity and environment cuts.

For a packet of rows $i \in I_{\text{frame}}$, the directional residual is

$$\mathcal{Q}_{\text{frame}} = \sum_{i \in I_{\text{frame}}} (\mathbf{y}_i - \mathbf{m}_i)^T C_i^{-1} (\mathbf{y}_i - \mathbf{m}_i)$$

The frame-projection penalty is

$$\mathcal{P}_{\text{frame}} = \sum_{i < j} \sum_{a \in K_i \cap K_j} w_a [(\Pi_i \theta_{\text{frame}})_a - (\Pi_j \theta_{\text{frame}})_a]^2$$

and the combined frame score is

$$\mathcal{R}_{\text{frame}} = \mathcal{Q}_{\text{frame}} + \lambda_{\text{frame}} \mathcal{P}_{\text{frame}}$$

The packet also records a direction check for every nonzero row,

$$\alpha_i = \cos^{-1} \left(\frac{\mathbf{y}_i \cdot \mathbf{m}_i}{\|\mathbf{y}_i\| \|\mathbf{m}_i\|} \right)$$

Tolerances must be declared before fitting: maximum $\mathcal{Q}_{\text{frame}}$, maximum $\mathcal{P}_{\text{frame}}$, maximum $\mathcal{R}_{\text{frame}}$, minimum shared projection-key overlap, and maximum allowed α_i for nonzero vectors. These tolerances are not universal constants; they belong to the survey packet, covariance construction, redshift binning, and systematics budget.

The falsifiers are concrete:

Failure code	Meaning
frame-split-coverage-open	At least one required family from $\mathcal{F}_{\text{frame}}$ is absent.
frame-split-residual-open	The directional residual total exceeds the declared tolerance.
frame-split-projection-open	Families can fit their own vectors only by using incompatible frame-state projections.

Failure code	Meaning
frame-split-projection-overlap-open	The packet does not share enough projection coordinates to test a common frame.
frame-split-angle-open	A measured vector points too far away from its expected shared-frame vector.
frame-split-shared-open	The combined residual-plus-projection score exceeds tolerance.

Any of these failures activates the witness code `cosmology.frame_split`. Passing the mock gate means only that the packet shape is coherent; a real packet must replace the mock vectors with survey-derived dipoles, covariance matrices, redshift-bin definitions, and nuisance records.

Runtime Artifact

The first scaffold is:

```
node scripts/cosmology/shared-residual-fit.mjs --pretty
```

It consumes:

```
scripts/cosmology/shared-residual-mock.json
```

and emits a JSON result with this shape:

Output field	Meaning
residual_terms	computed $\mathcal{R}_X$ for each observable family
projection_penalties	all pairwise $\mathcal{P}_{XY}$ terms, including shared keys and per-key contributions
totals.observable_residual	$\sum_X \mathcal{R}_X$
totals.projection_penalty_raw	$\sum_{X<Y} \mathcal{P}_{XY}$
totals.projection_penalty_weighted	$\lambda \sum_{X<Y} \mathcal{P}_{XY}$
totals.shared_residual	full $\mathcal{R}_{\text{shared}}$
gates	coverage, residual, projection, overlap, and total shared-residual pass/fail records
failure_code	observable-coverage-open, residual-total-open, projection-penalty-open, projection-overlap-open, shared-residual-open, or null
frame_split	optional directional frame-consistency result with vector rows, projection penalties, gates, and <code>cosmology.frame_split</code> witness status

The mock packet is deliberately small enough to inspect by hand. A real packet should replace the dimensionless residual entries with survey-derived residual vectors and covariance matrices, but it should keep the same gate shape unless this protocol is explicitly revised.

Acceptance Boundary

Passing the mock packet means only that the scaffold computes the intended residual and gate structure. It does not validate dark energy, H_0 , S_8 , BBN, CMB, or growth claims.

A real shared-state packet becomes promotable only if:

1. every required observable family is present exactly once;
2. residual vectors and covariance models are stated before fitting;
3. $\Pi_X \theta_{\text{sea}}$ projections share enough coordinates to test compatibility;
4. ordinary residuals stay inside declared tolerance;
5. the projection penalty stays inside declared tolerance;
6. any included `frame_split` packet passes coverage, residual, projection, angle, and shared-score gates;
7. the same θ_{sea} also remains compatible with the cosmology sector predicate in [Failure Criteria](#).

Failure is informative. If the ordinary residual passes but the projection penalty fails, the candidate has fit the data products while splitting the Noether sea state record. If the projection penalty passes but an observable residual fails, the shared state is coherent but not yet accurate. If coverage fails, the packet is not a cosmology closure artifact.

Redshift-Budget Toy Model

This protocol documents the first redshift-budget simulation fixture for the cosmology branch. The fixture is a bookkeeping replay of the factorized redshift record in [Expansion Mechanism](#), not an empirical distance-ladder fit.

Its purpose is narrow: verify that endpoint cadence, source-branch state, launch geometry, and Noether sea path-history remain separable in a machine-readable packet before any survey-facing cosmology comparison is attempted. The current packet also exposes the continuity-

disciplined path-rate law, so source loading, equilibration, frequency-space current, flow divergence, and anisotropic response are not hidden as unrelated fitted terms.

Runtime Artifact

Run the default mock packet with:

```
node scripts/cosmology/redshift-budget-toy-model.mjs --pretty
```

The script consumes:

```
scripts/cosmology/redshift-budget-mock.json
```

and emits one result row per scenario. The packet is deliberately dimensionless except for declared line frequencies, Euclidean path distance in megaparsecs, and the comparison constants c_0 and h .

Replay Equation

For a line family X , the path record is divided into segments of length Δs_j . The propagation bookkeeping variable starts at

$$Y_{X,0} = 0$$

and advances by

$$Y_{X,j+1} = Y_{X,j} + \alpha_{\text{prop},X,j} \Delta s_j$$

The fixture then reconstructs the logarithmic redshift budget

$$Z_X \equiv \ln(1 + z_X) = \ln \Gamma_{N,E} - \ln \Gamma_{N,R} + Y_{X,N} - \ln B_X(E) - \ln D_v$$

The observed receiver-facing frequency and photon energy are

$$\nu_{\text{obs},X} = \nu_{X,0} \exp(-Z_X), \quad E_{\text{obs},X} = h\nu_{\text{obs},X}$$

This is not an untracked photon-energy loss model. $Y_{X,N}$ is the path-history phase-cadence stretch left after endpoint cadence, source-branch shift, and launch geometry have been declared.

The path-history term is signed. A positive increment in Y_X is a redward frequency depletion relative to the clean emitted line, while a negative increment is a blueward frequency boost. For a segment-level exchange row,

$$\Delta Y_{X,j}^{\text{ex}} = -\ln \frac{\nu_{X,j}^+}{\nu_{X,j}^-}$$

with $\nu_{X,j}^-$ and $\nu_{X,j}^+$ measured in the same local comparison convention before and after the exchange. Sunyaev-Zeldovich-like mock rows should therefore be represented as signed exchange events rather than as a new expansion variable: a hot or coherently moving intervening medium may produce $\Delta Y_{X,j}^{\text{ex}} < 0$, while a lower-energy absorbing or relaxing segment may produce $\Delta Y_{X,j}^{\text{ex}} > 0$.

Each exchange row should also carry the local energy residual

$$R_{\nu\text{-ex},j} = \frac{\left| h(\nu_{X,j}^+ - \nu_{X,j}^-) + \Delta E_{\text{med},j} + \Delta E_{\text{recoil},j} + \Delta E_{\text{rem},j} \right|}{\epsilon_E}$$

The signs of the ΔE terms are ledger signs, not assumptions about the outcome. A photon boost is allowed only when the intervening medium or target record supplies the energy; a photon depletion is allowed only when the lost photon energy is routed into a named medium, recoil, remnant, or thermalization row.

For cosmology-facing packets, the same replay should expose whether the redshift channel also supplies the standard time-dilation and flux factors. The comparison target is

$$\frac{\Delta t_{\text{obs}}}{\Delta t_{\text{emit}}} = 1 + z_X, \quad F = \frac{L}{4\pi D_A^2 (1 + z_X)^2}, \quad d_L = (1 + z_X)^2 D_A$$

These are observer-level distance-ladder diagnostics. A path law that shifts line frequencies but does not dilate packet cadence, or that loses flux without the two redshift factors and angular-distance reciprocity, is not an acceptable cosmological redshift replacement.

Input Packet

Each scenario supplies:

Field	Meaning
line_family	spectral family whose reference frequency is replayed
comparison_line_family	optional clean comparison family used for chromaticity diagnostics
distance_mpc	corrected Euclidean path length used for the local transfer slope
B_X_E	source-branch factor $B_X(E)$
D_v	launch or relative-motion factor D_v
Gamma_N_E	emitter endpoint Noether sea cadence factor $\Gamma_{N,E}$
Gamma_N_R	receiver endpoint Noether sea cadence factor $\Gamma_{N,R}$
endpoint_records	optional endpoint records from which $\Gamma_{N,E}$ and $\Gamma_{N,R}$ are extracted
launch_record	optional source/receiver velocity record from which D_v is extracted
segments	path segments carrying Δs_j and propagation coefficients
continuity_transport_by_line	optional segment-level continuity packet for $\mathbf{p}_X \cdot D_\gamma \theta_{\text{sea}}, C_N[f_N]$, flow divergence, and anisotropic response
transport_terms_by_line	optional segment-level decomposition of $\alpha_{\text{prop},X}$ into named source, relaxation, or perturbation terms
transport_terms_cadence_by_line	optional cadence-channel version of the same decomposition for time-dilation checks
dark_energy_transport_by_line	optional coefficient packet that computes $\alpha_{\text{prop},X}^{\text{DE}}$ from a declared λ_X row and $\mathbf{q}_{\text{DE}}$ record
frequency_exchange_events_by_line	optional signed exchange rows with before/after photon frequency, medium energy change, recoil/remnant terms, and $R_{\nu\text{-ex}}$

Segment records may provide separate coefficient arrays for frequency, packet cadence, line-family comparison, and image-bundle beams. This is intentional: the first validation target is to expose when those channels agree and when they split.

Endpoint records may declare Γ_N directly or provide a cadence measurement from which the same factor is computed:

$$\Gamma_N = \frac{T_N}{T_{N0}} = \frac{\Omega_{N0}}{\Omega_N}$$

In JSON, this is supplied as Gamma_N, T_N_over_T_N0, Omega_N_over_Omega_N0, or the weak-field proxy Phi_N_over_c0_squared, for which the fixture uses $\Gamma_N \approx 1 - \Phi_N/c_0^2$. Scalar Gamma_N_E and Gamma_N_R values remain valid fallbacks for older or hand-written scenarios.

Launch records compute the low-speed source/receiver geometry factor from the radial endpoint velocity,

$$\beta_r = \frac{(\mathbf{v}_R - \mathbf{v}_E) \cdot \hat{\mathbf{k}}}{c_0}, \quad D_v = \sqrt{\frac{1 - \beta_r}{1 + \beta_r}}$$

where $\hat{\mathbf{k}}$ points from emitter to receiver and $v_r > 0$ means the endpoint separation is increasing. A packet may provide beta_r, radial_velocity_km_s, or the triple emitter_velocity_km_s, receiver_velocity_km_s, and line_of_sight. Scalar D_v remains the fallback.

The continuity-transport extension uses the segment packet

$$\alpha_{\text{prop},X,j} = \mathbf{p}_X \cdot \mathbf{d}_{\theta,j} + p_{\nu,X} \frac{S_{\text{BH},j} + S_{\text{GW},j} - R_{\text{eq},j} - \partial_\nu J_{\nu,j}}{f_{N,j} + \epsilon_f} + p_{u,X} \delta_{u,j} + p_{\sigma,X} \sigma_{X,j} + \mathcal{R}_{\text{coh},X,j}$$

In JSON, continuity_transport_by_line supplies p_theta_row, D_gamma_theta, p_nu, f_N, S_BH, S_GW, R_eq, partial_nu_J_nu, p_u, div_u_sea, p_sigma, sigma_projection, and R_coh as needed. The fixture logs the resulting pieces as continuity.theta_gradient, continuity.cadence_residual, continuity.flow_divergence, continuity.anisotropic_response, and continuity.coherence_residue. Legacy named transport_terms_by_line values are still accepted as explicit additions, but a promotable transport scenario should prefer the continuity packet whenever it is claiming to test Noether sea equilibrium transport.

Coefficient-Row Validation Notes

The fixture now reads each scenario as a restriction of the same coefficient-row map, not as a separate explanation for each redshift class. The endpoint extraction tests the cadence row

$$\mathbf{b}_N = (b_n, b_\chi, b_\lambda, 1, b_R)$$

with the weak static condition $b_n a_n + b_\chi a_\chi + b_\lambda a_\lambda + b_R a_R = 1$, or $b_n a_n + b_\chi (1 + \gamma_{\text{eff}}) + b_\lambda a_\lambda + b_R a_R = 1$ when the shared clock/signal delay closure is imposed. This fixture does not determine the individual endpoint coefficients; it checks whether endpoint records are replayed as endpoint cadence rather than hidden inside propagation or source factors.

The launch extraction tests the separate relative-motion term. In a homogeneous record with no source-branch or path-history contribution, the replay must reduce to

$$Z_X = -\ln D_v, \quad Y_{X,N} = 0$$

The scalar launch fallback and `launch_record` extractor therefore validate the sign and ownership of the motion term. A scenario fails the coefficient-row reading if it needs a nonzero propagation packet to recover a clean relative-motion redshift.

The continuity packet tests only the path row

$$(\mathbf{p}_X, p_{\nu,X}, p_{u,X}, p_{\sigma,X})$$

After endpoint, source-branch, and launch corrections have been subtracted, the residual must be

$$Z_{\text{prop},X} = \sum_j [\mathbf{p}_X \cdot \mathbf{d}_{\theta,j} + p_{\nu,X} \mathcal{C}_{N,j} + p_{u,X} \delta_{u,j} + p_{\sigma,X} \sigma_{X,j} + \mathcal{R}_{\text{coh},X,j}] \Delta s_j$$

The current mock rows constrain products of coefficients with declared segment records; they do not by themselves fix $\mathbf{p}_X$, $p_{\nu,X}$, $p_{u,X}$, or $p_{\sigma,X}$ individually. Those freedoms are falsified by the diagnostics already exposed here: chromaticity residuals, image-bundle variance, time-dilation residuals, nonzero laboratory residuals, or a need to replace the continuity packet with unrelated named terms.

The dark-energy coefficient extension uses

$$\alpha_{\text{prop},X}^{\text{DE}} = \frac{1}{c_\gamma} (\lambda_\rho^X q_\rho + \lambda_w^X q_w + \lambda_{\text{sea}}^X q_{\text{sea}} + \lambda_{\text{BH}}^X q_{\text{BH}})$$

In JSON, `lambda_row` supplies the four dimensionless coefficients and `q_DE_per_s` supplies the corresponding rate entries in inverse seconds. The script divides by the declared photon-channel speed, using `c_gamma_km_s` when present and otherwise `c0_km_s`, to convert the result into a path coefficient in Mpc^{-1} . A packet may instead supply `q_DE_per_mpc` when the rate has already been converted into path units.

Output Diagnostics

The fixture reports:

Output field	Meaning
<code>diagnostics.Z_prop_X</code>	corrected propagation residual $Y_{X,N}$
<code>diagnostics.Z_total_X</code>	total reconstructed logarithmic redshift Z_X
<code>diagnostics.redshift_z</code>	observed redshift $z_X = \exp(Z_X) - 1$
<code>diagnostics.inferred_H_eff_km_s_Mpc</code>	short-path slope proxy $c_0 Y_{X,N}/D$
<code>diagnostics.chromaticity_residual</code>	$ Y_{X,N} - Y_{Y,N} $ for two clean lines
<code>diagnostics.image_bundle_variance</code>	variance of beam-specific Y values
<code>diagnostics.time_dilation_residual</code>	split between frequency and packet-cadence propagation
<code>diagnostics.luminosity_factor_residual</code>	mismatch between the replayed flux factor and the expected $(1+z)^2$ distance-ladder factor
<code>diagnostics.distance_reciprocity_residual</code>	mismatch in the observer-level $d_L = (1+z)^2 D_A$ relation
<code>diagnostics.frequency_exchange_residual</code>	maximum or norm of the signed exchange energy-ledger residuals $R_{\nu\text{-ex},j}$
<code>diagnostics.path_transfer_sign</code>	whether the corrected path term is net redward, net blueward, or balanced after endpoint, source, and launch terms are removed
<code>observables.nu_obs_hz</code>	receiver-facing observed frequency
<code>observables.E_obs_j</code>	receiver-facing photon energy
<code>component_logs</code>	endpoint, propagation, source-branch, and launch contributions to Z_X
<code>transport_term_logs</code>	integrated named contributions to $Y_{X,N}$ for frequency and cadence channels
<code>extraction_logs</code>	endpoint and launch extraction methods, including scalar fallback versus record-derived values

The diagnostics are not pass/fail cosmology claims. They are failure witnesses for the factorization itself.

Expected Mock Behavior

The default mock packet has six hand-checkable rows.

Scenario	Expected behavior
<code>clean_laboratory_line</code>	All factors are unity or zero, so $Z_{\text{prop},X} = 0$, $z = 0$, and $H_{\text{eff}} = 0$.
<code>endpoint_launch_record_extraction</code>	Endpoint and launch factors are extracted from records: $\Gamma_{N,E} = 1/0.995$, $\Gamma_{N,R} = 1$, and $D_e \approx 0.998501$. The path residual remains $Z_{\text{prop},X} = 0$, so the total redshift comes only from endpoint cadence plus launch geometry.
<code>clean_galaxy_path</code>	Path history dominates the corrected residual: $Z_{\text{prop},X} = 0.02812$, giving a local slope near $70.25 \text{ km s}^{-1} \text{ Mpc}^{-1}$ while chromaticity, beam variance, and time-dilation residuals remain small.

Scenario	Expected behavior
equilibrium_transport_smooth_h_step	The continuity packet supplies $Z_{\text{prop},X} = 0.02800$, giving a local slope near $69.95 \text{ km s}^{-1} \text{ Mpc}^{-1}$ with source and gravitational-wave contributions logged inside the source-balanced cadence residual.
dark_energy_coefficient_packet	The propagation coefficient is computed from <code>lambda_row</code> and <code>q_DE_per_s</code> , giving $Z_{\text{prop},X} \approx 0.02788$ and a local slope near $69.66 \text{ km s}^{-1} \text{ Mpc}^{-1}$.
strong_source_near_black_hole	Endpoint cadence and source-branch terms dominate the total redshift. The path residual is only $Z_{\text{prop},X} = 0.00201$, so a propagation-only distance estimate would be invalid without the endpoint and source corrections.

These numbers are fixture expectations only. They validate arithmetic, packet shape, and diagnostic separation, not an observed cosmological model.

Failure Reading

The first failure modes are concrete:

Diagnostic pattern	Meaning
large chromaticity_residual on clean lines	the path law is behaving like a line-dependent loss process rather than a shared transport law
large image_bundle_variance	neighboring beams accumulate incompatible Y values, which threatens image sharpness
large time_dilation_residual	frequency shift and packet-cadence stretch no longer share one propagation record
large dark_energy.* dominance with failed chromaticity or cadence checks	the dark-energy handoff is acting like a fitted redshift source rather than a shared Noether sea transport coefficient
continuity packet replaced by unrelated named source terms	the run is not testing the no-case-switch transport law because $\partial_\nu J_\nu$, source loading, equilibration, and flow response have been separated into free fit parameters
large total Z_X with small $Z_{\text{prop},X}$	endpoint cadence, source branch, or launch geometry dominate, so distance cannot be inferred from propagation alone
nonzero laboratory residual after local corrections	the factorization leaks local calibration or source-branch effects into the propagation channel

A promotable redshift-distance packet must keep these diagnostics attached to the same Noether sea state record that later feeds supernova, BAO, CMB, growth, and local-ladder comparisons.

Closure Scorecard

This chapter is the reusable assessment surface for closure progress across the theory stack. Its purpose is to keep evaluation criteria stable from one scoring cycle to the next so that changes in score reflect actual progress or regression rather than drift in the assessment lens itself.

It is meant to be used with [Failure Criteria](#), [Validation Protocols](#), [No-Go Theorems](#), and [Parameter Ledger](#).

Reusable Assessment Prompt

Use this prompt for each new assessment cycle:

```

Perform a full validated-closure assessment of theory, mathematics, and geometry of modern physics vs. architrino theory.
Requirements:
1) Do a full read of all markdown documents in content/markdown/aaa (including subdirectories).
2) Evaluate each existing scorecard category in closure-scorecard.md on a 0-100 scale.
3) Use the validated-closure lens: certified equations, derivation depth, coefficient recovery, parameter determination, empirical precision, geometry/dynamics consistency, unresolved placeholders, and falsification-readiness.
4) Do not let architectural coherence, explanatory logic, or ontology compensate for missing equations, missing coefficients, unfixed parameters, or unvalidated benchmark recovery.
5) Add or populate the next dated assessment column in closure-scorecard.md with raw numeric scores, placing it after existing dated AAA columns and before Δ; preserve previous assessment columns unless explicitly told to replace or remove them.
6) Recompute each Δ value as the latest dated AAA score minus max(Modern Physics Operational, Modern Physics Mechanism).
7) Recompute the TOTAL row as the weighted arithmetic mean using the Weight column; round the displayed TOTAL only after computing the weighted mean from the raw row scores.
8) Add a dated assessment-notes section for the new column, naming concrete gains, regressions, and remaining blockers. If an assessment column or note is removed, remove or rewrite stale date references that pointed to it.
9) Keep all TeX intact and preserve category definitions unless explicitly asked to revise them.
```

Scale: 0-100 (standard numeric grading scale).
Total score rule: weighted arithmetic mean using the Weight column.

Scoring Lens

The scorecard now weights highly validated mathematical closure. A high score requires not only a coherent theory route, but also explicit equations, coefficient-level derivations, parameter fixing, and contact with tested benchmark physics.

This lens scores accepted closure, not the presence of a plan for closure. Protocols, ledgers, mock packets, replay fixtures, and negative controls can raise Falsification Gates, Coverage+Interface Readiness, or adjacent readiness rows. They should raise Formula+Coefficient Recovery, Parameter+Scale Closure, or Empirical Precision+Benchmark Validation only when they produce retained branch-derived coefficients, fixed parameters, or benchmark passes under declared tolerances.

Shared-record discipline is part of the score. A result that works only after changing the branch record, Noether sea state, coefficient row, apparatus kernel, or calibration context per observable remains local; it should not be scored as cross-regime or empirical closure. Negative and no-go diagnostics can improve auditability and falsification readiness, but they do not by themselves recover target formulas, constants, or benchmark data.

Score bands:

- 90–100: equation-level closure with derived coefficients or theorems, fixed parameters where relevant, and strong empirical or formal validation.
- 70–89: validated or mathematically mature closure in a broad regime, but with known interface limits, fitted quantities, or incomplete foundational mechanism.
- 50–69: coherent formal route with substantial equations or models, but missing key derivations, coefficients, or validation passes.
- 30–49: developed architecture or proof program with major mathematical targets still open.
- 0–29: hypothesis, placeholder, or early scaffold without certified mathematical closure.

Architectural coherence and ontic logic remain explicit criteria because they matter to theory quality. They carry limited weight so that a strong **AAA** architecture can score high as architecture without inflating the validated-closure total.

Assessment Table

Modern physics columns use the same categories for the effective-theory stack (GR, QM, QED, QFT, QCD, SM, LCDM): one operational/effective score and one mechanism/foundational score. The operational column measures validated mathematical and empirical closure of the effective theories. The mechanism column measures how far the same stack supplies a unified underlying mechanism rather than a collection of successful effective descriptions.

The Δ column is computed as the latest dated **AAA** score minus $\max(\text{Modern Physics Operational}, \text{Modern Physics Mechanism})$; negative values mark current **AAA** deficits against the stronger modern-physics comparator.

Category	Weight	Description	Modern Physics Operational	Modern Physics Mechanism	AAA 2026-05-16	AAA 2026-06-20	Δ
Empirical Precision+Benchmark Validation	14	Agreement with direct observation, precision tests, benchmark experiments, simulations, and quantitative pass/fail thresholds.	98	78	20	39	-59
Formula+Coefficient Recovery	13	Explicit recovery of target formulas and coefficients: Lorentz behavior, clock/redshift laws, PPN terms, mass formulas, quantum probabilities, and Standard Model mappings.	96	78	28	48	-48
Parameter+Scale Closure	10	Determination status of constants, couplings, scales, constitutive coefficients, and renormalization or calibration freedom.	70	42	25	42	-28
Potential+Action Closure	9	Action, potential, variational, and force/acceleration closure, including whether the central dynamics derive from a stable mathematical principle.	98	86	45	71	-27
Conservation+Invariant Closure	7	Energy, momentum, angular momentum, charge, quantum-number, and symmetry-invariant closure, including no-go consistency.	98	92	50	71	-27
UV/IR+Regularization Completion	6	Ultraviolet and infrared completion quality, including cutoff dependence, singular behavior, regularization limits, horizon/singularity issues, and asymptotics.	70	35	30	49	-21
Master EOM+Local Dynamics	10	Certified closure of the core equations of motion: local field/effective equations in modern physics and delayed path-history dynamics in AAA .	96	75	60	80	-16
Cross-Regime Bridge	8	Mathematical consistency across regimes: micro to macro, quantum to classical, particle to cosmology, weak to strong gravity, and thermodynamics.	82	48	42	69	-13
Internal Constituent Dynamics	5	Detailed closure of internal constituent regimes: bound-state/composite dynamics in modern physics and nested shell braid/Noether braid dynamics in AAA .	82	50	55	78	-4

Category	Weight	Description	Modern Physics Operational	Modern Physics Mechanism	AAA 2026-05- 16	AAA 2026-06- 20	Δ
Falsification Gates	4	Explicitness and enforceability of falsification thresholds, stop conditions, validation gates, and failure criteria.	98	88	80	95	-3
Coverage+Interface Readiness	2	Coverage completeness across mathematics/geometry-relevant domains, including interface consistency and minimally developed sections.	99	82	72	96	-3
Axiom+Notation	4	Canonical symbols, definitions, and cross-chapter mathematical language consistency.	96	82	92	99	3
Theory Architecture+Ontic Logic	8	Unified theoretical architecture, explanatory parsimony, substrate logic, and avoidance of ad-hoc patching, scored separately from validated formula recovery.	58	35	96	99	41
TOTAL	100	Weighted mean across all categories.	88	67	46	65	-23

2026-06-20 Assessment Notes

The 2026-06-20 assessment records a weighted AAA score of 65 after a full read of the 163 markdown files under content/markdown/aaa. The score is concentrated in mathematical scaffolding, validation discipline, and interface coverage rather than in final recovery of observed coefficients. The corpus now has a much stronger causal-action and energy/conservation spine: the scalar causal-hit functional has a regularized theorem spine and finite-memory bounds, the energy chapter separates finite-window wake-history balances from particle-only conservation, and nested shell braid dynamics states a shared causal-closure certificate target that ties causal-root ledgers, Jacobian floors, mass response, observer exports, event ledgers, and stability rows to the same retained branch.

The score increase is deliberately limited by the validated-closure lens. Many of the strongest new artifacts are still explicitly theorem targets, mock packets, replay fixtures, or rejection diagnostics. The hydrogen Γ_N spectral scan now keeps density, Noether sea delay, scale, envelope, and core rows separate and uses a shared coefficient row, but it does not yet derive hydrogen envelope gaps, real observer frequencies, or the static response vector from the master dynamics. The cosmology shared-residual fit, Bell-family record-measure harness, radiation ledgers, massive-superposition gravity packet, and thermodynamic residual protocol improve falsification-readiness and benchmark shape, but they do not yet supply empirical joint fits or accepted branch-derived coefficients.

Formula, parameter, and empirical rows remain the main drag on the total. The corpus still lacks a single accepted native record that supplies $E_{\text{internal}}(A_0)$, $\zeta(A_0)$, M_{sea}^{ab} , Lorentz/PPN coefficients, photon-channel coefficients, Born/Bell measures, weak-mixing and CKM/PMNS values, Standard Model mass formulas, radiation benchmarks, and shared cosmology residual fits. The Parameter Ledger improves the bookkeeping of primitive constants, geometric closure targets, constitutive closure targets, CODATA benchmark rows, and null-result discipline, but most decisive symbols remain open or branch-dependent rather than fixed outputs.

Falsification and coverage now score near modern-operational levels because the corpus contains explicit sector acceptance sets, null-result residuals, failure conditions, benchmark protocols, and cross-regime packet schemas. That does not make the total near modern physics. Architecture and ontology remain very strong, but their limited score weight prevents coherence from compensating for missing derivations, missing coefficients, unfixed parameters, and unvalidated benchmark recovery.

2026-05-22 Assessment Notes

The 2026-05-22 assessment raises the weighted AAA score from 59 to 61. The increase is concentrated in notation, internal constituent dynamics, cross-regime bridge quality, and falsification discipline. It is not a coefficient-recovery jump: the central benchmark rows still lack a retained branch that recovers masses, Lorentz / PPN coefficients, photon-channel coefficients, Born/Bell measures, weak mixing, Standard Model masses, or cosmological residuals from one accepted native record.

The largest corpus-side improvement is the Noether braid taxonomy. The corpus now separates the broad neutral braid, shell braid, and nested shell braid cases; treats exact binaries as a proof assumption rather than a naming axiom; and routes dynamic exclusion-envelope geometry into a dedicated nested shell braid geometry chapter. That chapter adds a computable assembly/Noether sea interface diagnostic,

$$D_{a,X}(\mathbf{x}, t) = \frac{\|\mathcal{W}_{a,X}^{\text{locked}}(\mathbf{x}, t)\|}{\|\mathcal{W}_{a,X}^{\text{locked}}(\mathbf{x}, t)\| + \|\mathcal{W}_{\text{sea},X}^{\text{ambient}}(\mathbf{x}, t)\|}$$

with locked and ambient contributions built from the same causal-root kernel, Jacobian floors, branch records, channel projections, and ledger-derived tolerance scales. This justifies raising Axiom+Notation, Cross-Regime Bridge, Internal Constituent Dynamics, and Coverage+Interface Readiness, while keeping the claim below full closure because the interface diagnostic is still a recovery target rather than a validated medium-response theorem.

The Noether sea branch embedding also improves the master-equation bridge. Local assembly branches are now stated as retained branches inside a surrounding Noether sea state and nearby-assembly record:

$$\mathcal{R}_{\text{branch}}(B; \Theta_{\text{sea}}, \Theta_{\text{asm}}, \mathcal{H}_{\partial\Omega}) = 0$$

with the force-ledger split

$$F_i = F_{i,\text{internal}} + F_{i,\text{sea}} + F_{i,\text{asm}} + F_{i,\partial\Omega}$$

This is a concrete mathematical advance because it prevents isolated seed charts from being read as physical branch closure unless Noether sea, assembly, and boundary residuals are statused. It supports modest increases in Master EOM+Local Dynamics, Potential+Action Closure, Conservation+Invariant Closure, Parameter+Scale Closure, and UV/IR+Regularization Completion.

Executable neutral-braid diagnostics add negative evidence and sharper first-failure semantics. The current sampled octahedral root-ledger diagnostic passes the all-pairs sampled root/Jacobian screen, while the rigid zero-offset fixed-speed row is rejected by a nonzero tangential residual witness and an ordinary same-source positive-delay no-go. These artifacts improve falsification readiness and empirical/simulation discipline because they report `not_retained` rather than converting a failed seed into branch evidence. The score increase is deliberately small because sampled diagnostics, no-go witnesses for one rigid seed, and finite-mode search schemas do not yet replace an interval-certified all-pairs root ledger, action/Noether row, event ledger, stability certificate, or observer-export recovery.

The total remains far below modern operational closure for the same reason as the prior assessments. The theory stack has stronger taxonomy, residual surfaces, and fail-closed diagnostics, but not the decisive retained branch. Until a single native record supplies $E_{\text{internal}}(A_0)$, $\zeta(A_0)$, $\mathcal{M}_{\text{sea}}^{ab}$, Lorentz/PPN recovery, photon-channel recovery, quantum source measures, Standard Model mapping coefficients, and shared cosmology fits, architecture and auditability must not inflate the validated-closure total.

2026-05-19 Assessment Notes

The 2026-05-19 assessment records a weighted AAA score of 55. The gain is broad but still pre-closure: the corpus now carries more explicit proof scaffolds, branch-certificate packet schemas, CODATA benchmark discipline, Standard Model mapping targets, quantum record-measure residuals, and shared cosmology residual gates. These changes improve mathematical auditability and executable validation readiness, but they do not yet close the first accepted branch, derive the central constants, or pass precision benchmark rows.

The largest score changes come from the proof and validation surfaces. The collinear-breather and Master EOM material now contain stronger dual-mollified branch-chart, finite-certificate, fold-layer, impulse-bound, continuity, and self-map structures. The A_0 branch-certificate protocol and run protocols now specify machine-checkable residual vectors, gate semantics, artifact lists, hidden-tuning failures, and promotion boundaries. These additions justify higher scores for Master EOM+Local Dynamics, Potential+Action Closure, UV/IR+Regularization Completion, Falsification Gates, and Empirical Precision+Benchmark Validation.

Formula, parameter, and cross-regime scores also rise because the corpus now separates exact SI conventions from adjusted CODATA benchmark rows, states the high-pressure roles of α , m_p/m_e , R_∞ , particle masses, and G , and gives the hydrogen Γ_N spectral row an executable shared-row scaffold rather than a per-line fit. The electroweak, weak-mixing, CKM/PMNS, Higgs, mass-map, Noether sea, and cosmology files now expose more of the required shared-record structure across particle, atomic, gravitational, thermodynamic, and cosmological regimes.

The total remains far below modern operational closure because the decisive derivations are still open. The first certified A_0 branch has not passed; $\zeta(A_0)$, $E_{\text{internal}}(A_0)$, and $\mathcal{M}_{\text{sea}}^{ab}$ are not accepted outputs; Lorentz, PPN, redshift, and photon-channel coefficients still lack one accepted Noether sea constitutive map; Born/Bell closure still has negative controls and measure targets rather than a positive pair-provenance theorem; Standard Model mixing and mass formulas remain shared-record theorem targets; and cosmology has a shared residual scaffold but not a fit to SN, BAO, CMB, growth, BBN, and pre-BBN rows with one θ_{sea} .

2026-05-16 Assessment Notes

The 2026-05-16 assessment is rescored under the validated-closure lens. The previous AAA columns were removed because they used a softer equal-weight closure lens that allowed architecture, coverage, and auditability to dominate the total.

AAA still scores very high in Theory Architecture+Ontic Logic because the corpus has a coherent substrate-first architecture, explicit causal-wake ontology, delayed Master Equation of Motion, Noether sea bridge program, and strong cross-document logic. That score is intentionally preserved rather than diluted.

The total is much lower because the central tested-physics closures remain open. The first certified A_0 branch is still absent, $\zeta(A)$ and $E_{\text{internal}}(A)$ are not extracted for a mass map, Lorentz and PPN coefficients are not yet derived from accepted attractors, Born-rule and Bell closure remain source-measure targets, weak V-A and CKM/PMNS quantitative closure are open, cosmology lacks an empirical shared-state fit, and UV/IR completion still depends on terminal-alignment, singularity, horizon-entropy, and effective-GR recovery proofs.

Modern physics now scores higher in the operational column because the revised lens rewards validated mathematical closure: GR, QFT, QED, QCD, the Standard Model, and LCDM-era phenomenology carry many precise equations, coefficients, and benchmark tests. Its mechanism/foundational score remains lower because the inherited stack does not supply a single ontic mechanism for quantum measurement, gauge/matter origin, gravity/quantum unification, parameter values, or cosmological initial conditions.